

Lightning Book of BOLTs

The Lightning BOLT authors

Compiled by Adam Shannon

Abstract

The Lightning Book of BOLTs is a readable compilation of the Lightning Network Specifications. The BOLTs themselves are unchanged. They are grouped here in the order a node actually uses them — connect, open a channel, route a payment, settle on-chain — with short chapter introductions so you know why you’re looking at a given document.

Credit for the protocol belongs to the original BOLT authors. This book is a reading order, not a rewrite.

Contents

Introduction

How to Read This Book

Source snapshot

Overview

- BOLT #0: Introduction and Index

Connecting Nodes

- BOLT #8: Encrypted and Authenticated Transport

- BOLT #1: Base Protocol

- BOLT #9: Assigned Feature Flags

Opening and Running Channels

- BOLT #2: Peer Protocol for Channel Management

- BOLT #3: Bitcoin Transaction and Script Formats

Routing Payments

- BOLT #4: Onion Routing Protocol

- BOLT #7: P2P Node and Channel Discovery

Paying People

- BOLT #11: Invoice Protocol for Lightning Payments

- BOLT #12: Negotiation Protocol for Lightning Payments

When Channels Close

- BOLT #5: Recommendations for On-chain Transaction Handling

Finding the Network

- BOLT #10: DNS Bootstrap and Assisted Node Location

Newer Work

- Extension BOLT XX: Simple Taproot Channels

Conclusion

Introduction

Welcome to the Lightning Book of BOLTs. This is a reading companion to the [Lightning Network Specifications \(BOLTs\)](#) — the documents that Lightning implementations actually implement.

I didn't write the BOLTs. Credit belongs to the original authors and the many contributors who have argued, patched, and shipped this protocol over the years. My job here is to put those specs in an order you can actually read, and to add a little context at the start of each chapter so you know why you're looking at a given BOLT.

The official repository numbers the BOLTs 00, 01, 02, and so on. That's a fine catalog. It's a rough way to learn the protocol. This book groups them the way a node actually uses them: connect, open a channel, route a payment, settle on-chain.

Each chapter starts with a short introduction, then the BOLT text itself, unchanged. If something in a spec looks dense, that's because it is — these documents are written for implementers. The wrapping text is here so you don't have to guess which document to open next.

Thanks for picking this up. I hope it makes Lightning a little less intimidating, and that it sends you back to the real BOLTs (and to contributing upstream) when you're ready.

How to Read This Book

You don't have to read this cover to cover. The BOLTs are a specification, not a novel. Still, there is a path through them that matches how Lightning actually works, and that's the path this book follows.

Start with the overview and BOLT 00. That document is the table of contents, the glossary, and the design goals. If a later BOLT uses a word you don't recognize — `htlc_id`, `cltv_expiry`, `channel_ready` — BOLT 00 is where it was defined.

Then walk the stack. First you connect two nodes (transport, messages, features). Then you open a channel and learn the Bitcoin transactions that back it. Then you route a payment through other people's channels. Then you look at invoices and offers, which are how humans actually pay each other. Closures and on-chain recovery come after you understand a live channel. DNS bootstrap and Simple Taproot Channels are extras: useful, but they aren't the core loop.

Skip freely. If you only care about invoice format, jump to the payments chapter. If you're writing a gossip crawler, start at routing. Come back to the earlier chapters when a type or a message doesn't make sense.

One more thing: this book does not edit the BOLTs. Typos, ambiguities, and proposed changes belong in [lightning/bolts](#), not here. The git commit page near the front tells you which snapshot of the specs you're reading.

Source snapshot

This book was built from the following commit of [lightning/bolts](#). If something here disagrees with upstream, upstream wins.

```
commit 1aadb719b4007c4cea0ba6e36b08c4fb53788dee
Merge: 1528972 807aaf8
Author: Olaoluwa Osuntokun <laolu32@gmail.com>
Date: Mon Sep 21 10:24:56 2026 -0700
```

```
Merge pull request #1355 from SomberNight/202608_bolt04_total_msat_assumed
```

```
BOLT 4: clarify final onion payload must always contain total_msat
```

Overview

Lightning is a network of payment channels sitting on top of Bitcoin. Two nodes lock coins into a shared output, then update a balance off-chain by exchanging signatures. When they are done, they broadcast the latest agreed state and walk away. Payments to people you don't have a channel with hop through other nodes, each hop wrapped in an onion.

The BOLTs — Basis of Lightning Technology — are the documents that make this interoperable. If your node speaks the BOLTs, it can open channels, route payments, and gossip with any other implementation. They are not a tutorial and they are not a whitepaper. They are the contract between implementations.

BOLT 00 is the front door. It explains the goals of the protocol, the keywords (MUST, SHOULD, MAY), and a glossary of terms you'll see everywhere else. Read it once for orientation, then keep it as a dictionary.

After that we'll stop going in numerical order. A node doesn't "do BOLT 01" in isolation. It first sets up an encrypted connection, then it talks, then it negotiates features. That's the next chapter.

In this chapter

- BOLT 00 — Introduction, glossary, and design goals

BOLT #0: Introduction and Index

Welcome, friend! These Basis of Lightning Technology (BOLT) documents describe a layer-2 protocol for off-chain bitcoin transfer by mutual cooperation, relying on on-chain transactions for enforcement if necessary.

Some requirements are subtle; we have tried to highlight motivations and reasoning behind the results you see here. I'm sure we've fallen short; if you find any part confusing or wrong, please contact us and help us improve.

This is version 0.

1. [BOLT #1](#): Base Protocol
2. [BOLT #2](#): Peer Protocol for Channel Management
3. [BOLT #3](#): Bitcoin Transaction and Script Formats

4. [BOLT #4](#): Onion Routing Protocol
5. [BOLT #5](#): Recommendations for On-chain Transaction Handling
6. [BOLT #7](#): P2P Node and Channel Discovery
7. [BOLT #8](#): Encrypted and Authenticated Transport
8. [BOLT #9](#): Assigned Feature Flags
9. [BOLT #10](#): DNS Bootstrap and Assisted Node Location
10. [BOLT #11](#): Invoice Protocol for Lightning Payments
11. [BOLT #12](#): Negotiation Protocol for Lightning Payments

The Spark: A Short Introduction to Lightning

Lightning is a protocol for making fast payments with Bitcoin using a network of channels.

Channels

Lightning works by establishing *channels*: two participants create a Lightning payment channel that contains some amount of bitcoin (e.g., 0.1 bitcoin) that they've locked up on the Bitcoin network. It is spendable only with both their signatures.

Initially they each hold a bitcoin transaction that sends all the bitcoin (e.g. 0.1 bitcoin) back to one party. They can later sign a new bitcoin transaction that splits these funds differently, e.g. 0.09 bitcoin to one party, 0.01 bitcoin to the other, and invalidate the previous bitcoin transaction so it won't be spent.

See [BOLT #2: Channel Establishment](#) for more on channel establishment and [BOLT #3: Funding Transaction Output](#) for the format of the bitcoin transaction that creates the channel. See [BOLT #5: Recommendations for On-chain Transaction Handling](#) for the requirements when participants disagree or fail, and the cross-signed bitcoin transaction must be spent.

Conditional Payments

A Lightning channel only allows payment between two participants, but channels can be connected together to form a network that allows payments between all members of the network. This requires the technology of a conditional payment, which can be added to a channel, e.g. "you get 0.01 bitcoin if you reveal the secret within 6 hours". Once the recipient presents the secret, that bitcoin transaction is replaced with one lacking the conditional payment and adding the funds to that recipient's output.

See [BOLT #2: Adding an HTLC](#) for the commands a participant uses to add a conditional payment, and [BOLT #3: Commitment Transaction](#) for the complete format of the bitcoin transaction.

Forwarding

Such a conditional payment can be safely forwarded to another participant with a lower time limit, e.g. "you get 0.01 bitcoin if you reveal the secret within 5 hours". This allows channels to be chained into a network without trusting the intermediaries.

See [BOLT #2: Forwarding HTLCs](#) for details on forwarding payments, [BOLT #4: Packet Structure](#) for how payment instructions are transported.

Network Topology

To make a payment, a participant needs to know what channels it can send through. Participants tell each other about channel and node creation, and updates.

See [BOLT #7: P2P Node and Channel Discovery](#) for details on the communication protocol, and [BOLT #10: DNS Bootstrap and Assisted Node Location](#) for initial network bootstrap.

Payment Invoicing

A participant receives invoices that tell them what payments to make.

See [BOLT #11: Invoice Protocol for Lightning Payments](#) for the protocol describing the destination and purpose of a payment such that the payer can later prove successful payment.

Glossary and Terminology Guide

- **Announcement:**

- A gossip message sent between [peers](#) intended to aid the discovery of a [channel](#) or a [node](#).

- **chain_hash:**

- The uniquely identifying hash of the target blockchain (usually the genesis hash). This allows [nodes](#) to create and reference *channels* on several blockchains. Nodes are to ignore any messages that reference a chain_hash that are unknown to them. Unlike bitcoin-cli, the hash is not reversed but is used directly.

For the main chain Bitcoin blockchain, the chain_hash value MUST be (encoded in hex):

6fe28c0ab6f1b372c1a6a246ae63f74f931e8365e15a089c68d6190000000000.

- **Channel:**

- A fast, off-chain method of mutual exchange between two [peers](#). To transact funds, peers exchange signatures to create an updated [commitment transaction](#).
- See closure methods: [mutual close](#), [revoked transaction close](#), [unilateral close](#)
- See related: [route](#)

- **Closing transaction:**

- A transaction generated as part of a [mutual close](#). A closing transaction is similar to a *commitment transaction*, but with no pending payments.
- See related: [commitment transaction](#), [funding transaction](#), [penalty transaction](#)

- **Commitment number:**

- A 48-bit incrementing counter for each [commitment transaction](#); counters are independent for each *peer* in the *channel* and start at 0.
- See container: [commitment transaction](#)

- See related: [closing transaction](#), [funding transaction](#), [penalty transaction](#)

- **Commitment revocation private key:**

- Every [commitment transaction](#) has a unique commitment revocation private-key value that allows the other *peer* to spend all outputs immediately: revealing this key is how old commitment transactions are revoked. To support revocation, each output of the commitment transaction refers to the commitment revocation public key.
- See container: [commitment transaction](#)
- See originator: [per-commitment secret](#)

- **Commitment transaction:**

- A transaction that spends the [funding transaction](#). Each *peer* holds the other peer's signature for this transaction, so that each always has a commitment transaction that it can spend. After a new commitment transaction is negotiated, the old one is *revoked*.
- See parts: [commitment number](#), [commitment revocation private key](#), [HTLC](#), [per-commitment secret](#), [outpoint](#)
- See related: [closing transaction](#), [funding transaction](#), [penalty transaction](#)
- See types: [revoked commitment transaction](#)

- **Fail the channel:**

- This is a forced close of the channel. Very early on (before opening), this may not require any action but forgetting the existence of the channel. Usually it requires signing and broadcasting the latest commitment transaction, although during mutual close it can also be performed by signing and broadcasting a mutual close transaction. See [BOLT #5](#).

- **Close the connection:**

- This means closing communication with the peer (such as closing the TCP socket). It does not imply closing any channels with the peer, but does cause the discarding of uncommitted state for connections with channels: see [BOLT #2](#).

- **Final node:**

- The final recipient of a packet that is routing a payment from an [origin node](#) through some number of [hops](#). It is also the final [receiving peer](#) in a chain.
- See category: [node](#)
- See related: [origin node](#), [processing node](#)

- **Funding transaction:**

- An irreversible on-chain transaction that pays to both [peers](#) on a [channel](#). It can only be spent by mutual consent.
- See related: [closing transaction](#), [commitment transaction](#), [penalty transaction](#)

- **Hop:**

- A [node](#). Generally, an intermediate node lying between an [origin node](#) and a [final node](#).
- See category: [node](#)

- **HTLC: Hashed Time Locked Contract.**

- A conditional payment between two [peers](#): the recipient can spend the payment by presenting its signature and a *payment preimage*, otherwise the payer can cancel the contract by spending it after a given time. These are implemented as outputs from the [commitment transaction](#).
- See container: [commitment transaction](#)
- See parts: [Payment hash](#), [Payment preimage](#)

- **Invoice: A request for funds on the Lightning Network, possibly**

including payment type, payment amount, expiry, and other information. This is how payments are made on the Lightning Network, rather than using Bitcoin-style addresses.

- **It's ok to be odd:**

- A rule applied to some numeric fields that indicates either optional or compulsory support for features. Even numbers indicate that both endpoints MUST support the feature in question, while odd numbers indicate that the feature MAY be disregarded by the other endpoint.

- **MSAT:**

- A millisatoshi, often used as a field name.

- **Mutual close:**

- A cooperative close of a [channel](#), accomplished by broadcasting an unconditional spend of the [funding transaction](#) with an output to each *peer* (unless one output is too small, and thus is not included).
- See related: [revoked transaction close](#), [unilateral close](#)

- **Node:**

- A computer or other device that is part of the Lightning network.
- See related: [peers](#)
- See types: [final node](#), [hop](#), [origin node](#), [processing node](#), [receiving node](#), [sending node](#)

- **Origin node:**

- The [node](#) that originates a packet that will route a payment through some number of [hops](#) to a [final node](#). It is also the first [sending peer](#) in a chain.
- See category: [node](#)
- See related: [final node](#), [processing node](#)

- **Outpoint:**

- A transaction hash and output index that uniquely identify an unspent transaction output. Needed to compose a new transaction, as an input.
- See related: [funding transaction](#), [commitment transaction](#)

- **Payment hash:**

- The [HTLC](#) contains the payment hash, which is the hash of the [payment preimage](#).
- See container: [HTLC](#)
- See originator: [Payment preimage](#)

- **Payment preimage:**

- Proof that payment has been received, held by the final recipient, who is the only person who knows this secret. The final recipient releases the preimage in order to release funds. The payment preimage is hashed as the [payment hash](#) in the [HTLC](#).
- See container: [HTLC](#)
- See derivation: [payment hash](#)

- **Peers:**

- Two [nodes](#) that are in communication with each other.
 - Two peers may gossip with each other prior to setting up a channel.
 - Two peers may establish a [channel](#) through which they transact.
- See related: [node](#)

- **Penalty transaction:**

- A transaction that spends all outputs of a [revoked commitment transaction](#), using the *commitment revocation private key*. A [peer](#) uses this if the other peer tries to “cheat” by broadcasting a [revoked commitment transaction](#).
- See related: [closing transaction](#), [commitment transaction](#), [funding transaction](#)

- **Per-commitment secret:**

- Every [commitment transaction](#) derives its keys from a per-commitment secret, which is generated such that the series of per-commitment secrets for all previous commitments can be stored compactly.
- See container: [commitment transaction](#)
- See derivation: [commitment revocation private key](#)

- **Processing node:**

- A [node](#) that is processing a packet that originated with an [origin node](#) and that is being sent toward a [final node](#) in order to route a payment. It acts as a [receiving peer](#) to receive the message, then a [sending peer](#) to send on the packet.
- See category: [node](#)

- See related: [final node](#), [origin node](#)

- **Receiving node:**

- A [node](#) that is receiving a message.
- See category: [node](#)
- See related: [sending node](#)

- **Receiving peer:**

- A [node](#) that is receiving a message from a directly connected *peer*.
- See category: [peer](#)
- See related: [sending peer](#)

- **Revoked commitment transaction:**

- An old [commitment transaction](#) that has been revoked because a new commitment transaction has been negotiated.
- See category: [commitment transaction](#)

- **Revoked transaction close:**

- An invalid close of a [channel](#), accomplished by broadcasting a *revoked commitment transaction*. Since the other *peer* knows the *commitment revocation secret key*, it can create a [penalty transaction](#).
- See related: [mutual close](#), [unilateral close](#)

- **Route:**

- A path across the Lightning Network that enables a payment from an *origin node* to a [final node](#) across one or more [hops](#).
- See related: [channel](#)

- **Sending node:**

- A [node](#) that is sending a message.
- See category: [node](#)
- See related: [receiving node](#)

- **Sending peer:**

- A [node](#) that is sending a message to a directly connected *peer*.
- See category: [peer](#)
- See related: [receiving peer](#).

- **Unilateral close:**

- An uncooperative close of a [channel](#), accomplished by broadcasting a [commitment transaction](#). This transaction is larger (i.e. less efficient) than a [closing transaction](#), and the [peer](#) whose commitment is broadcast cannot access its own outputs for some previously-negotiated duration.

◦ See related: [mutual close](#), [revoked transaction close](#)

Theme Song

Why this network could be democratic...
Numismatic...
Cryptographic!
Why it could be released Lightning!
(Release Lightning!)

We'll have some timelocked contracts with hashed pubkeys, oh yeah.
(Keep talking, whoa keep talkin')
We'll segregate the witness for trustless starts, oh yeah.
(I'll get the money, I've got to get the money)
With dynamic onion routes, they'll be shakin' in their boots;
You know that's just the truth, we'll be scaling through the roof.
Release Lightning!
(Go, go, go, go; go, go, go, go, go, go)

[Chorus:]
Oh released Lightning, it's better than a debit card..
(Release Lightning, go release Lightning!)
With released Lightning, micropayments just ain't hard...
(Release Lightning, go release Lightning!)
Then kaboom: we'll hit the moon -- release Lightning!
(Go, go, go, go; go, go, go, go, go, go)

We'll have QR codes, and smartphone apps, oh yeah.
(Ooo ooo ooo ooo ooo ooo ooo)
P2P messaging, and passive incomes, oh yeah.
(Ooo ooo ooo ooo ooo ooo ooo)
Outsourced closure watch, gives me feelings in my crotch.
You'll know it's not a brag when the repo gets a tag:
Released Lightning.

[Chorus]
[Instrumental, ~1m10s]
[Chorus]
(Lightning! Lightning! Lightning! Lightning!
Lightning! Lightning! Lightning! Lightning!)

C'mon guys, let's get to work!

– Anthony Towns aj@erisian.com.au

Authors

[FIXME: Insert Author List]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Connecting Nodes

Before two Lightning nodes can open a channel, they have to talk. That means a TCP connection, an encrypted handshake, and a shared idea of which messages exist.

BOLT 08 is the transport. Lightning uses a Noise protocol handshake (Noise_XK) so that by the time the first application message is sent, both sides have authenticated keys and an encrypted session. If you've ever wondered why Lightning node IDs are public keys, this is why: the transport identity *is* the node identity.

BOLT 01 is the messaging layer that rides on that session. It defines the binary framing, the type system, and the messages every peer needs — init, error, ping / pong. Everything later in the book is a BOLT 01 message with a more interesting type.

BOLT 09 is the feature bits. Lightning is still changing, so peers advertise what they support during init and inside invoices, node announcements, and other gossip. If a message in a later BOLT is optional, BOLT 09 is where that option got a number.

Read these three in order. Transport, then messages, then the flags that say which messages you're allowed to send.

In this chapter

- BOLT 08 — Encrypted and authenticated transport
- BOLT 01 — Base messaging
- BOLT 09 — Feature flags

BOLT #8: Encrypted and Authenticated Transport

All communications between Lightning nodes is encrypted in order to provide confidentiality for all transcripts between nodes and is authenticated in order to avoid malicious interference. Each node has a known long-term identifier that is a public key on Bitcoin's secp256k1 curve. This long-term public key is used within the protocol to establish an encrypted and authenticated connection with peers, and also to authenticate any information advertised on behalf of a node.

Table of Contents

- [Cryptographic Messaging Overview](#)
 - [Authenticated Key Agreement Handshake](#)
 - [Handshake Versioning](#)
 - [Noise Protocol Instantiation](#)
- [Authenticated Key Exchange Handshake Specification](#)
 - [Handshake State](#)
 - [Handshake State Initialization](#)
 - [Handshake Exchange](#)

- [Lightning Message Specification](#)
 - [Encrypting and Sending Messages](#)
 - [Receiving and Decrypting Messages](#)
- [Lightning Message Key Rotation](#)
- [Security Considerations](#)
- [Appendix A: Transport Test Vectors](#)
 - [Initiator Tests](#)
 - [Responder Tests](#)
 - [Message Encryption Tests](#)
- [Acknowledgments](#)
- [References](#)
- [Authors](#)

Cryptographic Messaging Overview

Prior to sending any Lightning messages, nodes MUST first initiate the cryptographic session state that is used to encrypt and authenticate all messages sent between nodes. The initialization of this cryptographic session state is completely distinct from any inner protocol message header or conventions.

The transcript between two nodes is separated into two distinct segments:

1. Before any actual data transfer, both nodes participate in an authenticated key agreement handshake, which is based on the Noise Protocol Framework².
2. If the initial handshake is successful, then nodes enter the Lightning message exchange phase. In the Lightning message exchange phase, all messages are Authenticated Encryption with Associated Data (AEAD) ciphertexts.

Authenticated Key Agreement Handshake

The handshake chosen for the authenticated key exchange is Noise_XK. As a pre-message, the initiator must know the identity public key of the responder. This provides a degree of identity hiding for the responder, as its static public key is *never* transmitted during the handshake. Instead, authentication is achieved implicitly via a series of Elliptic-Curve Diffie-Hellman (ECDH) operations followed by a MAC check.

The authenticated key agreement (Noise_XK) is performed in three distinct steps (acts). During each act of the handshake the following occurs: some (possibly encrypted) keying material is sent to the other party; an ECDH is performed, based on exactly which act is being executed, with the result mixed into the current set of encryption keys (ck the chaining key and k the encryption key); and an AEAD payload with a zero-length cipher text is sent. As this payload has no length, only a MAC is sent across. The mixing of ECDH outputs into a hash digest forms an incremental TripleDH handshake.

Using the language of the Noise Protocol, e and s (both public keys with e being the ephemeral key and s being the static key which in our case is usually the nodeid) indicate possibly encrypted keying material, and es, ee, and se each indicate an ECDH operation between two keys. The handshake is laid out as follows:

```
Noise_XK(s, rs):
  <- s
  ...
  -> e, es
```

<- e, ee
-> s, se

All of the handshake data sent across the wire, including the keying material, is incrementally hashed into a session-wide “handshake digest”, *h*. Note that the handshake state *h* is never transmitted during the handshake; instead, *digest* is used as the Associated Data within the zero-length AEAD messages.

Authenticating each message sent ensures that a man-in-the-middle (MITM) hasn’t modified or replaced any of the data sent as part of a handshake, as the MAC check would fail on the other side if so.

A successful check of the MAC by the receiver indicates implicitly that all authentication has been successful up to that point. If a MAC check ever fails during the handshake process, then the connection is to be immediately terminated.

Handshake Versioning

Each message sent during the initial handshake starts with a single leading byte, which indicates the version used for the current handshake. A version of 0 indicates that no change is necessary, while a non-zero version indicate that the client has deviated from the protocol originally specified within this document.

Clients **MUST** reject handshake attempts initiated with an unknown version.

Noise Protocol Instantiation

Concrete instantiations of the Noise Protocol require the definition of three abstract cryptographic objects: the hash function, the elliptic curve, and the AEAD cipher scheme. For Lightning, SHA-256 is chosen as the hash function, secp256k1 as the elliptic curve, and ChaChaPoly-1305 as the AEAD construction.

The composition of ChaCha20 and Poly1305 that are used **MUST** conform to RFC 8439¹.

The official protocol name for the Lightning variant of Noise is `Noise_XK_secp256k1_ChaChaPoly_SHA256`. The ASCII string representation of this value is hashed into a digest used to initialize the starting handshake state. If the protocol names of two endpoints differ, then the handshake process fails immediately.

Authenticated Key Exchange Handshake Specification

The handshake proceeds in three acts, taking 1.5 round trips. Each handshake is a *fixed* sized payload without any header or additional meta-data attached. The exact size of each act is as follows:

- **Act One:** 50 bytes
- **Act Two:** 50 bytes
- **Act Three:** 66 bytes

Handshake State

Throughout the handshake process, each side maintains these variables:

- **ck:** the **chaining key**. This value is the accumulated hash of all previous ECDH outputs. At the end of the handshake, *ck* is used to derive the encryption keys for Lightning messages.

- **h**: the **handshake hash**. This value is the accumulated hash of *all* handshake data that has been sent and received so far during the handshake process.
- **temp_k1, temp_k2, temp_k3**: the **intermediate keys**. These are used to encrypt and decrypt the zero-length AEAD payloads at the end of each handshake message.
- **e**: a party's **ephemeral keypair**. For each session, a node **MUST** generate a new ephemeral key with strong cryptographic randomness.
- **s**: a party's **static keypair** (**ls** for local, **rs** for remote)

The following functions will also be referenced:

- **ECDH(k, rk)**: performs an Elliptic-Curve Diffie-Hellman operation using **k**, which is a valid secp256k1 private key, and **rk**, which is a valid public key
 - The returned value is the SHA256 of the compressed format of the generated point.
- **HKDF(salt, ikm)**: a function defined in RFC 5869³, evaluated with a zero-length info field
 - All invocations of HKDF implicitly return 64 bytes of cryptographic randomness using the extract-and-expand component of the HKDF.
- **encryptWithAD(k, n, ad, plaintext)**: outputs **encrypt(k, n, ad, plaintext)**
 - Where **encrypt** is an evaluation of ChaCha20-Poly1305 (IETF variant) with the passed arguments, with nonce **n** encoded as 32 zero bits, followed by a *little-endian* 64-bit value. Note: this follows the Noise Protocol convention, rather than our normal endian.
- **decryptWithAD(k, n, ad, ciphertext)**: outputs **decrypt(k, n, ad, ciphertext)**
 - Where **decrypt** is an evaluation of ChaCha20-Poly1305 (IETF variant) with the passed arguments, with nonce **n** encoded as 32 zero bits, followed by a *little-endian* 64-bit value.
- **generateKey()**: generates and returns a fresh secp256k1 keypair
 - Where the object returned by **generateKey** has two attributes:
 - **.pub**, which returns an abstract object representing the public key
 - **.priv**, which represents the private key used to generate the public key
 - Where the object also has a single method:
 - **.serializeCompressed()**
- **a || b** denotes the concatenation of two byte strings **a** and **b**

Handshake State Initialization

Before the start of Act One, both sides initialize their per-sessions state as follows:

1. **h = SHA-256(protocolName)**
 - where **protocolName** = "Noise_XK_secp256k1_ChaChaPoly_SHA256" encoded as an ASCII string
2. **ck = h**
3. **h = SHA-256(h || prologue)**
 - where **prologue** is the ASCII string: **lightning**

As a concluding step, both sides mix the responder's public key into the handshake digest:

- The initiating node mixes in the responding node's static public key serialized in Bitcoin's compressed format:
 - $h = \text{SHA-256}(h \parallel rs.\text{pub.serializeCompressed}())$
- The responding node mixes in their local static public key serialized in Bitcoin's compressed format:
 - $h = \text{SHA-256}(h \parallel ls.\text{pub.serializeCompressed}())$

Handshake Exchange

Act One

-> e, es

Act One is sent from initiator to responder. During Act One, the initiator attempts to satisfy an implicit challenge by the responder. To complete this challenge, the initiator must know the static public key of the responder.

The handshake message is *exactly* 50 bytes: 1 byte for the handshake version, 33 bytes for the compressed ephemeral public key of the initiator, and 16 bytes for the poly1305 tag.

Sender Actions:

1. $e = \text{generateKey}()$
2. $h = \text{SHA-256}(h \parallel e.\text{pub.serializeCompressed}())$
 - The newly generated ephemeral key is accumulated into the running handshake digest.
3. $es = \text{ECDH}(e.\text{priv}, rs)$
 - The initiator performs an ECDH between its newly generated ephemeral key and the remote node's static public key.
4. $ck, \text{temp_k1} = \text{HKDF}(ck, es)$
 - A new temporary encryption key is generated, which is used to generate the authenticating MAC.
5. $c = \text{encryptWithAD}(\text{temp_k1}, 0, h, \text{zero})$
 - where zero is a zero-length plaintext
6. $h = \text{SHA-256}(h \parallel c)$
 - Finally, the generated ciphertext is accumulated into the authenticating handshake digest.
7. Send $m = 0 \parallel e.\text{pub.serializeCompressed}() \parallel c$ to the responder over the network buffer.

Receiver Actions:

1. Read *exactly* 50 bytes from the network buffer.
2. Parse the read message (m) into v, re, and c:
 - where v is the *first* byte of m, re is the next 33 bytes of m, and c is the last 16 bytes of m
 - The raw bytes of the remote party's ephemeral public key (re) are to be deserialized into a point on the curve using affine coordinates as encoded by the key's serialized composed format.
3. If v is an unrecognized handshake version, then the responder **MUST** abort the connection attempt.
4. $h = \text{SHA-256}(h \parallel re.\text{serializeCompressed}())$
 - The responder accumulates the initiator's ephemeral key into the authenticating handshake digest.

5. `es = ECDH(s.priv, re)`
 - The responder performs an ECDH between its static private key and the initiator's ephemeral public key.
6. `ck, temp_k1 = HKDF(ck, es)`
 - A new temporary encryption key is generated, which will shortly be used to check the authenticating MAC.
7. `p = decryptWithAD(temp_k1, 0, h, c)`
 - If the MAC check in this operation fails, then the initiator does *not* know the responder's static public key. If this is the case, then the responder **MUST** terminate the connection without any further messages.
8. `h = SHA-256(h || c)`
 - The received ciphertext is mixed into the handshake digest. This step serves to ensure the payload wasn't modified by a MITM.

Act Two

`<- e, ee`

Act Two is sent from the responder to the initiator. Act Two will *only* take place if Act One was successful. Act One was successful if the responder was able to properly decrypt and check the MAC of the tag sent at the end of Act One.

The handshake is *exactly* 50 bytes: 1 byte for the handshake version, 33 bytes for the compressed ephemeral public key of the responder, and 16 bytes for the poly1305 tag.

Sender Actions:

1. `e = generateKey()`
2. `h = SHA-256(h || e.pub.serializeCompressed())`
 - The newly generated ephemeral key is accumulated into the running handshake digest.
3. `ee = ECDH(e.priv, re)`
 - where `re` is the ephemeral key of the initiator, which was received during Act One
4. `ck, temp_k2 = HKDF(ck, ee)`
 - A new temporary encryption key is generated, which is used to generate the authenticating MAC.
5. `c = encryptWithAD(temp_k2, 0, h, zero)`
 - where `zero` is a zero-length plaintext
6. `h = SHA-256(h || c)`
 - Finally, the generated ciphertext is accumulated into the authenticating handshake digest.
7. Send `m = 0 || e.pub.serializeCompressed() || c` to the initiator over the network buffer.

Receiver Actions:

1. Read *exactly* 50 bytes from the network buffer.
2. Parse the read message (`m`) into `v`, `re`, and `c`:
 - where `v` is the *first* byte of `m`, `re` is the next 33 bytes of `m`, and `c` is the last 16 bytes of `m`.
3. If `v` is an unrecognized handshake version, then the responder **MUST** abort the connection attempt.
4. `h = SHA-256(h || re.serializeCompressed())`

5. $ee = \text{ECDH}(e.\text{priv}, re)$
 - where re is the responder's ephemeral public key
 - The raw bytes of the remote party's ephemeral public key (re) are to be deserialized into a point on the curve using affine coordinates as encoded by the key's serialized composed format.
6. $ck, \text{temp_k2} = \text{HKDF}(ck, ee)$
 - A new temporary encryption key is generated, which is used to generate the authenticating MAC.
7. $p = \text{decryptWithAD}(\text{temp_k2}, 0, h, c)$
 - If the MAC check in this operation fails, then the initiator MUST terminate the connection without any further messages.
8. $h = \text{SHA-256}(h \parallel c)$
 - The received ciphertext is mixed into the handshake digest. This step serves to ensure the payload wasn't modified by a MITM.

Act Three

-> s, se

Act Three is the final phase in the authenticated key agreement described in this section. This act is sent from the initiator to the responder as a concluding step. Act Three is executed *if and only if* Act Two was successful. During Act Three, the initiator transports its static public key to the responder encrypted with *strong* forward secrecy, using the accumulated HKDF derived secret key at this point of the handshake.

The handshake is *exactly* 66 bytes: 1 byte for the handshake version, 33 bytes for the static public key encrypted with the ChaCha20 stream cipher, 16 bytes for the encrypted public key's tag generated via the AEAD construction, and 16 bytes for a final authenticating tag.

Sender Actions:

1. $c = \text{encryptWithAD}(\text{temp_k2}, 1, h, s.\text{pub.serializeCompressed}())$
 - where s is the static public key of the initiator
2. $h = \text{SHA-256}(h \parallel c)$
3. $se = \text{ECDH}(s.\text{priv}, re)$
 - where re is the ephemeral public key of the responder
4. $ck, \text{temp_k3} = \text{HKDF}(ck, se)$
 - The final intermediate shared secret is mixed into the running chaining key.
5. $t = \text{encryptWithAD}(\text{temp_k3}, 0, h, \text{zero})$
 - where zero is a zero-length plaintext
6. $sk, rk = \text{HKDF}(ck, \text{zero})$
 - where zero is a zero-length plaintext, sk is the key to be used by the initiator to encrypt messages to the responder, and rk is the key to be used by the initiator to decrypt messages sent by the responder
 - The final encryption keys, to be used for sending and receiving messages for the duration of the session, are generated.
7. $rn = 0, sn = 0$
 - The sending and receiving nonces are initialized to 0.
8. $rck = sck = ck$
 - The sending and receiving chaining keys are initialized the same.

9. Send $m = 0 \parallel c \parallel t$ over the network buffer.

Receiver Actions:

1. Read *exactly* 66 bytes from the network buffer.
2. Parse the read message (m) into v , c , and t :
 - where v is the *first* byte of m , c is the next 49 bytes of m , and t is the last 16 bytes of m
3. If v is an unrecognized handshake version, then the responder MUST abort the connection attempt.
4. $rs = \text{decryptWithAD}(\text{temp_k2}, 1, h, c)$
 - At this point, the responder has recovered the static public key of the initiator.
 - If the MAC check in this operation fails, then the responder MUST terminate the connection without any further messages.
5. $h = \text{SHA-256}(h \parallel c)$
6. $se = \text{ECDH}(e.\text{priv}, rs)$
 - where e is the responder's original ephemeral key
7. $ck, \text{temp_k3} = \text{HKDF}(ck, se)$
8. $p = \text{decryptWithAD}(\text{temp_k3}, 0, h, t)$
 - If the MAC check in this operation fails, then the responder MUST terminate the connection without any further messages.
9. $rk, sk = \text{HKDF}(ck, \text{zero})$
 - where zero is a zero-length plaintext, rk is the key to be used by the responder to decrypt the messages sent by the initiator, and sk is the key to be used by the responder to encrypt messages to the initiator
 - The final encryption keys, to be used for sending and receiving messages for the duration of the session, are generated.
10. $rn = 0, sn = 0$
 - The sending and receiving nonces are initialized to 0.
11. $rck = sck = ck$
 - The sending and receiving chaining keys are initialized the same.

Lightning Message Specification


At the conclusion of Act Three, both sides have derived the encryption keys, which will be used to encrypt and decrypt messages for the remainder of the session.

The actual Lightning protocol messages are encapsulated within AEAD ciphertexts. Each message is prefixed with another AEAD ciphertext, which encodes the total length of the following Lightning message (not including its MAC).

The *maximum* size of *any* Lightning message MUST NOT exceed 65535 bytes. A maximum size of 65535 simplifies testing, makes memory management easier, and helps mitigate memory-exhaustion attacks.

In order to make traffic analysis more difficult, the length prefix for all encrypted Lightning messages is also encrypted. Additionally a 16-byte Poly-1305 tag is added to the encrypted length prefix in order to ensure that the packet length hasn't been modified when in-flight and also to avoid creating a decryption oracle.

The structure of packets on the wire resembles the following:



The prefixed message length is encoded as a 2-byte big-endian integer, for a total maximum packet length of $2 + 16 + 65535 + 16 = 65569$ bytes.

Encrypting and Sending Messages

In order to encrypt and send a Lightning message (m) to the network stream, given a sending key (sk) and a nonce (sn), the following steps are completed:

1. Let $l = \text{len}(m)$.
 - where len obtains the length in bytes of the Lightning message
2. Serialize l into 2 bytes encoded as a big-endian integer.
3. Encrypt l (using ChaChaPoly-1305, sn , and sk), to obtain lc (18 bytes)
 - The nonce sn is encoded as a 96-bit little-endian number. As the decoded nonce is 64 bits, the 96-bit nonce is encoded as: 32 bits of leading 0s followed by a 64-bit value.
 - The nonce sn MUST be incremented after this step.
 - A zero-length byte slice is to be passed as the AD (associated data).
4. Finally, encrypt the message itself (m) using the same procedure used to encrypt the length prefix. Let encrypted ciphertext be known as c .
 - The nonce sn MUST be incremented after this step.
5. Send $lc || c$ over the network buffer.

Receiving and Decrypting Messages

In order to decrypt the *next* message in the network stream, the following steps are completed:

1. Read *exactly* 18 bytes from the network buffer.
2. Let the encrypted length prefix be known as lc .
3. Decrypt lc (using ChaCha20-Poly1305, rn , and rk), to obtain the size of the encrypted packet l .
 - A zero-length byte slice is to be passed as the AD (associated data).
 - The nonce rn MUST be incremented after this step.
4. Read *exactly* $l+16$ bytes from the network buffer, and let the bytes be known as c .
5. Decrypt c (using ChaCha20-Poly1305, rn , and rk), to obtain decrypted plaintext packet p .
 - The nonce rn MUST be incremented after this step.

Lightning Message Key Rotation

Changing keys regularly and forgetting previous keys is useful to prevent the decryption of old messages, in the case of later key leakage (i.e. backwards secrecy).

Key rotation is performed for *each* key (sk and rk) *individually*, using sck and rck respectively. A key is to be rotated after a party encrypts or decrypts 1000 times with it (i.e. every 500 messages). This can be properly accounted for by rotating the key once the nonce dedicated to it reaches 1000.

Key rotation for a key k is performed according to the following steps:

1. Let ck be the chaining key (i.e. rck for rk or sck for sk)
2. $ck', k' = \text{HKDF}(ck, k)$
3. Reset the nonce for the key to $n = \emptyset$.
4. $k = k'$
5. $ck = ck'$

Security Considerations

It is strongly recommended that existing, commonly-used, validated libraries be used for encryption and decryption, to avoid the many possible implementation pitfalls.

Appendix A: Transport Test Vectors

To make a repeatable test handshake, the following specifies what `generateKey()` will return (i.e. the value for `e.priv`) for each side. Note that this is a violation of the spec, which requires randomness.

Initiator Tests

The initiator SHOULD produce the given output when fed this input. The comments reflect internal states, for debugging purposes.

[illegible]

[illegible]

[illegible]

```
output: ERROR (ACT3_BAD_CIPHertext)
```

```
name: transport-responder act3 bad rs test
```

```
ls.priv=2121212121212121212121212121212121212121212121212121212121212121
```

ls.pub=028d7500dd4c12685d1f568b4c2b5048e8534b873319f3a8daa612b469132ec7f7

[illegible]

e.pub=0x02466d7fcae563e5cb09a0d1870bb580344804617879a14949cf22285f1bae3f27

Act One

input:

```
0x00036360e856310ce5d294e8be33fc807077dc56ac80d95d9cd4ddbd21325eff73f70df6086551151f58b8afe6c195782c6a
```

Act Two

output:

```
0x0002466d7fcae563e5cb09a0d1870bb580344804617879a14949cf22285f1bae3f276e2470b93aac583c9ef6eafca3f730ae
```

Act Three

input:

```
0x00bfe3a702e93e3a9948c2ed6e5fd7590a6e1c3a0344cfc9d5b57357049aa2235536ad09a8ee351870c2bb7f78b754a26c6cef79a98d25139c856d7efd252c2ae73c
```

```
# decryptWithAD(0x908b166535c01a935cf1e130a5fe895ab4e6f3ef8855d87e9b7581c4ab663ddc,
0x00000000000000000000000000000001, 0x90578e247e98674e661013da3c5c1ca6a8c8f48c90b485c0dfa149a23d56d72,
0xd7fedc211450dd9602b41081c9bd05328b8bf8c0238880f7b7cb8a34bb6d8354081e8d4b81887fae47a74fe8aab300
8653)
```

```
# rs=0x044f355bdbcb7cc0af728ef3cceb9615d90684bb5b2ca5f859ab0f0b704075871aa
```

```
output: ERROR (ACT3_BAD_PUBKEY)
```

```
name: transport-responder act3 bad MAC test
```

```
ls.priv=2121212121212121212121212121212121212121212121212121212121212121
```

ls.pub=028d7500dd4c12685d1f568b4c2b5048e8534b873319f3a8daa612b469132ec7f7

[illegible]

e.pub=0x02466d7fcae563e5cb09a0d1870bb580344804617879a14949cf22285f1bae3f27

Act One

input:

```
0x00036360e856310ce5d294e8be33fc807077dc56ac80d95d9cd4ddbd21325eff73f70df6086551151f58b8afe6c195782c6a
```

Act Two

output:

0x0002466d7fcae563e5cb09a0d1870bb580344804617879a14949cf22285f1bae3f276e2470b93aac583c9ef6eafca3f730ae

Act Three

input:

```
0x00b9e3a702e93e3a9948c2ed6e5fd7590a6e1c3a0344cfc9d5b57357049aa22355361aa02e55a8fc28fef5bd6d71ad
0c38228dc68b1c466263b47fdf31e560e139bb
```

```
output: ERROR (ACT3 BAD TAG)
```

Message Encryption Tests

In this test, the initiator sends length 5 messages containing “hello” 1001 times. Only six example outputs are shown, for brevity and to test two key rotations:

```
name: transport-message test
```

ck=0x919219dbb2920afa8db80f9a51787a840bcf111ed8d588caf9ab4be716e42b01

```
sk=0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9
```

rk=0xbb9020b8965f4df047e07f955f3c4b88418984aad5c5cdb35096b9ea8fa5c3442

```
# encrypt l: cleartext=0x0005, AD=NULL, sn=0x00000000000000000000000000000000,
```

```

sk=0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9 =>
0xcf2b30ddf0cf3f80e7c35a6e6730b59fe802
# encrypt m: cleartext=0x68656c6c6f, AD=NULL, sn=0x000000000100000000000000,
sk=0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9 =>
0x473180f396d88a8fb0db8cbcf25d2f214cf9ea1d95
output 0: 0xcf2b30ddf0cf3f80e7c35a6e6730b59fe802473180f396d88a8fb0db8cbcf25d2f214cf9ea1d95
# encrypt l: cleartext=0x0005, AD=NULL, sn=0x000000000200000000000000,
sk=0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9 =>
0x72887022101f0b6753e0c7de21657d35a4cb
# encrypt m: cleartext=0x68656c6c6f, AD=NULL, sn=0x000000000300000000000000,
sk=0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9 =>
0x2a1f5cde2650528bbc8f837d0f0d7ad833b1a256a1
output 1: 0x72887022101f0b6753e0c7de21657d35a4cb2a1f5cde2650528bbc8f837d0f0d7ad833b1a256a1
# 0xcc2c6e467efc8067720c2d09c139d1f77731893aad1defa14f9bf3c48d3f1d31,
0x3fbd101abd1132ca3a0ae34a669d8d9ba69a587e0bb4ddd59524541cf4813d8 =
HKDF(0x919219dbb2920afa8db80f9a51787a840bcf111ed8d588caf9ab4be716e42b01,
0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9)
# 0xcc2c6e467efc8067720c2d09c139d1f77731893aad1defa14f9bf3c48d3f1d31,
0x3fbd101abd1132ca3a0ae34a669d8d9ba69a587e0bb4ddd59524541cf4813d8 =
HKDF(0x919219dbb2920afa8db80f9a51787a840bcf111ed8d588caf9ab4be716e42b01,
0x969ab31b4d288cedf6218839b27a3e2140827047f2c0f01bf5c04435d43511a9)
output 500: 0x178cb9d7387190fa34db9c2d50027d21793c9bc2d40b1e14dcf30ebbbb220f48364f7a4c68bf8
output 501: 0x1b186c57d44eb6de4c057c49940d79bb838a145cb528d6e8fd26dbe50a60ca2c104b56b60e45bd
# 0x728366ed68565dc17cf6dd97330a859a6a56e87e2beef3bd828a4c4a54d8df06,
0x9e0477f9850dca41e42db0e4d154e3a098e5a000d995e421849fcd5df27882bd =
HKDF(0xcc2c6e467efc8067720c2d09c139d1f77731893aad1defa14f9bf3c48d3f1d31,
0x3fbd101abd1132ca3a0ae34a669d8d9ba69a587e0bb4ddd59524541cf4813d8)
# 0x728366ed68565dc17cf6dd97330a859a6a56e87e2beef3bd828a4c4a54d8df06,
0x9e0477f9850dca41e42db0e4d154e3a098e5a000d995e421849fcd5df27882bd =
HKDF(0xcc2c6e467efc8067720c2d09c139d1f77731893aad1defa14f9bf3c48d3f1d31,
0x3fbd101abd1132ca3a0ae34a669d8d9ba69a587e0bb4ddd59524541cf4813d8)
output 1000: 0x4a2f3cc3b5e78ddb83dcb426d9863d9d9a723b0337c89dd0b005d89f8d3c05c52b76b29b740f09
output 1001: 0x2ecd8c8a5629d0d02ab457a0fdd0f7b90a192cd46be5ecb6ca570bfc5e268338b1a16cf4ef2d36

```

Acknowledgments

TODO(roasbeef); fin

References

1. <https://tools.ietf.org/html/rfc8439>
2. <http://noiseprotocol.org/noise.html>
3. <https://tools.ietf.org/html/rfc5869>

Authors

FIXME



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

BOLT #1: Base Protocol

Overview

This protocol assumes an underlying authenticated and ordered transport mechanism that takes care of framing individual messages. [BOLT #8](#) specifies the canonical transport layer used in Lightning, though it can be replaced by any transport that fulfills the above guarantees.

The default TCP port depends on the network used. The most common networks are:

- Bitcoin mainnet with port number 9735 or the corresponding hexadecimal `0x2607`;
- Bitcoin testnet with port number 19735 (`0x4D17`);
- Bitcoin signet with port number 39735 (`0x9B37`).

The Unicode code point for LIGHTNING [1](#), and the port convention try to follow the Bitcoin Core convention.

All data fields are unsigned big-endian unless otherwise specified.

Table of Contents

- [Connection Handling and Multiplexing](#)
- [Lightning Message Format](#)
- [Type-Length-Value Format](#)
- [Fundamental Types](#)
- [Setup Messages](#)
 - [The init Message](#)
 - [The error and warning Messages](#)
- [Control Messages](#)
 - [The ping and pong Messages](#)
- [Peer Storage](#)
 - [The peer_storage and peer_storage_retrieval Messages](#)
- [Appendix A: BigSize Test Vectors](#)
- [Appendix B: Type-Length-Value Test Vectors](#)
- [Appendix C: Message Extension](#)
- [Appendix D: Signed Integers Test Vectors](#)
- [Acknowledgments](#)
- [References](#)
- [Authors](#)

Connection Handling and Multiplexing

Implementations MUST use a single connection per peer; channel messages (which include a channel ID) are multiplexed over this single connection.

Lightning Message Format

After decryption, all Lightning messages are of the form:

1. type: a 2-byte big-endian field indicating the type of message
2. payload: a variable-length payload that comprises the remainder of the message and that conforms to a format matching the type
3. extension: an optional [TLV stream](#)

The type field indicates how to interpret the payload field. The format for each individual type is defined by a specification in this repository. The type follows the *it's ok to be odd* rule, so nodes MAY send *odd*-numbered types without ascertaining that the recipient understands it.

The messages are grouped logically into five groups, ordered by the most significant bit that is set:

- Setup & Control (types 0-31): messages related to connection setup, control, supported features, and error reporting (described below)
- Channel (types 32-127): messages used to setup and tear down micropayment channels (described in [BOLT #2](#))
- Commitment (types 128-255): messages related to updating the current commitment transaction, which includes adding, revoking, and settling HTLCs as well as updating fees and exchanging signatures (described in [BOLT #2](#))
- Routing (types 256-511): messages containing node and channel announcements, as well as any active route exploration (described in [BOLT #7](#))
- Custom (types 32768-65535): experimental and application-specific messages

The size of the message is required by the transport layer to fit into a 2-byte unsigned int; therefore, the maximum possible size is 65535 bytes.

A sending node: - MUST NOT send an evenly-typed message not listed here without prior negotiation. - MUST NOT send evenly-typed TLV records in the extension without prior negotiation. - that negotiates an option in this specification: - MUST include all the fields annotated with that option. - When defining custom messages: - SHOULD pick a random type to avoid collision with other custom types. - SHOULD pick a type that doesn't conflict with other experiments listed in [this issue](#). - SHOULD pick an odd type identifiers when regular nodes should ignore the additional data. - SHOULD pick an even type identifiers when regular nodes should reject the message and close the connection.

A receiving node: - upon receiving a message of *odd*, unknown type: - MUST ignore the received message. - upon receiving a message of *even*, unknown type: - MUST close the connection. - MAY fail the channels. - upon receiving a known message with insufficient length for the contents: - MUST close the connection. - MAY fail the channels. - upon receiving a message with an extension: - MAY ignore the extension. - Otherwise, if the extension is invalid: - MUST close the connection. - MAY fail the channels.

Rationale

By default SHA2 and Bitcoin public keys are both encoded as big endian, thus it would be unusual to use a different endian for other fields.

Length is limited to 65535 bytes by the cryptographic wrapping, and messages in the protocol are never more than that length anyway.

The *it's ok to be odd* rule allows for future optional extensions without negotiation or special coding in clients. The *extension* field similarly allows for future expansion by letting senders include additional TLV data. Note that an *extension* field can only be added when the message payload doesn't already fill the 65535 bytes maximum length.

Implementations may prefer to have message data aligned on an 8-byte boundary (the largest natural alignment requirement of any type here); however, adding a 6-byte padding after the type field was considered wasteful: alignment may be achieved by decrypting the message into a buffer with 6-bytes of pre-padding.

Type-Length-Value Format

Throughout the protocol, a TLV (Type-Length-Value) format is used to allow for the backwards-compatible addition of new fields to existing message types.

A `tlv_record` represents a single field, encoded in the form:

- [bigsize: type]
- [bigsize: length]
- [length: value]

A `tlv_stream` is a series of (possibly zero) `tlv_records`, represented as the concatenation of the encoded `tlv_records`. When used to extend existing messages, a `tlv_stream` is typically placed after all currently defined fields.

The type is encoded using the BigSize format. It functions as a message-specific, 64-bit identifier for the `tlv_record` determining how the contents of value should be decoded. type identifiers below 2^{16} are reserved for use in this specification. type identifiers greater than or equal to 2^{16} are available for custom records. Any record not defined in this specification is considered a custom record. This includes experimental and application-specific messages.

The length is encoded using the BigSize format signaling the size of value in bytes.

The value depends entirely on the type, and should be encoded or decoded according to the message-specific format determined by type.

Requirements

The sending node: - MUST order `tlv_records` in a `tlv_stream` by strictly-increasing type, hence MUST not produce more than a single TLV record with the same type - MUST minimally encode type and length. - When defining custom record type identifiers: - SHOULD pick random type identifiers to avoid collision with other custom types. - SHOULD pick odd type identifiers when regular nodes should ignore the additional data. - SHOULD pick even type identifiers when regular nodes should reject the full tlv stream containing the custom record. - SHOULD NOT use redundant, variable-length encodings in a `tlv_record`.

The receiving node: - if zero bytes remain before parsing a type: - MUST stop parsing the `tlv_stream`. - if a type or length is not minimally encoded: - MUST fail to parse the `tlv_stream`. - if decoded types are not strictly-increasing (including situations when two or more occurrences of the same type are met): - MUST fail

to parse the `tlv_stream`. - if length exceeds the number of bytes remaining in the message: - MUST fail to parse the `tlv_stream`. - if type is known: - MUST decode the next length bytes using the known encoding for type. - if length is not exactly equal to that required for the known encoding for type: - MUST fail to parse the `tlv_stream`. - if variable-length fields within the known encoding for type are not minimal: - MUST fail to parse the `tlv_stream`. - otherwise, if type is unknown: - if type is even: - MUST fail to parse the `tlv_stream`. - otherwise, if type is odd: - MUST discard the next length bytes.

Rationale

The primary advantage in using TLV is that a reader is able to ignore new fields that it does not understand, since each field carries the exact size of the encoded element. Without TLV, even if a node does not wish to use a particular field, the node is forced to add parsing logic for that field in order to determine the offset of any fields that follow.

The strict monotonicity constraint ensures that all types are unique and can appear at most once. Fields that map to complex objects, e.g. vectors, maps, or structs, should do so by defining the encoding such that the object is serialized within a single `tlv_record`. The uniqueness constraint, among other things, enables the following optimizations: - canonical ordering is defined independent of the encoded values. - canonical ordering can be known at compile-time, rather than being determined dynamically at the time of encoding. - verifying canonical ordering requires less state and is less-expensive. - variable-size fields can reserve their expected size up front, rather than appending elements sequentially and incurring double-and-copy overhead.

The use of a `bigsize` for type and length permits a space savings for small types or short values. This potentially leaves more space for application data over the wire or in an onion payload.

All types must appear in increasing order to create a canonical encoding of the underlying `tlv_records`. This is crucial when computing signatures over a `tlv_stream`, as it ensures verifiers will be able to recompute the same message digest as the signer. Note that the canonical ordering over the set of fields can be enforced even if the verifier does not understand what the fields contain.

Writers should avoid using redundant, variable-length encodings in a `tlv_record` since this results in encoding the length twice and complicates computing the outer length. As an example, when writing a variable length byte array, the value should contain only the raw bytes and forgo an additional internal length since the `tlv_record` already carries the number of bytes that follow. On the other hand, if a `tlv_record` contains multiple, variable-length elements then this would not be considered redundant, and is needed to allow the receiver to parse individual elements from value.

Fundamental Types

Various fundamental types are referred to in the message specifications:

- `byte`: an 8-bit byte
- `s8`: an 8-bit signed integer
- `u16`: a 2 byte unsigned integer
- `s16`: a 2 byte signed integer
- `u32`: a 4 byte unsigned integer
- `s32`: a 4 byte signed integer
- `u64`: an 8 byte unsigned integer
- `s64`: an 8 byte signed integer

Signed integers use standard big-endian two's complement representation (see test vectors [below](#)).

For the final value in TLV records, truncated integers may be used. Leading zeros in truncated integers MUST be omitted:

- tu16: a 0 to 2 byte truncated unsigned integer
- tu32: a 0 to 4 byte truncated unsigned integer
- tu64: a 0 to 8 byte truncated unsigned integer

When used to encode amounts, the previous fields MUST comply with the upper bound of 21 million BTC:

- satoshi amounts MUST be at most `0x000775f05a074000`
- milli-satoshi amounts MUST be at most `0x1d24b2dfac520000`

The following convenience types are also defined:

- chain_hash: a 32-byte chain identifier (see [BOLT #0](#))
- channel_id: a 32-byte channel_id (see [BOLT #2](#))
- sha256: a 32-byte SHA2-256 hash
- signature: a 64-byte bitcoin Elliptic Curve signature
- bip340sig: a 64-byte bitcoin Elliptic Curve Schnorr signature as per [BIP-340](#)
- point: a 33-byte Elliptic Curve point (compressed encoding as per [SEC 1 standard](#))
- short_channel_id: an 8 byte value identifying a channel (see [BOLT #7](#))
- sciddir_or_pubkey: either 9 or 33 bytes referencing or identifying a node, respectively
 - if the first byte is 0 or 1, then an 8-byte short_channel_id follows for a total of 9 bytes
 - 0 for the first byte indicates this refers to node_id_1 in the channel_announcement for short_channel_id
 - 1 for the first byte indicates this refers to node_id_2 in the channel_announcement for short_channel_id (see [BOLT #7](#))
 - if the first byte is 2 or 3, then the value is a 33-byte point
- bigsize: a variable-length, unsigned integer similar to Bitcoin's CompactSize encoding, but big-endian. Described in [BigSize](#).
- utf8: a byte as part of a UTF-8 string. A writer MUST ensure an array of these is a valid UTF-8 string, a reader MAY reject any messages containing an array of these which is not a valid UTF-8 string.

Setup Messages

The init Message

Once authentication is complete, the first message reveals the features supported or required by this node, even if this is a reconnection.

[BOLT #9](#) specifies lists of features. Each feature is generally represented by 2 bits. The least-significant bit is numbered 0, which is *even*, and the next most significant bit is numbered 1, which is *odd*. For historical reasons, features are divided into global and local feature bitmasks.

A feature is *offered* if a peer set it in the `init` message for the current connection (as either even or odd). A feature is *negotiated* if either both peers offered it, or the local node offered it as even: it can assume the peer supports it, as it did not disconnect as it would be required to do.

The features field MUST be padded to bytes with 0s.

1. type: 16 (init)
2. data:
 - [u16:gflen]
 - [gflen*byte:globalfeatures]
 - [u16:flen]
 - [flen*byte:features]
 - [init_tlvs:tlvs]
3. tlv_stream: init_tlvs
4. types:
 1. type: 1 (networks)
 2. data:
 - [...*chain_hash:chains]
 3. type: 3 (remote_addr)
 4. data:
 - [...*byte:data]

The optional networks indicates the chains the node is interested in. The optional remote_addr can be used to circumvent NAT issues.

Requirements

The sending node: - MUST send init as the first Lightning message for any connection. - MUST set feature bits as defined in [BOLT #9](#). - MUST set any undefined feature bits to 0. - SHOULD NOT set features greater than 13 in globalfeatures. - SHOULD use the minimum length required to represent the features field. - SHOULD set networks to all chains it will gossip or open channels for. - SHOULD set remote_addr to reflect the remote IP address (and port) of an incoming connection, if the node is the receiver and the connection was done via IP. - if it sets remote_addr: - MUST set it to a valid address descriptor (1 byte type and data) as described in [BOLT 7](#). - SHOULD NOT set private addresses as remote_addr.

The receiving node: - MUST wait to receive init before sending any other messages. - MUST combine (bitwise OR) the two feature bitmaps into one logical features map. - MUST respond to known feature bits as specified in [BOLT #9](#). - upon receiving unknown *odd* feature bits that are non-zero: - MUST ignore the bit. - upon receiving unknown *even* feature bits that are non-zero: - MUST close the connection. - upon receiving networks containing no common chains - MAY close the connection. - if the feature vector does not set all known, transitive dependencies: - MUST close the connection. - MAY use the remote_addr to update its node_announcement

Rationale

There used to be two feature bitfields here, but for backwards compatibility they're now combined into one.

This semantic allows both future incompatible changes and future backward compatible changes. Bits should generally be assigned in pairs, in order that optional features may later become compulsory.

Nodes wait for receipt of the other's features to simplify error diagnosis when features are incompatible.

Since all networks share the same port, but most implementations only support a single network, the networks fields avoids nodes erroneously believing they will receive updates about their preferred network, or that they can open channels.

The error and warning Messages

For simplicity of diagnosis, it's often useful to tell a peer that something is incorrect.

1. type: 17 (error)
2. data:
 - [channel_id:channel_id]
 - [u16:len]
 - [len*byte:data]
3. type: 1 (warning)
4. data:
 - [channel_id:channel_id]
 - [u16:len]
 - [len*byte:data]

Requirements

The channel is referred to by channel_id, unless channel_id is 0 (i.e. all bytes are 0), in which case it refers to all channels.

The funding node using channel establishment v1 (open_channel): - for all error messages sent before (and including) the funding_created message: - MUST use temporary_channel_id in lieu of channel_id.

The fundee node using channel establishment v1 (accept_channel): - for all error messages sent before (and not including) the funding_signed message: - MUST use temporary_channel_id in lieu of channel_id.

The opener node using channel establishment v2 (open_channel2): - for all error messages sent before the accept_channel2 message is received: - MUST use temporary_channel_id in lieu of channel_id.

The accepter node using channel establishment v2 (open_channel2): - for all error messages sent before (and including) the accept_channel2 message: - MUST use temporary_channel_id in lieu of channel_id.

A sending node: - SHOULD send error for protocol violations or internal errors that make channels unusable or that make further communication unusable. - SHOULD send error with the unknown channel_id in reply to messages of type 32-255 related to unknown channels. - when sending error: - MUST fail the channel(s) referred to by the error message. - MAY set channel_id to all zero to indicate all channels. - when sending warning: - MAY set channel_id to all zero if the warning is not related to a specific channel. - MAY send an empty data field. - when failure was caused by an invalid signature check: - SHOULD include the raw, hex-encoded transaction in reply to a funding_created, funding_signed, closing_signed, or commitment_signed message.

The receiving node: - upon receiving error: - if channel_id is all zero: - MUST fail all channels with the sending node. - otherwise: - MUST fail the channel referred to by channel_id, if that channel is with the sending node. - upon receiving warning: - SHOULD log the message for later diagnosis. - MAY attempt shutdown if permitted at this point. - if no existing channel is referred to by channel_id: - MUST ignore the

message. - if data is not composed solely of printable ASCII characters (For reference: the printable character set includes byte values 32 through 126, inclusive): - SHOULD NOT print out data verbatim.

Rationale

There are unrecoverable errors that require an abort of conversations; if the connection is simply dropped, then the peer may retry the connection. It's also useful to describe protocol violations for diagnosis, as this indicates that one peer has a bug.

On the other hand, overuse of error messages has lead to implementations ignoring them (to avoid an otherwise expensive channel break), so the "warning" message was added to allow some degree of retry or recovery for spurious errors.

It may be wise not to distinguish errors in production settings, lest it leak information — hence, the optional data field.

Control Messages

The ping and pong Messages

In order to allow for the existence of long-lived TCP connections, at times it may be required that both ends keep alive the TCP connection at the application level. Such messages also allow obfuscation of traffic patterns.

1. type: 18 (ping)
2. data:
 - [u16:num_pong_bytes]
 - [u16:byteslen]
 - [byteslen*byte:ignored]

The pong message is to be sent whenever a ping message is received. It serves as a reply and also serves to keep the connection alive, while explicitly notifying the other end that the receiver is still active. Within the received ping message, the sender will specify the number of bytes to be included within the data payload of the pong message.

1. type: 19 (pong)
2. data:
 - [u16:byteslen]
 - [byteslen*byte:ignored]

Requirements

A node sending a ping message: - SHOULD set ignored to 0s. - MUST NOT set ignored to sensitive data such as secrets or portions of initialized memory. - if it doesn't receive a corresponding pong: - MAY close the network connection, - and MUST NOT fail the channels in this case.

A node sending a pong message: - SHOULD set ignored to 0s. - MUST NOT set ignored to sensitive data such as secrets or portions of initialized memory.

A node receiving a ping message: - if num_pong_bytes is less than 65532: - MUST respond by sending a pong message, with bytes len equal to num_pong_bytes. - otherwise (num_pong_bytes is **not** less than 65532): - MUST ignore the ping.

A node receiving a pong message: - if bytes len does not correspond to any ping's num_pong_bytes value it has sent: - MAY close the connection.

Rationale

The largest possible message is 65535 bytes; thus, the maximum sensible bytes len is 65531 — in order to account for the type field (pong) and the bytes len itself. This allows a convenient cutoff for num_pong_bytes to indicate that no reply should be sent.

Connections between nodes within the network may be long lived, as payment channels have an indefinite lifetime. However, it's likely that no new data will be exchanged for a significant portion of a connection's lifetime. Also, on several platforms it's possible that Lightning clients will be put to sleep without prior warning. Hence, a distinct ping message is used, in order to probe for the liveness of the connection on the other side, as well as to keep the established connection active.

Additionally, the ability for a sender to request that the receiver send a response with a particular number of bytes enables nodes on the network to create *synthetic* traffic. Such traffic can be used to partially defend against packet and timing analysis — as nodes can fake the traffic patterns of typical exchanges without applying any true updates to their respective channels.

When combined with the onion routing protocol defined in [BOLT #4](#), careful statistically driven synthetic traffic can serve to further bolster the privacy of participants within the network.

Limited precautions are recommended against ping flooding, however some latitude is given because of network delays. Note that there are other methods of incoming traffic flooding (e.g. sending *odd* unknown message types, or padding every message maximally).

Finally, the usage of periodic ping messages serves to promote frequent key rotations as specified within [BOLT #8](#).

Peer storage

The peer_storage and peer_storage_retrieval Messages

Nodes that advertise the option_provide_storage feature offer storing arbitrary data for their peers. The data stored must not exceed 65531 bytes, which lets it fit in lightning messages.

Nodes can verify that their option_provide_storage peers correctly store their data at each reconnection, by comparing the contents of the retrieved data with the last one they sent. However, nodes should not expect their peers to always have their latest data available.

Nodes ask their peers to store data using the peer_storage message and expect peers to return the latest data to them using the peer_storage_retrieval message:

1. type: 7 (peer_storage)

2. data:
 - [u16:length]
 - [length*byte:blob]
3. type: 9 (peer_storage_retrieval)
4. data:
 - [u16:length]
 - [length*byte:blob]

Requirements:

The sender of peer_storage: - MAY send peer_storage whenever necessary. - MUST limit its blob to 65531 bytes. - MUST encrypt the data in a manner that ensures its integrity upon receipt. - SHOULD pad the blob to ensure its length is always exactly 65531 bytes.

The receiver of peer_storage: - If it offered option_provide_storage: - if it has an open channel with the sender: - MUST store the message. - MAY store the message anyway.

- If it does store the message:
 - MAY delay storage to ratelimit peer to no more than one update per minute.
 - MUST replace the old blob with the latest received.
 - MUST send peer_storage_retrieval again after reconnection, after exchanging init messages.

The sender of peer_storage_retrieval: - MUST include the last blob it stored for that peer. - when all channels with that peer are closed: - SHOULD wait at least 2016 blocks before deleting the blob.

The receiver of peer_storage_retrieval: - when it receives peer_storage_retrieval with an outdated or irrelevant data: - MAY send a warning.

Rationale:

The peer_storage and peer_storage_retrieval messages enable nodes to securely store and share data with other nodes in the network, serving as a backup mechanism for important information. By utilizing them, nodes can safeguard crucial data, enhancing the network's resilience and reliability. Additionally, even if we don't have an open channel, some nodes might provide this service in exchange for some sats so they may store peer_storage.

Nodes should pad the blob to obscure its actual size, enhancing privacy by making size-based analysis more difficult for the receiver.

peer_storage_retrieval should not be sent after channel_reestablish because then the user wouldn't have an option to recover the node and update its state in case they lost data.

Nodes should send a peer_storage message whenever they wish to update the blob stored with their peers. This blob can be used to distribute encrypted data, which could be helpful in restoring the node.

Appendix A: BigSize Test Vectors

The following test vectors can be used to assert the correctness of a BigSize implementation used in the TLV format. The format is identical to the CompactSize encoding used in bitcoin, but replaces the little-endian encoding of multi-byte values with big-endian.

Values encoded with `BigSize` will produce an encoding of either 1, 3, 5, or 9 bytes depending on the size of the integer. The encoding is a piece-wise function that takes a `uint64` value `x` and produces:

<code>uint8(x)</code>	if <code>x < 0xfd</code>
<code>0xfd + be16(uint16(x))</code>	if <code>x < 0x10000</code>
<code>0xfe + be32(uint32(x))</code>	if <code>x < 0x100000000</code>
<code>0xff + be64(x)</code>	otherwise.

Here `+` denotes concatenation and `be16`, `be32`, and `be64` produce a big-endian encoding of the input for 16, 32, and 64-bit integers, respectively.

A value is said to be *minimally encoded* if it could not be encoded using fewer bytes. For example, a `BigSize` encoding that occupies 5 bytes but whose value is less than `0x10000` is not minimally encoded. All values decoded with `BigSize` should be checked to ensure they are minimally encoded.

BigSize Decoding Tests

The following is an example of how to execute the `BigSize` decoding tests.

```
func testReadBigSize(t *testing.T, test bigSizeTest) {
    var buf [8]byte
    r := bytes.NewReader(test.Bytes)
    val, err := tlv.ReadBigSize(r, &buf)
    if err != nil && err.Error() != test.ExpErr {
        t.Fatalf("expected decoding error: %v, got: %v",
            test.ExpErr, err)
    }

    // If we expected a decoding error, there's no point checking the value.
    if test.ExpErr != "" {
        return
    }

    if val != test.Value {
        t.Fatalf("expected value: %d, got %d", test.Value, val)
    }
}
```

A correct implementation should pass against these test vectors:

```
[
  {
    "name": "zero",
    "value": 0,
    "bytes": "00"
  },
  {
    "name": "one byte high",
    "value": 252,
    "bytes": "fc"
  },
  {
    "name": "two byte low",
    "value": 253,
```

```

    "bytes": "fd00fd"
  },
  {
    "name": "two byte high",
    "value": 65535,
    "bytes": "fdffff"
  },
  {
    "name": "four byte low",
    "value": 65536,
    "bytes": "fe00010000"
  },
  {
    "name": "four byte high",
    "value": 4294967295,
    "bytes": "feffffffff"
  },
  {
    "name": "eight byte low",
    "value": 4294967296,
    "bytes": "ff0000000100000000"
  },
  {
    "name": "eight byte high",
    "value": 18446744073709551615,
    "bytes": "ffffffffffffffffffff"
  },
  {
    "name": "two byte not canonical",
    "value": 0,
    "bytes": "fd00fc",
    "exp_error": "decoded bigsize is not canonical"
  },
  {
    "name": "four byte not canonical",
    "value": 0,
    "bytes": "fe0000ffff",
    "exp_error": "decoded bigsize is not canonical"
  },
  {
    "name": "eight byte not canonical",
    "value": 0,
    "bytes": "ff00000000ffffffff",
    "exp_error": "decoded bigsize is not canonical"
  },
  {
    "name": "two byte short read",
    "value": 0,
    "bytes": "fd00",
    "exp_error": "unexpected EOF"
  },
  {
    "name": "four byte short read",
    "value": 0,
    "bytes": "feffff",

```

```

        "exp_error": "unexpected EOF"
    },
    {
        "name": "eight byte short read",
        "value": 0,
        "bytes": "ffffffff",
        "exp_error": "unexpected EOF"
    },
    {
        "name": "one byte no read",
        "value": 0,
        "bytes": "",
        "exp_error": "EOF"
    },
    {
        "name": "two byte no read",
        "value": 0,
        "bytes": "fd",
        "exp_error": "unexpected EOF"
    },
    {
        "name": "four byte no read",
        "value": 0,
        "bytes": "fe",
        "exp_error": "unexpected EOF"
    },
    {
        "name": "eight byte no read",
        "value": 0,
        "bytes": "ff",
        "exp_error": "unexpected EOF"
    }
}
]

```

BigSize Encoding Tests

The following is an example of how to execute the BigSize encoding tests.

```

func testWriteBigSize(t *testing.T, test bigSizeTest) {
    var (
        w    bytes.Buffer
        buf [8]byte
    )
    err := tlv.WriteBigSize(&w, test.Value, &buf)
    if err != nil {
        t.Fatalf("unable to encode %d as bigsize: %v",
            test.Value, err)
    }

    if bytes.Compare(w.Bytes(), test.Bytes) != 0 {
        t.Fatalf("expected bytes: %v, got %v",
            test.Bytes, w.Bytes())
    }
}

```


A correct implementation should pass against the following test vectors:

```
[
  {
    "name": "zero",
    "value": 0,
    "bytes": "00"
  },
  {
    "name": "one byte high",
    "value": 252,
    "bytes": "fc"
  },
  {
    "name": "two byte low",
    "value": 253,
    "bytes": "fd00fd"
  },
  {
    "name": "two byte high",
    "value": 65535,
    "bytes": "fdffff"
  },
  {
    "name": "four byte low",
    "value": 65536,
    "bytes": "fe00010000"
  },
  {
    "name": "four byte high",
    "value": 4294967295,
    "bytes": "feffffffff"
  },
  {
    "name": "eight byte low",
    "value": 4294967296,
    "bytes": "ff0000000100000000"
  },
  {
    "name": "eight byte high",
    "value": 18446744073709551615,
    "bytes": "ffffffffffffffffffff"
  }
]
```

Appendix B: Type-Length-Value Test Vectors

The following tests assume that two separate TLV namespaces exist: n1 and n2.

The n1 namespace supports the following TLV types:

1. tlv_stream: n1
2. types:
 1. type: 1 (tlv1)

- 2. data:
 - [tu64:amount_msat]
- 1. type: 2 (tlv2)
- 2. data:
 - [short_channel_id:scid]
- 1. type: 3 (tlv3)
- 2. data:
 - [point:node_id]
 - [u64:amount_msat_1]
 - [u64:amount_msat_2]
- 1. type: 254 (tlv4)
- 2. data:
 - [u16:cltv_delta]

The n2 namespace supports the following TLV types:

- 1. tlv_stream: n2
- 2. types:
 - 1. type: 0 (tlv1)
 - 2. data:
 - [tu64:amount_msat]
 - 1. type: 11 (tlv2)
 - 2. data:
 - [tu32:cltv_expiry]

TLV Decoding Failures

The following TLV streams in any namespace should trigger a decoding failure:

- 1. Invalid stream: 0xfd
- 2. Reason: type truncated
- 3. Invalid stream: 0xfd01
- 4. Reason: type truncated
- 5. Invalid stream: 0xfd0001 00
- 6. Reason: not minimally encoded type
- 7. Invalid stream: 0xfd0101
- 8. Reason: missing length
- 9. Invalid stream: 0x0f fd
- 10. Reason: (length truncated)
- 11. Invalid stream: 0x0f fd26

- [illegible]

The following TLV streams in either namespace should trigger a decoding failure:

1. Invalid stream: 0x12 00
2. Reason: unknown even type.
3. Invalid stream: 0xfd0102 00
4. Reason: unknown even type.
5. Invalid stream: 0xfe01000002 00
6. Reason: unknown even type.
7. Invalid stream: 0xff0100000000000002 00
8. Reason: unknown even type.

The following TLV streams in namespace n1 should trigger a decoding failure:

1. Invalid stream: 0x01 09 ffffffff
2. Reason: greater than encoding length for n1s tlv1.
3. Invalid stream: 0x01 01 00
4. Reason: encoding for n1s tlv1s amount_msat is not minimal
5. Invalid stream: 0x01 02 0001
6. Reason: encoding for n1s tlv1s amount_msat is not minimal
7. Invalid stream: 0x01 03 000100
8. Reason: encoding for n1s tlv1s amount_msat is not minimal
9. Invalid stream: 0x01 04 00010000
10. Reason: encoding for n1s tlv1s amount_msat is not minimal

11. Invalid stream: 0x01 05 0001000000
12. Reason: encoding for n1s tlv1s amount_msat is not minimal
13. Invalid stream: 0x01 06 000100000000
14. Reason: encoding for n1s tlv1s amount_msat is not minimal
15. Invalid stream: 0x01 07 00010000000000
16. Reason: encoding for n1s tlv1s amount_msat is not minimal
17. Invalid stream: 0x01 08 0001000000000000
18. Reason: encoding for n1s tlv1s amount_msat is not minimal
19. Invalid stream: 0x02 07 01010101010101
20. Reason: less than encoding length for n1s tlv2.
21. Invalid stream: 0x02 09 0101010101010101
22. Reason: greater than encoding length for n1s tlv2.
23. Invalid stream: 0x03 21 023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb
24. Reason: less than encoding length for n1s tlv3.
25. Invalid stream: 0x03 29
023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb0000000000000001
26. Reason: less than encoding length for n1s tlv3.
27. Invalid stream: 0x03 30
023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb00000000000000010000000000000001
28. Reason: less than encoding length for n1s tlv3.
29. Invalid stream: 0x03 31
043da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb00000000000000010000000000000002
30. Reason: n1s node_id is not a valid point.
31. Invalid stream: 0x03 32
023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb00000000000000010000000000000001
32. Reason: greater than encoding length for n1s tlv3.
33. Invalid stream: 0xfd00fe 00
34. Reason: less than encoding length for n1s tlv4.
35. Invalid stream: 0xfd00fe 01 01

36. Reason: less than encoding length for n1s tlv4.
37. Invalid stream: 0xfd00fe 03 010101
38. Reason: greater than encoding length for n1s tlv4.
39. Invalid stream: 0x00 00
40. Reason: unknown even field for n1s namespace.

TLV Decoding Successes

The following TLV streams in either namespace should correctly decode, and be ignored:

1. Valid stream: 0x
2. Explanation: empty message
3. Valid stream: 0x21 00
4. Explanation: Unknown odd type.
5. Valid stream: 0xfd0201 00
6. Explanation: Unknown odd type.
7. Valid stream: 0xfd00fd 00
8. Explanation: Unknown odd type.
9. Valid stream: 0xfd00ff 00
10. Explanation: Unknown odd type.
11. Valid stream: 0xfe02000001 00
12. Explanation: Unknown odd type.
13. Valid stream: 0xff02000000000001 00
14. Explanation: Unknown odd type.

The following TLV streams in n1 namespace should correctly decode, with the values given here:

1. Valid stream: 0x01 00
2. Values: tlv1 amount_msat=0
3. Valid stream: 0x01 01 01
4. Values: tlv1 amount_msat=1
5. Valid stream: 0x01 02 0100
6. Values: tlv1 amount_msat=256

7. Valid stream: 0x01 03 010000
8. Values: tlv1 amount_msat=65536
9. Valid stream: 0x01 04 01000000
10. Values: tlv1 amount_msat=16777216
11. Valid stream: 0x01 05 0100000000
12. Values: tlv1 amount_msat=4294967296
13. Valid stream: 0x01 06 010000000000
14. Values: tlv1 amount_msat=1099511627776
15. Valid stream: 0x01 07 01000000000000
16. Values: tlv1 amount_msat=281474976710656
17. Valid stream: 0x01 08 0100000000000000
18. Values: tlv1 amount_msat=72057594037927936
19. Valid stream: 0x02 08 0000000000000226
20. Values: tlv2 scid=0x0x550
21. Valid stream: 0x03 31
023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb00000000000000010000000000000002
22. Values: tlv3 node_id=023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb
amount_msat_1=1 amount_msat_2=2
23. Valid stream: 0xfd00fe 02 0226
24. Values: tlv4 cltv_delta=550

TLV Stream Decoding Failure

Any appending of an invalid stream to a valid stream should trigger a decoding failure.

Any appending of a higher-numbered valid stream to a lower-numbered valid stream should not trigger a decoding failure.

In addition, the following TLV streams in namespace n1 should trigger a decoding failure:

1. Invalid stream: 0x02 08 0000000000000226 01 01 2a
2. Reason: valid TLV records but invalid ordering
3. Invalid stream: 0x02 08 0000000000000231 02 08 0000000000000451
4. Reason: duplicate TLV type

5. Invalid stream: 0x1f 00 0f 01 2a
6. Reason: valid (ignored) TLV records but invalid ordering
7. Invalid stream: 0x1f 00 1f 01 2a
8. Reason: duplicate TLV type (ignored)

The following TLV stream in namespace n2 should trigger a decoding failure:

1. Invalid stream: 0xffffffffffffff 00 00 00
2. Reason: valid TLV records but invalid ordering

Appendix C: Message Extension

This section contains examples of valid and invalid extensions on the init message. The base init message (without extensions) for these examples is 0x001000000000 (all features turned off).

The following init messages are valid:

- 0x001000000000: no extension provided
- 0x001000000000c9012acb0104: the extension contains two unknown *odd* TLV records (with types 0xc9 and 0xcb)

The following init messages are invalid:

- 0x00100000000001: the extension is present but truncated
- 0x001000000000ca012a: the extension contains unknown *even* TLV records (assuming that TLV type 0xca is unknown)
- 0x001000000000c90101c90102: the extension TLV stream is invalid (duplicate TLV record type 0xc9)

Note that when messages are signed, the *extension* is part of the signed bytes. Nodes should store the *extension* bytes even if they don't understand them to be able to correctly verify signatures.

Appendix D: Signed Integers Test Vectors

The following test vector show how signed integers (s8, s16, s32 and s64) are encoded using big-endian two's complement.

```
[
  {
    "value": 0,
    "bytes": "00"
  },
  {
    "value": 42,
    "bytes": "2a"
  },
  {
    "value": -42,
    "bytes": "d6"
  },
  {
```

```
    "value": 127,  
    "bytes": "7f"  
  },  
  {  
    "value": -128,  
    "bytes": "80"  
  },  
  {  
    "value": 128,  
    "bytes": "0080"  
  },  
  {  
    "value": -129,  
    "bytes": "ff7f"  
  },  
  {  
    "value": 15000,  
    "bytes": "3a98"  
  },  
  {  
    "value": -15000,  
    "bytes": "c568"  
  },  
  {  
    "value": 32767,  
    "bytes": "7fff"  
  },  
  {  
    "value": -32768,  
    "bytes": "8000"  
  },  
  {  
    "value": 32768,  
    "bytes": "00008000"  
  },  
  {  
    "value": -32769,  
    "bytes": "ffff7fff"  
  },  
  {  
    "value": 21000000,  
    "bytes": "01406f40"  
  },  
  {  
    "value": -21000000,  
    "bytes": "febf90c0"  
  },  
  {  
    "value": 2147483647,  
    "bytes": "7fffffff"  
  },  
  {  
    "value": -2147483648,  
    "bytes": "80000000"  
  },  
}
```



```

{
  "value": 2147483648,
  "bytes": "0000000080000000"
},
{
  "value": -2147483649,
  "bytes": "ffffffff7fffffff"
},
{
  "value": 500000000000,
  "bytes": "000000746a528800"
},
{
  "value": -500000000000,
  "bytes": "ffffff8b95ad7800"
},
{
  "value": 9223372036854775807,
  "bytes": "7fffffffffffffff"
},
{
  "value": -9223372036854775808,
  "bytes": "8000000000000000"
}
]

```

Acknowledgments

[TODO: (roasbeef); fin]

References

1. <http://www.unicode.org/charts/PDF/U2600.pdf>

Authors

[FIXME: Insert Author List]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

BOLT #9: Assigned Feature Flags

This document tracks the assignment of features flags in the init message ([BOLT #1](#)), as well as features fields in the channel_announcement and node_announcement messages ([BOLT #7](#)). The flags are tracked separately, since new flags will likely be added over time.

Some features were introduced and became so widespread they are ASSUMED to be present by all nodes, and can be safely ignored (and the semantics are only defined in prior revisions of this spec).

Flags are numbered from the least-significant bit, at bit 0 (i.e. 0x1, an *even* bit). They are generally assigned in pairs so that features can be introduced as optional (*odd* bits) and later upgraded to be compulsory (*even* bits), which will be refused by outdated nodes: see [BOLT #1: The init Message](#).

Some features don't make sense on a per-channels or per-node basis, so each feature defines how it is presented in those contexts. Some features may be required for opening a channel, but not a requirement for use of the channel, so the presentation of those features depends on the feature itself.

The Context column decodes as follows:

- I: presented in the init message.
- N: presented in the node_announcement messages
- C: presented in the channel_announcement message.
- C-: presented in the channel_announcement message, but always odd (optional).
- C+: presented in the channel_announcement message, but always even (required).
- 9: presented in [BOLT 11](#) invoices.
- B: presented in the allowed_features field of a blinded path.
- T: used in the channel_type field [when opening channels](#).

Bits	Name	Description	Context	Dependencies	Link
0/1	option_data_loss_protect	ASSUMED			
4/5	option_upfront_shutdown_script	Commits to a shutdown scriptpubkey when opening channel	IN		BOLT #2
6/7	gossip_queries	Peer has useful gossip to share			
8/9	var_onion_optin	ASSUMED			
10/11	gossip_queries_ex	Gossip queries can include additional information	IN		BOLT #7
12/13	option_static_remotekey	ASSUMED			
14/15	payment_secret	ASSUMED			[Routing Onion Specifica [bolt04]
16/17	basic_mpp	Node can receive basic multi-part payments	IN9	payment_secret	BOLT #4
18/19	option_support_large_channel	Can create large channels	IN		BOLT #2
22/23	option_anchors	Anchor commitment type with zero fee HTLC transactions	INT		BOLT #3 lightning dev

Bits	Name	Description	Context	Dependencies	Link
24/25	option_route_blinding	Node supports blinded paths	IN9		BOLT #4
26/27	option_shutdown_anysegwit	Future segwit versions allowed in shutdown	IN		BOLT #2
28/29	option_dual_fund	Use v2 of channel open, enables dual funding	IN		BOLT #2
34/35	option_quiesce	Support for stfu message	IN		BOLT #2
36/37	option_attribution_data	Can generate/relay attribution data and fulfillment_payload	IN9		BOLT #4 successful payment
38/39	option_onion_messages	Can forward onion messages	IN		BOLT #4
40/41	zero_fee_commitments	Zero-fee commitment and HTLC transactions	IN	option_channel_type	BOLT #3
42/43	option_provide_storage	Can store other nodes' encrypted backup data	IN		BOLT #1
44/45	option_channel_type	ASSUMED			
46/47	option_scid_alias	Supply channel aliases for routing	INT		BOLT #2
48/49	option_payment_metadata	Payment metadata in tlv record	9		BOLT #1
50/51	option_zeroconf	Understands zeroconf channel types	INT	option_scid_alias	BOLT #2
60/61	option_simple_close	Simplified closing negotiation	IN	option_shutdown_anysegwit	BOLT #2
62/63	option_splice	Allows replacing the funding transaction with a new one	IN		BOLT #2
66/67	option_onion_messages_only_channels	Only accepts onion messages from peers with a channel	IN	option_onion_messages	BOLT #4

Requirements

The origin node: * If it supports a feature above, SHOULD set the corresponding odd bit in all feature fields indicated by the Context column unless indicated that it must set the even feature bit instead. * If it requires a feature above, MUST set the corresponding even feature bit in all feature fields indicated by the Context column, unless indicated that it must set the odd feature bit instead. * MUST NOT set feature bits it does not support. * MUST NOT set feature bits in fields not specified by the table above. * MUST NOT set both the optional and mandatory bits. * MUST set all transitive feature dependencies. * MUST support: * `var_onion_optin`

The receiving node: * if both the optional and the mandatory feature bits in a pair are set, the feature should be treated as mandatory.

The requirements for receiving specific bits are defined in the linked sections in the table above. The requirements for feature bits that are not defined above can be found in [BOLT #1: The init Message](#).

Rationale

Note that for feature flags which are available in both the `node_announcement` and [BOLT 11](#) invoice contexts, the features as set in the [BOLT 11](#) invoice should override those set in the `node_announcement`. This keeps things consistent with the unknown features behavior as specified in [BOLT 7](#).

The origin must set all transitive feature dependencies in order to create a well-formed feature vector. By validating all known dependencies up front, this simplifies logic gated on a single feature bit; the feature's dependencies are known to be set, and do not need to be validated at every feature gate.



This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

Opening and Running Channels

A Lightning channel is a Bitcoin output that two people share, plus a set of rules for updating who owns how much without going back to the chain every time.

BOLT 02 is the peer protocol for that. It covers the whole channel lifecycle: opening (`open_channel` / `accept_channel` / `funding_created` / `funding_signed`), going live (`channel_ready`), sending payments with HTLCs (`update_add_htlc` and friends), committing a new state (`commitment_signed` / `revoke_and_ack`), and cooperative close. If you're implementing a node, you'll live in this document.

BOLT 03 is the Bitcoin underneath. It specifies the funding transaction, the commitment transactions, HTLC-timeout and HTLC-success transactions, and the scripts and signature formats that make the penalty model work. BOLT 02 is the conversation; BOLT 03 is the money.

You can read BOLT 02 first for the flow, then BOLT 03 when you need to know exactly what gets signed. Or keep them side by side. They were meant to be a pair.

In this chapter

- BOLT 02 — Peer protocol for channel management
- BOLT 03 — Bitcoin transactions and scripts

BOLT #2: Peer Protocol for Channel Management

The peer channel protocol has three phases: establishment, normal operation, and closing.

Table of Contents

- [Channel](#)
 - [Definition of channel_id](#)
 - [Interactive Transaction Construction](#)
 - [Set-Up and Vocabulary](#)
 - [Fee Responsibility](#)
 - [Overview](#)
 - [The tx_add_input Message](#)
 - [The tx_add_output Message](#)
 - [The tx_remove_input and tx_remove_output Messages](#)
 - [The tx_complete Message](#)
 - [The tx_signatures Message](#)
 - [The tx_init_rbf Message](#)
 - [The tx_ack_rbf Message](#)
 - [The tx_abort Message](#)
 - [Channel Establishment v1](#)
 - [The open_channel Message](#)
 - [The accept_channel Message](#)
 - [The funding_created Message](#)
 - [The funding_signed Message](#)
 - [The channel_ready Message](#)
 - [Channel Establishment v2](#)
 - [The open_channel2 Message](#)
 - [The accept_channel2 Message](#)
 - [Funding Composition](#)
 - [The commitment_signed Message](#)
 - [Sharing funding signatures: tx_signatures](#)
 - [Fee bumping: tx_init_rbf and tx_ack_rbf](#)
 - [Channel Quiescence](#)
 - [Channel Splicing](#)
 - [The splice_init Message](#)
 - [The splice_ack Message](#)
 - [Splice Transaction Construction](#)
 - [Splice Completion](#)
 - [Channel Close](#)
 - [Closing Initiation: shutdown](#)
 - [Closing Negotiation: closing_complete and closing_sig](#)

- [Legacy Closing Negotiation: closing_signed](#)
- [Normal Operation](#)
 - [Forwarding HTLCs](#)
 - [cltv_expiry_delta Selection](#)
 - [Adding an HTLC: update_add_htlc](#)
 - [Removing an HTLC: update_fulfill_htlc, update_fail_htlc, and update_fail_malformed_htlc](#)
 - [Batching channel messages](#)
 - [Committing Updates So Far: commitment_signed](#)
 - [Completing the Transition to the Updated State: revoke_and_ack](#)
 - [Updating Fees: update_fee](#)
- [Message Retransmission: channel_reestablish_message](#)
- [Authors](#)

Channel

Definition of channel_id

Some messages use a `channel_id` to identify the channel. It's derived from the funding transaction by combining the `funding_txid` and the `funding_output_index`, using big-endian exclusive-OR (i.e. `funding_output_index` alters the last 2 bytes).

Prior to channel establishment, a `temporary_channel_id` is used, which is a random nonce.

Note that as duplicate `temporary_channel_ids` may exist from different peers, APIs which reference channels by their channel id before the funding transaction is created are inherently unsafe. The only protocol-provided identifier for a channel before `funding_created` has been exchanged is the (`source_node_id`, `destination_node_id`, `temporary_channel_id`) tuple. Note that any such APIs which reference channels by their channel id before the funding transaction is confirmed are also not persistent - until you know the script pubkey corresponding to the funding output nothing prevents duplicative channel ids.

`channel_id`, v2

For channels established using the v2 protocol, the `channel_id` is the `SHA256(lesser-revocation-basepoint || greater-revocation-basepoint)`, where the lesser and greater is based off the order of the basepoint.

When sending `open_channel2`, the peer's revocation basepoint is unknown. A `temporary_channel_id` must be computed by using a zeroed out basepoint for the non-initiator.

When sending `accept_channel2`, the `temporary_channel_id` from `open_channel2` must be used, to allow the initiator to match the response to its request.

Rationale

The revocation basepoints must be remembered by both peers for correct operation anyway. They're known after the first exchange of messages, obviating the need for a `temporary_channel_id` in subsequent messages. By mixing information from both sides, they avoid `channel_id` collisions, and they remove the dependency on the funding txid.

Interactive Transaction Construction

Interactive transaction construction allows two peers to collaboratively build a transaction for broadcast. This protocol is the foundation for dual-funded channels establishment (v2).

Set-Up and Vocabulary

There are two parties to a transaction construction: an *initiator* and a *non-initiator*. The *initiator* is the peer which initiates the protocol, e.g. for channel establishment v2 the *initiator* would be the peer which sends `open_channel2`.

The protocol makes the following assumptions:

- The feerate for the transaction is known.
- The `dust_limit` for the transaction is known.
- The `nLocktime` for the transaction is known.
- The `nVersion` for the transaction is known.

Fee Responsibility

The *initiator* is responsible for paying the fees for the following fields, to be referred to as the `common` fields.

- `version`
- `segwit` marker + flag
- `input` count
- `output` count
- `locktime`

The rest of the transaction bytes' fees are the responsibility of the peer who contributed that input or output via `tx_add_input` or `tx_add_output`, at the agreed upon feerate.

Overview

The *initiator* initiates the interactive transaction construction protocol with `tx_add_input`. The *non-initiator* responds with any of `tx_add_input`, `tx_add_output`, `tx_remove_input`, `tx_remove_output`, or `tx_complete`. The protocol continues with the synchronous exchange of interactive transaction protocol messages until both nodes have sent and received a consecutive `tx_complete`. This is a turn-based protocol.


Once peers have exchanged consecutive `tx_completes`, the interactive transaction construction protocol is considered concluded. Both peers should construct the transaction and fail the negotiation if an error is discovered.

This protocol is expressly designed to allow for parallel, multi-party sessions to collectively construct a single transaction. This preserves the ability to open multiple channels in a single transaction. While `serial_ids` are generally chosen randomly, to maintain consistent transaction ordering across all peer sessions, it is simplest to reuse received `serial_ids` when forwarding them to other peers, inverting the bottom bit as necessary to satisfy the parity requirement.

Here are a few example exchanges.

initiator only


A, the *initiator*, has two inputs and an output (the funding output). B, the *non-initiator* has nothing to contribute.



initiator and non-initiator

A, the *initiator*, contributes 2 inputs and an output that they then remove. B, the *non-initiator*, contributes 1 input and an output, but waits until A adds a second input before contributing.

Note that if A does not send a second input, the negotiation will end without B's contributions.



The tx_add_input Message

This message contains a transaction input.

1. type: 66 (tx_add_input)
2. data:
 - [channel_id:channel_id]
 - [u64:serial_id]
 - [u16:prevtx_len]
 - [prevtx_len*byte:prevtx]
 - [u32:prevtx_vout]
 - [u32:sequence]
 - [tx_add_input_tlvs:tlvs]
3. tlv_stream: tx_add_input_tlvs
4. types:
 1. type: 0 (shared_input_txid)

2. data:
- [sha256:funding_txid]

Requirements

The sending node: - MUST add all sent inputs to the transaction - MUST use a unique `serial_id` for each input currently added to the transaction - MUST set sequence to be less than or equal to 4294967293 (0xFFFFFFFF) - MUST NOT re-transmit inputs it has received from the peer - if is the *initiator*: - MUST send even `serial_ids` - if is the *non-initiator*: - MUST send odd `serial_ids`

The receiving node: - MUST add all received inputs to the transaction - MUST fail the negotiation if: - sequence is set to 0xFFFFFFFF or 0xFFFFFFFF - if `prevtx_len` is 0: - `shared_input_txid` is not set - `shared_input_txid` and `prevtx_vout` don't match the previous funding output - a previously added (and not removed) input already exists with `shared_input_txid` set - if `prevtx_len` is not 0: - `prevtx` and `prevtx_vout` are identical to a previously added (and not removed) input - `prevtx` is not a valid transaction - `prevtx_vout` is greater or equal to the number of outputs on `prevtx` - the `scriptPubKey` of the `prevtx_vout` output of `prevtx` is not exactly a 1-byte push opcode (for the numeric values 0 to 16) followed by a data push between 2 and 40 bytes - the `serial_id` is already included in the transaction - the `serial_id` has the wrong parity - if has received 4096 `tx_add_input` messages during this negotiation

Rationale

Each node must know the set of the transaction inputs. The *non-initiator* MAY omit this message.

`serial_id` is a randomly chosen number which uniquely identifies this input. Inputs in the constructed transaction MUST be sorted by `serial_id`.

`prevtx` is the serialized transaction that contains the output this input spends, used to verify that the input is non-malleable. It can be omitted (`prevtx_len` set to 0) when both peers already know that the input is non-malleable (e.g. when it is the previous funding output).

`prevtx_vout` is the index of the output being spent.

sequence is the sequence number of this input: it must signal replaceability, and the same value should be used across implementations to avoid on-chain fingerprinting.

Liquidity griefing

When sending `tx_add_input`, senders have no guarantee that the remote node will complete the protocol in a timely manner. Malicious remote nodes could delay messages or stop responding, which can result in a partially created transaction that cannot be broadcast by the honest node. If the honest node is locking the corresponding UTXO exclusively for this remote node, this can be exploited to lock up the honest node's liquidity.

It is thus recommended that implementations keep UTXOs unlocked and actively reuse them in concurrent sessions, which guarantees that transactions created with honest nodes double-spend pending transactions with malicious nodes at no additional cost for the honest node.

Unfortunately, this will also create conflicts between concurrent sessions with honest nodes. This is a reasonable trade-off though because:

- on-chain funding attempts are relatively infrequent operations
- honest nodes should complete the protocol quickly, reducing the risk of conflicts
- failed attempts can simply be retried at no cost

The tx_add_output Message

This message adds a transaction output.

1. type: 67 (tx_add_output)
2. data:
 - [channel_id:channel_id]
 - [u64:serial_id]
 - [u64:sats]
 - [u16:scriptlen]
 - [scriptlen*byte:script]

Requirements

Either node: - MAY omit this message

The sending node: - MUST add all sent outputs to the transaction - if is the *initiator*: - MUST send even serial_ids - if is the *non-initiator*: - MUST send odd serial_ids

The receiving node: - MUST add all received outputs to the transaction - MUST accept P2WSH, P2WPKH, P2TR scripts - MAY fail the negotiation if script is non-standard - MUST fail the negotiation if: - the serial_id is already included in the transaction - the serial_id has the wrong parity - it has received 4096 tx_add_output messages during this negotiation - the sats amount is less than the dust_limit - the sats amount is greater than 2,100,000,000,000,000 (MAX_MONEY)

Rationale

Each node must know the set of the transaction outputs.

serial_id is a randomly chosen number which uniquely identifies this output. Outputs in the constructed transaction MUST be sorted by serial_id.

sats is the satoshi value of the output.

script is the scriptPubKey for the output (with its length omitted). scripts are not required to follow standardness rules: non-standard scripts such as OP_RETURN may be accepted, but the corresponding transaction may fail to relay across the network.

The tx_remove_input and tx_remove_output Messages

This message removes an input from the transaction.

1. type: 68 (tx_remove_input)

2. data:
 - [channel_id:channel_id]
 - [u64:serial_id]

This message removes an output from the transaction.

1. type: 69 (tx_remove_output)
2. data:
 - [channel_id:channel_id]
 - [u64:serial_id]

Requirements

The sending node: - MUST NOT send a tx_remove with a serial_id it did not add to the transaction or has already been removed

The receiving node: - MUST remove the indicated input or output from the transaction - MUST fail the negotiation if: - the input or output identified by the serial_id was not added by the sender - the serial_id does not correspond to a currently added input (or output)

The tx_complete Message

This message signals the conclusion of a peer's transaction contributions.

1. type: 70 (tx_complete)
2. data:
 - [channel_id:channel_id]

Requirements

The nodes: - MUST send this message in succession to conclude this protocol

The receiving node: - MUST use the negotiated inputs and outputs to construct a transaction - MUST fail the negotiation if: - the peer's total input satoshis is less than their outputs. One MUST account for the peer's portion of the funding output when verifying compliance with this requirement. - the peer's paid feerate does not meet or exceed the agreed feerate (based on the minimum fee). - if is the *non-initiator*: - the *initiator*'s fees do not cover the common fields - there are more than 252 inputs - there are more than 252 outputs - the estimated weight of the tx is greater than 400,000 (MAX_STANDARD_TX_WEIGHT)

Rationale

To signal the conclusion of exchange of transaction inputs and outputs.

Upon successful exchange of tx_complete messages, both nodes should construct the transaction and proceed to the next portion of the protocol. For channel establishment v2, exchanging commitment transactions.

For the minimum fee calculation see [BOLT #3](#).

The maximum inputs and outputs are capped at 252. This effectively fixes the byte size of the input and output counts on the transaction to one (1).

The tx_signatures Message

1. type: 71 (tx_signatures)
2. data:
 - [channel_id:channel_id]
 - [sha256:txid]
 - [u16:num_witnesses]
 - [num_witnesses*witness:witnesses]
 - [tx_signatures_tlvs:tlvs]
3. subtype: witness
4. data:
 - [u16:len]
 - [len*byte:witness_data]
5. tlv_stream: tx_signatures_tlvs
6. types:
 1. type: 0 (shared_input_signature)
 2. data:
 - [signature:signature]

Requirements

The sending node: - if it has the lowest total satoshis contributed, as defined by total tx_add_input values, or both peers have contributed equal amounts but it has the lowest node_id (sorted lexicographically): - MUST transmit their tx_signatures first - MUST order the witnesses by the serial_id of the input they correspond to - num_witnesses MUST equal the number of inputs they added - MUST use the SIGHASH_ALL (0x01) flag on each signature

The receiving node: - MUST fail the negotiation if: - the message contains an empty witness - the number of witnesses does not equal the number of inputs added by the sending node - the txid does not match the txid of the transaction - the witnesses are non-standard - a signature uses a flag that is not SIGHASH_ALL (0x01) - SHOULD apply the witnesses to the transaction and broadcast it - MUST reply with their tx_signatures if not already transmitted

Rationale

A strict ordering is used to decide which peer sends tx_signatures first. This prevents deadlocks where each peer is waiting for the other peer to send tx_signatures, and enables multiparty tx collaboration.

The witness_data is encoded as per bitcoin's wire protocol (a CompactSize number of elements, with each element a CompactSize length and that many bytes following).

While the minimum fee is calculated and verified at tx_complete conclusion, it is possible for the fee for the exchanged witness data to be underpaid. It is the responsibility of the sending peer to correctly account for the required fee.

The tx_init_rbf Message

This message initiates a replacement of the transaction after it's been completed.

1. type: 72 (tx_init_rbf)
2. data:
 - [channel_id:channel_id]
 - [u32:locktime]
 - [u32:feerate]
 - [tx_init_rbf_tlvs:tlvs]
3. tlv_stream: tx_init_rbf_tlvs
4. types:
 1. type: 0 (funding_output_contribution)
 2. data:
 - [s64:satoshis]
 3. type: 2 (require_confirmed_inputs)

Requirements

The sender: - MUST set feerate greater than or equal to the maximum of: - 25/24 times the feerate of the previously constructed transaction, rounded down. - 25 sat per kw greater than the feerate of the previously constructed transaction. - If it contributes to the transaction's funding output: - MUST set funding_output_contribution - If it requires the receiving node to only use confirmed inputs: - MUST set require_confirmed_inputs - If it contributed to previous transactions: - MUST ensure that the new transaction double-spends all other attempts, by sending tx_add_input with at least one input from each previous transaction construction attempt.

The recipient: - MUST respond either with tx_abort or with tx_ack_rbf - MUST respond with tx_abort if: - the feerate is not greater than or equal to the maximum of: - 25/24 times the feerate of the last successfully constructed transaction, rounded down. - 25 sat per kw greater than the feerate of the last successfully constructed transaction - MAY send tx_abort for any reason - MUST fail the negotiation if: - require_confirmed_inputs is set but it cannot provide confirmed inputs

Rationale

feerate is the feerate this transaction will pay. It must be at least the maximum of 25/24 times the previous feerate (rounded down) and 25 sat per kw above the previous feerate. The multiplicative rule ensures progress at higher feerates, while the additive floor of 25 sat per kw (0.1 sat/vB) ensures that the replacement transaction meets BIP125's minimum relay fee requirement (default incrementalRelayFee of 0.1 sat/vB since Bitcoin Core v30.0), which the multiplicative increase alone could fail to satisfy at low feerates.

E.g. if the last feerate was 520, the next sent feerate must be at least 545 ($\max(520 * 25 / 24, 520 + 25) = \max(541, 545) = 545$).

If the previous transaction confirms in the middle of an RBF attempt, the attempt MUST be abandoned.

funding_output_contribution is the amount of satoshis that this peer will contribute to the funding output of the transaction, when there is such an output. Note that it may be different from the contribution made in the previously completed transaction. If omitted, the sender is not contributing to the funding output.

The tx_ack_rbf Message

1. type: 73 (tx_ack_rbf)
2. data:
 - [channel_id:channel_id]
 - [tx_ack_rbf_tlvs:tlvs]
3. tlv_stream: tx_ack_rbf_tlvs
4. types:
 1. type: 0 (funding_output_contribution)
 2. data:
 - [s64:satoshis]
 3. type: 2 (require_confirmed_inputs)

Requirements

The sender: - If it contributes to the transaction's funding output: - MUST set funding_output_contribution - If it requires the receiving node to only use confirmed inputs: - MUST set require_confirmed_inputs - If it contributed to previous transactions: - MUST ensure that the new transaction double-spends all other attempts, by sending tx_add_input with at least one input from each previous transaction construction attempt.

The recipient: - MUST respond with tx_abort or with a tx_add_input message, restarting the interactive tx collaboration protocol. - MUST fail the negotiation if: - require_confirmed_inputs is set but it cannot provide confirmed inputs

Rationale

funding_output_contribution is the amount of satoshis that this peer will contribute to the funding output of the transaction, when there is such an output. Note that it may be different from the contribution made in the previously completed transaction. If omitted, the sender is not contributing to the funding output.

It's recommended that a peer, rather than fail the RBF negotiation due to a large feerate change, instead stop contributing to the funding output, and decline to participate further in the transaction (by not contributing, they may obtain incoming liquidity at no cost).

The tx_abort Message

1. type: 74 (tx_abort)
2. data:
 - [channel_id:channel_id]
 - [u16:len]
 - [len*byte:data]

Requirements

A sending node: - MUST NOT have already transmitted tx_signatures - SHOULD forget the current negotiation and reset their state. - MAY send an empty data field. - when failure was caused by an invalid signature check: - SHOULD include the raw, hex-encoded transaction in reply to a tx_signatures or commitment_signed message.

A receiving node: - if they have already sent tx_signatures to the peer: - MUST NOT forget the channel until any inputs to the negotiated tx have been spent. - if they have not sent tx_signatures: - SHOULD forget the current negotiation and reset their state. - if they have not sent tx_abort: - MUST echo back tx_abort - if data is not composed solely of printable ASCII characters (For reference: the printable character set includes byte values 32 through 126, inclusive): - SHOULD NOT print out data verbatim.

Rationale

A receiving node, if they've already sent their tx_signatures has no guarantee that the transaction won't be signed and published by their peer. They must remember the transaction and channel (if appropriate) until the transaction is no longer eligible to be spent (i.e. any input has been spent in a different transaction).

The tx_abort message allows for the cancellation of an in progress negotiation, and a return to the initial starting state. It is distinct from the error message, which triggers a channel close.

Echoing back tx_abort allows the peer to ack that they've seen the abort message, permitting the originating peer to terminate the in-flight process without worrying about stale messages.


Channel Establishment v1

After authenticating and initializing a connection ([BOLT #8](#) and [BOLT #1](#), respectively), channel establishment may begin.

There are two pathways for establishing a channel, a legacy version presented here, and a second version ([below](#)). Which channel establishment protocols are available for use is negotiated in the init message.

This consists of the funding node (funder) sending an open_channel message, followed by the responding node (fundee) sending accept_channel. With the channel parameters locked in, the funder is able to create the funding transaction and both versions of the commitment transaction, as described in [BOLT #3](#). The funder then sends the outpoint of the funding output with the funding_created message, along with the signature for the fundee's version of the commitment transaction. Once the fundee learns the funding outpoint, it's able to generate the signature for the funder's version of the commitment transaction and send it over using the funding_signed message.

Once the channel funder receives the funding_signed message, it must broadcast the funding transaction to the Bitcoin network. After the funding_signed message is sent/received, both sides should wait for the funding transaction to enter the blockchain and reach the specified depth (number of confirmations). After both sides have sent the channel_ready message, the channel is established and can begin normal operation. The channel_ready message includes information that will be used to construct channel authentication proofs.



- where node A is 'funder' and node B is 'fundee'

If this fails at any stage, or if one node decides the channel terms offered by the other node are not suitable, the channel establishment fails.

Note that multiple channels can operate in parallel, as all channel messages are identified by either a `temporary_channel_id` (before the funding transaction is created) or a `channel_id` (derived from the funding transaction).

The `open_channel` Message

This message contains information about a node and indicates its desire to set up a new channel. This is the first step toward creating the funding transaction and both versions of the commitment transaction.

1. type: 32 (`open_channel`)
2. data:
 - [chain_hash:chain_hash]
 - [32*byte:temporary_channel_id]
 - [u64:funding_satoshis]
 - [u64:push_msat]
 - [u64:dust_limit_satoshis]
 - [u64:max_htlc_value_in_flight_msat]
 - [u64:channel_reserve_satoshis]
 - [u64:htlc_minimum_msat]
 - [u32:feerate_per_kw]
 - [u16:to_self_delay]
 - [u16:max_accepted_htlcs]
 - [point:funding_pubkey]
 - [point:revocation_basepoint]
 - [point:payment_basepoint]
 - [point:delayed_payment_basepoint]
 - [point:htlc_basepoint]
 - [point:first_per_commitment_point]
 - [byte:channel_flags]
 - [open_channel_tlvs:tlvs]
3. tlv_stream: open_channel_tlvs
4. types:
 1. type: 0 (`upfront_shutdown_script`)
 2. data:
 - [...*byte:shutdown_scriptpubkey]
 3. type: 1 (`channel_type`)
 4. data:
 - [...*byte:type]

The `chain_hash` value denotes the exact blockchain that the opened channel will reside within. This is usually the genesis hash of the respective blockchain. The existence of the `chain_hash` allows nodes to open channels

across many distinct blockchains as well as have channels within multiple blockchains opened to the same peer (if it supports the target chains).

The `temporary_channel_id` is used to identify this channel on a per-peer basis until the funding transaction is established, at which point it is replaced by the `channel_id`, which is derived from the funding transaction.

`funding_satoshis` is the amount the sender is putting into the channel. `push_msat` is an amount of initial funds that the sender is unconditionally giving to the receiver. `dust_limit_satoshis` is the threshold below which outputs should not be generated for this node's commitment or HTLC transactions (i.e. HTLCs below this amount plus HTLC transaction fees are not enforceable on-chain). This reflects the reality that tiny outputs are not considered standard transactions and will not propagate through the Bitcoin network.

`channel_reserve_satoshis` is the minimum amount that the other node is to keep as a direct payment.

`htlc_minimum_msat` indicates the smallest value HTLC this node will accept.

`max_htlc_value_in_flight_msat` is a cap on total value of outstanding HTLCs offered by the remote node, which allows the local node to limit its exposure to HTLCs; similarly, `max_accepted_htlcs` limits the number of outstanding HTLCs the remote node can offer.

`feerate_per_kw` indicates the initial fee rate in satoshi per 1000-weight (i.e. 1/4 the more normally-used 'satoshi per 1000 vbytes') that this side will pay for commitment and HTLC transactions, as described in [BOLT #3](#) (this can be adjusted later with an `update_fee` message).

`to_self_delay` is the number of blocks that the other node's to-self outputs must be delayed, using `OP_CHECKSEQUENCEVERIFY` delays; this is how long it will have to wait in case of breakdown before redeeming its own funds.

`funding_pubkey` is the public key in the 2-of-2 multisig script of the funding transaction output.

The various `_basepoint` fields are used to derive unique keys as described in [BOLT #3](#) for each commitment transaction. Varying these keys ensures that the transaction ID of each commitment transaction is unpredictable to an external observer, even if one commitment transaction is seen; this property is very useful for preserving privacy when outsourcing penalty transactions to third parties.

`first_per_commitment_point` is the per-commitment point to be used for the first commitment transaction,

Only the least-significant bit of `channel_flags` is currently defined: `announce_channel`. This indicates whether the initiator of the funding flow wishes to advertise this channel publicly to the network, as detailed within [BOLT #7](#).

The `shutdown_scriptpubkey` allows the sending node to commit to where funds will go on mutual close, which the remote node should enforce even if a node is compromised later.

The `option_support_large_channel` is a feature used to let everyone know this node will accept `funding_satoshis` greater than or equal to 2^{24} . Since it's broadcast in the `node_announcement` message other nodes can use it to identify peers willing to accept large channel even before exchanging the `init` message with them.

Defined Channel Types

Channel types are an explicit enumeration: for convenience of future definitions they reuse even feature bits, but they are not an arbitrary combination (they represent the persistent features which affect the channel operation).

The currently defined basic types are: - `option_static_remotekey` (bit 12) - `option_anchors` and `option_static_remotekey` (bits 22 and 12) - `zero_fee_commitments` (bit 40)

Each basic type has the following variations allowed: - `option_scid_alias` (bit 46) - `option_zeroconf` (bit 50)

Requirements

The sending node: - MUST ensure the `chain_hash` value identifies the chain it wishes to open the channel within. - MUST ensure `temporary_channel_id` is unique from any other channel ID with the same peer. - if both nodes advertised `option_support_large_channel`: - MAY set `funding_satoshis` greater than or equal to 2^{24} satoshi. - otherwise: - MUST set `funding_satoshis` to less than 2^{24} satoshi. - MUST set `push_msat` to equal or less than $1000 * \text{funding_satoshis}$. - MUST set `funding_pubkey`, `revocation_basepoint`, `htlc_basepoint`, `payment_basepoint`, and `delayed_payment_basepoint` to valid secp256k1 pubkeys in compressed format. - MUST set `first_per_commitment_point` to the per-commitment point to be used for the initial commitment transaction, derived as specified in [BOLT #3](#). - MUST set `channel_reserve_satoshis` greater than or equal to `dust_limit_satoshis`. - MUST set undefined bits in `channel_flags` to 0. - if both nodes advertised the `option_upfront_shutdown_script` feature: - MUST include `upfront_shutdown_script` with either a valid `shutdown_scriptpubkey` as required by `shutdown_scriptpubkey`, or a zero-length `shutdown_scriptpubkey` (ie. `0x0000`). - otherwise: - MAY include `upfront_shutdown_script`. - if it includes `open_channel_tlvs`: - MUST include `upfront_shutdown_script`. - MUST set `channel_type`: - MUST set it to a defined type representing the type it wants. - MUST use the smallest bitmap possible to represent the channel type. - SHOULD NOT set it to a type containing a feature which was not negotiated. - if `channel_type` includes `zero_fee_commitments`: - MUST set `feerate_per_kw` to 0. - if `announce_channel` is true (not 0): - MUST NOT send `channel_type` with the `option_scid_alias` bit set.

The sending node SHOULD: - set `to_self_delay` sufficient to ensure the sender can irreversibly spend a commitment transaction output, in case of misbehavior by the receiver. - set `dust_limit_satoshis` to a sufficient value to allow commitment transactions to propagate through the Bitcoin network. - set `htlc_minimum_msat` to the minimum value HTLC it's willing to accept from this peer. - when using `zero_fee_commitments` or `option_anchors`: - set `dust_limit_satoshis` to a sufficient value to ensure that it is economical to spend, taking the cost of HTLC transactions into account. - otherwise (not using `zero_fee_commitments` or `option_anchors`): - set `feerate_per_kw` to at least the rate it estimates would cause the transaction to be immediately included in a block.

The receiving node MUST: - ignore undefined bits in `channel_flags`. - if the message doesn't include a `channel_type`: - fail the channel. - if the connection has been re-established after receiving a previous `open_channel`, BUT before receiving a `funding_created` message: - accept a new `open_channel` message. - discard the previous `open_channel` message. - if `option_dual_fund` has been negotiated: - fail the channel.

The receiving node MAY fail the channel if: - `announce_channel` is false (0), yet it wishes to publicly announce the channel. - `funding_satoshis` is too small. - it considers `htlc_minimum_msat` too large. - it considers `max_htlc_value_in_flight_msat` too small. - it considers `channel_reserve_satoshis` too large. - it considers `max_accepted_htlcs` too small. - it considers `dust_limit_satoshis` too large (see [BOLT 3](#)).

The receiving node MUST fail the channel if: - the `chain_hash` value is set to a hash of a chain that is unknown to the receiver. - `push_msat` is greater than `funding_satoshis * 1000`. - `to_self_delay` is unreasonably large. - `channel_type` includes `zero_fee_commitments` and: - `max_accepted_htlcs` is greater than 114. - `feerate_per_kw` is not 0. - `channel_type` does not include `zero_fee_commitments` and: - `max_accepted_htlcs` is greater than 483. - it considers `feerate_per_kw` too small for timely processing or unreasonably large. - `funding_pubkey`, `revocation_basepoint`, `htlc_basepoint`, `payment_basepoint`, or `delayed_payment_basepoint` are not valid secp256k1 pubkeys in compressed format. - `dust_limit_satoshis` is greater than `channel_reserve_satoshis`. - `dust_limit_satoshis` is smaller than 354 satoshis (see [BOLT 3](#)). - the funder's amount for the initial commitment transaction is not sufficient for full [fee payment](#). - both `to_local` and `to_remote` amounts for the initial commitment transaction are less than or equal to `channel_reserve_satoshis` (see [BOLT 3](#)). - `funding_satoshis` is greater than or equal to 2^{24} and the receiver does not support `option_support_large_channel`. - the `channel_type` is not suitable. - the `channel_type` includes `option_zeroconf` and it does not trust the sender to open an unconfirmed channel.

The receiving node MUST NOT: - consider funds received, using `push_msat`, to be received until the funding transaction has reached sufficient depth.

Rationale

The requirement for `funding_satoshis` to be less than 2^{24} satoshi was a temporary self-imposed limit while implementations were not yet considered stable, it can be lifted by advertising `option_support_large_channel`.

The *channel reserve* is specified by the peer's `channel_reserve_satoshis`: 1% of the channel total is suggested. Each side of a channel maintains this reserve so it always has something to lose if it were to try to broadcast an old, revoked commitment transaction. Initially, this reserve may not be met, as only one side has funds; but the protocol ensures that there is always progress toward meeting this reserve, and once met, it is maintained.

The sender can unconditionally give initial funds to the receiver using a non-zero `push_msat`, but even in this case we ensure that the funder has sufficient remaining funds to pay fees and that one side has some amount it can spend (which also implies there is at least one non-dust output). Note that, like any other on-chain transaction, this payment is not certain until the funding transaction has been confirmed sufficiently (with a danger of double-spend until this occurs) and may require a separate method to prove payment via on-chain confirmation.

The `feerate_per_kw` is generally only of concern to the sender (who pays the fees), but there is also the fee rate paid by HTLC transactions; thus, unreasonably large fee rates can also penalize the recipient.

Separating the `htlc_basepoint` from the `payment_basepoint` improves security: a node needs the secret associated with the `htlc_basepoint` to produce HTLC signatures for the protocol, but the secret for the `payment_basepoint` can be in cold storage.

The requirement that `channel_reserve_satoshis` is not considered dust according to `dust_limit_satoshis` eliminates cases where all outputs would be eliminated as dust. The similar requirements in `accept_channel` ensure that both sides' `channel_reserve_satoshis` are above both `dust_limit_satoshis`.

The receiver should not accept large `dust_limit_satoshis`, as this could be used in griefing attacks, where the peer publishes its commitment with a lot of dust htlcs, which effectively become miner fees. But it must allow

values higher than the standard Bitcoin Core dust limits, since HTLC outputs need to be spent by a second-stage transaction at a feerate matching the current on-chain feerate.

Details for how to handle a channel failure can be found in [BOLT 5:Failing a Channel](#).

The `accept_channel` Message

This message contains information about a node and indicates its acceptance of the new channel. This is the second step toward creating the funding transaction and both versions of the commitment transaction.

1. type: 33 (`accept_channel`)
2. data:
 - [32*byte:temporary_channel_id]
 - [u64:dust_limit_satoshis]
 - [u64:max_htlc_value_in_flight_msat]
 - [u64:channel_reserve_satoshis]
 - [u64:htlc_minimum_msat]
 - [u32:minimum_depth]
 - [u16:to_self_delay]
 - [u16:max_accepted_htlcs]
 - [point:funding_pubkey]
 - [point:revocation_basepoint]
 - [point:payment_basepoint]
 - [point:delayed_payment_basepoint]
 - [point:htlc_basepoint]
 - [point:first_per_commitment_point]
 - [accept_channel_tlvs:tlvs]
3. tlv_stream: `accept_channel_tlvs`
4. types:
 1. type: 0 (`upfront_shutdown_script`)
 2. data:
 - [...*byte:shutdown_scriptpubkey]
 3. type: 1 (`channel_type`)
 4. data:
 - [...*byte:type]

Requirements

The `temporary_channel_id` MUST be the same as the `temporary_channel_id` in the `open_channel` message.

The sender: - if `channel_type` includes `option_zeroconf`: - MUST set `minimum_depth` to zero. - otherwise: - SHOULD set `minimum_depth` to a number of blocks it considers reasonable to avoid double-spending of the funding transaction. - MUST set `channel_reserve_satoshis` greater than or equal to `dust_limit_satoshis` from the `open_channel` message. - MUST set `dust_limit_satoshis` less than or equal to `channel_reserve_satoshis` from the `open_channel` message. - MUST set `channel_type` to the `channel_type` from `open_channel`.

The receiver: - if `minimum_depth` is unreasonably large: - MAY fail the channel. - if `channel_reserve_satoshis` is less than `dust_limit_satoshis` within the `open_channel` message: - MUST fail the channel. - if `channel_reserve_satoshis` from the `open_channel` message is less than `dust_limit_satoshis`: - MUST fail the channel. - if the message doesn't include a `channel_type`: - MUST fail the channel. - if `channel_type` does not match the `channel_type` from `open_channel`: - MUST fail the channel.

Other fields have the same requirements as their counterparts in `open_channel`.

The `funding_created` Message

This message describes the outpoint which the funder has created for the initial commitment transactions. After receiving the peer's signature, via `funding_signed`, it will broadcast the funding transaction.

1. type: 34 (`funding_created`)
2. data:
 - [32*byte:temporary_channel_id]
 - [sha256:funding_txid]
 - [u16:funding_output_index]
 - [signature:signature]

Requirements

The sender MUST set: - `temporary_channel_id` the same as the `temporary_channel_id` in the `open_channel` message. - `funding_txid` to the transaction ID of a non-malleable transaction, - and MUST NOT broadcast this transaction. - `funding_output_index` to the output number of that transaction that corresponds the funding transaction output, as defined in [BOLT #3](#). - `signature` to the valid signature using its `funding_pubkey` for the initial commitment transaction, as defined in [BOLT #3](#).

The sender: - when creating the funding transaction: - SHOULD use only BIP141 (Segregated Witness) inputs. - SHOULD ensure the funding transaction confirms in the next 2016 blocks.

The recipient: - if `signature` is incorrect OR non-compliant with LOW-S-standard rule [LWS](#): - MUST send a warning and close the connection, or send an error and fail the channel.

Rationale

The `funding_output_index` can only be 2 bytes, since that's how it's packed into the `channel_id` and used throughout the gossip protocol. The limit of 65535 outputs should not be overly burdensome.

A transaction with all Segregated Witness inputs is not malleable, hence the funding transaction recommendation.

The funder may use CPFP on a change output to ensure that the funding transaction confirms before 2016 blocks, otherwise the fundee may forget that channel.

The `funding_signed` Message

This message gives the funder the signature it needs for the first commitment transaction, so it can broadcast the transaction knowing that funds can be redeemed, if need be.

This message introduces the `channel_id` to identify the channel. It's derived from the funding transaction by combining the `funding_txid` and the `funding_output_index`, using big-endian exclusive-OR (i.e. `funding_output_index` alters the last 2 bytes).

1. type: 35 (`funding_signed`)
2. data:
 - [`channel_id:channel_id`]
 - [`signature:signature`]

Requirements

Both peers: - MUST use the negotiated `channel_type` for all commitment transactions.

The sender MUST set: - `channel_id` by exclusive-OR of the `funding_txid` and the `funding_output_index` from the `funding_created` message. - `signature` to the valid signature, using its `funding_pubkey` for the initial commitment transaction, as defined in [BOLT #3](#).

The recipient: - if `signature` is incorrect OR non-compliant with LOW-S-standard rule [LWS](#): - MUST send a warning and close the connection, or send an error and fail the channel. - MUST NOT broadcast the funding transaction before receipt of a valid `funding_signed`. - on receipt of a valid `funding_signed`: - SHOULD broadcast the funding transaction.

Rationale

We generate the commitment transaction at this point, using the `channel_type` that was communicated in the `open_channel` and `accept_channel` messages. This `channel_type` determines the channel commitment format for the total lifetime of the channel.

The `channel_ready` Message

This message (which used to be called `funding_locked`) indicates that the funding transaction has sufficient confirms for channel use. Once both nodes have sent this, the channel enters normal operating mode.

Note that the opener is free to send this message at any time (since it presumably trusts itself), but the acceptor would usually wait until the funding has reached the `minimum_depth` asked for in `accept_channel`.

1. type: 36 (`channel_ready`)
2. data:
 - [`channel_id:channel_id`]
 - [`point:second_per_commitment_point`]
 - [`channel_ready_tlvs:tlvs`]
3. `tlv_stream`: `channel_ready_tlvs`
4. types:
 1. type: 1 (`short_channel_id`)
 2. data:
 - [`short_channel_id:alias`]

Requirements

The sender: - MUST NOT send `channel_ready` unless outpoint of given by `funding_txid` and `funding_output_index` in the `funding_created` message pays exactly `funding_satoshis` to the scriptpubkey specified in [BOLT #3](#). - if it is not the node opening the channel: - SHOULD wait until the funding transaction has reached `minimum_depth` before sending this message. - MUST wait for at least 100 blocks if the funding transaction is the coinbase transaction. - MUST set `second_per_commitment_point` to the per-commitment point to be used for commitment transaction #1, derived as specified in [BOLT #3](#). - if `option_scid_alias` was negotiated: - MUST set `short_channel_id alias`. - otherwise: - MAY set `short_channel_id alias`. - if it sets `alias`: - if the `announce_channel` bit was set in `open_channel`: - SHOULD initially set `alias` to value not related to the real `short_channel_id`. - otherwise: - MUST set `alias` to a value not related to the real `short_channel_id`. - MUST NOT send the same `alias` for multiple peers or use an `alias` which collides with a `short_channel_id` of a channel on the same node. - MUST always recognize the `alias` as a `short_channel_id` for incoming HTLCs to this channel. - if `channel_type` has `option_scid_alias` set: - MUST NOT allow incoming HTLCs to this channel using the real `short_channel_id` - MAY send multiple `channel_ready` messages to the same peer with different `alias` values. - otherwise: - MUST wait until the funding transaction has reached `minimum_depth` before sending this message.

The sender:

A non-funding node (fundee): - SHOULD forget the channel if it does not see the correct funding transaction after a timeout of 2016 blocks.

The receiver: - MAY use any of the `alias` it received, in BOLT 11 `r` fields. - if `channel_type` has `option_scid_alias` set: - MUST NOT use the real `short_channel_id` in BOLT 11 `r` fields.

From the point of waiting for `channel_ready` onward, either node MAY send an error and fail the channel if it does not receive a required response from the other node after a reasonable timeout.

Rationale

The non-funder can simply forget the channel ever existed, since no funds are at risk. If the fundee were to remember the channel forever, this would create a Denial of Service risk; therefore, forgetting it is recommended (even if the promise of `push_msat` is significant).

If the fundee forgets the channel before it was confirmed, the funder will need to broadcast the commitment transaction to get his funds back and open a new channel. To avoid this, the funder should ensure the funding transaction confirms in the next 2016 blocks.


The `alias` here is required for two distinct use cases. The first one is for routing payments through channels that are not confirmed yet (since the real `short_channel_id` is unknown until confirmation). The second one is to provide one or more aliases to use for private channels (even once a real `short_channel_id` is available).

While a node can send multiple `alias`, it must remember all of the ones it has sent so it can use them should they be requested by incoming HTLCs. The recipient only need remember one, for use in `r` route hints in BOLT 11 invoices.

If an RBF negotiation is in progress when a `channel_ready` message is exchanged, the negotiation must be abandoned.

Channel Establishment v2

This is a revision of the channel establishment protocol. It changes the previous protocol to allow the `accept_channel2` peer (the *accepter* / *non-initiator*) to contribute inputs to the funding transaction, via the interactive transaction construction protocol.



- where node A is **opener*/***initiator** and node B is **accepter*/***non-initiator**

The `open_channel2` Message

This message initiates the v2 channel establishment workflow.

1. type: 64 (`open_channel2`)
2. data:
 - `[chain_hash:chain_hash]`
 - `[channel_id:temporary_channel_id]`
 - `[u32:funding_feerate_perkw]`
 - `[u32:commitment_feerate_perkw]`
 - `[u64:funding_satoshis]`
 - `[u64:dust_limit_satoshis]`
 - `[u64:max_htlc_value_in_flight_msat]`
 - `[u64:htlc_minimum_msat]`
 - `[u16:to_self_delay]`
 - `[u16:max_accepted_htlcs]`

- [u32:locktime]
 - [point:funding_pubkey]
 - [point:revocation_basepoint]
 - [point:payment_basepoint]
 - [point:delayed_payment_basepoint]
 - [point:htlc_basepoint]
 - [point:first_per_commitment_point]
 - [point:second_per_commitment_point]
 - [byte:channel_flags]
 - [opening_tlvs:tlvs]
3. tlv_stream: opening_tlvs
4. types:
1. type: 0 (upfront_shutdown_script)
 2. data:
 - [...*byte:shutdown_scriptpubkey]
 3. type: 1 (channel_type)
 4. data:
 - [...*byte:type]
 5. type: 2 (require_confirmed_inputs)

Rationale and Requirements are the same as for [open_channel](#), with the following additions.

Requirements

If nodes have negotiated option_dual_fund: - the opening node: - MUST NOT send open_channel

The sending node: - MUST set channel_type - MUST set funding_feerate_perkw to the feerate for this transaction - if channel_type includes zero_fee_commitments: - MUST set commitment_feerate_perkw to 0. - If it requires the receiving node to only use confirmed inputs: - MUST set require_confirmed_inputs

The receiving node: - MAY fail the negotiation if: - the locktime is unacceptable - the funding_feerate_perkw is unacceptable - MUST fail the negotiation if: - channel_type includes zero_fee_commitments and commitment_feerate_perkw is not 0. - require_confirmed_inputs is set but it cannot provide confirmed inputs - channel_type is not set

Rationale

temporary_channel_id MUST be derived using a zeroed out basepoint for the peer's revocation basepoint. This allows the peer to return channel-assignable errors before the *accepter's* revocation basepoint is known.

funding_feerate_perkw indicates the fee rate that the opening node will pay for the funding transaction in satoshi per 1000-weight, as described in [BOLT-3, Appendix F](#).

locktime is the locktime for the funding transaction.

The receiving node, if the locktime or funding_feerate_perkw is considered out of an acceptable range, may fail the negotiation. However, it is recommended that the *accepter* permits the channel open to proceed without their participation in the channel's funding.

Note that `open_channel`'s `channel_reserve_satoshi` has been omitted. Instead, the channel reserve is fixed at 1% of the total channel balance (`open_channel2.funding_satoshis + accept_channel2.funding_satoshis`) rounded down to the nearest whole satoshi or the `dust_limit_satoshis`, whichever is greater.

Note that `push_msat` has been omitted.

`second_per_commitment_point` is now sent here (as well as in `channel_ready`) as a convenience for implementations.

The sending node may require the other participant to only use confirmed inputs. This ensures that the sending node doesn't end up paying the fees of a low feerate unconfirmed ancestor of one of the other participant's inputs.

The `accept_channel2` Message

This message contains information about a node and indicates its acceptance of the new channel.

1. type: 65 (`accept_channel2`)
2. data:
 - `[channel_id:temporary_channel_id]`
 - `[u64:funding_satoshis]`
 - `[u64:dust_limit_satoshis]`
 - `[u64:max_htlc_value_in_flight_msat]`
 - `[u64:htlc_minimum_msat]`
 - `[u32:minimum_depth]`
 - `[u16:to_self_delay]`
 - `[u16:max_accepted_htlcs]`
 - `[point:funding_pubkey]`
 - `[point:revocation_basepoint]`
 - `[point:payment_basepoint]`
 - `[point:delayed_payment_basepoint]`
 - `[point:htlc_basepoint]`
 - `[point:first_per_commitment_point]`
 - `[point:second_per_commitment_point]`
 - `[accept_tlvs:tlvs]`
3. `tlv_stream: accept_tlvs`
4. types:
 1. type: 0 (`upfront_shutdown_script`)
 2. data:
 - `[...*byte:shutdown_scriptpubkey]`
 3. type: 1 (`channel_type`)
 4. data:
 - `[...*byte:type]`
 5. type: 2 (`require_confirmed_inputs`)

Rationale and Requirements are the same as listed above, for [accept_channel](#) with the following additions.

Requirements

The accepting node: - MUST use the `temporary_channel_id` of the `open_channel2` message. - MUST set `channel_type` to the `channel_type` from `open_channel2`. - MAY respond with a `funding_satoshis` value of zero. - If it requires the opening node to only use confirmed inputs: - MUST set `require_confirmed_inputs`.

The receiving node: - MUST fail the negotiation if: - `require_confirmed_inputs` is set but it cannot provide confirmed inputs. - `channel_type` is not set.

Rationale

The `funding_satoshis` is the amount of bitcoin in satoshis the *accepter* will be contributing to the channel's funding transaction.

Note that `accept_channel's channel_reserve_satoshi` has been omitted. Instead, the channel reserve is fixed at 1% of the total channel balance (`open_channel2.funding_satoshis + accept_channel2.funding_satoshis`) rounded down to the nearest whole satoshi or the `dust_limit_satoshis`, whichever is greater.

Funding Composition

Funding composition for channel establishment v2 makes use of the [Interactive Transaction Construction](#) protocol, with the following additional caveats.

The `tx_add_input` Message

Requirements

The sending node: - if the receiver set `require_confirmed_inputs` in `open_channel2`, `accept_channel2`, `tx_init_rbf` or `tx_ack_rbf`: - MUST NOT send a `tx_add_input` that contains an unconfirmed input

The `tx_add_output` Message

Requirements

The sending node: - if is the *opener*: - MUST send at least one `tx_add_output`, which contains the channel's funding output

Rationale

The channel funding output must be added by the *opener*, who pays its fees.

The `tx_complete` Message

Upon receipt of consecutive `tx_completes`, the receiving node: - if is the *accepter*: - MUST fail the negotiation if: - no funding output was received - the value of the funding output is not equal to the sum of `open_channel2.funding_satoshis` and `accept_channel2.funding_satoshis` - the value of the funding output is less than the `dust_limit` - if this is an RBF attempt: - MUST fail the negotiation if: - the transaction's total fees is less than the last successfully negotiated transaction's fees - the transaction does not share at least one input with each previous funding transaction - if it has sent `require_confirmed_inputs` in `open_channel2`,

accept_channel2, tx_init_rbf or tx_ack_rbf: - MUST fail the negotiation if: - one of the inputs added by the other peer is unconfirmed

The commitment_signed Message

This message is exchanged by both peers. It contains the signatures for the first commitment transaction, which uses a format determined by the channel_type sent in open_channel2 and accept_channel2.

Rationale and Requirements are the same as listed below, for [commitment_signed](#) with the following additions.

Requirements

The sending node: - MUST send zero HTLCs. - MUST remember the details of this funding transaction.

The receiving node: - if the message has one or more HTLCs: - MUST fail the negotiation - if it has not already transmitted its commitment_signed: - MUST send commitment_signed - Otherwise: - MUST send tx_signatures if it should sign first, as specified in the [tx_signatures_requirements](#)

Rationale

The first commitment transaction has no HTLCs.

Once peers are ready to exchange commitment signatures, they must remember the details of the funding transaction to allow resuming the signatures exchange if a disconnection happens.

Sharing funding signatures: tx_signatures

After a valid commitment_signed has been received from the peer and a commitment_signed has been sent, a peer: - MUST transmit tx_signatures with their signatures for the funding transaction, following the order specified in the [tx_signatures_requirements](#)

Requirements

The sending node: - MUST verify it has received a valid commitment signature from its peer - MUST remember the details of this funding transaction - if it has NOT received a valid commitment_signed message: - MUST NOT send a tx_signatures message

The receiving node: - if has already sent or received a channel_ready message for this channel: - MUST ignore this message - if the witness weight lowers the effective feerate below the *opener's* feerate for the funding transaction and the effective feerate is determined by the receiving node to be insufficient for getting the transaction confirmed in a timely manner: - SHOULD broadcast their commitment transaction, closing the channel - SHOULD double-spend their channel inputs when there is a productive opportunity to do so; effectively canceling this channel open - SHOULD apply witnesses to the funding transaction and broadcast it

Rationale

A peer sends their tx_signatures after receiving a valid commitment_signed message, following the order specified in the [tx_signatures_section](#).

In the case where a peer provides valid witness data that causes their paid feerate to fall beneath the `open_channel2.funding_feerate_perkw`, the channel should be considered failed and the channel should be double-spent when there is a productive opportunity to do so. This should disincentivize peers from underpaying fees.

Fee bumping: `tx_init_rbf` and `tx_ack_rbf`

After the funding transaction has been broadcast, it can be replaced by a transaction paying more fees to make the channel confirm faster.

Requirements

The sender of `tx_init_rbf`: - MAY be either the *initiator* or the *accepter* - If the sender is the accepter, it becomes the initiator of the interactive-tx session and thus: - MUST send `tx_add_output` for the channel output - MUST pay the fees for the shared transaction fields - MUST NOT have sent or received a `channel_ready` message.

The recipient: - MUST fail the negotiation if they have already sent or received `channel_ready` - MAY fail the negotiation for any reason

Rationale

If a valid `channel_ready` message is received in the middle of an RBF attempt, the attempt MUST be abandoned.

Peers can use different values in `tx_init_rbf.funding_output_contribution` and `tx_ack_rbf.funding_output_contribution` from the amounts transmitted in `open_channel2` and `accept_channel2`: they are allowed to change how much they wish to commit to the funding output.

It's recommended that a peer, rather than fail the RBF negotiation due to a large feerate change, instead sets their sats to zero, and decline to participate further in the channel funding: by not contributing, they may obtain incoming liquidity at no cost.

We allow both nodes to initiate RBF, because any one of them may want to take this opportunity to contribute additional funds to the channel without waiting for the initial funding transaction to confirm.

Channel Quiescence

Various fundamental changes, in particular protocol upgrades, are easiest on channels where both commitment transactions match, and no pending updates are in flight. We define a protocol to quiesce the channel by indicating that "Something Fundamental is Underway".

stfu

1. type: 2 (stfu)
2. data:
 - [channel_id:channel_id]
 - [u8:initiator]

Requirements

The sender of stfu: - MUST NOT send stfu unless option_quiesce is negotiated. - MUST NOT send stfu if any of the sender's htlc additions, htlc removals or fee updates are pending for either peer. - MUST NOT send stfu twice. - if it is replying to an stfu: - MUST set initiator to 0 - otherwise: - MUST set initiator to 1 - MUST set channel_id to the id of the channel to quiesce. - MUST now consider the channel to be quiescing. - MUST NOT send an update message after stfu.

The receiver of stfu: - if it has sent stfu then: - MUST now consider the channel to be quiescent - otherwise: - SHOULD NOT send any more update messages. - MUST reply with stfu once it can do so.

Both nodes: - MUST disconnect after 60 seconds of quiescence if the HTLCs are pending.

Upon disconnection: - the channel is no longer considered quiescent.

Dependent Protocols: - MUST specify all states that terminate quiescence. - NOTE: this prevents batching executions of protocols that depend on quiescence.

Rationale

The normal use would be to cease sending updates, then wait for all the current updates to be acknowledged by both peers, then start quiescence. For some protocols, choosing the initiator matters, so this flag is sent.

If both sides send stfu simultaneously, they will both set initiator to 1, in which case the "initiator" is arbitrarily considered to be the channel funder (the sender of open_channel). The quiescence effect is exactly the same as if one had replied to the other.

Dependent protocols have to specify termination conditions to prevent the need for disconnection to resume channel traffic. An explicit resume message was [considered but rejected](#) since it introduces a number of edge cases that make bilateral consensus of channel state significantly more complex to maintain. This introduces the derivative property that it is impossible to batch multiple downstream protocols in the same quiescence session.

Channel Splicing

Splicing is the term given for replacing the funding transaction with a new one. For simplicity, splicing takes place once a channel is [quiescent](#).

Operation returns to normal once the splice transaction has been signed (while waiting for one of the splice transactions to confirm), at which point the channel isn't quiescent anymore.


The splice is finally terminated when both sides send splice_locked to indicate that one of the splice transactions reached acceptable depth.

```
+-----+
|      |--- splice_init ----->|      |
|  A    |<----- splice_ack ---|  B    |
|      |                         |      |
|      |--- tx_add_input ----->|      |
|      |<----- tx_add_input ---|      |
|      |--- tx_add_input ----->|      |
```

```

|<----- tx_add_output ---|
|--- tx_add_output ----->|
|<----- tx_complete ---|
|--- tx_add_output ----->|
|<----- tx_complete ---|
|--- tx_complete ----->|
|
|--- commit_sig ----->|
|<----- commit_sig ---|
|--- tx_signatures ----->|
|<----- tx_signatures ---|
|
|<RESUME CHANNEL>
|
|--- update_add_htlc ----->|
|--- commit_sig ----->|
|--- commit_sig ----->|
|<----- revoke_and_ack ---|
|<----- commit_sig ---|
|<----- commit_sig ---|
|--- revoke_and_ack ----->|
|
|<RBF>
|
|<----- tx_init_rbf ---|
|--- tx_ack_rbf ----->|
|<----- tx_add_input ---|
|--- tx_add_input ----->|
|<----- tx_add_input ---|
|--- tx_add_output ----->|
|<----- tx_add_output ---|
|--- tx_complete ----->|
|<----- tx_add_output ---|
|--- tx_complete ----->|
|<----- tx_complete ---|
|
|<----- commit_sig ---|
|--- commit_sig ----->|
|--- tx_signatures ----->|
|<----- tx_signatures ---|
|
|<RESUME CHANNEL>
|
|--- update_add_htlc ----->|
|--- commit_sig ----->|
|--- commit_sig ----->|
|--- commit_sig ----->|
|<----- revoke_and_ack ---|
|<----- commit_sig ---|
|<----- commit_sig ---|
|<----- commit_sig ---|
|--- revoke_and_ack ----->|
|
|<SPICE COMPLETION>

```



The splice_init Message

1. type: 80 (splice_init)
2. data:
 - [channel_id:channel_id]
 - [s64:funding_contribution_satoshis]
 - [u32:funding_feerate_perkw]
 - [u32:locktime]
 - [point:funding_pubkey]
 - [splice_init_tlvs:tlvs]
3. tlv_stream: splice_init_tlvs
4. types:
 1. type: 2 (require_confirmed_inputs)

funding_contribution_satoshis is the amount the sender is adding to their channel balance (splice-in) or removing from their channel balance (splice-out).

Requirements

The sending node: - MUST NOT send splice_init if the channel is not quiescent. - MUST NOT send splice_init if it is not the quiescence initiator. - MUST NOT send splice_init before sending and receiving channel_ready. - MUST NOT send splice_init while another splice is being negotiated. - MUST NOT send splice_init if another splice has been negotiated but splice_locked has not been sent and received. - MUST NOT send splice_init if it has previously sent shutdown. - MUST set funding_feerate_perkw to the feerate for the splice transaction. - If it is splicing funds out of the channel: - MUST set funding_contribution_satoshis to a negative value matching the amount that will be subtracted from its current channel balance. - If it is splicing funds into the channel: - MUST set funding_contribution_satoshis to a positive value matching the amount that will be added to its current channel balance. - If it requires the receiving node to only use confirmed inputs: - MUST set require_confirmed_inputs. - SHOULD use a different funding_pubkey than the one used for the previous funding transaction.

The receiving node: - If the channel is not quiescent: - MUST send a warning and close the connection or send an error and fail the channel. - If the sending node is not the quiescence initiator: - MUST send a warning and close the connection or send an error and fail the channel. - If another splice is already being negotiated: - MUST send a warning and close the connection or send an error and fail the channel. - If another splice has been negotiated but isn't locked yet: - MUST send a warning and close the connection or send an error and fail the channel. - If it has received shutdown: - MUST send a warning and close the connection or send an error

and fail the channel. - If the `funding_feerate_perkw` is unacceptable: - MUST respond with `tx_abort`. - If `funding_contribution_satoshis` is negative and its absolute value is greater than the sending node's current channel balance: - MUST send a warning and close the connection or send an error and fail the channel. - If it accepts the splice attempt: - MUST respond with `splice_ack`. - Otherwise (it rejects the splice): - MUST respond with `tx_abort`.

The `splice_ack` Message

1. type: 81 (`splice_ack`)
2. data:
 - `[channel_id:channel_id]`
 - `[s64:funding_contribution_satoshis]`
 - `[point:funding_pubkey]`
 - `[splice_ack_tlvs:tlvs]`
3. `tlv_stream`: `splice_ack_tlvs`
4. types:
 1. type: 2 (`require_confirmed_inputs`)

Requirements

The sending node: - SHOULD use a different `funding_pubkey` than the one used for the previous funding transaction. - MAY set `funding_contribution_satoshis` to 0 if they don't want to contribute to the splice. - If it requires the receiving node to only use confirmed inputs: - MUST set `require_confirmed_inputs`.

The receiving node: - If it has sent `splice_init`: - If `funding_contribution_satoshis` is negative and its absolute value is greater than the sending node's current channel balance: - MUST send a warning and close the connection or send an error and fail the channel. - If it accepts the splice attempt: - MUST start an interactive-tx session to create the splice transaction. - Otherwise: - MUST reject the splice attempt by sending `tx_abort`. - Otherwise (it has not sent `splice_init`): - MUST send a warning and close the connection or send an error and fail the channel.

Splice Transaction Construction

The splice transaction is created using the [Interactive Transaction Construction](#) protocol, with the following additional requirements.

The `tx_add_input` Message

Requirements

The sending node: - If it is the splice initiator: - MUST add the current channel input to the splice transaction by sending `tx_add_input` with `shared_input_txid` containing the txid of the previous funding transaction. - MUST NOT include `prevtx` for that shared input. - MUST set `prevtx_vout` to the previous funding output index. - If the receiver set `require_confirmed_inputs` in `splice_init`, `splice_ack`, `tx_init_rbf` or `tx_ack_rbf`: - MUST NOT send a `tx_add_input` that contains an unconfirmed input.

The receiving node: - If `shared_input_txid` is set: - If it doesn't match the `txid` of the previous funding transaction: - MUST fail the negotiation by sending `tx_abort`. - If `prevtx_vout` doesn't match the previous funding output index: - MUST fail the negotiation by sending `tx_abort`.

Rationale

The splice transaction must spend the current channel funding output. The splice initiator is responsible for adding that input to the transaction, and pay the fees for its weight. It would be wasteful to transmit the previous funding transaction in the `prevtx` field, and wouldn't even be possible for funding transactions that exceed 65kB, so we only transmit its `txid` using the `shared_input_txid` field.

The **tx_add_output** Message

Requirements

The sending node: - If it is the splice initiator: - MUST send at least one `tx_add_output`, which contains the new channel's funding output based on the `funding_pubkeys` from `splice_init` and `splice_ack`. - MUST set the amount of that `tx_add_output` to the previous channel capacity with the `funding_contribution_satoshis` from `splice_init` and `splice_ack` applied.

Rationale

The splice initiator is responsible for adding the new channel funding output to the transaction and paying the fees for its weight.

The **tx_complete** Message

Requirements

The receiving node: - MUST compute the channel balance for each side by adding their respective `funding_contribution_satoshis` to their previous channel balance. - MUST fail the negotiation by sending `tx_abort` if: - There is not exactly one input spending the current funding transaction. - There is not exactly one channel funding output using the funding public keys and funding contributions from `splice_init` and `splice_ack`. - This is an RBF attempt and the transaction's total fees is less than the last successfully negotiated splice transaction's fees. - Either side has added an output other than the channel funding output and the balance for that side is less than the channel reserve that matches the new channel capacity.

Rationale

If a side does not meet the reserve requirements, that's OK: but if they take funds out of the channel, they must ensure that they do meet them. If your peer adds a massive amount to the channel, then you only have to add more reserve if you want to contribute to the splice (and you can use `tx_remove_output` and/or `tx_remove_input` part-way through if this happens).

The **commitment_signed** Message

After exchanging `tx_complete`, both peers send `commitment_signed` to commit to the splice transaction by creating a commitment transaction spending the new channel funding output.

The usual [commitment_signed](#) requirements apply with the following additions.

Requirements

The sending node: - MUST create a commitment transaction that spends the splice funding output and: - Adds `funding_contribution_satoshis` from `splice_init` and `splice_ack` to the main balance of their respective sender. - Uses the same feerate as the existing commitment transaction. - Uses the same `commitment_number` as the existing commitment transaction. - MUST send signatures for pending HTLCs. - MUST remember the details of this splice transaction.

The receiving node: - MUST NOT respond with `revoke_and_ack`. - If it has not already transmitted its `commitment_signed`: - MUST send `commitment_signed`. - If it should sign first, as specified in the [tx_signatures requirements](#): - MUST send `tx_signatures`. - Note that since the initiator sends `tx_add_input` for the shared input (corresponding to the previous channel output), 100% of the previous channel capacity is attributed to the initiator when computing who must send `tx_signatures` first (instead of using each node's previous balance).

On reconnection: - If `next_funding` matches the splice transaction: - MUST retransmit `commitment_signed`.

Rationale

Once peers are ready to exchange commitment signatures, they must remember the details of the splice transaction to allow resuming the signatures exchange if a disconnection happens.

The `tx_signatures` Message

Requirements

The sending node: - MUST set `shared_input_signature` to a valid ECDSA signature for the `tx_add_input` spending the previous channel funding output using the `funding_pubkey` that matches this input.

The receiving node: - If `shared_input_signature` is not set: - MUST send an error and fail the channel. - If `shared_input_signature` is not valid or non-compliant with the LOW-S-standard rule [LWS](#): - MUST send an error and fail the channel. - MUST consider the channel no longer quiescent.

On reconnection: - If `next_funding` matches the splice transaction: - MUST retransmit `tx_signatures`.

Rationale

Spending the channel funding output requires a signature from both peers. Each peer transmits its own signature, which allows creating a valid witness for the shared input without adding an additional message.

Once `tx_signatures` have been exchanged, the splice transaction can be broadcast. The channel is no longer quiescent: normal operation can resume while waiting for the transaction to confirm and `splice_locked` messages to be exchanged.

The **tx_init_rbf** Message

Requirements

The sending node: - MUST NOT send **tx_init_rbf** if the channel is not quiescent. - MUST NOT send **tx_init_rbf** if it is not the quiescence initiator. - MAY send **tx_init_rbf** even if it is not the splice initiator. - If there are more than 10 pending RBF attempts: - MUST set a high enough feerate to ensure quick confirmation. - MUST NOT send **tx_init_rbf** if it has previously sent **splice_locked**. - MUST NOT send **tx_init_rbf** if **option_zeroconf** has been negotiated. - MAY set **funding_output_contribution** to a different value than the **funding_contribution_satoshis** used in **splice_init** or **splice_ack**, or in previous RBF attempts.

The receiving node: - If the channel is not quiescent: - MUST send a warning and close the connection or send an error and fail the channel. - If the sending node is not the quiescence initiator: - MUST send a warning and close the connection or send an error and fail the channel. - If another RBF attempt has been created recently: - SHOULD send **tx_abort** to reject this RBF attempt and wait for the previous RBF attempt to confirm. - If there are more than 10 pending RBF attempts and the feerate is not high enough to ensure quick confirmation: - SHOULD send **tx_abort** to reject the RBF attempt. - If the sender previously sent **splice_locked**: - MUST send a warning and close the connection or send an error and fail the channel. - If **option_zeroconf** has been negotiated: - MUST send a warning and close the connection or send an error and fail the channel. - If **funding_output_contribution** is negative and its absolute value is greater than the sending node's current channel balance: - MUST send a warning and close the connection or send an error and fail the channel.

Rationale

Splice transactions can be RBF-ed to react to changes in the mempool feerate. We allow both nodes to initiate RBF, because any one of them may want to take this opportunity to splice additional funds into or out of the channel without waiting for the initial splice transaction to confirm.

We limit the number of pending RBF attempts, otherwise we may reach the **batch_size** limit defined in [start_batch](#). We require using a large enough feerate when we've already created many RBF attempts, and we wait between RBF attempts to allow a previous attempt to confirm.

Since splice transactions always spend the current channel funding output, the RBF attempts automatically double-spend each other. We thus disallow RBF when **option_zeroconf** has been negotiated, because that creates a risk of losing funds.

The **tx_ack_rbf** Message

Requirements


The sending node: - MAY set **funding_output_contribution** to a different value than the **funding_contribution_satoshis** used in **splice_init** or **splice_ack**, or in previous RBF attempts.

The receiving node: - If **funding_output_contribution** is negative and its absolute value is greater than the sending node's current channel balance: - MUST send a warning and close the connection or send an error and fail the channel.

Splice Completion

Once splice transactions have been signed but haven't reached acceptable depth, channel operations go back to normal and HTLCs can be exchanged, with the caveat that payments must be valid for all splice transactions.

Nodes keep track of multiple commitment transactions (one for the current funding transaction and one for each splice transaction) and exchange signatures for each of these commitment transactions.



The splice completes by exchanging `splice_locked` messages, at which point the locked transaction replaces the previous funding transaction.

The `splice_locked` Message

1. type: 77 (`splice_locked`)
2. data:
 - `[channel_id:channel_id]`
 - `[sha256:splice_txid]`

Requirements

Each node: - If any splice transaction reaches acceptable depth: - MUST send `splice_locked` with the txid of that transaction. - If `option_zeroconf` has been negotiated: - SHOULD send `splice_locked` immediately after exchanging `tx_signatures`.

The receiving node: - If `splice_txid` doesn't match any of its pending splice transactions: - MUST send a warning and close the connection, or send an error and fail the channel.

Once a node has sent and received `splice_locked`: - If the `splice_txids` match: - MUST stop sending `commitment_signed` for RBF attempts and ancestors of this splice transaction. - MAY discard RBF attempts and ancestor transactions. - If `announce_channel` is set for this channel: - MUST send `announcement_signatures` with `short_channel_id` matching this splice transaction. - If the `splice_txids` are for different RBF candidates: - SHOULD ignore the message. - MAY send an error and fail the channel.

Rationale


If nodes are on a different fork of the blockchain, they may disagree on which RBF attempt has been confirmed: in that case nodes can either close the channel or simply ignore `splice_locked` and wait for one of the forks to


eventually replace the other, at which point both nodes should agree on which RBF attempt confirmed and exchange splice_locked for the same splice_txid to complete the splice.

Channel Close

Nodes can negotiate a mutual close of the connection, which unlike a unilateral close, allows them to access their funds immediately and can be negotiated with lower fees.

Closing happens in two stages: 1. one side indicates it wants to clear the channel (and thus will accept no new HTLCs) 2. once all HTLCs are resolved, the final channel close negotiation begins.





Closing Initiation: shutdown

Either node (or both) can send a shutdown message to initiate closing, along with the scriptpubkey it wants to be paid to.

1. type: 38 (shutdown)
2. data:
 - [channel_id:channel_id]
 - [u16:len]
 - [len*byte:scriptpubkey]

Requirements

A sending node: - if it hasn't sent a `funding_created` (if it is a funder) or a `funding_signed` (if it is a fundee): - MUST NOT send a shutdown - MAY send a shutdown before a `channel_ready`, i.e. before the funding transaction has reached `minimum_depth`. - if there are updates pending on the receiving node's commitment transaction: - MUST NOT send a shutdown. - MUST NOT send multiple shutdown messages. - MUST NOT send shutdown if there is a splice transaction that isn't locked yet. - MUST NOT send an `update_add_htlc` after a shutdown. - if no HTLCs remain in either commitment transaction (including dust HTLCs) and neither side has a pending `revoke_and_ack` to send: - MUST NOT send any update message after that point. - SHOULD fail to route any HTLC added after it has sent shutdown. - if it sent a non-zero-length `shutdown_scriptpubkey` in `open_channel` or `accept_channel`: - MUST send the same value in `scriptpubkey`. - MUST set `scriptpubkey` in one of the following forms:

1. ``OP_0` `20`` 20-bytes (version 0 pay to witness pubkey hash), OR
2. ``OP_0` `32`` 32-bytes (version 0 pay to witness script hash), OR
3. if (and only if) ``option_shutdown_anysegwit`` is negotiated:
 - * ``OP_1`` through ``OP_16`` inclusive, followed by a single push of 2 to 40 bytes (witness program versions 1 through 16)
4. if (and only if) ``option_simple_close`` is negotiated:
 - * ``OP_RETURN`` followed by one of:
 - * ``6`` to ``75`` inclusive followed by exactly that many bytes
 - * ``76`` followed by ``76`` to ``80`` followed by exactly that many bytes

A receiving node: - if it hasn't received a `funding_signed` (if it is a funder) or a `funding_created` (if it is a fundee): - SHOULD send an error and fail the channel. - if the `scriptpubkey` is not in one of the above forms: - SHOULD send a warning. - if it hasn't sent a `channel_ready` yet: - MAY reply to a shutdown message with a shutdown - once there are no outstanding updates on the peer, UNLESS it has already sent a shutdown: - MUST reply to a shutdown message with a shutdown - if both nodes advertised the `option_upfront_shutdown_script` feature, and the receiving node received a non-zero-length `shutdown_scriptpubkey` in `open_channel` or `accept_channel`, and that `shutdown_scriptpubkey` is not equal to `scriptpubkey`: - MAY send a warning. - MUST fail the connection.

Rationale

If channel state is always "clean" (no pending changes) when a shutdown starts, the question of how to behave if it wasn't is avoided: the sender always sends a `commitment_signed` first.

As shutdown implies a desire to terminate, it implies that no new HTLCs will be added or accepted. Once any HTLCs are cleared, there are no commitments for which a revocation is owed, and all updates are included on both commitment transactions, the peer may immediately begin closing negotiation, so we ban further updates to the commitment transaction (in particular, `update_fee` would be possible otherwise). However, while there are HTLCs on the commitment transaction, the initiator may find it desirable to increase the feerate as there may be pending HTLCs on the commitment which could timeout.

The `scriptpubkey` forms include only standard segwit forms accepted by the Bitcoin network, which ensures the resulting transaction will propagate to miners. However old nodes may send non-segwit scripts, which may be accepted for backwards-compatibility (with a caveat to force-close if this output doesn't meet dust relay requirements).

The `option_upfront_shutdown_script` feature means that the node wanted to pre-commit to `shutdown_scriptpubkey` in case it was compromised somehow. This is a weak commitment (a malevolent implementation tends to ignore specifications like this one!), but it provides an incremental improvement in security by requiring the cooperation of the receiving node to change the `scriptpubkey`.

The shutdown response requirement implies that the node sends `commitment_signed` to commit any outstanding changes before replying; however, it could theoretically reconnect instead, which would simply erase all outstanding uncommitted changes.

`OP_RETURN` is only standard if followed by `PUSH` opcodes, and the total script is 83 bytes or less. We are slightly stricter, to only allow a single `PUSH`, but there are two forms in script: one which pushes up to 75 bytes, and a longer one (`OP_PUSHDAT1`) which is needed for 76-80 bytes.

Closing Negotiation: `closing_complete` and `closing_sig`

Once shutdown is complete, the channel is empty of HTLCs, there are no commitments for which a revocation is owed, and all updates are included on both commitments, the final current commitment transactions will have no HTLCs.

If `option_simple_close` is not negotiated, see [Legacy Closing Negotiation](#) below.

Each peer creates their own closing transaction where they pay the fee, and sends `closing_complete` to the other peer with the transaction details. The other peer simply signs that transaction and sends back `closing_sig`. Each peer will thus independently send `closing_complete` and receive `closing_sig`, resulting in two independent (but conflicting) closing transactions being created.

The lesser-paid peer (if either is) can opt to omit their own output from the closing transaction.

This process can be repeated multiple times by sending `closing_complete` again, which allows increasing the fees and changing the output script.

1. type: 40 (`closing_complete`)
2. data:
 - `[channel_id:channel_id]`
 - `[u16:closer_scriptpubkey_len]`
 - `[closer_scriptpubkey_len*byte:closer_scriptpubkey]`
 - `[u16:closee_scriptpubkey_len]`
 - `[closee_scriptpubkey_len*byte:closee_scriptpubkey]`
 - `[u64:fee_satoshis]`
 - `[u32:locktime]`
 - `[closing_tlvs:tlvs]`
3. type: 41 (`closing_sig`)
4. data:
 - `[channel_id:channel_id]`
 - `[u16:closer_scriptpubkey_len]`
 - `[closer_scriptpubkey_len*byte:closer_scriptpubkey]`
 - `[u16:closee_scriptpubkey_len]`
 - `[closee_scriptpubkey_len*byte:closee_scriptpubkey]`
 - `[u64:fee_satoshis]`

- [u32:locktime]
- [closing_tlvs:tlvs]
- 5. tlv_stream: closing_tlvs
- 6. types:
 1. type: 1 (closer_output_only)
 2. data:
 - [signature:sig]
 3. type: 2 (closee_output_only)
 4. data:
 - [signature:sig]
 5. type: 3 (closer_and_closee_outputs)
 6. data:
 - [signature:sig]

Requirements

Note: the details and requirements for the transaction being signed are in [BOLT 3](#).

An output is *dust* if the amount is less than the [Bitcoin Core Dust Thresholds](#).

Note: These requirements only apply if option_simple_close is negotiated, otherwise the requirements are in [Legacy Closing Negotiation](#).

Both nodes: - After a shutdown has been sent and received, AND no HTLCs remain in either commitment transaction: - SHOULD send a closing_complete message.

The sender of closing_complete (aka. “the closer”): - MUST set fee_satoshis to a fee less than or equal to its outstanding balance, rounded down to whole satoshis. - MUST set fee_satoshis so that at least one output is not dust. - MUST set closer_scriptpubkey to its desired output script. - MUST set closee_scriptpubkey to the last script it received from its peer (from closing_complete or from the initial shutdown). - MUST set locktime to the desired nLockTime of the closing transaction. - If the local outstanding balance (in millisatoshi) is less than the remote outstanding balance: - MUST NOT set closer_output_only. - MUST set closee_output_only if the local output amount is dust. - MAY set closee_output_only if it considers the local output amount uneconomical AND its closer_scriptpubkey is not OP_RETURN. - Otherwise (not lesser amount, cannot remove its own output): - MUST NOT set closee_output_only. - If it considers the local output amount uneconomical: - MAY send a closer_scriptpubkey that is a valid OP_RETURN script. - If it does, the output value MUST be set to zero so that all funds go to fees, as specified in [BOLT #3](#). - If the closee’s output amount is dust: - MUST set closer_output_only. - MUST NOT set closer_and_closee_outputs. - Otherwise: - MUST set both closer_output_only and closer_and_closee_outputs. - MUST generate its closing transaction as specified in [BOLT #3](#). - MUST set signature fields as valid signature using its funding_pubkey of: - closer_output_only: closing transaction with only the local (“closer”) output. - closee_output_only: closing transaction with only the remote (“closee”) output. - closer_and_closee_outputs: closing transaction with both the closer and closee outputs. - If it wants to send another closing_complete (e.g. with a different fee_satoshis or closer_scriptpubkey): - MUST wait until it has received closing_sig first. - SHOULD close the connection if it doesn’t receive closing_sig.

The receiver of closing_complete (aka. “the closee”): - If fee_satoshis is greater than the closer’s outstanding balance: - MUST either send a warning and close the connection, or send an error and fail the channel. - If

closee_scriptpubkey does not match the last script it sent (from closing_complete or from the initial shutdown): - SHOULD ignore closing_complete. - SHOULD send a warning. - SHOULD close the connection. - If closer_scriptpubkey is invalid (as detailed in the [shutdown requirements](#)): - SHOULD ignore closing_complete. - SHOULD send a warning. - SHOULD close the connection. - If closer_scriptpubkey is a valid OP_RETURN script: - MUST set the closer's output amount to zero so that all funds go to fees, as specified in [BOLT #3](#). - MUST generate the remote closing transaction as specified in [BOLT #3](#). - Select a signature for validation: - If the local output amount is dust: - MUST use closer_output_only. - Otherwise, if it considers the local output amount uneconomical AND its closee_scriptpubkey is not OP_RETURN: - MUST use closer_output_only. - Otherwise, if closer_and_closee_outputs is present: - MUST use closer_and_closee_outputs. - Otherwise: - MUST use closee_output_only. - If the selected signature field does not exist: - MUST either send a warning and close the connection, or send an error and fail the channel. - If the signature field is not valid for the corresponding closing transaction specified in [BOLT #3](#): - MUST either send a warning and close the connection, or send an error and fail the channel. - If the signature field is non-compliant with LOW-S-standard rule [LWS](#): - MUST either send a warning and close the connection, or send an error and fail the channel. - MUST sign and broadcast the corresponding closing transaction. - MUST send closing_sig with a single valid signature in the same TLV field as the closing_complete. - MUST use closer_scriptpubkey for its own future closing_complete messages.

The receiver of closing_sig: - If closer_scriptpubkey, closee_scriptpubkey, fee_satoshis or locktime don't match what was sent in closing_complete: - MUST either send a warning and close the connection, or send an error and fail the channel. - If tlvs does not contain exactly one signature: - MUST either send a warning and close the connection, or send an error and fail the channel. - If tlvs does not contain one of the TLV fields sent in closing_complete: - MUST either send a warning and close the connection, or send an error and fail the channel. - If the signature field is not valid for the corresponding closing transaction specified in [BOLT #3](#): - MUST either send a warning and close the connection, or send an error and fail the channel. - If the signature field is non-compliant with LOW-S-standard rule [LWS](#): - MUST either send a warning and close the connection, or send an error and fail the channel. - otherwise: - MUST broadcast the corresponding closing transaction. - MAY send another closing_complete (e.g. with a different fee_satoshis or closer_scriptpubkey).

Rationale

The close protocol is designed to avoid any failure scenarios caused by fee disagreement, since each side offers to pay its own desired fee.

If one side has less funds than the other, it may choose to omit its own output, and in this case dust MUST be omitted, to ensure that the resulting transaction can be broadcast.

The corner case where fees are so high that both outputs are dust is addressed in two ways: paying a low fee to avoid the problem, or using an OP_RETURN (which is never "dust"). If one side chooses to use an OP_RETURN output, its amount must be 0 to ensure that the resulting transaction can be broadcast.

Note that there is usually no reason to pay a high fee for rapid processing, since an urgent child could pay the fee on the closing transactions' behalf. If rapid processing is desired and CPFP is not an option, the closer can RBF its previous closing transactions by sending closing_complete again.

Sending a new closing_complete message overrides previous ones, so you can negotiate again (even changing the output address if upfront_shutdown_script was not negotiated). This creates a rare race condition if both

nodes send `closing_complete` to change their `closer_scriptpubkey` at the same time: when that happens, the `closing_complete` they receive will be using their previous output script, so they shouldn't sign the corresponding transaction. When that happens, we simply reconnect, which provides the opportunity for both nodes to send their latest output script in shutdown and restart the signing flow. We include the closer and closer scripts in `closing_sig` to allow our peer to detect a script mismatch and correctly ignore the signatures, which also helps debugging race conditions.

If the closer proposes a transaction which will not relay (an output is dust, or the fee rate it proposes is too low), it doesn't harm the closer to sign the transaction.

Similarly, if the closer proposes a high fee, it doesn't harm the closer to sign the transaction, as the closer is paying.

There's a slight game where each side would prefer the other side pay the fee, and proposes a minimal fee. If neither side proposes a fee which will relay, the negotiation can occur again, or the final commitment transaction can be spent. In practice, the opener has an incentive to offer a reasonable closing fee, as they would pay the fee for the commitment transaction, which also costs more to spend.

Legacy Closing Negotiation: `closing_signed`

Once shutdown is complete, the channel is empty of HTLCs, there are no commitments for which a revocation is owed, and all updates are included on both commitments, the final current commitment transactions will have no HTLCs, and closing fee negotiation begins: if `option_simple_close` is negotiated, the section above applies, otherwise this legacy section applies.

The funder chooses a fee it thinks is fair, and signs the closing transaction with the `scriptpubkey` fields from the shutdown messages (along with its chosen fee) and sends the signature; the other node then replies similarly, using a fee it thinks is fair. This exchange continues until both agree on the same fee or when one side fails the channel.

In the modern method, the funder sends its permissible fee range, and the non-funder has to pick a fee in this range. If the non-funder chooses the same value, negotiation is complete after two messages, otherwise the funder will reply with the same value (completing after three messages).

1. type: 39 (`closing_signed`)
2. data:
 - `[channel_id:channel_id]`
 - `[u64:fee_satoshis]`
 - `[signature:signature]`
 - `[closing_signed_tlvs:tlvs]`
3. `tlv_stream: closing_signed_tlvs`
4. types:
 1. type: 1 (`fee_range`)
 2. data:
 - `[u64:min_fee_satoshis]`
 - `[u64:max_fee_satoshis]`

Requirements

Note: These requirements only apply if `option_simple_close` is NOT negotiated, otherwise the requirements [Closing Negotiation: closing_complete and closing_sig](#) apply.

The funding node: - after shutdown has been received, AND no HTLCs remain in either commitment transaction: - SHOULD send a `closing_signed` message.

The sending node: - SHOULD set the initial `fee_satoshis` according to its estimate of cost of inclusion in a block. - SHOULD set `fee_range` according to the minimum and maximum fees it is prepared to pay for a close transaction. - if it doesn't receive a `closing_signed` response after a reasonable amount of time: - MUST fail the channel - if it is not the funder: - SHOULD set `max_fee_satoshis` to at least the `max_fee_satoshis` received - SHOULD set `min_fee_satoshis` to a fairly low value - MUST set `signature` to the Bitcoin signature of the close transaction, as specified in [BOLT #3](#).

The receiving node: - if the `signature` is not valid for either variant of closing transaction specified in [BOLT #3](#) OR non-compliant with LOW-S-standard rule [LOW_S](#): - MUST send a warning and close the connection, or send an error and fail the channel. - if `fee_satoshis` is equal to its previously sent `fee_satoshis`: - SHOULD sign and broadcast the final closing transaction. - MAY close the connection. - if `fee_satoshis` matches its previously sent `fee_range`: - SHOULD use `fee_satoshis` to sign and broadcast the final closing transaction - SHOULD reply with a `closing_signed` with the same `fee_satoshis` value if it is different from its previously sent `fee_satoshis` - MAY close the connection. - if the message contains a `fee_range`: - if there is no overlap between that and its own `fee_range`: - SHOULD send a warning - MUST fail the channel if it doesn't receive a satisfying `fee_range` after a reasonable amount of time - otherwise: - if it is the funder: - if `fee_satoshis` is not in the overlap between the sent and received `fee_range`: - MUST fail the channel - otherwise: - MUST reply with the same `fee_satoshis`. - otherwise (it is not the funder): - if it has already sent a `closing_signed`: - if `fee_satoshis` is not the same as the value it sent: - MUST fail the channel - otherwise: - MUST propose a `fee_satoshis` in the overlap between received and (about-to-be) sent `fee_range`. - otherwise, if `fee_satoshis` is not strictly between its last-sent `fee_satoshis` and its previously-received `fee_satoshis`, UNLESS it has since reconnected: - SHOULD send a warning and close the connection, or send an error and fail the channel. - otherwise, if the receiver agrees with the fee: - SHOULD reply with a `closing_signed` with the same `fee_satoshis` value. - otherwise: - MUST propose a value "strictly between" the received `fee_satoshis` and its previously-sent `fee_satoshis`.

The receiving node: - if one of the outputs in the closing transaction is below the dust limit for its `scriptpubkey` (see [BOLT 3](#)): - MUST fail the channel

Rationale

When `fee_range` is not provided, the "strictly between" requirement ensures that forward progress is made, even if only by a single satoshi at a time. To avoid keeping state and to handle the corner case, where fees have shifted between disconnection and reconnection, negotiation restarts on reconnection.

Note there is limited risk if the closing transaction is delayed, but it will be broadcast very soon; so there is usually no reason to pay a premium for rapid processing.


Note that the non-funder is not paying the fee, so there is no reason for it to have a maximum feerate. It may want a minimum feerate, however, to ensure that the transaction propagates. It can always use CPFP later to speed up confirmation if necessary, so that minimum should be low.

It may happen that the closing transaction doesn't meet bitcoin's default relay policies (e.g. when using a non-segwit shutdown script for an output below 546 satoshis, which is possible if `dust_limit_satoshis` is below 546 satoshis). No funds are at risk when that happens, but the channel must be force-closed as the closing transaction will likely never reach miners.

Normal Operation

Once both nodes have exchanged `channel_ready` (and optionally [announcement_signatures](#)), the channel can be used to make payments via Hashed Time Locked Contracts.

Changes are sent in batches: one or more `update_` messages are sent before a `commitment_signed` message, as in the following diagram:



Counter-intuitively, these updates apply to the *other node's* commitment transaction; the node only adds those updates to its own commitment transaction when the remote node acknowledges it has applied them via `revoke_and_ack`.

Thus each update traverses through the following states:

1. pending on the receiver
2. in the receiver's latest commitment transaction
3. ... and the receiver's previous commitment transaction has been revoked, and the update is pending on the sender
4. ... and in the sender's latest commitment transaction
5. ... and the sender's previous commitment transaction has been revoked

As the two nodes' updates are independent, the two commitment transactions may be out of sync indefinitely. This is not concerning: what matters is whether both sides have irrevocably committed to a particular update or not (the final state, above).

Forwarding HTLCs

In general, a node offers HTLCs for two reasons: to initiate a payment of its own, or to forward another node's payment. In the forwarding case, care must be taken to ensure the *outgoing* HTLC cannot be redeemed unless the *incoming* HTLC can be redeemed. The following requirements ensure this is always true.

The respective **addition/removal** of an HTLC is considered *irrevocably committed* when:

1. The commitment transaction **with/without** it is committed to by both nodes, and any previous commitment transaction **without/with** it has been revoked, OR
2. The commitment transaction **with/without** it has been irreversibly committed to the blockchain.

Requirements

A node: - until an incoming HTLC has been irrevocably committed: - MUST NOT offer the corresponding outgoing HTLC (`update_add_htlc`) in response to that incoming HTLC. - until the removal of an outgoing HTLC is irrevocably committed, OR until the outgoing on-chain HTLC output has been spent via the HTLC-timeout transaction (with sufficient depth): - MUST NOT fail the incoming HTLC (`update_fail_htlc`) that corresponds to that outgoing HTLC. - once the `cltv_expiry` of an incoming HTLC has been reached, OR if `cltv_expiry` minus `current_height` is less than `cltv_expiry_delta` for the corresponding outgoing HTLC: - MUST fail that incoming HTLC (`update_fail_htlc`). - if an incoming HTLC's `cltv_expiry` is unreasonably far in the future: - SHOULD fail that incoming HTLC (`update_fail_htlc`). - upon receiving an `update_fulfill_htlc` for an outgoing HTLC, OR upon discovering the `payment_preimage` from an on-chain HTLC spend: - MUST fulfill the incoming HTLC that corresponds to that outgoing HTLC.

Rationale

In general, one side of the exchange needs to be dealt with before the other. Fulfilling an HTLC is different: knowledge of the preimage is, by definition, irrevocable and the incoming HTLC should be fulfilled as soon as possible to reduce latency.

An HTLC with an unreasonably long expiry is a denial-of-service vector and therefore is not allowed. Note that the exact value of “unreasonable” is currently unclear and may depend on network topology.

`cltv_expiry_delta` Selection

Once an HTLC has timed out, it can either be fulfilled or timed-out; care must be taken around this transition, both for offered and received HTLCs.

Consider the following scenario, where A sends an HTLC to B, who forwards to C, who delivers the goods as soon as the payment is received.

1. C needs to be sure that the HTLC from B cannot time out, even if B becomes unresponsive; i.e. C can fulfill the incoming HTLC on-chain before B can time it out on-chain.
2. B needs to be sure that if C fulfills the HTLC from B, it can fulfill the incoming HTLC from A; i.e. B can get the preimage from C and fulfill the incoming HTLC on-chain before A can time it out on-chain.

The critical settings here are the `cltv_expiry_delta` in [BOLT #7](#) and the related `min_final_cltv_expiry_delta` in [BOLT #11](#). `cltv_expiry_delta` is the minimum difference in HTLC CLTV timeouts, in the forwarding case (B). `min_final_cltv_expiry_delta` is the minimum difference between HTLC CLTV timeout and the current block height, for the terminal case (C).

Note that a node is at risk if it accepts an HTLC in one channel and offers an HTLC in another channel with too small of a difference between the CLTV timeouts. For this reason, the `cltv_expiry_delta` for the *outgoing* channel is used as the delta across a node.

The worst-case number of blocks between outgoing and incoming HTLC resolution can be derived, given a few assumptions:

- a worst-case reorganization depth R blocks
- a grace-period G blocks after HTLC timeout before giving up on an unresponsive peer and dropping to chain
- a number of blocks S between transaction broadcast and the transaction being included in a block

The worst case is for a forwarding node (B) that takes the longest possible time to spot the outgoing HTLC fulfillment and also takes the longest possible time to redeem it on-chain:

1. The B->C HTLC times out at block N , and B waits G blocks until it gives up waiting for C. B or C commits to the blockchain, and B spends HTLC, which takes S blocks to be included.
2. Bad case: C wins the race (just) and fulfills the HTLC, B only sees that transaction when it sees block $N+G+S+1$.
3. Worst case: There's reorganization R deep in which C wins and fulfills. B only sees transaction at $N+G+S+R$.
4. B now needs to fulfill the incoming A->B HTLC, but A is unresponsive: B waits G more blocks before giving up waiting for A. A or B commits to the blockchain.
5. Bad case: B sees A's commitment transaction in block $N+G+S+R+G+1$ and has to spend the HTLC output, which takes S blocks to be mined.
6. Worst case: there's another reorganization R deep which A uses to spend the commitment transaction, so B sees A's commitment transaction in block $N+G+S+R+G+R$ and has to spend the HTLC output, which takes S blocks to be mined.
7. B's HTLC spend needs to be at least R deep before it times out, otherwise another reorganization could allow A to timeout the transaction.

Thus, the worst case is $3R+2G+2S$, assuming R is at least 1. Note that the chances of three reorganizations in which the other node wins all of them is low for R of 2 or more. Since high fees are used (and HTLC spends can use almost arbitrary fees), S should be small during normal operation; although, given that block times are irregular, empty blocks still occur, fees may vary greatly, and the fees cannot be bumped on HTLC transactions, $S=12$ should be considered a minimum. S is also the parameter that may vary the most under attack, so a higher value may be desirable when non-negligible amounts are at risk. The grace period G can be low (1 or 2), as nodes are required to timeout or fulfill as soon as possible; but if G is too low it increases the risk of unnecessary channel closure due to networking delays.

There are four values that need be derived:

1. the `cltv_expiry_delta` for channels, $3R+2G+2S$: if in doubt, a `cltv_expiry_delta` of at least 34 is reasonable ($R=2$, $G=2$, $S=12$).
2. the deadline for offered HTLCs: the deadline after which the channel has to be failed and timed out on-chain. This is G blocks after the HTLC's `cltv_expiry`: 1 or 2 blocks is reasonable.
3. the deadline for received HTLCs this node has fulfilled: the deadline after which the channel has to be failed and the HTLC fulfilled on-chain before its `cltv_expiry`. See steps 4-7 above, which imply a deadline of $2R+G+S$ blocks before `cltv_expiry`: 18 blocks is reasonable.

4. the minimum `cltv_expiry` accepted for terminal payments: the worst case for the terminal node C is $2R+G+S$ blocks (as, again, steps 1-3 above don't apply). The default in [BOLT #11](#) is 18, which matches this calculation.

Requirements

An offering node: - MUST estimate a timeout deadline for each HTLC it offers. - MUST NOT offer an HTLC with a timeout deadline before its `cltv_expiry`. - if an HTLC which it offered is in either node's current commitment transaction, AND is past this timeout deadline: - SHOULD send an error to the receiving peer (if connected). - MUST fail the channel.

A fulfilling node: - for each HTLC it is attempting to fulfill: - MUST estimate a fulfillment deadline. - MUST fail (and not forward) an HTLC whose fulfillment deadline is already past. - if an HTLC it has fulfilled is in either node's current commitment transaction, AND is past this fulfillment deadline: - SHOULD send an error to the offering peer (if connected). - MUST fail the channel.

Bounding exposure to trimmed in-flight HTLCs: `max_dust_htlc_exposure_msat`

When an HTLC in a channel is below the "trimmed" threshold in [BOLT3 #3](#), the HTLC cannot be claimed on-chain, instead being turned into additional miner fees if either party unilaterally closes the channel. Because the threshold is per-HTLC, the total exposure to such HTLCs may be substantial if there are many dust HTLCs committed when the channel is force-closed.

This can be exploited in grieving attacks or even in miner-extractable-value attacks, if the malicious entity wins [mining capabilities](#).

The total exposure is given by the following back-of-the-envelope computation:

```
remote `max_accepted_htlcs` * (`HTLC-success-kiloweight` * `feerate_per_kw` + remote
`dust_limit_satoshis`)
+ local `max_accepted_htlcs` * (`HTLC-timeout-kiloweight` * `feerate_per_kw` + remote
`dust_limit_satoshis`)
```

To mitigate this scenario, a `max_dust_htlc_exposure_msat` threshold can be applied when sending, forwarding and receiving HTLCs.

A node: - when receiving an HTLC: - if the HTLC's `amount_msat` is smaller than the remote `dust_limit_satoshis` plus the HTLC-timeout fee at `feerate_per_kw`: - if the `amount_msat` plus the dust balance of the remote transaction is greater than `max_dust_htlc_exposure_msat`: - SHOULD fail this HTLC once it's committed - SHOULD NOT reveal a preimage for this HTLC - if the HTLC's `amount_msat` is smaller than the local `dust_limit_satoshis` plus the HTLC-success fee at `feerate_per_kw`: - if the `amount_msat` plus the dust balance of the local transaction is greater than `max_dust_htlc_exposure_msat`: - SHOULD fail this HTLC once it's committed - SHOULD NOT reveal a preimage for this HTLC - when offering an HTLC: - if the HTLC's `amount_msat` is smaller than the remote `dust_limit_satoshis` plus the HTLC-success fee at `feerate_per_kw`: - if the `amount_msat` plus the dust balance of the remote transaction is greater than `max_dust_htlc_exposure_msat`: - SHOULD NOT send this HTLC - SHOULD fail the corresponding incoming HTLC (if any) - if the HTLC's `amount_msat` is inferior to the holder's `dust_limit_satoshis` plus the HTLC-timeout fee at the `feerate_per_kw`: - if the `amount_msat` plus the dust balance of the local transaction is greater

than `max_dust_htlc_exposure_msat`: - SHOULD NOT send this HTLC - SHOULD fail the corresponding incoming HTLC (if any)

The `max_dust_htlc_exposure_msat` is an upper bound on the trimmed balance from dust exposure. The exact value used is a matter of node policy.

For channels that don't use `zero_fee_commitments` or `option_anchors`, an increase of the `feerate_per_kw` may trim multiple htlcs from commitment transactions, which could create a large increase in dust exposure.

Adding an HTLC: `update_add_htlc`

Either node can send `update_add_htlc` to offer an HTLC to the other, which is redeemable in return for a payment preimage. Amounts are in millisatoshi, though on-chain enforcement is only possible for whole satoshi amounts greater than the dust limit (in commitment transactions these are rounded down as specified in [BOLT #3](#)).

The format of the `onion_routing_packet` portion, which indicates where the payment is destined, is described in [BOLT #4](#).

1. type: 128 (`update_add_htlc`)
2. data:
 - `[channel_id:channel_id]`
 - `[u64:id]`
 - `[u64:amount_msat]`
 - `[sha256:payment_hash]`
 - `[u32:cltv_expiry]`
 - `[1366*byte:onion_routing_packet]`
3. tlv_stream: `update_add_htlc_tlvs`
4. types:
 1. type: 0 (`blinded_path`)
 2. data:
 - `[point:path_key]`

Requirements

A sending node: - if `zero_fee_commitments` has not been negotiated: - if it is *responsible* for paying the Bitcoin fee: - MUST NOT offer `amount_msat` if, after adding that HTLC to its commitment transaction, it cannot pay the fee for either the local or remote commitment transaction at the current `feerate_per_kw` while maintaining its channel reserve (see [Updating Fees](#)). - if `option_anchors` applies to this commitment transaction and the sending node is the funder: - MUST be able to additionally pay for `to_local_anchor` and `to_remote_anchor` above its reserve. - SHOULD NOT offer `amount_msat` if, after adding that HTLC to its commitment transaction, its remaining balance doesn't allow it to pay the commitment transaction fee when receiving or sending a future additional non-dust HTLC while maintaining its channel reserve. It is recommended that this "fee spike buffer" can handle twice the current `feerate_per_kw` to ensure predictability between implementations. - if it is *not responsible* for paying the Bitcoin fee: - SHOULD NOT offer `amount_msat` if, once the remote node adds that HTLC to its commitment transaction, it cannot pay the fee for the updated local or remote transaction at the current `feerate_per_kw` while maintaining its channel reserve. - MUST offer `amount_msat` greater than 0. - MUST NOT offer `amount_msat` below the receiving node's `htlc_minimum_msat` - MUST set `cltv_expiry` less

than 500000000. - if result would be offering more than the remote's `max_accepted_htlcs` HTLCs, in the remote commitment transaction: - MUST NOT add an HTLC. - if the total value of offered HTLCs would exceed the remote's `max_htlc_value_in_flight_msat`: - MUST NOT add an HTLC. - for the first HTLC it offers: - MUST set `id` to 0. - MUST increase the value of `id` by 1 for each successive offer. - if it is relaying a payment inside a blinded route: - MUST set `path_key` (see [Route Blinding](#)) - if a splice is pending: - MUST ensure that requirements are met for all commitment transactions.

`id` MUST NOT be reset to 0 after the update is complete (i.e. after `revoke_and_ack` has been received). It MUST continue incrementing instead.

A receiving node: - receiving an `amount_msat` equal to 0, OR less than its own `htlc_minimum_msat`: - SHOULD send a warning and close the connection, or send an error and fail the channel. - receiving an `amount_msat` that the sending node cannot afford at the current `feerate_per_kw` (while maintaining its channel reserve and any `to_local_anchor` and `to_remote_anchor` costs): - SHOULD send a warning and close the connection, or send an error and fail the channel. - if a sending node adds more than receiver `max_accepted_htlcs` HTLCs to its local commitment transaction, OR adds more than receiver `max_htlc_value_in_flight_msat` worth of offered HTLCs to its local commitment transaction: - SHOULD send a warning and close the connection, or send an error and fail the channel. - if sending node sets `cltv_expiry` to greater or equal to 500000000: - SHOULD send a warning and close the connection, or send an error and fail the channel. - MUST allow multiple HTLCs with the same `payment_hash`. - if the sender did not previously acknowledge the commitment of that HTLC: - MUST ignore a repeated `id` value after a reconnection. - if other `id` violations occur: - MAY send a warning and close the connection, or send an error and fail the channel. - MUST decrypt `onion_routing_packet` as described in [Onion Decryption](#) to extract a payload. - MUST use `path_key` (if specified). - MUST use `payment_hash` as `associated_data`. - If decryption fails, the result is not a valid payload TLV, or it contains unknown even types: - MUST respond with an error as detailed in [Failure Messages](#) - Otherwise: - MUST follow the requirements for the reader of payload in [Payload Format](#) - If a splice is pending: - MUST ensure that requirements are met for all commitment transactions.

The `onion_routing_packet` contains an obfuscated list of hops and instructions for each hop along the path. It commits to the HTLC by setting the `payment_hash` as associated data, i.e. includes the `payment_hash` in the computation of HMACs. This prevents replay attacks that would reuse a previous `onion_routing_packet` with a different `payment_hash`.

Rationale

Invalid amounts are a clear protocol violation and indicate a breakdown.

If a node did not accept multiple HTLCs with the same payment hash, an attacker could probe to see if a node had an existing HTLC. This requirement, to deal with duplicates, leads to the use of a separate identifier; it's assumed a 64-bit counter never wraps.

Retransmissions of unacknowledged updates are explicitly allowed for reconnection purposes; allowing them at other times simplifies the recipient code (though strict checking may help debugging).

`max_accepted_htlcs` is limited to 483 to ensure that, even if both sides send the maximum number of HTLCs, the `commitment_signed` message will still be under the maximum message size. It also ensures that a single penalty transaction can spend the entire commitment transaction, as calculated in [BOLT #5](#). When using `zero_fee_commitments`, we further limit `max_accepted_htlcs` to 114 because v3 transactions are limited to 10kB, which decreases the number of outputs the commitment transaction can have.

cltv_expiry values equal to or greater than 500000000 would indicate a time in seconds, and the protocol only supports an expiry in blocks.

When zero_fee_commitments is not used, the node *responsible* for paying the Bitcoin fee should maintain a “fee spike buffer” on top of its reserve to accommodate a future fee increase. Without this buffer, the node *responsible* for paying the Bitcoin fee may reach a state where it is unable to send or receive any non-dust HTLC while maintaining its channel reserve (because of the increased weight of the commitment transaction), resulting in a degraded channel. See [#728](#) for more details.

If splicing is supported, there can be more than one commitment transaction at a time: proposed changes must be valid for all of them.

Removing an HTLC: update_fulfill_htlc, update_fail_htlc, and update_fail_malformed_htlc

For simplicity, a node can only remove HTLCs added by the other node. There are four reasons for removing an HTLC: the payment preimage is supplied, it has timed out, it has failed to route, or it is malformed.

To supply the preimage:

1. type: 130 (update_fulfill_htlc)
2. data:
 - [channel_id:channel_id]
 - [u64:id]
 - [32*byte:payment_preimage]
3. tlv_stream: update_fulfill_htlc_tlvs
4. types:
 1. type: 1 (attribution_data)
 2. data:
 - [20*u32:htlc_hold_times]
 - [210*sha256[..4]:truncated_hmacs]
 3. type: 3 (fulfillment_payload)
 4. data:
 - [...*byte:fulfillment_payload]

The fulfillment_payload field is an opaque encrypted blob for the benefit of the original HTLC initiator, as defined in [BOLT #4](#). It mirrors the reason field of update_fail_htlc: the final node originates it, and intermediate nodes obfuscate and relay it.

For a timed out or route-failed HTLC:

1. type: 131 (update_fail_htlc)
2. data:
 - [channel_id:channel_id]
 - [u64:id]
 - [u16:len]
 - [len*byte:reason]
3. tlv_stream: update_fail_htlc_tlvs
4. types:
 1. type: 1 (attribution_data)

2. data:

- [20*sha256:htlc_hold_times]
- [210*sha256[...4]:truncated_hmacs]

The reason field is an opaque encrypted blob for the benefit of the original HTLC initiator, as defined in [BOLT #4](#); however, there's a special malformed failure variant for the case where the peer couldn't parse it: in this case the current node instead takes action, encrypting it into a `update_fail_htlc` for relaying.

For an unparsable HTLC:

1. type: 135 (`update_fail_malformed_htlc`)

2. data:

- [channel_id:channel_id]
- [u64:id]
- [sha256:sha256_of_onion]
- [u16:failure_code]

Requirements

A node: - SHOULD remove an HTLC as soon as it can. - SHOULD fail an HTLC which has timed out. - until the corresponding HTLC is irrevocably committed in both sides' commitment transactions: - MUST NOT send an `update_fulfill_htlc`, `update_fail_htlc`, or `update_fail_malformed_htlc`. - When failing an incoming HTLC: - If `current_path_key` is set in the onion payload and it is not the final node: - MUST send an `update_fail_htlc` error using the `invalid_onion_blinding` failure code for any local or downstream errors. - SHOULD use the `sha256_of_onion` of the onion it received. - MAY use an all zero `sha256_of_onion`. - SHOULD add a random delay before sending `update_fail_htlc`. - If `path_key` is set in the incoming `update_add_htlc`: - MUST send an `update_fail_malformed_htlc` error using the `invalid_onion_blinding` failure code for any local or downstream errors. - SHOULD use the `sha256_of_onion` of the onion it received. - MAY use an all zero `sha256_of_onion`. - When supporting `option_attribution_data`: - if `path_key` is not set in the incoming `update_add_htlc`: - MUST initialize `attribution_data` and include it in `update_fail_htlc` and `update_fulfill_htlc`. See [BOLT04](#). - if a `fulfillment_payload` is present in the `update_fulfill_htlc` received from the downstream node: - MUST obfuscate and relay it in the outgoing `update_fulfill_htlc`, including in a blinded path. See [BOLT04](#).

A receiving node: - if the `id` does not correspond to an HTLC in its current commitment transaction: - MUST send a warning and close the connection, or send an error and fail the channel. - if the `payment_preimage` value in `update_fulfill_htlc` doesn't SHA256 hash to the corresponding HTLC `payment_hash`: - MUST send a warning and close the connection, or send an error and fail the channel. - if the `fulfillment_payload` in `update_fulfill_htlc` is longer than 32768 bytes (32 KiB): - MUST send an error and fail the channel. - if the `BADONION` bit in `failure_code` is not set for `update_fail_malformed_htlc`: - MUST send a warning and close the connection, or send an error and fail the channel. - if the `sha256_of_onion` in `update_fail_malformed_htlc` doesn't match the onion it sent and is not all zero: - MAY retry or choose an alternate error response. - otherwise, a receiving node which has an outgoing HTLC canceled by `update_fail_malformed_htlc`: - MUST return an error in the `update_fail_htlc` sent to the link which originally sent the HTLC, using the `failure_code` given and setting the data to `sha256_of_onion`.

Rationale

A node that doesn't time out HTLCs risks channel failure (see [cltv_expiry_delta Selection](#)).

A node that sends `update_fulfill_htlc`, before the sender, is also committed to the HTLC and risks losing funds.

If the onion is malformed, the upstream node won't be able to extract the shared key to generate a response — hence the special failure message, which makes this node do it.

The node can check that the SHA256 that the upstream is complaining about does match the onion it sent, which may allow it to detect random bit errors. However, without re-checking the actual encrypted packet sent, it won't know whether the error was its own or the remote's; so such detection is left as an option.

Nodes inside a blinded route must use `invalid_onion_blinding` to avoid leaking information to senders trying to probe the blinded route.

Batching channel messages

Multiple channel messages can be grouped together as a single logical message using the `start_batch` message.

1. type: 127 (`start_batch`)
2. data:
 - `[channel_id:channel_id]`
 - `[u16:batch_size]`
3. tlv_stream: `start_batch_tlvs`
4. types:
 1. type: 1 (`message_type`)
 2. data:
 - `[u16:message_type]`

Requirements

A sending node: - MUST set `batch_size` to a value strictly greater than 1. - MUST set `batch_size` to a value lower than or equal to 20. - MUST set `message_type` to 132 (`commitment_signed` message type as defined in [Bolt 1](#)). - After sending `start_batch`: - MUST send `batch_size` `commitment_signed` messages for the same `channel_id`, without any other unrelated messages in-between.

A receiving node: - If `batch_size` is not strictly greater than 1: - MUST ignore the `start_batch` message. - SHOULD send a warning. - If `batch_size` is strictly greater than 20: - MUST send a warning and close the connection, or send an error and fail the channel. - MUST group the next `batch_size` messages and process them together. - If one of those messages is not for the specified `channel_id`: - MUST send a warning and close the connection, or send an error and fail the channel. - If `message_type` is missing or not set to the type for `commitment_signed`: - MUST ignore the `start_batch` message and process the following messages sequentially.

Rationale

The `start_batch` message is only used when sending multiple `commitment_signed` messages during splicing. We thus narrow the requirements for this specific scenario: they can be relaxed in the future if we use `start_batch` for other features.

We limit the `batch_size` to 20 elements to protect against excessive queuing that could be abused to DoS receiving nodes. The `start_batch` message is only used for splice RBF attempts so far, which shouldn't need that many attempts to get transactions confirmed.

Committing Updates So Far: `commitment_signed`

When a node has changes for the remote commitment, it can apply them, sign the resulting transaction (as defined in [BOLT #3](#)), and send a `commitment_signed` message.

1. type: 132 (`commitment_signed`)
2. data:
 - `[channel_id:channel_id]`
 - `[signature:signature]`
 - `[u16:num_htlcs]`
 - `[num_htlcs*signature:htlc_signature]`
 - `[commitment_signed_tlvs:tlvs]`
3. `tlv_stream`: `commitment_signed_tlvs`
4. types:
 1. type: 1 (`funding_txid`)
 2. data:
 - `[sha256:funding_txid]`

Requirements

A sending node: - MUST NOT send a `commitment_signed` message that does not include any updates. - MAY send a `commitment_signed` message that only alters the fee. - MAY send a `commitment_signed` message that doesn't change the commitment transaction aside from the new revocation number (due to dust, identical HTLC replacement, or insignificant or multiple fee changes). - MUST include one `htlc_signature` for every HTLC transaction corresponding to the ordering of the commitment transaction (see [BOLT #3](#)). - MUST set `funding_txid` to the funding transaction spent by this commitment transaction. - if it has not recently received a message from the remote node: - SHOULD use ping and await the reply pong before sending `commitment_signed`. - If there are N pending splice transactions (with N greater than 0): - MUST send `start_batch` first with `batch_size` set to N + 1 and `message_type` set to 132 (`commitment_signed`). - MUST send `commitment_signed` for the current channel funding output. - MUST send `commitment_signed` for each of the splice transactions. - MUST set `funding_txid` in each `commitment_signed` message to match the funding transaction spent by that commitment transaction. - MUST NOT send any other message before it has sent the batch of `commitment_signed` messages.

A receiving node: - once all pending updates are applied: - if signature is not valid for its local commitment transaction OR non-compliant with LOW-S-standard rule [LWS](#): - MUST send a warning and close the connection, or send an error and fail the channel. - if `num_htlcs` is not equal to the number of HTLC outputs in the local commitment transaction: - MUST send a warning and close the connection, or send an error and fail

the channel. - if any `htlc_signature` is not valid for the corresponding HTLC transaction OR non-compliant with LOW-S-standard rule [LOWS](#): - MUST send a warning and close the connection, or send an error and fail the channel. - If there are pending splice transactions and the sending node did not send `start_batch` followed by a batch of `commitment_signed` messages: - MUST send an error and fail the channel. - If the sending node sent `start_batch` and we are processing a batch of `commitment_signed` messages: - If `funding_txid` is missing in one of the `commitment_signed` messages: - MUST send an error and fail the channel. - If there are pending splice transactions: - MUST validate each `commitment_signed` based on `funding_txid`. - If `commitment_signed` is missing for a funding transaction: - MUST send an error and fail the channel. - Otherwise: - MUST respond with a `revoke_and_ack` message. - Otherwise (no pending splice transactions): - MUST ignore `commitment_signed` where `funding_txid` does not match the current funding transaction. - If `commitment_signed` is missing for the current funding transaction: - MUST send an error and fail the channel. - Otherwise: - MUST respond with a `revoke_and_ack` message. - Otherwise: - MUST respond with a `revoke_and_ack` message.

Rationale

There's little point offering spam updates: it implies a bug.

The `num_htlcs` field is redundant, but makes the packet length check fully self-contained.

The recommendation to require recent messages recognizes the reality that networks are unreliable: nodes might not realize their peers are offline until after sending `commitment_signed`. Once `commitment_signed` is sent, the sender considers itself bound to those HTLCs, and cannot fail the related incoming HTLCs until the output HTLCs are fully resolved.

Note that the `htlc_signature` implicitly enforces the time-lock mechanism in the case of offered HTLCs being timed out or received HTLCs being spent. This is done to reduce fees by creating smaller scripts compared to explicitly stating time-locks on HTLC outputs.

Splicing requires us to send and receive additional signatures, as we don't know which (if any) of the splice transactions will end up being the new channel funding transaction. We use `start_batch` to send a batch of `commitment_signed` messages containing signatures for each of the pending splice transactions and for the current funding transaction. After sending `splice_locked`, we may receive a batch containing obsolete `commitment_signed` messages from our peer that they started sending before receiving our `splice_locked`: we can safely ignore them by filtering on `funding_txid`.

Completing the Transition to the Updated State: `revoke_and_ack`

Once the recipient of `commitment_signed` checks the signature and knows it has a valid new commitment transaction, it replies with the commitment preimage for the previous commitment transaction in a `revoke_and_ack` message.

This message also implicitly serves as an acknowledgment of receipt of the `commitment_signed`, so this is a logical time for the `commitment_signed` sender to apply (to its own commitment) any pending updates it sent before that `commitment_signed`.

The description of key derivation is in [BOLT #3](#).

1. type: 133 (`revoke_and_ack`)

2. data:

- [channel_id:channel_id]
- [32*byte:per_commitment_secret]
- [point:next_per_commitment_point]

Requirements

A sending node: - MUST set `per_commitment_secret` to the secret used to generate keys for the previous commitment transaction. - MUST set `next_per_commitment_point` to the values for its next commitment transaction. - MUST send a single `revoke_and_ack` message, even if it is responding to a batch of `commitment_signed` messages.

A receiving node: - if `per_commitment_secret` is not a valid secret key or does not generate the previous `per_commitment_point`: - MUST send an error and fail the channel. - if the `per_commitment_secret` was not generated by the protocol in [BOLT #3](#): - MAY send a warning and close the connection, or send an error and fail the channel.

A node: - MUST NOT broadcast old (revoked) commitment transactions, - Note: doing so will allow the other node to seize all channel funds. - SHOULD NOT sign commitment transactions, unless it's about to broadcast them (due to a failed connection), - Note: this is to reduce the above risk.

Updating Fees: `update_fee`

For channels that don't use `zero_fee_commitments`, an `update_fee` message is sent by the node which is paying the Bitcoin fee. Like any update, it's first committed to the receiver's commitment transaction and then (once acknowledged) committed to the sender's. Unlike an HTLC, `update_fee` is never closed but simply replaced.

There is a possibility of a race, as the recipient can add new HTLCs before it receives the `update_fee`. Under this circumstance, the sender may not be able to afford the fee on its own commitment transaction, once the `update_fee` is finally acknowledged by the recipient. In this case, the fee will be less than the fee rate, as described in [BOLT #3](#).

The exact calculation used for deriving the fee from the fee rate is given in [BOLT #3](#).

1. type: 134 (`update_fee`)

2. data:

- [channel_id:channel_id]
- [u32:feerate_per_kw]

Requirements

When using `zero_fee_commitments`, nodes: - MUST NOT send `update_fee`.

The node *responsible* for paying the Bitcoin fee: - SHOULD send `update_fee` to ensure the current fee rate is sufficient (by a significant margin) for timely processing of the commitment transaction.

The node *not responsible* for paying the Bitcoin fee: - MUST NOT send `update_fee`.

A sending node: - if `option_anchors` was not negotiated: - if the `update_fee` increases `feerate_per_kw`: - if the dust balance of the remote transaction at the updated `feerate_per_kw` is greater than

max_dust_htlc_exposure_msat: - MAY NOT send update_fee - MAY fail the channel - if the dust balance of the local transaction at the updated feerate_per_kw is greater than max_dust_htlc_exposure_msat: - MAY NOT send update_fee - MAY fail the channel - if a splice is pending: - MUST ensure that requirements are met for all commitment transactions.

A receiving node: - if zero_fee_commitments is used: - MUST send an error and fail the channel. - if the update_fee is too low for timely processing, OR is unreasonably large: - MUST send a warning and close the connection, or send an error and fail the channel. - if the sender is not responsible for paying the Bitcoin fee: - MUST send a warning and close the connection, or send an error and fail the channel. - if the sender cannot afford the new fee rate on the receiving node's current commitment transaction: - SHOULD send a warning and close the connection, or send an error and fail the channel. - but MAY delay this check until the update_fee is committed. - if option_anchors was not negotiated: - if the update_fee increases feerate_per_kw: - if the dust balance of the remote transaction at the updated feerate_per_kw is greater than max_dust_htlc_exposure_msat: - MAY fail the channel - if the dust balance of the local transaction at the updated feerate_per_kw is greater than max_dust_htlc_exposure_msat: - MAY fail the channel - if a splice is pending: - MUST ensure that requirements are met for all commitment transactions.

Rationale

Bitcoin fees are required for unilateral closes to be effective. With option_anchors, feerate_per_kw is not as critical anymore to guarantee confirmation as it was in the legacy commitment format, but it still needs to be enough to be able to enter the mempool (satisfy min relay fee and mempool min fee). With zero_fee_commitments, transactions can enter the mempool as a package with a child transaction paying the fees: channel transactions thus don't need to directly include a fee, so this message isn't used.

For the legacy commitment format, there is no general method for the broadcasting node to use child-pays-for-parent to increase its effective fee.

Given the variance in fees, and the fact that the transaction may be spent in the future, it's a good idea for the fee payer to keep a good margin (say 5x the expected fee requirement) for legacy commitment txes; but, due to differing methods of fee estimation, an exact value is not specified.

Since the fees are currently one-sided (the party which requested the channel creation always pays the fees for the commitment transaction), it's simplest to only allow it to set fee levels; however, as the same fee rate applies to HTLC transactions, the receiving node must also care about the reasonableness of the fee.

If on-chain fees increase while commitments contain many HTLCs that will be trimmed at the updated feerate, this could overflow the configured max_dust_htlc_exposure_msat. Whether to close the channel preemptively or not is left as a matter of node policy.

If splicing is supported, there can be more than one commitment transaction at a time: proposed changes must be valid for all of them.

Message Retransmission

Because communication transports are unreliable, and may need to be re-established from time to time, the design of the transport has been explicitly separated from the protocol.

Nonetheless, it's assumed our transport is ordered and reliable. Reconnection introduces doubt as to what has been received, so there are explicit acknowledgments at that point.

This is fairly straightforward in the case of channel establishment and close, where messages have an explicit order, but during normal operation, acknowledgments of updates are delayed until the `commitment_signed` / `revoke_and_ack` exchange; so it cannot be assumed that the updates have been received. This also means that the receiving node only needs to store updates upon receipt of `commitment_signed`.

Note that messages described in [BOLT #7](#) are independent of particular channels; their transmission requirements are covered there, and besides being transmitted after `init` (as all messages are), they are independent of requirements here.

1. type: 136 (`channel_reestablish`)
2. data:
 - [`channel_id:channel_id`]
 - [`u64:next_commitment_number`]
 - [`u64:next_revocation_number`]
 - [`32*byte:your_last_per_commitment_secret`]
 - [`point:my_current_per_commitment_point`]
3. tlv_stream: `channel_reestablish_tlvs`
4. types:
 1. type: 1 (`next_funding`)
 2. data:
 - [`sha256:next_funding_txid`]
 - [`byte:retransmit_flags`]
 3. type: 5 (`my_current_funding_locked`)
 4. data:
 - [`sha256:my_current_funding_locked_txid`]
 - [`byte:retransmit_flags`]

`next_commitment_number`: A commitment number is a 48-bit incrementing counter for each commitment transaction; counters are independent for each peer in the channel and start at 0. They're only explicitly relayed to the other node in the case of re-establishment, otherwise they are implicit.

The `next_funding.retransmit_flags` bitfield is used to let the receiving peer know which messages they must retransmit for the corresponding `next_funding_txid` after the reconnection:

Bit Position	Name
0	<code>commitment_signed</code>

The `my_current_funding_locked.retransmit_flags` bitfield is used to let our peer know which messages we expect them to retransmit after the reconnection:

Bit Position	Name
0	<code>announcement_signatures</code>

Requirements

A funding node: - upon disconnection: - if it has broadcast the funding transaction: - MUST remember the channel for reconnection. - otherwise: - SHOULD NOT remember the channel for reconnection.

A non-funding node: - upon disconnection: - if it has sent the funding_signed message: - MUST remember the channel for reconnection. - otherwise: - SHOULD NOT remember the channel for reconnection.

A node: - MUST handle continuation of a previous channel on a new encrypted transport. - upon disconnection: - MUST reverse any uncommitted updates sent by the other side (i.e. all messages beginning with update_ for which no commitment_signed has been received). - Note: a node MAY have already used the payment_preimage value from the update_fulfill_htlc, so the effects of update_fulfill_htlc are not completely reversed. - upon reconnection: - if a channel is in an error state: - SHOULD retransmit the error packet and ignore any other packets for that channel. - otherwise: - MUST transmit channel_reestablish for each channel. - MUST wait to receive the other node's channel_reestablish message before sending any other messages for that channel.

The sending node: - MUST set next_commitment_number to the commitment number of the next commitment_signed it expects to receive. - MUST set next_revocation_number to the commitment number of the next revoke_and_ack message it expects to receive. - MUST set my_current_per_commitment_point to a valid point. - if next_revocation_number equals 0: - MUST set your_last_per_commitment_secret to all zeroes - otherwise: - MUST set your_last_per_commitment_secret to the last per_commitment_secret it received - if it has sent commitment_signed for an interactive transaction construction but it has not received tx_signatures: - MUST include the next_funding TLV. - MUST set next_funding_txid to the txid of that interactive transaction. - if it has not received commitment_signed for this next_funding_txid: - MUST set the commitment_signed bit in retransmit_flags. - otherwise: - MUST NOT include the next_funding TLV. - if option_splice was negotiated: - if a splice transaction reached acceptable depth while disconnected: - MUST include my_current_funding_locked with the txid of the latest such transaction. - otherwise, if it has already sent splice_locked for any transaction: - MUST include my_current_funding_locked with the txid of the last splice_locked it sent. - otherwise, if it has already sent channel_ready: - MUST include my_current_funding_locked with the txid of the channel funding transaction. - otherwise (it has never sent channel_ready or splice_locked): - MUST NOT include my_current_funding_locked. - if my_current_funding_locked is included: - if announce_channel is set for this channel: - if it has not received announcement_signatures for that transaction: - MUST set the announcement_signatures bit to 1 in retransmit_flags. - otherwise: - MUST set the announcement_signatures bit to 0 in retransmit_flags.

A node: - if next_commitment_number is zero: - MUST immediately fail the channel and broadcast any relevant latest commitment transaction. - if next_commitment_number is 1 in both the channel_reestablish it sent and received, and none of those channel_reestablish messages contain my_current_funding_locked or next_funding for a splice transaction: - MUST retransmit channel_ready. - otherwise: - MUST NOT retransmit channel_ready, but MAY send channel_ready with a different short_channel_id alias field. - upon reconnection: - MUST ignore any redundant channel_ready it receives. - if next_commitment_number is equal to the commitment number of the last commitment_signed message the receiving node has sent: - MUST reuse the same commitment number for its next commitment_signed. - otherwise: - if next_commitment_number is not equal to the commitment number of the next commitment_signed that the receiving node would send: - SHOULD send an error and fail the channel. - if next_revocation_number is equal to the commitment number of the last revoke_and_ack the receiving node sent, AND the receiving node hasn't already received a closing_signed: - MUST re-send the revoke_and_ack. - if it has previously sent a commitment_signed that

needs to be retransmitted: - MUST retransmit `revoke_and_ack` and `commitment_signed` in the same relative order as initially transmitted. - otherwise: - if `next_revocation_number` is not equal to 1 greater than the commitment number of the last `revoke_and_ack` the receiving node has sent: - SHOULD send an error and fail the channel. - if it has not sent `revoke_and_ack`, AND `next_revocation_number` is not equal to 0: - SHOULD send an error and fail the channel.

A receiving node: - MUST ignore `my_current_per_commitment_point`, but MAY require it to be a valid point. - if `next_revocation_number` is greater than expected above, AND `your_last_per_commitment_secret` is correct for that `next_revocation_number` minus 1: - MUST NOT broadcast its commitment transaction. - SHOULD send an error to request the peer to fail the channel. - otherwise: - if `your_last_per_commitment_secret` does not match the expected values: - SHOULD send an error and fail the channel.

A receiving node: - if the `next_funding` TLV is set: - if `next_funding_txid` matches the latest interactive funding transaction: - if it has not received `tx_signatures` for that funding transaction: - if the `commitment_signed` bit is set in `retransmit_flags`: - MUST retransmit its `commitment_signed` for that funding transaction. - if it has already received `commitment_signed` and it should sign first, as specified in the [tx_signatures requirements](#): - MUST send its `tx_signatures` for that funding transaction. - if it has already received `tx_signatures` for that funding transaction: - MUST send its `tx_signatures` for that funding transaction. - if it also sets `next_funding` in its own `channel_reestablish`, but the values don't match: - MUST send an error and fail the channel. - otherwise: - MUST send `tx_abort` to let the sending node know that they can forget this funding transaction.

A receiving node: - if splice transactions are pending and `my_current_funding_locked` matches one of those splice transactions, for which it hasn't received `splice_locked` yet: - MUST process `my_current_funding_locked` as if it was receiving `splice_locked` for this txid. - if `my_current_funding_locked` is included with the `announcement_signatures` bit set in the `retransmit_flags`: - if `announce_channel` is set for this channel and the receiving node is ready to send `announcement_signatures` for the corresponding splice transaction: - MUST retransmit `announcement_signatures`.

A node: - MUST NOT assume that previously-transmitted messages were lost, - if it has sent a previous `commitment_signed` message: - MUST handle the case where the corresponding commitment transaction is broadcast at any time by the other side, - Note: this is particularly important if the node does not simply retransmit the exact `update_` messages as previously sent. - upon reconnection: - if it has sent a previous shutdown: - MUST retransmit shutdown.

Rationale

The requirements above ensure that the opening phase is nearly atomic: if it doesn't complete, it starts again. The only exception is if the `funding_signed` message is sent but not received. In this case, the funder will forget the channel, and presumably open a new one upon reconnection; meanwhile, the other node will eventually forget the original channel, due to never receiving `channel_ready` or seeing the funding transaction on-chain.

There's no acknowledgment for error, so if a reconnect occurs it's polite to retransmit before disconnecting again; however, it's not a MUST, because there are also occasions where a node can simply forget the channel altogether.

`closing_signed` also has no acknowledgment so must be retransmitted upon reconnection (though negotiation restarts on reconnection, so it needs not be an exact retransmission). The only acknowledgment for shutdown is `closing_signed`, so one or the other needs to be retransmitted.

The handling of updates is similarly atomic: if the commit is not acknowledged (or wasn't sent) the updates are re-sent. However, it's not insisted they be identical: they could be in a different order, involve different fees, or even be missing HTLCs which are now too old to be added. Requiring they be identical would effectively mean a write to disk by the sender upon each transmission, whereas the scheme here encourages a single persistent write to disk for each `commitment_signed` sent or received. But if you need to retransmit both a `commitment_signed` and a `revoke_and_ack`, the relative order of these two must be preserved, otherwise it will lead to a channel closure.

A re-transmittal of `revoke_and_ack` should never be asked for after a `closing_signed` has been received, since that would imply a shutdown has been completed — which can only occur after the `revoke_and_ack` has been received by the remote node.

Note that the `next_commitment_number` starts at 1, since commitment number 0 is created during opening. `next_revocation_number` will be 0 until the `commitment_signed` for commitment number 1 is sent and then the revocation for commitment number 0 is received.

`channel_ready` is implicitly acknowledged by the start of normal operation, which is known to have begun after a `commitment_signed` has been received (hence, the test for a `next_commitment_number` greater than 1) or after a splice transaction has been created.

A node, which has somehow fallen behind (e.g. has been restored from old backup), can detect that it has fallen behind. A fallen-behind node must know it cannot broadcast its current commitment transaction — which would lead to total loss of funds — as the remote node can prove it knows the revocation preimage. The error returned by the fallen-behind node should make the other node drop its current commitment transaction to the chain. The other node should wait for that error to give the fallen-behind node an opportunity to fix its state first (e.g. by restarting with a different backup).

The `next_funding` TLV allows peers to finalize the signing steps of an interactive transaction construction, or safely abort that transaction if it was not signed by one of the peers, who has thus already removed it from its state.

`my_current_funding_locked` is equivalent to sending `splice_locked`, but is handled atomically during `channel_reestablish` (instead of requiring a retransmission of the `splice_locked` message). This is useful to avoid race conditions with channel updates (for more details about this race condition, see [this example](#)). It handles the case where the `splice_locked` message was lost during the disconnection, or when a splice transaction reaches acceptable depth while peers are disconnected. It also allows requesting a retransmission of the `announcement_signatures` message for the latest splice transaction, in case it wasn't received before disconnecting.

Authors

[FIXME: Insert Author List]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

BOLT #3: Bitcoin Transaction and Script Formats

This details the exact format of on-chain transactions, which both sides need to agree on to ensure signatures are valid. This consists of the funding transaction output script, the commitment transactions, and the HTLC transactions.

Table of Contents

- [Transactions](#)
 - [Transaction Output Ordering](#)
 - [Use of Segwit](#)
 - [Funding Transaction Output](#)
 - [Commitment Transaction](#)
 - [Commitment Transaction Outputs](#)
 - [to_local Output](#)
 - [to_remote Output](#)
 - [to_local_anchor and to_remote_anchor](#)
 - [shared_anchor](#)
 - [Offered HTLC Outputs](#)
 - [Received HTLC Outputs](#)
 - [Trimmed Outputs](#)
 - [HTLC-timeout and HTLC-success Transactions](#)
 - [Legacy Closing Transaction](#)
 - [Closing Transaction](#)
 - [Fees](#)
 - [Fee Calculation](#)
 - [Fee Payment](#)
 - [Calculating Fees for Collaborative Transactions](#)
 - [Dust Limits](#)
 - [Commitment Transaction Construction](#)
- [Keys](#)
 - [Key Derivation](#)
 - [localpubkey, local_htlcpubkey, remote_htlcpubkey, local_delayedpubkey, and remote_delayedpubkey Derivation](#)
 - [remotepubkey Derivation](#)
 - [revocationpubkey Derivation](#)
 - [Per-commitment Secret Requirements](#)
 - [Efficient Per-commitment Secret Storage](#)
- [Appendix A: Expected Weights](#)
 - [Expected Weight of the Funding Transaction \(v2 Channel Establishment\)](#)
 - [Expected Weight of the Commitment Transaction](#)
 - [Expected Weight of HTLC-timeout and HTLC-success Transactions](#)
- [Appendix B: Funding Transaction Test Vectors](#)
- [Appendix C: Commitment and HTLC Transaction Test Vectors](#)
- [Appendix D: Per-commitment Secret Generation Test Vectors](#)
 - [Generation Tests](#)
 - [Storage Tests](#)

- [Appendix E: Key Derivation Test Vectors](#)
- [Appendix F: Commitment and HTLC Transaction Test Vectors \(anchors\)](#)
- [Appendix G: Dual Funded Transaction Test Vectors](#)
- [Appendix H: Commitment and HTLC Transaction Test Vectors \(zero-fee-commitments\)](#)
- [References](#)
- [Authors](#)

Transactions

Transaction Output Ordering

Outputs in transactions are always sorted according to: * first according to their value, smallest first (in whole satoshis, note that for HTLC outputs, the millisatoshi part must be ignored) * followed by scriptpubkey, comparing the common-length prefix lexicographically as if by memcmp, then selecting the shorter script (if they differ in length), * finally, for HTLC outputs, in increasing cltv_expiry order.

Rationale

Two offered HTLCs which have the same amount (rounded from amount_msat) and payment_hash will have identical outputs, even if their cltv_expiry differs. This only matters because the same ordering is used to send htlc_signatures and the HTLC transactions themselves are different, thus the two peers must agree on the canonical ordering for this case.

Use of Segwit

Most transaction outputs used here are pay-to-witness-script-hash [BIP141](#) (P2WSH) outputs: the Segwit version of P2SH. To spend such outputs, the last item on the witness stack must be the actual script that was used to generate the P2WSH output that is being spent. This last item has been omitted for brevity in the rest of this document.

A <> designates an empty vector as required for compliance with MINIMALIF-standard rule. [MINIMALIF](#)

Funding Transaction Output

- The funding output script is a P2WSH to:

```
2 <pubkey1> <pubkey2> 2 OP_CHECKMULTISIG
```

- Where pubkey1 is the lexicographically lesser of the two funding_pubkey in compressed format, and where pubkey2 is the lexicographically greater of the two.

Commitment Transaction

- version: 3 when using zero_fee_commitments, otherwise 2
- locktime: upper 8 bits are 0x20, lower 24 bits are the lower 24 bits of the obscured commitment number
- txin count: 1
 - txin[0] output: txid and output_index from funding_created message

- txin[0] sequence: upper 8 bits are 0x80, lower 24 bits are upper 24 bits of the obscured commitment number
- txin[0] script bytes: 0
- txin[0] witness: 0 <signature_for_pubkey1> <signature_for_pubkey2>

The 48-bit commitment number is obscured by XOR with the lower 48 bits of:

```
SHA256(payment_basepoint from open_channel || payment_basepoint from accept_channel)
```

This obscures the number of commitments made on the channel in the case of unilateral close, yet still provides a useful index for both nodes (who know the payment_basepoints) to quickly find a revoked commitment transaction.

Commitment Transaction Outputs

To allow an opportunity for penalty transactions, in case of a revoked commitment transaction, all outputs that return funds to the owner of the commitment transaction (a.k.a. the “local node”) must be delayed for `to_self_delay` blocks. For HTLCs this delay is done in a second-stage HTLC transaction (HTLC-success for HTLCs accepted by the local node, HTLC-timeout for HTLCs offered by the local node).

The reason for the separate transaction stage for HTLC outputs is so that HTLCs can timeout or be fulfilled even though they are within the `to_self_delay` delay. Otherwise, the required minimum timeout on HTLCs is lengthened by this delay, causing longer timeouts for HTLCs traversing the network.

The amounts for each output MUST be rounded down to whole satoshis. If this amount, minus the fees for the HTLC transaction, is less than the `dust_limit_satoshis` set by the owner of the commitment transaction, the output MUST NOT be produced (thus the funds add to fees). If `zero_fee_commitments` is used, the amount may be partially added to the `shared_anchor` as detailed [here](#).

to_local Output

This output sends funds back to the owner of this commitment transaction and thus must be timelocked using `OP_CHECKSEQUENCEVERIFY`. It can be claimed, without delay, by the other party if they know the revocation private key. The output is a version-0 P2WSH, with a witness script:

```
OP_IF
  # Penalty transaction
  <revocationpubkey>
OP_ELSE
  `to_self_delay`
  OP_CHECKSEQUENCEVERIFY
  OP_DROP
  <local_delayedpubkey>
OP_ENDIF
OP_CHECKSIG
```

The output is spent by an input with `nSequence` field set to `to_self_delay` (which can only be valid after that duration has passed) and witness:

```
<local_delayedsig> <>
```

If a revoked commitment transaction is published, the other party can spend this output immediately with the following witness:

```
<revocation_sig> 1
```

to_remote Output

If `option_anchors` applies to the commitment transaction, the `to_remote` output is encumbered by a one block csv lock.

```
<remotepubkey> OP_CHECKSIGVERIFY 1 OP_CHECKSEQUENCEVERIFY
```

The output is spent by an input with `nSequence` field set to 1 and witness:

```
<remote_sig>
```

Otherwise, this output is a simple P2WPKH to `remotepubkey`.

to_local_anchor and to_remote_anchor Output (option_anchors)

This output can be spent by the local and remote nodes respectively to provide incentive to mine the transaction, using child-pays-for-parent. Both anchor outputs are always added, except for the case where there are no htcs and one of the parties has a commitment output that is below the dust limit. In that case only an anchor is added for the commitment output that does materialize. This typically happens if the initiator closes right after opening (no `to_remote` output).

```
<local_funding_pubkey/remote_funding_pubkey> OP_CHECKSIG OP_IFDUP  
OP_NOTIF  
    OP_16 OP_CHECKSEQUENCEVERIFY  
OP_ENDIF
```

Each party has its own anchor output that locks to their funding key. This is to prevent a malicious peer from attaching child transactions with a low fee density to an anchor and thereby blocking the victim from getting the commit tx confirmed in time. This defense is supported by a change in Bitcoin core 0.19: <https://github.com/bitcoin/bitcoin/pull/15681>. This is also the reason that every non-anchor output on the commit tx is CSV locked.

To prevent utxo set pollution, any anchor that remains unspent can be spent by anyone after the commitment tx confirms. This is also the reason to lock the anchor outputs to the funding key. Third parties can observe this key and reconstruct the spend script, even if none of the commitment outputs would be spent. This does assume that at some point the fee market goes down to a level where sweeping the anchors is economical.

The amount of the output is fixed at 330 sats, the default dust limit for P2WSH.

Spending of the output requires the following witness:

```
<local_sig/remote_sig>
```

After 16 blocks, anyone can sweep the anchor with witness:

```
<>
```

shared_anchor Output (zero_fee_commitments)

This output can be spent by anyone to provide incentive to mine the transaction, using child-pays-for-parent. This output is only added when the commitment transaction is using version 3 (zero_fee_commitments) and prevents pinning attacks on the commitment transaction.

It is using the following standard P2A (pay-to-anchor) script:

```
OP_1 <0x4e73>
```

The standard dust limit for this type of script is 240 sat. However, if the parent transaction doesn't pay any on-chain fees, using values below this dust limit is allowed under the "ephemeral dust" rule.

If the difference between the commitment transaction input and the sum of all other commitment transaction outputs is smaller than 240 sat, we set the anchor amount to this value. This ensures that the commitment transaction doesn't pay any on-chain fees whenever the anchor amount is smaller than 240 sat.

Otherwise, we set the anchor amount to 240 sat, and the remaining difference between the commitment transaction input and its outputs directly contributes to the transaction's mining fees (which is standard since the shared_anchor has reached its dust limit).

Spending of the output requires the following (empty) witness:

```
<>
```

Offered HTLC Outputs

This output sends funds to either an HTLC-timeout transaction after the HTLC-timeout or to the remote node using the payment preimage or the revocation key. The output is a P2WSH, with a witness script (no option_anchors):

```
# To remote node with revocation key
OP_DUP OP_HASH160 <RIPEMD160(SHA256(revocationpubkey))> OP_EQUAL
OP_IF
  OP_CHECKSIG
OP_ELSE
  <remote_htlcpubkey> OP_SWAP OP_SIZE 32 OP_EQUAL
  OP_NOTIF
    # To local node via HTLC-timeout transaction (timelocked).
    OP_DROP 2 OP_SWAP <local_htlcpubkey> 2 OP_CHECKMULTISIG
  OP_ELSE
    # To remote node with preimage.
    OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY
    OP_CHECKSIG
  OP_ENDIF
OP_ENDIF
```

Or, with option_anchors:

```
# To remote node with revocation key
OP_DUP OP_HASH160 <RIPEMD160(SHA256(revocationpubkey))> OP_EQUAL
OP_IF
  OP_CHECKSIG
```

```

OP_ELSE
  <remote_htlcpubkey> OP_SWAP OP_SIZE 32 OP_EQUAL
  OP_NOTIF
    # To local node via HTLC-timeout transaction (timelocked).
    OP_DROP 2 OP_SWAP <local_htlcpubkey> 2 OP_CHECKMULTISIG
  OP_ELSE
    # To remote node with preimage.
    OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY
    OP_CHECKSIG
  OP_ENDIF
  1 OP_CHECKSEQUENCEVERIFY OP_DROP
OP_ENDIF

```

The remote node can redeem the HTLC with the witness:

```
<remotehtlcsig> <payment_preimage>
```

Note that if `option_anchors` applies, the `nSequence` field of the spending input must be 1.

If a revoked commitment transaction is published, the remote node can spend this output immediately with the following witness:

```
<revocation_sig> <revocationpubkey>
```

The sending node can use the HTLC-timeout transaction to timeout the HTLC once the HTLC is expired, as shown below. This is the only way that the local node can timeout the HTLC, and this branch requires `<remotehtlcsig>`, which ensures that the local node cannot prematurely timeout the HTLC since the HTLC-timeout transaction has `cltv_expiry` as its specified locktime. The local node must also wait `to_self_delay` before accessing these funds, allowing for the remote node to claim these funds if the transaction has been revoked.

Received HTLC Outputs

This output sends funds to either the remote node after the HTLC-timeout or using the revocation key, or to an HTLC-success transaction with a successful payment preimage. The output is a P2WSH, with a witness script (no `option_anchors`):

```

# To remote node with revocation key
OP_DUP OP_HASH160 <RIPEMD160(SHA256(revocationpubkey))> OP_EQUAL
OP_IF
  OP_CHECKSIG
OP_ELSE
  <remote_htlcpubkey> OP_SWAP OP_SIZE 32 OP_EQUAL
  OP_IF
    # To local node via HTLC-success transaction.
    OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY
    2 OP_SWAP <local_htlcpubkey> 2 OP_CHECKMULTISIG
  OP_ELSE
    # To remote node after timeout.
    OP_DROP <cltv_expiry> OP_CHECKLOCKTIMEVERIFY OP_DROP
    OP_CHECKSIG
  OP_ENDIF
OP_ENDIF

```

Or, with option_anchors:

```
# To remote node with revocation key
OP_DUP OP_HASH160 <RIPEMD160(SHA256(revocationpubkey))> OP_EQUAL
OP_IF
  OP_CHECKSIG
OP_ELSE
  <remote_htlcpubkey> OP_SWAP OP_SIZE 32 OP_EQUAL
  OP_IF
    # To local node via HTLC-success transaction.
    OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY
    2 OP_SWAP <local_htlcpubkey> 2 OP_CHECKMULTISIG
  OP_ELSE
    # To remote node after timeout.
    OP_DROP <cltv_expiry> OP_CHECKLOCKTIMEVERIFY OP_DROP
    OP_CHECKSIG
  OP_ENDIF
  1 OP_CHECKSEQUENCEVERIFY OP_DROP
OP_ENDIF
```

To timeout the HTLC, the remote node spends it with the witness:

```
<remotehtlcsig> <>
```

Note that if option_anchors applies, the nSequence field of the spending input must be 1.

If a revoked commitment transaction is published, the remote node can spend this output immediately with the following witness:

```
<revocation_sig> <revocationpubkey>
```

To redeem the HTLC, the HTLC-success transaction is used as detailed below. This is the only way that the local node can spend the HTLC, since this branch requires <remotehtlcsig>, which ensures that the local node must wait to_self_delay before accessing these funds allowing for the remote node to claim these funds if the transaction has been revoked.

Trimmed Outputs

Each peer specifies a dust_limit_satoshis below which outputs should not be produced; these outputs that are not produced are termed “trimmed”. A trimmed output is considered too small to be worth creating; it is instead either added to the commitment transaction fee or, if zero_fee_commitments applies, to the [shared anchor output](#).

For HTLCs, it needs to be taken into account that the second-stage HTLC transaction may also be below the limit. Note that when using option_anchors or zero_fee_commitments, HTLC transactions don’t include a fee and thus don’t contribute to trimming: setting a higher dust_limit_satoshis makes sense for those channels to ensure that outputs are economical to spend.

Requirements

Before the commitment transaction outputs are determined: - the base fee MUST be subtracted from the to_local or to_remote output, as specified in [Fee Calculation](#). - if option_anchors applies: - the

to_local_anchor output value MUST be subtracted from the funder's output. - the to_remote_anchor output value MUST be subtracted from the funder's output.

The commitment transaction: - if the amount of the commitment transaction to_local output would be less than dust_limit_satoshis set by the transaction owner: - MUST NOT contain that output. - otherwise: - MUST be generated as specified in [to_local_Output](#). - if the amount of the commitment transaction to_remote output would be less than dust_limit_satoshis set by the transaction owner: - MUST NOT contain that output. - otherwise: - MUST be generated as specified in [to_remote_Output](#). - for every offered HTLC: - if the HTLC amount minus the HTLC-timeout fee would be less than dust_limit_satoshis set by the transaction owner: - MUST NOT contain that output. - otherwise: - MUST be generated as specified in [Offered HTLC Outputs](#). - for every received HTLC: - if the HTLC amount minus the HTLC-success fee would be less than dust_limit_satoshis set by the transaction owner: - MUST NOT contain that output. - otherwise: - MUST be generated as specified in [Received HTLC Outputs](#). - if zero_fee_commitments applies: - MUST add a shared_anchor output with an amount computed as detailed in [the shared_anchor section](#).

HTLC-Timeout and HTLC-Success Transactions

These HTLC transactions are almost identical, except the HTLC-timeout transaction is timelocked. Both HTLC-timeout/HTLC-success transactions can be spent by a valid penalty transaction.

- version: 3 when using zero_fee_commitments, otherwise 2
- locktime: 0 for HTLC-success, cltv_expiry for HTLC-timeout
- txin count: 1
 - txin[0] outputpoint: txid of the commitment transaction and output_index of the matching HTLC output for the HTLC transaction
 - txin[0] sequence: 1 for option_anchors, 0 otherwise
 - txin[0] script bytes: 0
 - txin[0] witness stack: 0 <remotehtlcsig> <localhtlcsig> <payment_preimage> for HTLC-success, 0 <remotehtlcsig> <localhtlcsig> <> for HTLC-timeout
- txout count: 1
 - txout[0] amount: the HTLC amount_msat divided by 1000 (rounding down) minus fees in satoshis (see [Fee Calculation](#))
 - txout[0] script: version-0 P2WSH with witness script as shown below
- if option_anchors or zero_fee_commitments applies to this commitment transaction, SIGHASH_SINGLE | SIGHASH_ANYONECANPAY is used as described in [BOLT #5](#).

The witness script for the output is:

```
OP_IF
  # Penalty transaction
  <revocationpubkey>
OP_ELSE
  `to_self_delay`
  OP_CHECKSEQUENCEVERIFY
  OP_DROP
  <local_delayedpubkey>
OP_ENDIF
OP_CHECKSIG
```

To spend this via penalty, the remote node uses a witness stack `<revocationsig> 1`, and to collect the output, the local node uses an input with `nSequence to_self_delay` and a witness stack `<local_delayedsig> 0`.

Legacy Closing Transaction

This variant is used for `closing_signed` messages (i.e. where `option_simple_close` is not negotiated).

Note that there are two possible variants for each node.

- version: 2
- locktime: 0
- txin count: 1
 - `txin[0]` output: `txid` and `output_index` from `funding_created` message
 - `txin[0]` sequence: `0xFFFFFFFF`
 - `txin[0]` script bytes: 0
 - `txin[0]` witness: 0 `<signature_for_pubkey1> <signature_for_pubkey2>`
- txout count: 1 or 2
 - txout amount: final balance to be paid to one node (minus `fee_satoshis` from `closing_signed`, if this peer funded the channel)
 - txout script: as specified in that node's `scriptpubkey` in its shutdown message

Requirements

Each node offering a signature: - MUST round each output down to whole satoshis. - MUST subtract the fee given by `fee_satoshis` from the output to the funder. - MUST remove any output below its own `dust_limit_satoshis`. - MAY eliminate its own output.

Rationale

There is a possibility of irreparable differences on closing if one node considers the other's output too small to allow propagation on the Bitcoin network (a.k.a. "dust"), and that other node instead considers that output too valuable to discard. This is why each side uses its own `dust_limit_satoshis`, and the result can be a signature validation failure, if they disagree on what the closing transaction should look like.

However, if one side chooses to eliminate its own output, there's no reason for the other side to fail the closing protocol; so this is explicitly allowed. The signature indicates which variant has been used.

There will be at least one output, if the funding amount is greater than twice `dust_limit_satoshis`.

Closing Transaction

This variant is used for `closing_complete` and `closing_sig` messages (i.e. where `option_simple_close` is negotiated).

In this case, the node sending `closing_complete` ("the closer") pays the fees. The outputs are ordered as detailed in [Transaction Output Ordering](#).

The side with lesser funds can opt to omit their own output.

- version: 2

- locktime: locktime from the closing_complete message
- txin count: 1
 - txin[0] output: txid and output_index of the channel output
 - txin[0] sequence: 0xFFFFFFFF
 - txin[0] script bytes: 0
 - txin[0] witness: 0 <signature_for_pubkey1> <signature_for_pubkey2>
- txout count: 1 or 2
 - The closer output:
 - txout amount:
 - 0 if the scriptpubkey starts with OP_RETURN
 - otherwise the final balance for the closer, minus closing_complete.fee_satoshis, rounded down to whole satoshis
 - txout script: as specified in closer_scriptpubkey from the closing_complete message
 - The closer output:
 - txout amount:
 - 0 if the scriptpubkey starts with OP_RETURN
 - otherwise the final balance for the closer, rounded down to whole satoshis
 - txout script: as specified in closer_scriptpubkey from the closing_complete message

Requirements

Each node offering a signature: - MUST round each output down to whole satoshis. - MUST subtract the fee given by fee_satoshis from the closer output. - MUST set the output amount to 0 if the scriptpubkey is OP_RETURN.

Fees

Fee Calculation

The fee calculation for both commitment transactions and HTLC transactions is based on the current feerate_per_kw and the *expected weight* of the transaction. Note that feerate_per_kw is set to 0 when using zero_fee_commitments: in that case the commitment transactions and HTLC transactions don't pay any fee, which must be paid using child-pays-for-parent (for the commitment transactions) or by adding inputs (for the HTLC transactions).

The actual and expected weights vary for several reasons:

- Bitcoin uses DER-encoded signatures, which vary in size.
- Bitcoin also uses variable-length integers, so a large number of outputs will take 3 bytes to encode rather than 1.
- The to_remote output may be below the dust limit.
- The to_local output may be below the dust limit once fees are extracted.

Thus, a simplified formula for *expected weight* is used, which assumes:

- Signatures are 73 bytes long (the maximum length).
- There are a small number of outputs (thus 1 byte to count them).
- There are always both a to_local output and a to_remote output.

- (if option_anchors) there are always both a to_local_anchor and to_remote_anchor output.

This yields the following *expected weights* (details of the computation in [Appendix A](#)):

```
Commitment weight (no option_anchors): 724 + 172 * num-untrimmed-htlc-outputs
Commitment weight (option_anchors): 1124 + 172 * num-untrimmed-htlc-outputs
HTLC-timeout weight (no option_anchors): 663
HTLC-timeout weight (option_anchors): 666
HTLC-success weight (no option_anchors): 703
HTLC-success weight (option_anchors): 706
```

Note the reference to the “base fee” for a commitment transaction in the requirements below, which is what the funder pays. The actual fee may be higher than the amount calculated here, due to rounding and trimmed outputs.

Requirements

The fee for an HTLC-timeout transaction: - If option_anchors or zero_fee_commitments applies: 1. MUST be 0. - Otherwise, MUST be calculated to match: 1. Multiply feerate_per_kw by 663 and divide by 1000 (rounding down).

The fee for an HTLC-success transaction: - If option_anchors or zero_fee_commitments applies: 1. MUST be 0. - Otherwise, MUST be calculated to match: 1. Multiply feerate_per_kw by 703 and divide by 1000 (rounding down).

The base fee for a commitment transaction: - If zero_fee_commitments applies: 1. MUST be 0. - Otherwise, MUST be calculated to match: 1. Start with weight = 724 (1124 if option_anchors applies). 2. For each committed HTLC, if that output is not trimmed as specified in [Trimmed Outputs](#), add 172 to weight. 3. Multiply feerate_per_kw by weight, divide by 1000 (rounding down).

Example

For example, suppose there is a feerate_per_kw of 5000, a dust_limit_satoshis of 546 satoshis, and a commitment transaction with: * two offered HTLCs of 5000000 and 1000000 millisatoshis (5000 and 1000 satoshis) * two received HTLCs of 7000000 and 800000 millisatoshis (7000 and 800 satoshis)

The HTLC-timeout transaction weight is 663, and thus the fee is 3315 satoshis. The HTLC-success transaction weight is 703, and thus the fee is 3515 satoshis

The commitment transaction weight is calculated as follows:

- weight starts at 724.
- The offered HTLC of 5000 satoshis is above 546 + 3315 and results in:
 - an output of 5000 satoshi in the commitment transaction
 - an HTLC-timeout transaction of 5000 - 3315 satoshis that spends this output
 - weight increases to 896
- The offered HTLC of 1000 satoshis is below 546 + 3315 so it is trimmed.

- The received HTLC of 7000 satoshis is above $546 + 3515$ and results in:
 - an output of 7000 satoshi in the commitment transaction
 - an HTLC-success transaction of $7000 - 3515$ satoshis that spends this output
 - weight increases to 1068
- The received HTLC of 800 satoshis is below $546 + 3515$ so it is trimmed.

The base commitment transaction fee is 5340 satoshi; the actual fee (which adds the 1000 and 800 satoshi HTLCs that would make dust outputs) is 7140 satoshi. The final fee may be even higher if the `to_local` or `to_remote` outputs fall below `dust_limit_satoshis`.

Fee Payment

Base commitment transaction fees and amounts for `to_local_anchor` and `to_remote_anchor` outputs are extracted from the funder's amount.

When using `zero_fee_commitments`, the `shared_anchor` output amount is taken from [trimmed outputs](#), as detailed in the [shared_anchor section](#).

Note that after the fee amount is subtracted from the to-funder output, that output may be below `dust_limit_satoshis`, and thus will also contribute to fees (or, if `zero_fee_commitments` applies, to the `shared_anchor`).

A node: - if `zero_fee_commitments` is not used and the resulting fee rate is too low: - MAY send a warning and close the connection, or send an error and fail the channel.

Calculating Fees for Collaborative Transactions

For transactions constructed using the [interactive protocol](#), fees are paid by each party to the transaction, at a feerate determined during the initiation, with the initiator covering the fees for the common transaction fields.

Dust Limits

The `dust_limit_satoshis` parameter is used to configure the threshold below which nodes will not produce on-chain transaction outputs.

There is no consensus rule in Bitcoin that makes outputs below dust thresholds invalid or unspendable, but policy rules in popular implementations will prevent relaying transactions that contain such outputs.

Bitcoin Core defines the following dust thresholds:

- pay to pubkey hash (p2pkh): 546 satoshis
- pay to script hash (p2sh): 540 satoshis
- pay to witness pubkey hash (p2wpkh): 294 satoshis
- pay to witness script hash (p2wsh): 330 satoshis
- pay to anchor (p2a): 240 satoshis
- unknown segwit versions: 354 satoshis
- OP_RETURN outputs: these are never dust

The rationale of this calculation (implemented [here](#)) is explained in the following sections.

In all these sections, the calculations are done with a feerate of 3000 sat/kB as per Bitcoin Core's implementation.

Note that since the introduction of `option_anchors`, the second-stage HTLC transaction's weight is not taken into account when deciding whether outputs should be included in the commitment transaction or not. It thus makes sense to use a `dust_limit_satoshis` that takes into account the cost of those second-stage HTLC transactions, to ensure that outputs added to the commitment transaction can actually be claimed on-chain, otherwise they may pollute the utxo set indefinitely. At a minimum, nodes should allow their peer to use a `dust_limit_satoshis` that is higher than the values defined by Bitcoin Core. We cannot predict future feerates, so this will not always work and can still result in HTLC outputs that are unspendable if the on-chain fees are too high. We cannot use very large `dust_limit_satoshis` values either since it would create too much dust exposure in the commitment transaction (more details [here](#)).

Pay to pubkey hash (p2pkh)

A p2pkh output is 34 bytes:

- 8 bytes for the output amount
- 1 byte for the script length
- 25 bytes for the script (`OP_DUP OP_HASH160 20 20-bytes OP_EQUALVERIFY OP_CHECKSIG`)

A p2pkh input is at least 148 bytes:

- 36 bytes for the previous output (32 bytes hash + 4 bytes index)
- 4 bytes for the sequence
- 1 byte for the script sig length
- 107 bytes for the script sig:
 - 1 byte for the items count
 - 1 byte for the signature length
 - 71 bytes for the signature
 - 1 byte for the public key length
 - 33 bytes for the public key

The p2pkh dust threshold is then $(34 + 148) * 3000 / 1000 = 546$ satoshis

Pay to script hash (p2sh)

A p2sh output is 32 bytes:

- 8 bytes for the output amount
- 1 byte for the script length
- 23 bytes for the script (`OP_HASH160 20 20-bytes OP_EQUAL`)

A p2sh input doesn't have a fixed size, since it depends on the underlying script, so we use 148 bytes as a lower bound.

The p2sh dust threshold is then $(32 + 148) * 3000 / 1000 = 540$ satoshis

Pay to witness pubkey hash (p2wpkh)

A p2wpkh output is 31 bytes:

- 8 bytes for the output amount
- 1 byte for the script length
- 22 bytes for the script (0P_0 20 20-bytes)

A p2wpkh input is at least 67 bytes (depending on the signature length):

- 36 bytes for the previous output (32 bytes hash + 4 bytes index)
- 4 bytes for the sequence
- 1 byte for the script sig length
- 26 bytes for the witness (rounded down from 26.75, with the 75% segwit discount applied):
 - 1 byte for the items count
 - 1 byte for the signature length
 - 71 bytes for the signature
 - 1 byte for the public key length
 - 33 bytes for the public key

The p2wpkh dust threshold is then $(31 + 67) * 3000 / 1000 = 294$ satoshis

Pay to witness script hash (p2wsh)

A p2wsh output is 43 bytes:

- 8 bytes for the output amount
- 1 byte for the script length
- 34 bytes for the script (0P_0 32 32-bytes)

A p2wsh input doesn't have a fixed size, since it depends on the underlying script, so we use 67 bytes as a lower bound.

The p2wsh dust threshold is then $(43 + 67) * 3000 / 1000 = 330$ satoshis

Unknown segwit versions

Unknown segwit outputs are at most 51 bytes:

- 8 bytes for the output amount
- 1 byte for the script length
- 42 bytes for the script (0P_1 through 0P_16 inclusive, followed by a single push of 2 to 40 bytes)

The input doesn't have a fixed size, since it depends on the underlying script, so we use 67 bytes as a lower bound.

The unknown segwit version dust threshold is then $(51 + 67) * 3000 / 1000 = 354$ satoshis

Commitment Transaction Construction

This section ties the previous sections together to detail the algorithm for constructing the commitment transaction for one peer: given that peer's `dust_limit_satoshis`, the current `feerate_per_kw`, the amounts due to each peer (`to_local` and `to_remote`), and all committed HTLCs:

- Initialize the commitment transaction input and locktime, as specified in [Commitment Transaction](#).
- Calculate which committed HTLCs need to be trimmed (see [Trimmed Outputs](#)).
- Calculate the base [commitment transaction fee](#).
- Subtract this base fee from the funder (either `to_local` or `to_remote`). If `option_anchors` applies to the commitment transaction, also subtract two times the fixed anchor size of 330 sats from the funder (either `to_local` or `to_remote`).
- For every offered HTLC, if it is not trimmed, add an [offered HTLC output](#).
- For every received HTLC, if it is not trimmed, add a [received HTLC output](#).
- If the `to_local` amount is greater or equal to `dust_limit_satoshis`, add a [to_local output](#).
- If the `to_remote` amount is greater or equal to `dust_limit_satoshis`, add a [to_remote output](#).
- If `option_anchors` applies to the commitment transaction:
 - if `to_local` exists or there are untrimmed HTLCs, add a [to_local_anchor output](#).
 - if `to_remote` exists or there are untrimmed HTLCs, add a [to_remote_anchor output](#).
- If `zero_fee_commitments` applies to the commitment transaction:
 - add a [shared_anchor output](#).
- Sort the outputs into [BIP 69+CLTV order](#).

Keys

Key Derivation

Each commitment transaction uses a unique `localpubkey`. The HTLC-success and HTLC-timeout transactions use `local_delayedpubkey` and `revocationpubkey`. These are changed for every transaction based on the `per_commitment_point`.

The reason for key change is so that trustless watching for revoked transactions can be outsourced. Such a *watcher* should not be able to determine the contents of a commitment transaction — even if the *watcher* knows which transaction ID to watch for and can make a reasonable guess as to which HTLCs and balances may be included. Nonetheless, to avoid storage of every commitment transaction, a *watcher* can be given the `per_commitment_secret` values (which can be stored compactly) and the `revocation_basepoint` and `delayed_payment_basepoint` used to regenerate the scripts required for the penalty transaction; thus, a *watcher* need only be given (and store) the signatures for each penalty input.

Changing the `localpubkey` every time ensures that commitment transaction ID cannot be guessed except in the trivial case where there is no `to_local` output, as every commitment transaction uses an ID in its output script. Splitting the `local_delayedpubkey`, which is required for the penalty transaction, allows it to be shared with the *watcher* without revealing `localpubkey`; even if both peers use the same *watcher*, nothing is revealed.

Finally, even in the case of normal unilateral close, the HTLC-success and/or HTLC-timeout transactions do not reveal anything to the *watcher*, as it does not know the corresponding `per_commitment_secret` and cannot relate the `local_delayedpubkey` or `revocationpubkey` with their bases.

For efficiency, keys are generated from a series of per-commitment secrets that are generated from a single seed, which allows the receiver to compactly store them (see [below](#)).

localpubkey, local_htlcpubkey, remote_htlcpubkey, local_delayedpubkey, and remote_delayedpubkey Derivation

These pubkeys are simply generated by addition from their base points:

$$\text{pubkey} = \text{basepoint} + \text{SHA256}(\text{per_commitment_point} \parallel \text{basepoint}) * G$$

- The localpubkey uses the local node's payment_basepoint;
- The local_htlcpubkey uses the local node's htlc_basepoint;
- The remote_htlcpubkey uses the remote node's htlc_basepoint;
- The local_delayedpubkey uses the local node's delayed_payment_basepoint;
- The remote_delayedpubkey uses the remote node's delayed_payment_basepoint.

The corresponding private keys can be similarly derived, if the basepoint secrets are known (i.e. the private keys corresponding to localpubkey, local_htlcpubkey, and local_delayedpubkey only):

$$\text{privkey} = \text{basepoint_secret} + \text{SHA256}(\text{per_commitment_point} \parallel \text{basepoint})$$

remotepubkey Derivation

The remotepubkey is simply the remote node's payment_basepoint.

This means that a node can spend a commitment transaction even if it has lost data and doesn't know the corresponding per_commitment_point. A watchtower could correlate transactions given to it which only have a to_remote output if it sees one of them onchain, but such transactions do not need any enforcement and should not be handed to a watchtower.

revocationpubkey Derivation

The revocationpubkey is a blinded key: when the local node wishes to create a new commitment for the remote node, it uses its own revocation_basepoint and the remote node's per_commitment_point to derive a new revocationpubkey for the commitment. After the remote node reveals the per_commitment_secret used (thereby revoking that commitment), the local node can then derive the revocationprivkey, as it now knows the two secrets necessary to derive the key (revocation_basepoint_secret and per_commitment_secret).

The per_commitment_point is generated using elliptic-curve multiplication:

$$\text{per_commitment_point} = \text{per_commitment_secret} * G$$

And this is used to derive the revocation pubkey from the remote node's revocation_basepoint:

$$\text{revocationpubkey} = \text{revocation_basepoint} * \text{SHA256}(\text{revocation_basepoint} \parallel \text{per_commitment_point}) + \text{per_commitment_point} * \text{SHA256}(\text{per_commitment_point} \parallel \text{revocation_basepoint})$$

This construction ensures that neither the node providing the basepoint nor the node providing the per_commitment_point can know the private key without the other node's secret.

The corresponding private key can be derived once the per_commitment_secret is known:

```
revocationprivkey = revocation_basepoint_secret * SHA256(revocation_basepoint ||
per_commitment_point) + per_commitment_secret * SHA256(per_commitment_point ||
revocation_basepoint)
```

Per-commitment Secret Requirements

A node: - MUST select an unguessable 256-bit seed for each connection, - MUST NOT reveal the seed.

Up to $(2^{48} - 1)$ per-commitment secrets can be generated.

The first secret used: - MUST be index 281474976710655, - and from there, the index is decremented.

The I'th secret P: - MUST match the output of this algorithm:

```
generate_from_seed(seed, I):
    P = seed
    for B in 47 down to 0:
        if B set in I:
            flip(B) in P
        P = SHA256(P)
    return P
```

Where “flip(B)” alternates the $(B \bmod 8)$ bit of the $(B \div 8)$ byte of the value. So, “flip(0) in e3b0...” is “e2b0...”, and “flip(10) in e3b0...” is “e3b4...”.

The receiving node: - MAY store all previous per-commitment secrets. - MAY calculate them from a compact representation, as described below.

Efficient Per-commitment Secret Storage

The receiver of a series of secrets can store them compactly in an array of 49 (value,index) pairs. Because, for a given secret on a 2^X boundary, all secrets up to the next 2^X boundary can be derived; and secrets are always received in descending order starting at 0xFFFFFFFFFFFF.

In binary, it's helpful to think of any index in terms of a *prefix*, followed by some trailing 0s. You can derive the secret for any index that matches this *prefix*.

For example, secret 0xFFFFFFFFFFFF0 allows the secrets to be derived for 0xFFFFFFFFFFFF1 through 0xFFFFFFFFFFFFF, inclusive; and secret 0xFFFFFFFFFFFF8 allows the secrets to be derived for 0xFFFFFFFFFFFF9 through 0xFFFFFFFFFFFFF, inclusive.

This is done using a slight generalization of generate_from_seed above:

```
# Return I'th secret given base secret whose index has bits..47 the same.
derive_secret(base, bits, I):
    P = base
    for B in bits - 1 down to 0:
        if B set in I:
            flip(B) in P
        P = SHA256(P)
    return P
```

Only one secret for each unique prefix need be saved; in effect, the number of trailing 0s is counted, and this determines where in the storage array the secret is stored:

```
# a.k.a. count trailing 0s
where_to_put_secret(I):
    for B in 0 to 47:
        if testbit(I) in B == 1:
            return B
    # I = 0, this is the seed.
    return 48
```

A double-check, that all previous secrets derive correctly, is needed; if this check fails, the secrets were not generated from the same seed:

```
insert_secret(secret, I):
    B = where_to_put_secret(I)

    # This tracks the index of the secret in each bucket across the traversal.
    for b in 0 to B:
        if derive_secret(secret, B, known[b].index) != known[b].secret:
            error The secret for I is incorrect
            return

    # Assuming this automatically extends known[] as required.
    known[B].index = I
    known[B].secret = secret
```

Finally, if an unknown secret at index I needs be derived, it must be discovered which known secret can be used to derive it. The simplest method is iterating over all the known secrets, and testing if each can be used to derive the unknown secret:

```
derive_old_secret(I):
    for b in 0 to len(secrets):
        # Mask off the non-zero prefix of the index.
        MASK = ~((1 << b) - 1)
        if (I & MASK) == secrets[b].index:
            return derive_secret(known, i, I)
    error Index 'I' hasn't been received yet.
```

This looks complicated, but remember that the index in entry b has b trailing 0s; the mask and compare simply checks if the index at each bucket is a prefix of the desired index.

Appendix A: Expected Weights

Expected Weight of the Funding Transaction (v2 Channel Establishment)

The *expected weight* of a funding transaction is calculated as follows:

```
inputs: 41 bytes
- previous_out_point: 36 bytes
- hash: 32 bytes
- index: 4 bytes
- var_int: 1 byte
- script_sig: 0 bytes
```


- witness <---- Cost for "witness" data calculated separately.
- sequence: 4 bytes

non_funding_outputs: 9 bytes + `scriptlen`

- value: 8 bytes
- var_int: 1 byte <---- assuming a standard output script
- script: `scriptlen`

funding_output: 43 bytes

- value: 8 bytes
- var_int: 1 byte
- script: 34 bytes
 - OP_0: 1 byte
 - PUSHDATA(32-byte-hash): 33 bytes

Multiplying non-witness data by 4 results in a weight of:

```
// transaction_fields = 10 (version, input count, output count, locktime)
// segwit_fields = 2 (marker + flag)
// funding_transaction = 43 + num_inputs * 41 + num_outputs * 9 + sum(scriptlen)
funding_transaction_weight = 4 * (funding_transaction + transaction_fields) + segwit_fields

witness_weight = sum(witness_len)

overall_weight = funding_transaction_weight + witness_weight
```

Expected Weight of the Commitment Transaction

The *expected weight* of a commitment transaction is calculated as follows:

p2wsh: 34 bytes

- OP_0: 1 byte
- OP_DATA: 1 byte (witness_script_SHA256 length)
- witness_script_SHA256: 32 bytes

p2wpkh: 22 bytes

- OP_0: 1 byte
- OP_DATA: 1 byte (public_key_HASH160 length)
- public_key_HASH160: 20 bytes

multi_sig: 71 bytes

- OP_2: 1 byte
- OP_DATA: 1 byte (pub_key_alice length)
- pub_key_alice: 33 bytes
- OP_DATA: 1 byte (pub_key_bob length)
- pub_key_bob: 33 bytes
- OP_2: 1 byte
- OP_CHECKMULTISIG: 1 byte

witness: 222 bytes

- number_of_witness_elements: 1 byte
- nil_length: 1 byte
- sig_alice_length: 1 byte
- sig_alice: 73 bytes
- sig_bob_length: 1 byte

- sig_bob: 73 bytes
- witness_script_length: 1 byte
- witness_script (multi_sig)

funding_input: 41 bytes

- previous_out_point: 36 bytes
 - hash: 32 bytes
 - index: 4 bytes
- var_int: 1 byte (script_sig length)
- script_sig: 0 bytes
- witness <---- "witness" is used instead of "script_sig" for transaction validation; however, "witness" is stored separately, and the cost for its size is smaller. So, the calculation of ordinary data is separated from the witness data.
- sequence: 4 bytes

output_paying_to_local: 43 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wsh): 34 bytes

output_paying_to_remote (no option_anchors): 31 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wpkh): 22 bytes

output_paying_to_remote (option_anchors): 43 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wsh): 34 bytes

output_anchor (option_anchors): 43 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wsh): 34 bytes

htlc_output: 43 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wsh): 34 bytes

witness_header: 2 bytes

- flag: 1 byte
- marker: 1 byte

commitment_transaction (no option_anchors): $125 + 43 * \text{num-htlc-outputs}$ bytes

- version: 4 bytes
- witness_header <---- part of the witness data
- count_tx_in: 1 byte
- tx_in: 41 bytes
 - funding_input
- count_tx_out: 1 byte
- tx_out: $74 + 43 * \text{num-htlc-outputs}$ bytes
 - output_paying_to_remote,

- output_paying_to_local,
....htlc_output's...
- lock_time: 4 bytes

commitment_transaction (option_anchors): 225 + 43 * num-htlc-outputs bytes

- version: 4 bytes
- witness_header <---- part of the witness data
- count_tx_in: 1 byte
- tx_in: 41 bytes
 - funding_input
- count_tx_out: 3 byte
- tx_out: 172 + 43 * num-htlc-outputs bytes
 - output_paying_to_remote,
 - output_paying_to_local,
 - output_anchor,
 - output_anchor,
 -htlc_output's...
- lock_time: 4 bytes

Multiplying non-witness data by 4 results in a weight of:

```
// 500 + 172 * num-htlc-outputs weight (no option_anchors)
// 900 + 172 * num-htlc-outputs weight (option_anchors)
commitment_transaction_weight = 4 * commitment_transaction
```

```
// 224 weight
witness_weight = witness_header + witness
```

```
overall_weight (no option_anchors) = 500 + 172 * num-htlc-outputs + 224 weight
overall_weight (option_anchors) = 900 + 172 * num-htlc-outputs + 224 weight
```

Expected Weight of HTLC-timeout and HTLC-success Transactions

The *expected weight* of an HTLC transaction is calculated as follows:

accepted_htlc_script: 140 bytes (143 bytes with option_anchors)

- OP_DUP: 1 byte
- OP_HASH160: 1 byte
- OP_DATA: 1 byte (RIPEMD160(SHA256(revocationpubkey)) length)
- RIPEMD160(SHA256(revocationpubkey)): 20 bytes
- OP_EQUAL: 1 byte
- OP_IF: 1 byte
- OP_CHECKSIG: 1 byte
- OP_ELSE: 1 byte
- OP_DATA: 1 byte (remote_htlcpubkey length)
- remote_htlcpubkey: 33 bytes
- OP_SWAP: 1 byte
- OP_SIZE: 1 byte
- OP_DATA: 1 byte (32 length)
- 32: 1 byte
- OP_EQUAL: 1 byte
- OP_IF: 1 byte
- OP_HASH160: 1 byte
- OP_DATA: 1 byte (RIPEMD160(payment_hash) length)
- RIPEMD160(payment_hash): 20 bytes

- OP_EQUALVERIFY: 1 byte
- 2: 1 byte
- OP_SWAP: 1 byte
- OP_DATA: 1 byte (local_htlcpubkey length)
- local_htlcpubkey: 33 bytes
- 2: 1 byte
- OP_CHECKMULTISIG: 1 byte
- OP_ELSE: 1 byte
- OP_DROP: 1 byte
- OP_DATA: 1 byte (cltv_expiry length)
- cltv_expiry: 4 bytes
- OP_CHECKLOCKTIMEVERIFY: 1 byte
- OP_DROP: 1 byte
- OP_CHECKSIG: 1 byte
- OP_ENDIF: 1 byte
- OP_1: 1 byte (option_anchors)
- OP_CHECKSEQUENCEVERIFY: 1 byte (option_anchors)
- OP_DROP: 1 byte (option_anchors)
- OP_ENDIF: 1 byte

offered_htlc_script: 133 bytes (136 bytes with option_anchors)

- OP_DUP: 1 byte
- OP_HASH160: 1 byte
- OP_DATA: 1 byte (RIPEMD160(SHA256(revocationpubkey)) length)
- RIPEMD160(SHA256(revocationpubkey)): 20 bytes
- OP_EQUAL: 1 byte
- OP_IF: 1 byte
- OP_CHECKSIG: 1 byte
- OP_ELSE: 1 byte
- OP_DATA: 1 byte (remote_htlcpubkey length)
- remote_htlcpubkey: 33 bytes
- OP_SWAP: 1 byte
- OP_SIZE: 1 byte
- OP_DATA: 1 byte (32 length)
- 32: 1 byte
- OP_EQUAL: 1 byte
- OP_NOTIF: 1 byte
- OP_DROP: 1 byte
- 2: 1 byte
- OP_SWAP: 1 byte
- OP_DATA: 1 byte (local_htlcpubkey length)
- local_htlcpubkey: 33 bytes
- 2: 1 byte
- OP_CHECKMULTISIG: 1 byte
- OP_ELSE: 1 byte
- OP_HASH160: 1 byte
- OP_DATA: 1 byte (RIPEMD160(payment_hash) length)
- RIPEMD160(payment_hash): 20 bytes
- OP_EQUALVERIFY: 1 byte
- OP_CHECKSIG: 1 byte
- OP_ENDIF: 1 byte
- OP_1: 1 byte (option_anchors)
- OP_CHECKSEQUENCEVERIFY: 1 byte (option_anchors)
- OP_DROP: 1 byte (option_anchors)
- OP_ENDIF: 1 byte

timeout_witness: 285 bytes (288 bytes with option_anchors)

- number_of_witness_elements: 1 byte
- nil_length: 1 byte
- sig_alice_length: 1 byte
- sig_alice: 73 bytes
- sig_bob_length: 1 byte
- sig_bob: 73 bytes
- nil_length: 1 byte
- witness_script_length: 1 byte
- witness_script (offered_htlc_script)

success_witness: 324 bytes (327 bytes with option_anchors)

- number_of_witness_elements: 1 byte
- nil_length: 1 byte
- sig_alice_length: 1 byte
- sig_alice: 73 bytes
- sig_bob_length: 1 byte
- sig_bob: 73 bytes
- preimage_length: 1 byte
- preimage: 32 bytes
- witness_script_length: 1 byte
- witness_script (accepted_htlc_script)

commitment_input: 41 bytes

- previous_out_point: 36 bytes
 - hash: 32 bytes
 - index: 4 bytes
- var_int: 1 byte (script_sig length)
- script_sig: 0 bytes
- witness (success_witness or timeout_witness)
- sequence: 4 bytes

htlc_output: 43 bytes

- value: 8 bytes
- var_int: 1 byte (pk_script length)
- pk_script (p2wsh): 34 bytes

htlc_transaction:

- version: 4 bytes
- witness_header <---- part of the witness data
- count_tx_in: 1 byte
- tx_in: 41 bytes
 - commitment_input
- count_tx_out: 1 byte
- tx_out: 43
 - htlc_output
- lock_time: 4 bytes

Multiplying non-witness data by 4 results in a weight of 376. Adding the witness data for each case (285 or 288 + 2 for HTLC-timeout, 324 or 327 + 2 for HTLC-success) results in weights of:

663 (HTLC-timeout) (666 with option_anchors)

703 (HTLC-success) (706 with option_anchors)

- (really 702 and 705, but we use these numbers for historical reasons)

In the following:

- It's assumed that *local* is the funder.
- Private keys are displayed as 32 bytes plus a trailing 1 (Bitcoin's convention for "compressed" private keys, i.e. keys for which the public key is compressed).
- Transaction signatures are all deterministic, using RFC6979 (using HMAC-SHA256). A valid signature **MUST** sign all inputs and outputs of the relevant transaction (i.e. **MUST** be created with a `SIGHASH_ALL` [signature hash](#)), unless explicitly stated otherwise. Note that clients **MUST** send the signature in compact encoding and not in Bitcoin-script format, thus the signature hash byte is not transmitted.

[illegible]

```
local_funding_pubkey: 023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb
remote_funding_pubkey: 030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c39fa711c1
# funding witness script =
5221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d
66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae
```

```
input txid: fd2105607605d2302994ffea703b09f66b6351816ee737a93e42a841ea20bbad
input index: 0
input satoshis: 5000000000
funding satoshis: 10000000
# funding witness script =
5221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d
66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae
# feerate_per_kw: 15000
change satoshis: 4989986080
funding output: 0
```

The resulting funding transaction is:

```
0200000001adbb20ea41a8423ea937e76e8151636bf6093b70eaff942930d20576600521fd000000006b483045022100
90587b6201e166ad6af0227d3036a9454223d49a1f11839c1a362184340ef0240220577f7cd5cca78719405cbf1de741
4ac027f0239ef6e214c90fcaab0454d84b3b012103535b32d5eb0a6ed0982a0479bbadc9868d9836f6ba94dd5a63be16
d875069184fffffffff028096980000000000220020c015c4a6be010e21657068fc2e6a9d02b27ebe4d490a25846f7237
f104d1a3cd20256d29010000001600143ca33c2e4446f4a305f23c80df8ad1afdcf652f900000000
# txid: 8984484a580b825b9972d7adb15050b3ab624ccd731946b3eeddb92f4e7ef6be
```

In the following: - *local* transactions are considered, which implies that all payments to *local* are delayed. - It's assumed that *local* is the opener. - Private keys are displayed as 32 bytes plus a trailing 1 (Bitcoin's convention for "compressed" private keys, i.e. keys for which the public key is compressed). - Transaction signatures are all deterministic, using RFC6979 (using HMAC-SHA256).

[illegible]

Here are the points used to derive the obscuring factor for the commitment number:

```
local_payment_basepoint: 034f355bdc7cc0af728ef3cceb9615d90684bb5b2ca5f859ab0f0b704075871aa
remote_payment_basepoint: 032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991
# obscured commitment number = 0x2bb038521914 ^ 42
```

We use the same values for the HTLC basepoints:

```
local_htlc_basepoint: 034f355bdc7cc0af728ef3cceb9615d90684bb5b2ca5f859ab0f0b704075871aa
remote_htlc_basepoint: 032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991
```

And, here are the keys needed to create the transactions:

```
local_funding_privkey: 30ff4956bbdd3222d44cc5e8a1261dab1e07957bdac5ae88fe3261ef321f374901
local_funding_pubkey: 023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb
remote_funding_pubkey: 030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c39fa711c1
local_privkey: bb13b121cdc357cd2e608b0aea294afca36e2b34cf958e2e6451a2f27469449101
local_htlcpubkey: 030d417a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e7
remote_htlcpubkey: 0394854aa6eab5b2a8122cc726e9dded053a2184d88256816826d6231c068d4a5b
local_delayedpubkey: 03fd5960528dc152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c
local_revocation_pubkey: 0212a140cd0c6539d07cd08dfe09984dec3251ea808b892efecac3ede9402bf2b19
# funding wscript =
5221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d
66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae
```

And, here are the test vectors themselves:

```
name: simple commitment tx with no HTLCs
to_local_msat: 7000000000
to_remote_msat: 3000000000
local_feerate_per_kw: 15000
# base commitment transaction fee = 10860
# actual commitment transaction fee = 10860
# to_local amount 6989140 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efecac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
3045022100c3127b33dcc741dd6b05b1e63cbd1a9a7d816f37af9b6756fa2376b056f032370220408b96279808fe57eb
7e463710804cdf4f108388bc5cf722d8c848d2c7f9f3b00
# local_signature =
30440220616210b2cc4d3afb601013c373bbd8aac54febd9f15400379a8cb65ce7deca60022034236c010991beb7ff77
0510561ae8dc885b8d38d1947248c38f2ae055647142
output commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80
02c0c62d0000000000160014cc1b07838e387deacd0e5232e1e8b49f4c29e48454a56a0000000002200204adb4e2f00
643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e04004730440220616210b2cc4d3afb601013c373bb
d8aac54febd9f15400379a8cb65ce7deca60022034236c010991beb7ff770510561ae8dc885b8d38d1947248c38f2ae0
5564714201483045022100c3127b33dcc741dd6b05b1e63cbd1a9a7d816f37af9b6756fa2376b056f032370220408b96
279808fe57eb7e463710804cdf4f108388bc5cf722d8c848d2c7f9f3b001475221023da092f6980e58d2c037173180e9
a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c3
9fa711c152ae3e195220
num_htlcs: 0
```

name: commitment tx with all five HTLCs untrimmed (minimum feerate)


```
to_local_msat: 6988000000
to_remote_msat: 3000000000
local_feerate_per_kw: 0
# base commitment transaction fee = 0
# actual commitment transaction fee = 0
# HTLC #2 offered amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a914b43e1b38138a41b37f7cd9a1d274bc63e3a9b5d188ac6868
# HTLC #3 offered amount 3000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a9148a486ff2e31d6158bf39e2608864d63fef09d5b88ac6868
# HTLC #0 received amount 1000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a914b8bcb07f6344b42ab04250c86a6e8b75d3fdbbc688527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f401b175ac6868
# HTLC #1 received amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a9144b6b2e5444c2639cc0fb7bcea5afba3f3cdce23988527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f501b175ac6868
# HTLC #4 received amount 4000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a91418bc1a114ccf9c052d3d23e28d3b0a9d1227434288527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f801b175ac6868
# to_local amount 6988000 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efeac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
3044022009b048187705a8cbc9ad73adbe5af148c3d012e1f067961486c822c7af08158c022006d66f3704cfab3eb2dc
49dae24e4aa22a6910fc9b424007583204e3621af2e5
# local_signature =
304402206fc2d1f10ea59951eefac0b4b7c396a3c3d87b71ff0b019796ef4535beaf36f902201765b0181e514d04f4c8
ad75659d7037be26cdb3f8bb6f78fe61decef484c3ea
output_commit_tx:
0200000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80
07e8030000000000022002052bfef0479d7b293c27e0f1eb294bea154c63a3294ef092c19af51409bce0e2ad0070000
00000000220020403d394747cae42e98ff01734ad5c08f82ba123d3d9a620abda88989651e2ab5d0070000000000022
0020748eba944fedc8827f6b06bc44678f93c0f9e6078b35c6331ed31e75f8ce0c2db80b00000000000220020c20b5d
1f8584fd90443e7b7b720136174fa4b9333c261d04dbbd012635c0f419a00f000000000002200208c48d15160397c97
31df9bc3b236656efb6665fbfe92b4a6878e88a499f741c4c0c62d0000000000160014cc1b07838e387deacd0e5232e1
e8b49f4c29e484e0a06a0000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7d
d80e040047304402206fc2d1f10ea59951eefac0b4b7c396a3c3d87b71ff0b019796ef4535beaf36f902201765b0181e
514d04f4c8ad75659d7037be26cdb3f8bb6f78fe61decef484c3ea01473044022009b048187705a8cbc9ad73adbe5af1
48c3d012e1f067961486c822c7af08158c022006d66f3704cfab3eb2dc49dae24e4aa22a6910fc9b424007583204e362
1af2e501475221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2c
cc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae3e195220
num_htlcs: 5
# signature for output #0 (htlc-success for htlc #0)
remote_htlc_signature =
3045022100d9e29616b8f3959f1d3d7f7ce893ffedcdc407717d0de8e37d808c91d3a7c50d022078c3033f6d00095c87
20a4bc943c1b45727818c082e4e3ddbc6d3116435b624b
# local_htlc_signature =
```

[illegible]

[illegible]

[illegible]

[illegible]

```
# HTLC #2 offered amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a914b43e1b38138a41b37f7cd9a1d274bc63e3a9b5d188ac6868
# HTLC #3 offered amount 3000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a9148a486ff2e31d6158bf39e2608864d63fef09d5b88ac6868
# HTLC #1 received amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a9144b6b2e5444c2639cc0fb7bcea5afba3f3cdce23988527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f501b175ac6868
# HTLC #4 received amount 4000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a91418bc1a114ccf9c052d3d23e28d3b0a9d1227434288527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f801b175ac6868
# to_local amount 6987086 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efeac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304402203948f900a5506b8de36a4d8502f94f21dd84fd9c2314ab427d52feaa7a0a19f2022059b6a37a4adaa2c5419d
c8aea63c6e2a2ec4c4bde46207f6dc1fcd22152fc6e5
# local_signature =
3045022100b15f72908ba3382a34ca5b32519240a22300cc6015b6f9418635fb41f3d01d8802207adb331b9ed1575383
dca0f2355e86c173802feecf8298fba53b9d4610583e9
output_commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80
06d0070000000000000220020403d394747cae42e98ff01734ad5c08f82ba123d3d9a620abda88989651e2ab5d0070000
00000000220020748eba944fedc8827f6b06bc44678f93c0f9e6078b35c6331ed31e75f8ce0c2db80b0000000000022
0020c20b5d1f8584fd90443e7b7b720136174fa4b9333c261d04dbbd012635c0f419a00f0000000000002200208c48d1
5160397c9731df9bc3b236656efb6665fbfe92b4a6878e88a499f741c4c0c62d0000000000160014cc1b07838e387dea
cd0e5232e1e8b49f4c29e4844e9d6a0000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857ac
cc862d6b7dd80e0400483045022100b15f72908ba3382a34ca5b32519240a22300cc6015b6f9418635fb41f3d01d8802
207adb331b9ed1575383dca0f2355e86c173802feecf8298fba53b9d4610583e90147304402203948f900a5506b8de3
6a4d8502f94f21dd84fd9c2314ab427d52feaa7a0a19f2022059b6a37a4adaa2c5419dc8aea63c6e2a2ec4c4bde46207
f6dc1fcd22152fc6e501475221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb2103
0e9f7b623d2ccc7c9bd44d66d5ce21ce504c0ac6f385a132cec6d3c39fa711c152ae3e195220
num_htlcs: 4
# signature for output #0 (htlc-timeout for htlc #2)
remote_htlc_signature =
3045022100a031202f3be94678f0e998622ee95ebb6ada8da1e9a5110228b5e04a747351e4022010ca6a21e18314ed53
cfaae3b1f51998552a61a468e596368829a50ce40110e0
# local_htlc_signature =
304502210097e1873b57267730154595187a34949d3744f52933070c74757005e61ce2112e02204ecfba2aa42d4f14bd
f8bad4206bb97217b702e6c433e0e1b0ce6587e6d46ec6
htlc_timeout_tx (htlc #2):
020000000001010f44041fd9b175987cf4e6135ba2a154e3b7fb96483dc0ed5efc0678e5b6bf10000000000000000
01230600000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e05004830
45022100a031202f3be94678f0e998622ee95ebb6ada8da1e9a5110228b5e04a747351e4022010ca6a21e18314ed53cf
aae3b1f51998552a61a468e596368829a50ce40110e00148304502210097e1873b57267730154595187a34949d3744f5
2933070c74757005e61ce2112e02204ecfba2aa42d4f14bdf8bad4206bb97217b702e6c433e0e1b0ce6587e6d46ec601
008576a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a
2184d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4da
```

[illegible]

```
local_feerate_per_kw: 2069
# base commitment transaction fee = 2921
# actual commitment transaction fee = 3921
# HTLC #2 offered amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a914b43e1b38138a41b37f7cd9a1d274bc63e3a9b5d188ac6868
# HTLC #3 offered amount 3000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a9148a486ff2e31d6158bf39e2608864d63fef09d5b88ac6868
# HTLC #1 received amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a9144b6b2e5444c2639cc0fb7bcea5afba3f3cdce23988527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f501b175ac6868
# HTLC #4 received amount 4000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a91418bc1a114ccf9c052d3d23e28d3b0a9d1227434288527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f801b175ac6868
# to_local amount 6985079 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efeac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304502210090b96a2498ce0c0f2fadbec2aab278fed54c1a7838df793ec4d2c78d96ec096202204fdd439c50f90d483b
aa7b68feef4bd33bc277695405447bcd0bfb2ca34d7bc
# local_signature =
3045022100ad9a9bbbb75d506ca3b716b336ee3cf975dd7834fcf129d7dd188146eb58a8b4022061a759ee417339f7fe
2ea1e8deb83abb6a74db31a09b7648a932a639cda23e33
output_commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80
06d007000000000000220020403d394747cae42e98ff01734ad5c08f82ba123d3d9a620abda88989651e2ab5d0070000
00000000220020748eba944fedc8827f6b06bc44678f93c0f9e6078b35c6331ed31e75f8ce0c2db80b00000000000022
0020c20b5d1f8584fd90443e7b7b720136174fa4b9333c261d04dbbd012635c0f419a00f0000000000002200208c48d1
5160397c9731df9bc3b236656efb6665fbfe92b4a6878e88a499f741c4c0c62d0000000000160014cc1b07838e387dea
cd0e5232e1e8b49f4c29e48477956a00000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857ac
cc862d6b7dd80e0400483045022100ad9a9bbbb75d506ca3b716b336ee3cf975dd7834fcf129d7dd188146eb58a8b402
2061a759ee417339f7fe2ea1e8deb83abb6a74db31a09b7648a932a639cda23e330148304502210090b96a2498ce0c0f
2fadbec2aab278fed54c1a7838df793ec4d2c78d96ec096202204fdd439c50f90d483baa7b68feef4bd33bc27769540
5447bcd0bfb2ca34d7bc01475221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f54eb21
030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acff6385a132cec6d3c39fa711c152ae3e195220
num_htlcs: 4
# signature for output #0 (htlc-timeout for htlc #2)
remote_htlc_signature =
3045022100f33513ee38abf1c582876f921f8fddc06acff48e04515532a32d3938de938fffd02203aa308a2c1863b7d6f
df53159a1465bf2e115c13152546cc5d74483ceaa7f699
# local_htlc_signature =
3045022100a637902a5d4c9ba9e7c472a225337d5aac9e2e3f6744f76e237132e7619ba0400220035c60d784a031c0d9
f6df66b7eab8726a5c25397399ee4aa960842059eb3f9d
htlc_timeout_tx (htlc #2):
02000000000101adbe717a63fb658add30ada1e6e12ed257637581898abe475c11d7bbcd65bd4d000000000000000000
0175020000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e05004830
45022100f33513ee38abf1c582876f921f8fddc06acff48e04515532a32d3938de938fffd02203aa308a2c1863b7d6fdf
53159a1465bf2e115c13152546cc5d74483ceaa7f69901483045022100a637902a5d4c9ba9e7c472a225337d5aac9e2e
```


[illegible]

```
name: commitment tx with five outputs untrimmed (minimum feerate)
to_local_msat: 6988000000
to_remote_msat: 3000000000
local_feerate_per_kw: 2070
# base commitment transaction fee = 2566
# actual commitment transaction fee = 5566
# HTLC #2 offered amount 2000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a914b43e1b38138a41b37f7cd9a1d274bc63e3a9b5d188ac6868
# HTLC #3 offered amount 3000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca81
3e21734b140639e752ae67a9148a486ff2e31d6158bf39e2608864d63fef09d5b88ac6868
# HTLC #4 received amount 4000 wscript
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
d88256816826d6231c068d4a5b7c8201208763a91418bc1a114ccf9c052d3d23e28d3b0a9d1227434288527c21030d41
7a46946384f88d5f3337267c5e579765875dc4daca813e21734b140639e752ae677502f801b175ac6868
# to_local amount 6985434 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efeac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304402204ca1ba260dee913d318271d86e10ca0f5883026fb5653155cff600fb40895223022037b145204b7054a40e08
bb1fefbd826f827b40838d3e501423bcc57924bcb50c
# local_signature =
3044022001014419b5ba00e083ac4e0a85f19afc848aacac2d483b4b525d15e2ae5adbfe022015ebddad6ee1e72b47cb
09f3e78459da5be01cccd95dceca0e056a00cc773c1
output commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80
05d0070000000000000220020403d394747cae42e98ff01734ad5c08f82ba123d3d9a620abda88989651e2ab5b80b0000
00000000220020c20b5d1f8584fd90443e7b7b720136174fa4b9333c261d04dbbd012635c0f419a00f00000000000022
00208c48d15160397c9731df9bc3b236656efb6665fbfe92b4a6878e88a499f741c4c0c62d00000000000160014cc1b07
838e387deacd0e5232e1e8b49f4c29e484da966a00000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c1
1a40d857accc862d6b7dd80e0400473044022001014419b5ba00e083ac4e0a85f19afc848aacac2d483b4b525d15e2ae
5adbfe022015ebddad6ee1e72b47cb09f3e78459da5be01cccd95dceca0e056a00cc773c10147304402204ca1ba260d
ee913d318271d86e10ca0f5883026fb5653155cff600fb40895223022037b145204b7054a40e08bb1fefbd826f827b40
838d3e501423bcc57924bcb50c01475221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f
54eb21030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae3e195220
num_htlcs: 3
# signature for output #0 (htlc-timeout for htlc #2)
remote_htlc_signature =
304402205f6b6d12d8d2529fb24f4445630566cf4abbd0f9330ab6c2bdb94222d6a2a0c502202f556258ae6f05b19374
9e4c541dfcc13b525a5422f6291f073f15617ba8579b
# local_htlc_signature =
30440220150b11069454da70caf2492ded9e0065c9a57f25ac2a4c52657b1d15b6c6ed85022068a38833b603c8892717
206383611bad210f1cbb4b1f87ea29c6c65b9e1cb3e5
htlc_timeout_tx (htlc #2):
02000000000101403ad7602b43293497a3a2235a12ecefda4f3a1f1d06e49b1786d945685de1ff000000000000000000
01740200000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e05004730
4402205f6b6d12d8d2529fb24f4445630566cf4abbd0f9330ab6c2bdb94222d6a2a0c502202f556258ae6f05b193749e
4c541dfcc13b525a5422f6291f073f15617ba8579b014730440220150b11069454da70caf2492ded9e0065c9a57f25ac
2a4c52657b1d15b6c6ed85022068a38833b603c8892717206383611bad210f1cbb4b1f87ea29c6c65b9e1cb3e5010085
76a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a2184
```

[illegible]

```
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efead3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304402204bb3d6e279d71d9da414c82de42f1f954267c762b2e2eb8b76bc3be4ea07d4b0022014febc009c5edc8c3fc5
d94015de163200f780046f1c293bfed8568f08b70fb3
# local_signature =
3044022072c2e2b1c899b2242656a537dde2892fa3801be0d6df0a87836c550137acde8302201654aa1974d37a829083
c3ba15088689f30b56d6a4f6cb14c7bad0ee3116d398
output_commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80
05d0070000000000000220020403d394747cae42e98ff01734ad5c08f82ba123d3d9a620abda88989651e2ab5b80b0000
00000000220020c20b5d1f8584fd90443e7b7b720136174fa4b9333c261d04dbbd012635c0f419a00f00000000000022
00208c48d15160397c9731df9bc3b236656efb6665fbfe92b4a6878e88a499f741c4c0c62d0000000000160014cc1b07
838e387deacd0e5232e1e8b49f4c29e48440966a00000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c1
1a40d857accc862d6b7dd80e0400473044022072c2e2b1c899b2242656a537dde2892fa3801be0d6df0a87836c550137
acde8302201654aa1974d37a829083c3ba15088689f30b56d6a4f6cb14c7bad0ee3116d3980147304402204bb3d6e279
d71d9da414c82de42f1f954267c762b2e2eb8b76bc3be4ea07d4b0022014febc009c5edc8c3fc5d94015de163200f780
046f1c293bfed8568f08b70fb301475221023da092f6980e58d2c037173180e9a465476026ee50f96695963e8efe436f
54eb21030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c39fa711c152ae3e195220
num_htlcs: 3
# signature for output #0 (htlc-timeout for htlc #2)
remote_htlc_signature =
3045022100939726680351a7856c1bc386d4a1f422c7d29bd7b56afc139570f508474e6c40022023175a799ccf44c017
fbaadb924c40b2a12115a5b7d0dfd3228df803a2de8450
# local_htlc_signature =
304502210099c98c2edeeeee6ec0fb5f3bea8b79bb016a2717afa9b5072370f34382de281d302206f5e2980a995e045cf
90a547f0752a7ee99d48547bc135258fe7bc07e0154301
htlc_timeout_tx (htlc #2):
02000000000101153cd825fdb3aa624bfe513e8031d5d08c5e582fb3d1d1fe8faf27d3eed410cd00000000000000000
01220200000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e05004830
45022100939726680351a7856c1bc386d4a1f422c7d29bd7b56afc139570f508474e6c40022023175a799ccf44c017fb
aadb924c40b2a12115a5b7d0dfd3228df803a2de84500148304502210099c98c2edeeeee6ec0fb5f3bea8b79bb016a271
7afa9b5072370f34382de281d302206f5e2980a995e045cf90a547f0752a7ee99d48547bc135258fe7bc07e015430101
008576a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a
2184d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4da
ca813e21734b140639e752ae67a914b43e1b38138a41b37f7cd9a1d274bc63e3a9b5d188ac6868f6010000
# signature for output #1 (htlc-timeout for htlc #3)
remote_htlc_signature =
3044022021bb883bf324553d085ba2e821cad80c28ef8b303dbead8f98e548783c02d1600220638f9ef2a9bba25869af
c923f4b5dc38be3bb459f9efa5d869392d5f7779a4a0
# local_htlc_signature =
3045022100fd85bd7697b89c08ec12acc8ba89b23090637d83abd26ca37e01ae93e67c367302202b551fe69386116c47
f984aab9c8dfd25d864dcde5d3389cfbef2447a85c4b77
htlc_timeout_tx (htlc #3):
02000000000101153cd825fdb3aa624bfe513e8031d5d08c5e582fb3d1d1fe8faf27d3eed410cd010000000000000000
010a0600000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e05004730
44022021bb883bf324553d085ba2e821cad80c28ef8b303dbead8f98e548783c02d1600220638f9ef2a9bba25869afc9
23f4b5dc38be3bb459f9efa5d869392d5f7779a4a001483045022100fd85bd7697b89c08ec12acc8ba89b23090637d83
abd26ca37e01ae93e67c367302202b551fe69386116c47f984aab9c8dfd25d864dcde5d3389cfbef2447a85c4b770100
8576a91414011f7254d96b819c76986c277d115efce6f7b58763ac67210394854aa6eab5b2a8122cc726e9dded053a21
84d88256816826d6231c068d4a5b7c820120876475527c21030d417a46946384f88d5f3337267c5e579765875dc4daca
813e21734b140639e752ae67a9148a486ff2e31d6158bf39e2608864d63fef09d5b88ac6868f7010000
# signature for output #2 (htlc-success for htlc #4)
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
output commit_tx:
0200000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80
02c0c62d000000000160014cc1b07838e387deacd0e5232e1e8b49f4c29e484fa926a0000000002200204adb4e2f00
643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80e04004830450221008a953551f4d67cb4df3037207f
c082ddaf6be84d417b0bd14c80aab66f1b01a402207508796dc75034b2dee876fe01dc05a08b019f3e5d689ac8842ade
2f1befccf50147304402203a286936e74870ca1459c700c71202af0381910a6bfab687ef494ef1bc3e02c902202506c3
62d0e3bee15e802aa729bf378e051644648253513f1c085b264cc2a72001475221023da092f6980e58d2c037173180e9
a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c3
9fa711c152ae3e195220
num_htlcs: 0
```

```
name: commitment tx with two outputs untrimmed (maximum feerate)
to_local_msat: 6988000000
to_remote_msat: 3000000000
local_feerate_per_kw: 9651180
# base commitment transaction fee = 6987454
# actual commitment transaction fee = 6999454
# to_local amount 546 wscript
63210212a140cd0c6539d07cd08dfe09984dec3251ea808b892efecac3ede9402bf2b1967029000b2752103fd5960528d
c152014952efdb702a88f71e3c1653b2314431701ec77e57fde83c68ac
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304402200a8544eba1d216f5c5e530597665fa9bec56943c0f66d98fc3d028df52d84f7002201e45fa5c6bc3a506cc25
53e7d1c0043a9811313fc39c954692c0d47cfce2bbd3
# local_signature =
3045022100e11b638c05c650c2f63a421d36ef8756c5ce82f2184278643520311cdf50aa200220259565fb9c8e4a87cc
af17f27a3b9ca4f20625754a0920d9c6c239d8156a11de
output commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80
02220200000000000002200204adb4e2f00643db396dd120d4e7dc17625f5f2c11a40d857accc862d6b7dd80ec0c62d00
000000000160014cc1b07838e387deacd0e5232e1e8b49f4c29e4840400483045022100e11b638c05c650c2f63a421d36
ef8756c5ce82f2184278643520311cdf50aa200220259565fb9c8e4a87ccaf17f27a3b9ca4f20625754a0920d9c6c239
d8156a11de0147304402200a8544eba1d216f5c5e530597665fa9bec56943c0f66d98fc3d028df52d84f7002201e45fa
5c6bc3a506cc2553e7d1c0043a9811313fc39c954692c0d47cfce2bbd301475221023da092f6980e58d2c037173180e9
a465476026ee50f96695963e8efe436f54eb21030e9f7b623d2ccc7c9bd44d66d5ce21ce504c0acf6385a132cec6d3c3
9fa711c152ae3e195220
num_htlcs: 0
```

```
name: commitment tx with one output untrimmed (minimum feerate)
to_local_msat: 6988000000
to_remote_msat: 3000000000
local_feerate_per_kw: 9651181
# base commitment transaction fee = 6987455
# actual commitment transaction fee = 7000000
# to_remote amount 3000000
P2WPKH(032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991)
remote_signature =
304402202ade0142008309eb376736575ad58d03e5b115499709c6db0b46e36ff394b492022037b63d78d66404d6504d
4c4ac13be346f3d1802928a6d3ad95a6a944227161a2
# local_signature =
304402207e8d51e0c570a5868a78414f4e0cbfaed1106b171b9581542c30718ee4eb95ba02203af84194c97adf98898c
9afe2f2ed4a7f8dba05a2dfab28ac9d9c604aa49a379
output commit_tx:
02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80
```


[illegible]

Appendix D: Per-commitment Secret Generation Test Vectors

Generation Tests

```
name: generate_from_seed 01 last nontrivial node
seed: 0x0101010101010101010101010101010101010101010101010101010101010101
I: 1
output: 0x915c75942a26bb3a433a8ce2cb0427c29ec6c1775cfc78328b57f6ba7bfeaa9c
```

```
name: insert_secret correct sequence
I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518c8ae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
```

secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: OK
I: 281474976710651
secret: 0xc65716add7aa98ba7acb236352d665cab17345fe45b55fb879ff80e6bd0c41dd
output: OK
I: 281474976710650
secret: 0x969660042a28f32d9be17344e09374b379962d03db1574df5a8a5a47e19ce3f2
output: OK
I: 281474976710649
secret: 0xa5a64476122ca0925fb344bdc1854c1c0a59fc614298e50a33e331980a220f32
output: OK
I: 281474976710648
secret: 0x05cde6323d949933f7f7b78776bcc1ea6d9b31447732e3802e1f7ac44b650e17
output: OK

name: insert_secret #1 incorrect

I: 281474976710655
secret: 0x02a40c85b6f28da08dfdbe0926c53fab2de6d28c10301f8f7c4073d5e42e3148
output: OK
I: 281474976710654
secret: 0xc7518c8ae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: ERROR

name: insert_secret #2 incorrect (#1 derived from incorrect)

I: 281474976710655
secret: 0x02a40c85b6f28da08dfdbe0926c53fab2de6d28c10301f8f7c4073d5e42e3148
output: OK
I: 281474976710654
secret: 0xdddc3a8d14fddf2b68fa8c7fbad2748274937479dd0f8930d5ebb4ab6bd866a3
output: OK
I: 281474976710653
secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: ERROR

name: insert_secret #3 incorrect

I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518c8ae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
secret: 0xc51a18b13e8527e579ec56365482c62f180b7d5760b46e9477dae59e87ed423a
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: ERROR

name: insert_secret #4 incorrect (1,2,3 derived from incorrect)

I: 281474976710655
secret: 0x02a40c85b6f28da08dfdbe0926c53fab2de6d28c10301f8f7c4073d5e42e3148
output: OK
I: 281474976710654
secret: 0xdddc3a8d14fddf2b68fa8c7fbad2748274937479dd0f8930d5ebb4ab6bd866a3
output: OK
I: 281474976710653
secret: 0xc51a18b13e8527e579ec56365482c62f180b7d5760b46e9477dae59e87ed423a
output: OK
I: 281474976710652
secret: 0xba65d7b0ef55a3ba300d4e87af29868f394f8f138d78a7011669c79b37b936f4
output: OK
I: 281474976710651
secret: 0xc65716add7aa98ba7acb236352d665cab17345fe45b55fb879ff80e6bd0c41dd
output: OK
I: 281474976710650
secret: 0x969660042a28f32d9be17344e09374b379962d03db1574df5a8a5a47e19ce3f2
output: OK
I: 281474976710649
secret: 0xa5a64476122ca0925fb344bdc1854c1c0a59fc614298e50a33e331980a220f32
output: OK
I: 281474976710648
secret: 0x05cde6323d949933f7f7b78776bcc1ea6d9b31447732e3802e1f7ac44b650e17
output: ERROR

name: insert_secret #5 incorrect

I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518cae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: OK
I: 281474976710651
secret: 0x631373ad5f9ef654bb3dade742d09504c567edd24320d2fcd68e3cc47e2ff6a6
output: OK
I: 281474976710650
secret: 0x969660042a28f32d9be17344e09374b379962d03db1574df5a8a5a47e19ce3f2
output: ERROR

name: insert_secret #6 incorrect (5 derived from incorrect)

I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518cae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK

I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: OK
I: 281474976710651
secret: 0x631373ad5f9ef654bb3dade742d09504c567edd24320d2fcd68e3cc47e2ff6a6
output: OK
I: 281474976710650
secret: 0xb7e76a83668bde38b373970155c868a653304308f9896692f904a23731224bb1
output: OK
I: 281474976710649
secret: 0xa5a64476122ca0925fb344bdc1854c1c0a59fc614298e50a33e331980a220f32
output: OK
I: 281474976710648
secret: 0x05cde6323d949933f7f7b78776bcc1ea6d9b31447732e3802e1f7ac44b650e17
output: ERROR

name: insert_secret #7 incorrect

I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518c8ae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: OK
I: 281474976710651
secret: 0xc65716add7aa98ba7acb236352d665cab17345fe45b55fb879ff80e6bd0c41dd
output: OK
I: 281474976710650
secret: 0x969660042a28f32d9be17344e09374b379962d03db1574df5a8a5a47e19ce3f2
output: OK
I: 281474976710649
secret: 0xe7971de736e01da8ed58b94c2fc216cb1dca9e326f3a96e7194fe8ea8af6c0a3
output: OK
I: 281474976710648
secret: 0x05cde6323d949933f7f7b78776bcc1ea6d9b31447732e3802e1f7ac44b650e17
output: ERROR

name: insert_secret #8 incorrect

I: 281474976710655
secret: 0x7cc854b54e3e0dcdb010d7a3fee464a9687be6e8db3be6854c475621e007a5dc
output: OK
I: 281474976710654
secret: 0xc7518c8ae4660ed02894df8976fa1a3659c1a8b4b5bec0c4b872abeba4cb8964
output: OK
I: 281474976710653
secret: 0x2273e227a5b7449b6e70f1fb4652864038b1cbf9cd7c043a7d6456b7fc275ad8
output: OK
I: 281474976710652
secret: 0x27cddaa5624534cb6cb9d7da077cf2b22ab21e9b506fd4998a51d54502e99116
output: OK


```
I: 281474976710651
secret: 0xc65716add7aa98ba7acb236352d665cab17345fe45b55fb879ff80e6bd0c41dd
output: OK
I: 281474976710650
secret: 0x969660042a28f32d9be17344e09374b379962d03db1574df5a8a5a47e19ce3f2
output: OK
I: 281474976710649
secret: 0xa5a64476122ca0925fb344bdc1854c1c0a59fc614298e50a33e331980a220f32
output: OK
I: 281474976710648
secret: 0xa7efbc61aac46d34f77778bac22c8a20c6a46ca460addc49009bda875ec88fa4
output: ERROR
```

Appendix E: Key Derivation Test Vectors

These test the derivation for localpubkey, remotepubkey, local_htlcpubkey, remote_htlcpubkey, local_delayedpubkey, and remote_delayedpubkey (which use the same formula), as well as the revocationpubkey.

All of them use the following secrets (and thus the derived points):

```
base_secret: 0x000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
per_commitment_secret: 0x1f1e1d1c1b1a191817161514131211100f0e0d0c0b0a09080706050403020100
base_point: 0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2
per_commitment_point: 0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486
```

```
name: derivation of pubkey from basepoint and per_commitment_point
# SHA256(per_commitment_point || basepoint)
# => SHA256(0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486 ||
0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2)
# = 0xcbcdd70fcfad15ea8e9e5c5a12365cf00912504f08ce01593689dd426bca9ff0
# + basepoint (0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2)
# = 0x0235f2dbfaa89b57ec7b055afe29849ef7ddfeb1cefdb9ebdc43f5494984db29e5
localpubkey: 0x0235f2dbfaa89b57ec7b055afe29849ef7ddfeb1cefdb9ebdc43f5494984db29e5
```

```
name: derivation of private key from basepoint secret and per_commitment_secret
# SHA256(per_commitment_point || basepoint)
# => SHA256(0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486 ||
0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2)
# = 0xcbcdd70fcfad15ea8e9e5c5a12365cf00912504f08ce01593689dd426bca9ff0
# + basepoint_secret (0x000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f)
# = 0xcbced912d3b21bf196a766651e436aff192362621ce317704ea2f75d87e7be0f
localprivkey: 0xcbced912d3b21bf196a766651e436aff192362621ce317704ea2f75d87e7be0f
```

```
name: derivation of revocation pubkey from basepoint and per_commitment_point
# SHA256(revocation_basepoint || per_commitment_point)
# => SHA256(0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2 ||
0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486)
# = 0xefbf7ba5a074276701798376950a64a90f698997cce0dff4d24a6d2785d20963
# x revocation_basepoint = 0x02c00c4aad536290422a807250824a8d87f19d18da9d610d45621df22510db8ce
# SHA256(per_commitment_point || revocation_basepoint)
# => SHA256(0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486 ||
0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2)
# = 0xcbcdd70fcfad15ea8e9e5c5a12365cf00912504f08ce01593689dd426bca9ff0
```

```
# x_per_commitment_point = 0x0325ee7d3323ce52c4b33d4e0a73ab637711057dd8866e3b51202a04112f054c43
# 0x02c00c4aad536290422a807250824a8d87f19d18da9d610d45621df22510db8ce +
0x0325ee7d3323ce52c4b33d4e0a73ab637711057dd8866e3b51202a04112f054c43 =>
0x02916e326636d19c33f13e8c0c3a03dd157f332f3e99c317c141dd865eb01f8ff0
revocationpubkey: 0x02916e326636d19c33f13e8c0c3a03dd157f332f3e99c317c141dd865eb01f8ff0
```

```
name: derivation of revocation secret from basepoint_secret and per_commitment_secret
# SHA256(revocation_basepoint || per_commitment_point)
# => SHA256(0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2 ||
0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486)
# = 0xefbf7ba5a074276701798376950a64a90f698997cce0dff4d24a6d2785d20963
# * revocation_basepoint_secret
(0x000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f)# =
0x44bfd55f845f885b8e60b2dca4b30272d5343be048d79ce87879d9863dedc842
# SHA256(per_commitment_point || revocation_basepoint)
# => SHA256(0x025f7117a78150fe2ef97db7cfc83bd57b2e2c0d0dd25eaf467a4a1c2a45ce1486 ||
0x036d6caac248af96f6afa7f904f550253a0f3ef3f5aa2fe6838a95b216691468e2)
# = 0xcbbdd70fcfad15ea8e9e5c5a12365cf00912504f08ce01593689dd426bca9ff0
# * per_commitment_secret (0x1f1e1d1c1b1a191817161514131211100f0e0d0c0b0a09080706050403020100)#
= 0x8be02a96a97b9a3c1c9f59ebb718401128b72ec009d85ee1656319b52319b8ce
# => 0xd09ffff62ddb2297ab000cc85bcb4283fdeb6aa052affbc9dddcf33b61078110
revocationprivkey: 0xd09ffff62ddb2297ab000cc85bcb4283fdeb6aa052affbc9dddcf33b61078110
```

Appendix F: Commitment and HTLC Transaction Test Vectors (anchors)

The anchor test vectors are based on the test cases as defined in appendix C. Note that in appendix C, to_local_msat and to_remote_msat are balances before subtraction of: * Commit fee (funder only) * Anchor outputs (funder only) * In-flight htlcs

```
[
  {
    "Name": "simple commitment tx with no HTLCs",
    "LocalBalance": 7000000000,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 546,
    "FeePerKw": 15000,
    "UseTestHtlcs": false,
    "HtlcDescs": [],
    "ExpectedCommitmentTxHex":
      "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b8004",

    "RemoteSigHex":
      "3045022100f89034eba16b2be0e5581f750a0a6309192b75cce0f202f0ee2b4ec0cc394850022076c65dc507fe42276152b",

  },
  {
    "Name": "simple commitment tx with no HTLCs and single anchor",
    "LocalBalance": 10000000000,
    "RemoteBalance": 0,
    "DustLimitSatoshis": 546,
    "FeePerKw": 15000,
    "UseTestHtlcs": false,
    "HtlcDescs": [],
    "ExpectedCommitmentTxHex":
      "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b8002"
```

```

    "RemoteSigHex":
    "30440220655bf909fb6fa81d086f1336ac72c97906dce29d1b166e305c99152d810e26e1022051f577faa46412c46707aa",
  },
  {
    "Name": "commitment tx with seven outputs untrimmed",
    "LocalBalance": 6988000000,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 546,
    "FeePerKw": 644,
    "UseTestHtlcs": true,
    "HtlcDescs": [
      {
        "RemoteSigHex":
        "30440220746dc89a593e1b50915db63359b50c3c404f8324f78075c87708c866ccefd2502202a11062012dc8607b17e8c4",

        "ResolutionTxHex":
        "02000000000101b8cefef62ea66f5178b9361b2371be0759cbc8c689bcfa7a8e6746d497ec221a02000000000100000001",

        },
      {
        "RemoteSigHex":
        "3045022100a847bc3b8cf2441013725ae32dc589c419835d069afd6d7f3d9834d8be7cea6e02204ce9d35ec7f0788da80e4",

        "ResolutionTxHex":
        "02000000000101b8cefef62ea66f5178b9361b2371be0759cbc8c689bcfa7a8e6746d497ec221a03000000000100000001",

        },
      {
        "RemoteSigHex":
        "3044022035977f2aa6d6ae12f1dae3366440faa17558e9c88c3188bfc8df7e276c6c65410220659f1f9070c725d9d46e43b",

        "ResolutionTxHex":
        "02000000000101b8cefef62ea66f5178b9361b2371be0759cbc8c689bcfa7a8e6746d497ec221a04000000000100000001",

        },
      {
        "RemoteSigHex":
        "3044022002f3ff9a31270092c214d3d5b8b4f826599404bba64b87f7536ef6324d41551b022079dc4cb25f7ecd84b49f5cc",

        "ResolutionTxHex":
        "02000000000101b8cefef62ea66f5178b9361b2371be0759cbc8c689bcfa7a8e6746d497ec221a05000000000100000001",

        },
      {
        "RemoteSigHex":
        "30440220547937564288f64fb3e3c945a1348b912471f0c1c4cc7dc8ceca15a4cbd299b6022053c4f8e30832b13dfbe31e4",

        "ResolutionTxHex":
        "02000000000101b8cefef62ea66f5178b9361b2371be0759cbc8c689bcfa7a8e6746d497ec221a06000000000100000001",

        }
    ],
    "ExpectedCommitmentTxHex":
    "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b8009"
  }
}

```

```

    "RemoteSigHex":
    "3045022100e0106830467a558c07544a3de7715610c1147062e7d091deeebe8b5c661cda9402202ad049c1a6d04834317a7

},
{
    "Name": "commitment tx with six outputs untrimmed (minimum dust limit)",
    "LocalBalance": 6988000000,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 1001,
    "FeePerKw": 645,
    "UseTestHtlcs": true,
    "HtlcDescs": [
        {
            "RemoteSigHex":
            "3045022100e04d160a326432659fe9fb127304c1d348dfeaba840081bdc57d8efd902a48d8022008a824e7cf5492b97e4d9

            "ResolutionTxHex":
            "02000000000101104f394af4c4fad78337f95e3e9f802f4c0d86ab231853af09b285348561320002000000000100000001c

        },
        {
            "RemoteSigHex":
            "3045022100fbdc3c367ce3bf30796025cc590ee1f2ce0e72ae1ac19f5986d6d0a4fc76211f02207e45ae9267e8e820d1885

            "ResolutionTxHex":
            "02000000000101104f394af4c4fad78337f95e3e9f802f4c0d86ab231853af09b285348561320003000000000100000001c

        },
        {
            "RemoteSigHex":
            "3044022066c5ef625cee3ddd2bc7b6bfb354b5834cf1cc6d52dd972fb41b7b225437ae4a022066cb85647df65c6b87a54e4

            "ResolutionTxHex":
            "02000000000101104f394af4c4fad78337f95e3e9f802f4c0d86ab231853af09b285348561320004000000000100000001b

        },
        {
            "RemoteSigHex":
            "3044022022c7e11595c53ee89a57ca76baf0aed730da035952d6ab3fe6459f5eff3b337a022075e10cc5f5fd724a35ce408

            "ResolutionTxHex":
            "02000000000101104f394af4c4fad78337f95e3e9f802f4c0d86ab231853af09b285348561320005000000000100000001a

        }
    ],
    "ExpectedCommitmentTxHex":
    "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a488489000000000038b02b80084

    "RemoteSigHex":
    "3044022025d97466c8049e955a5afce28e322f4b34d2561118e52332fb400f9b908cc0a402205dc6fba3a0d67ee142c428c

},
{
    "Name": "commitment tx with four outputs untrimmed (minimum dust limit)",
    "LocalBalance": 6988000000,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 2001,

```

```
"FeePerKw": 2185,  
"UseTestHtlcs": true,  
"HtlcDescs": [  
  {  
    "RemoteSigHex":  
      "304402206870514a72ad6e723ff7f1e0370d7a33c1cd2a0b9272674143ebaf6a1d02dee102205bd953c34faf5e7322e9a1d...",  
  
    "ResolutionTxHex":  
      "020000000000101ac13a7715f80b8e52dda43c6929cade5521bdced3a405da02b443f1ffb1e33cc02000000000100000001...",  
  
  },  
  {  
    "RemoteSigHex":  
      "3045022100949e8dd938da56445b1cdfdebe1b7efea086edd05d89910d205a1e2e033ce47102202cbd68b5262ab144d9ec1...",  
  
    "ResolutionTxHex":  
      "020000000000101ac13a7715f80b8e52dda43c6929cade5521bdced3a405da02b443f1ffb1e33cc03000000000100000001...",  
  
  }  
],  
"ExpectedCommitmentTxHex":  
  "020000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80064...",  
  
"RemoteSigHex":  
  "3044022040f63a16148cf35c8d3d41827f5ae7f7c3746885bb64d4d1b895892a83812b3e02202fcf95c2bf02c466163b3fa...",  
  
},  
{  
  "Name": "commitment tx with three outputs untrimmed (minimum dust limit)",  
  "LocalBalance": 6988000000,  
  "RemoteBalance": 3000000000,  
  "DustLimitSatoshis": 3001,  
  "FeePerKw": 3687,  
  "UseTestHtlcs": true,  
  "HtlcDescs": [  
    {  
      "RemoteSigHex":  
        "3044022017b558a3cf5f0cb94269e2e927b29ed22bd2416abb8a7ce6de4d1256f359b93602202e9ca2b1a23ea3e69f433c7...",  
  
      "ResolutionTxHex":  
        "020000000000101542562b326c08e3a076d9cfca2be175041366591da334d8d513fff1686fd95a6002000000000100000001...",  
  
    }  
  ],  
  "ExpectedCommitmentTxHex":  
    "020000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b80064...",  
  
  "RemoteSigHex":  
    "3045022100ad6c71569856b2d7ff42e838b4abe74a713426b37f22fa667a195a4c88908c6902202b37272b02a42dc6d9f4f...",  
  
},  
{  
  "Name": "commitment tx with two outputs untrimmed (minimum dust limit)",  
  "LocalBalance": 6988000000,  
  "RemoteBalance": 3000000000,  
  "DustLimitSatoshis": 4001,  
  "FeePerKw": 4894,
```

```

    "UseTestHtlcs": true,
    "HtlcDescs": [],
    "ExpectedCommitmentTxHex":
    "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b8004a",

    "RemoteSigHex":
    "3045022100e784a66b1588575801e237d35e510fd92a81ae3a4a2a1b90c031ad803d07b3f3022021bc5f16501f167607d63",

  },
  {
    "Name": "commitment tx with one output untrimmed (minimum dust limit)",
    "LocalBalance": 6988000000,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 4001,
    "FeePerKw": 6216010,
    "UseTestHtlcs": true,
    "HtlcDescs": [],
    "ExpectedCommitmentTxHex":
    "02000000000101bef67e4e2fb9ddeeb3461973cd4c62abb35050b1add772995b820b584a48848900000000038b02b8004a",

    "RemoteSigHex":
    "30450221008fd5dbff02e4b59020d4cd23a3c30d3e287065fda75a0a09b402980adf68ccda022001e0b8b620cd915ddff1",

  },
  {
    "Name": "commitment tx with 3 htlc outputs, 2 offered having the same amount and
    preimage",
    "LocalBalance": 6987999999,
    "RemoteBalance": 3000000000,
    "DustLimitSatoshis": 546,
    "FeePerKw": 253,
    "UseTestHtlcs": true,
    "HtlcDescs": [
      {
        "RemoteSigHex":
        "30440220078fe5343dab88c348a3a8a9c1a9293259dbf35507ae971702cc39dd623ea9af022011ed0c0f35243cd0bb4d9ca",

        "ResolutionTxHex":
        "020000000001013d060d0305c9616eaabc21d41fae85bcb5477b5d7f1c92aa429cf15339bbe1c402000000000100000001",

      },
      {
        "RemoteSigHex":
        "304402202df6bf0f98a42cfd0172a16bde7d1b16c14f5f42ba23f5c54648c14b647531302200fe1508626817f23925bb50",

        "ResolutionTxHex":
        "020000000001013d060d0305c9616eaabc21d41fae85bcb5477b5d7f1c92aa429cf15339bbe1c403000000000100000001",

      },
      {
        "RemoteSigHex":
        "3045022100bd206b420c495f3aa714d3ea4766cbe95441deacb5d2f737f1913349aee7c2ae02200249d2c950dd3b15326bf",

        "ResolutionTxHex":
        "020000000001013d060d0305c9616eaabc21d41fae85bcb5477b5d7f1c92aa429cf15339bbe1c404000000000100000001",

      }
    ]
  }
}

```

Funding Transaction Construction

Genesis block 0:

[illegible][illegible][illegible]

Coinbase address **privkey**: cPxftfE1w3ptFnsZvvFeWji21kTArYa9GXwMkYsoQHdaJKrjUTek

Coinbase address **privkey**: cPxftfE1w3ptFnsZvvFeWji21kTArYa9GXwMkYsoQHdaJKrjUTek

```
020000000000101f86fd1d0db3ac5a72df968622f31e6b5e6566a09e29206d7c7a55df90e181de800000000171600141f
b9623ffd0d422eacc450fd1e967efc477b83ccfffd0f0580b2e60e00000000220020fd89acf65485df89797d9ba7ba
7a33624ac4452f00db08107f34257d33e5b94680b2e60e0000000017a9146a235d064786b49e7043e4a042d4cc429f7e
b6948780b2e60e00000000160014fbb4db9d85fba5e301f4399e3038928e44e37d3280b2e60e0000000017a9147ecd1b
519326bc13b0ec716e469b58ed02b112a087f0006bee0000000017a914f856a70093da3a5b5c4302ade033d4c2171705
d387024730440220696f6cee2929f1feb3fd6adf024ca0f9aa2f4920ed6d35fb9ec5b78c8408475302201641afae1124
2160101c6f9932aeb4fcd1f13a9c6df5d1386def000ea259a35001210381d7d5b1bc0d7600565d827242576d9cb793bf
e0754334af82289ee8b65d137600000000
```

```
02000000000101f86fd1d0db3ac5a72df968622f31e6b5e6566a09e29206d7c7a55df90e181de800000000171600141f
b9623ffd0d422eacc450fd1e967efc477b83ccfffd0f0580b2e60e00000000220020fd89acf65485df89797d9ba7ba
7a33624ac4452f00db08107f34257d33e5b94680b2e60e0000000017a9146a235d064786b49e7043e4a042d4cc429f7e
b6948780b2e60e00000000160014fbb4db9d85fba5e301f4399e3038928e44e37d3280b2e60e0000000017a9147ecd1b
519326bc13b0ec716e469b58ed02b112a087f0006bee0000000017a914f856a70093da3a5b5c4302ade033d4c2171705
d387024730440220696f6cee2929f1feb3fd6adf024ca0f9aa2f4920ed6d35fb9ec5b78c8408475302201641afae1124
2160101c6f9932aeb4fcd1f13a9c6df5d1386def000ea259a35001210381d7d5b1bc0d7600565d827242576d9cb793bf
e0754334af82289ee8b65d137600000000
```

Locktime: 120 Feerate: 253 sat/kiloweight Opener's funding_satoshi: 2 0000 0000 sat Acceptor's funding_satoshi: 2 0000 0000 sat

Locktime: 120 Feerate: 253 sat/kiloweight Opener's funding_satoshi: 2 0000 0000 sat Acceptor's funding_satoshi: 2 0000 0000 sat

4303ca8ff10c6c345b9299672a66f111c5b81ae027cc5b0d4d39d09c66b032b9 0

4303ca8ff10c6c345b9299672a66f111c5b81ae027cc5b0d4d39d09c66b032b9 0

```
preimage: 20 68656c6c6f2074686572652c2074686973206973206120626974636f6e212121
```

```
witness_script: 27
82012088a820add57dfe5277079d069ca4ad4893c96de91f88ffb981fdc6a2a34d5336c66aff87
scriptPubKey: 0020fd89acf65485df89797d9ba7ba7a33624ac4452f00db08107f34257d33e5b946
address: bcr1t1qlky6eaj5sh0cj7tanwnm573nvf9vg3f0qrdssyrlxsjh6vl9h9rq140v2g

4303ca8ff10c6c345b9299672a66f111c5b81ae027cc5b0d4d39d09c66b032b9 2
pubkey: 034695f5b7864c580bf11f9f8cb1a94eb336f2ce9ef872d2ae1a90ee276c772484
privkey: cUM8Dr33wK4uFmw3Tz8sbQ7BiBNgX5BthRurU7RkgXVvNUPcWrJf
witness_program: fbb4db9d85fba5e301f4399e3038928e44e37d32
scriptPubKey: 0014fbb4db9d85fba5e301f4399e3038928e44e37d32
address: bcr1qlw6dh8v9lwj7xq058x0rqwyj3ezwxfjxsy7er
```

Expected Opener's tx_add_input (input 0 above):

```
num_inputs: 1
tx_add_input:[
{
  channel_id: "xxx",
  serial_id: 20,
  prevtx_len: 353,
  prevtx:
"02000000000101f86fd1d0db3ac5a72df968622f31e6b5e6566a09e29206d7c7a55df90e181de800000000171600141
fb9623ffd0d422eacc450fd1e967efc477b83ccffffffffff0580b2e60e00000000220020fd89acf65485df89797d9ba7b
a7a33624ac4452f00db08107f34257d33e5b94680b2e60e0000000017a9146a235d064786b49e7043e4a042d4cc429f7
eb6948780b2e60e00000000160014fbb4db9d85fba5e301f4399e3038928e44e37d3280b2e60e0000000017a9147ecd1
b519326bc13b0ec716e469b58ed02b112a087f0006bee0000000017a914f856a70093da3a5b5c4302ade033d4c217170
5d387024730440220696f6cee2929f1feb3fd6adf024ca0f9aa2f4920ed6d35fb9ec5b78c8408475302201641afae112
42160101c6f9932aeb4fcd1f13a9c6df5d1386def000ea259a35001210381d7d5b1bc0d7600565d827242576d9cb793b
fe0754334af82289ee8b65d137600000000",
  prev_vout: 0,
  sequence: 4294967293
}
]
```

Expected Acceptor's tx_add_input (input 2 above):

```
num_inputs: 1
tx_add_input:[
{
  channel_id: "xxx",
  serial_id: 11,
  prevtx_len: 353,
  prevtx:
"02000000000101f86fd1d0db3ac5a72df968622f31e6b5e6566a09e29206d7c7a55df90e181de800000000171600141
fb9623ffd0d422eacc450fd1e967efc477b83ccffffffffff0580b2e60e00000000220020fd89acf65485df89797d9ba7b
a7a33624ac4452f00db08107f34257d33e5b94680b2e60e0000000017a9146a235d064786b49e7043e4a042d4cc429f7
eb6948780b2e60e00000000160014fbb4db9d85fba5e301f4399e3038928e44e37d3280b2e60e0000000017a9147ecd1
b519326bc13b0ec716e469b58ed02b112a087f0006bee0000000017a914f856a70093da3a5b5c4302ade033d4c217170
5d387024730440220696f6cee2929f1feb3fd6adf024ca0f9aa2f4920ed6d35fb9ec5b78c8408475302201641afae112
42160101c6f9932aeb4fcd1f13a9c6df5d1386def000ea259a35001210381d7d5b1bc0d7600565d827242576d9cb793b
fe0754334af82289ee8b65d137600000000",
  prev_vout: 2,
  sequence: 4294967293
}
```



```
}  
]
```

Outputs: (scriptPubKeys)

```
# opener's change address  
pubkey: 0206e626a4c6d4392d4030bc78bd93f728d1ba61214a77c63adc17d71e32ded3df  
# privkey: cSpC1KYEV1vsUFBwTdcuRkncbwfipY1m5zuQ9CjgAYwiVvbQ4fc1  
scriptPubKey: 00141ca1cca8855bad6bc1ea5436edd8cff10b7e448b  
address: bcrt1qrsue2y9twkkhs022smwmkx07y9hu3ytshgjmj  
  
# accepter's change address  
pubkey: 028f3978c211f4c0bf4d20674f345ae14e08871b25b2c957b4bdbd42e9726278fc  
privkey: cQ1HXnbAE4wGhuB2b9rJEydV5ayeEmMqxf1dvHPZmyMTPkwvZJyg  
scriptPubKey: 001444cb0c39f93ecc372b5851725bd29d865d333b10  
address: bcrt1qgn9scw0e8mxrw26c29e9h55asewnxcwdxdp50  
  
# the 2-of-2s  
pubkey1: 0292edb5f7bbf9e900f7e024be1c1339c6d149c11930e613af3a983d2565f4e41e  
pubkey2: 02e16172a41e928cbd78f761bd1c657c4afc7495a1244f7f30166b654fbf7661e3  
script_def:  
multi(2,0292edb5f7bbf9e900f7e024be1c1339c6d149c11930e613af3a983d2565f4e41e,02e16172a41e928cbd78f761bd1c657c4afc7495a1244f7f30166b654fbf7661e3)  
script:  
52210292edb5f7bbf9e900f7e024be1c1339c6d149c11930e613af3a983d2565f4e41e2102e16172a41e928cbd78f761bd1c657c4afc7495a1244f7f30166b654fbf7661e352ae  
scriptPubKey: 0020297b92c238163e820b82486084634b4846b86a3c658d87b9384192e6bea98ec5  
address: bcrt1q99ae9s3czclgyzuzfpgsggc6tfprts63uvkxc0wfcgxfwd04f3mzs3asq6l
```

Expected Opener's tx_add_output:

```
num_outputs: 2  
tx_add_output[  
  {  
    channel_id:"xxx",  
    serial_id: 30,  
    sats: 49999845,  
    scriptlen: 22,  
    script: "1600141ca1cca8855bad6bc1ea5436edd8cff10b7e448b"  
  },{  
    channel_id: "xxx",  
    serial_id: 44,  
    sats: 400000000,  
    scriptlen: 34,  
    script: "220020297b92c238163e820b82486084634b4846b86a3c658d87b9384192e6bea98ec5"  
  }  
]
```

Expected Acceptor's tx_add_output:

```
num_outputs: 1  
tx_add_output[  
  {  
    channel_id: xxx,
```

1

Expected Fee Calculation:

Opener's fees and change:

```

initiator_weight = (input_count, output_count, version, locktime) * 4
                  + segwit_fields (marker + flag)
                  + inputs (txid, vout, scriptSig, sequence) * number initiator inputs * 4
                  + funding_output * 4
                  + change_output * 4
                  + max(number initiator inputs
                      * minimum witness weight,
                      actual witness weight of all inputs)

```

```
initiator_weight = (1 + 1 + 4 + 4) * 4
                  + 2
                  + (32 + 4 + 1 + 4) * 1 * 4
                  + 43 * 4
                  + 31 * 4
                  + max(1 * 107, 71)
```

```
initiator_weight = 609
```

```
initiator_fees = initiator_weight * feerate
initiator_fees = 609 * 253 / 1000
initiator_fees = 155 sats
```

```
change = total_funding
        - funding_sats
        - fees
```

$$\begin{array}{r} \text{change} = 2\ 5000\ 0000 \\ - 2\ 0000\ 0000 \\ - \qquad\qquad 155 \end{array}$$

```
change = 4999 9845
```

```
as hex: e5effa0200000000
```

Acceptor's fees and change:

```
contributor_weight = inputs(txid, vout, scriptSig, sequence)
    * number contributor inputs * 4
    + change_output * 4
+ max(number contributor inputs
    * minimum witness weight,
    actual witness weight of all inputs)
```

```

contributor_weight = (32 + 4 + 1 + 4) * 1 * 4
                    + 31 * 4
                    + max(1 * 107, 99)

contributor_weight = 395

contributor_fees = contributor_weight * feerate
contributor_fees = 398 * 253 / 1000
contributor_fees = 100 sats

change = total_funding
        - funding_sats
        - fees

change = 2 5000 0000
        - 2 0000 0000
        -      100

change = 4999 9900

as hex: 1cf0fa0200000000

```

Unsigned Funding Transaction:

```

0200000002b932b0669cd0394d0d5bcc27e01ab8c511f1662a6799925b346c0cf18fca03430200000000fdffffffb932
b0669cd0394d0d5bcc27e01ab8c511f1662a6799925b346c0cf18fca03430000000000fdffffff03e5effa0200000000
1600141ca1cca8855bad6bc1ea5436edd8cff10b7e448b1cf0fa020000000016001444cb0c39f93ecc372b5851725bd2
9d865d333b100084d71700000000220020297b92c238163e820b82486084634b4846b86a3c658d87b9384192e6bea98e
c578000000

```

Expected Opener's tx_signatures:

```

tx_signatures{
  channel_id: xxx,
  txid: "5ca4e657c1aa9d069ea4a5d712045d233a7d7c52738cb02993637289e6386057",
  num_witnesses: 1,
  witness[{
    len: 74,
    witness_data:
"022068656c6c6f2074686572652c2074686973206973206120626974636f6e2121212782012088a820add57dfe52770
79d069ca4ad4893c96de91f88ffb981fdc6a2a34d5336c66aff87"
  ]}
}

```

Expected Acceptor's tx_signatures:

```

tx_signatures{
  channel_id: xxx,
  txid: "5ca4e657c1aa9d069ea4a5d712045d233a7d7c52738cb02993637289e6386057",
  num_witnesses: 1,
  witness[{
    len: 107,
    witness_data:
"0247304402207de9ba56bb9f641372e805782575ee840a899e61021c8b1572b3ec1d5b5950e9022069e9ba998915dae

```

```
193d3c25cb89b5e64370e6a3a7755e7f31cf6d7cbc2a49f6d0121034695f5b7864c580bf11f9f8cb1a94eb336f2ce9ef
872d2ae1a90ee276c772484"
    }
}
```

Signed Funding Transaction:

Note locktime is set to 120.

```
02000000000102b932b0669cd0394d0d5bcc27e01ab8c511f1662a6799925b346c0cf18fca03430200000000fdffffff
b932b0669cd0394d0d5bcc27e01ab8c511f1662a6799925b346c0cf18fca0343000000000fdffffff03e5effa020000
00001600141ca1cca8855bad6bc1ea5436edd8cff10b7e448b1cf0fa020000000016001444cb0c39f93ecc372b585172
5bd29d865d333b100084d71700000000220020297b92c238163e820b82486084634b4846b86a3c658d87b9384192e6be
a98ec50247304402207de9ba56bb9f641372e805782575ee840a899e61021c8b1572b3ec1d5b5950e9022069e9ba9989
15dae193d3c25cb89b5e64370e6a3a7755e7f31cf6d7cbc2a49f6d0121034695f5b7864c580bf11f9f8cb1a94eb336f2
ce9ef872d2ae1a90ee276c772484022068656c6c6f2074686572652c2074686973206973206120626974636f6e212121
2782012088a820add57dfe5277079d069ca4ad4893c96de91f88ffb981fdc6a2a34d5336c66aff8778000000
```

Appendix H: Commitment and HTLC Transaction Test Vectors (zero-fee-commitments)

Test vectors [are provided](#) that detail commitment and HTLC transactions created using `zero_fee_commitments` in various scenarios. All parameters are provided to ensure that implementations generate exactly the same signed transactions.

References

Authors

[FIXME:]



This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

Routing Payments

You don't open a channel with everyone you pay. Most Lightning payments bounce through other people's channels, and the sender is the one who picks the path.

BOLT 04 is onion routing. The sender builds a Sphinx onion: each hop can only decrypt its own layer, which tells it the next node, the amount, and the CLTV expiry. Inner hops don't learn the source, the destination, or the full route. Failures come back along the same onion, encrypted so that only the sender can make sense of them.

BOLT 07 is gossip — how you learn the graph that onions are built on. Nodes announce themselves and their channels, sign those announcements, and flood them through the network. Channel updates carry fees and expiry deltas. Without gossip, onion routing has nowhere to go; without onions, gossip is just a map.

Together these two BOLTs are how a payment finds its way. Invoices, in the next chapter, are how you know who to pay and what hash to use.

In this chapter

- BOLT 04 — Onion routing
- BOLT 07 — P2P node and channel discovery (gossip)

BOLT #4: Onion Routing Protocol

Overview

This document describes the construction of an onion routed packet that is used to route a payment from an *origin node* to a *final node*. The packet is routed through a number of intermediate nodes, called *hops*.

The routing schema is based on the [Sphinx](#) construction and is extended with a per-hop payload.

Intermediate nodes forwarding the message can verify the integrity of the packet and can learn which node they should forward the packet to. They cannot learn which other nodes, besides their predecessor or successor, are part of the packet's route; nor can they learn the length of the route or their position within it. The packet is obfuscated at each hop, to ensure that a network-level attacker cannot associate packets belonging to the same route (i.e. packets belonging to the same route do not share any correlating information). Notice that this does not preclude the possibility of packet association by an attacker via traffic analysis.

The route is constructed by the origin node, which knows the public keys of each intermediate node and of the final node. Knowing each node's public key allows the origin node to create a shared secret (using ECDH) for each intermediate node and for the final node. The shared secret is then used to generate a *pseudo-random stream* of bytes (which is used to obfuscate the packet) and a number of *keys* (which are used to encrypt the payload and compute the HMACs). The HMACs are then in turn used to ensure the integrity of the packet at each hop.

Each hop along the route only sees an ephemeral key for the origin node, in order to hide the sender's identity. The ephemeral key is blinded by each intermediate hop before forwarding to the next, making the onions unlinkable along the route.

This specification describes *version 0* of the packet format and routing mechanism.

A node: - upon receiving a higher version packet than it implements: - MUST report a route failure to the origin node. - MUST discard the packet.

Table of Contents

- [Conventions](#)
- [Key Generation](#)
- [Pseudo Random Byte Stream](#)
- [Packet Structure](#)
 - [Payload Format](#)
 - [Basic Multi-Part Payments](#)
- [Route Blinding](#)
 - [Inside encrypted_recipient_data: encrypted_data_tlv](#)
- [Accepting and Forwarding a Payment](#)
 - [Payload for the Last Node](#)

- [Non-strict Forwarding](#)
- [Shared Secret](#)
- [Blinding Ephemeral Onion Keys](#)
- [Packet Construction](#)
- [Onion Decryption](#)
- [Filler Generation](#)
- [Returning Errors](#)
 - [Failure Messages](#)
 - [Receiving Failure Codes](#)
- [Successful Payments](#)
- [Onion Messages](#)
- [max_htlc_cltv Selection](#)
- [Test Vector](#)
 - [Returning Errors](#)
- [References](#)
- [Authors](#)

Conventions

There are a number of conventions adhered to throughout this document:

- HMAC: the integrity verification of the packet is based on Keyed-Hash Message Authentication Code, as defined by the [FIPS 198 Standard](#) / [RFC 2104](#), and using a SHA256 hashing algorithm.
- Elliptic curve: for all computations involving elliptic curves, the Bitcoin curve is used, as specified in [secp256k1](#)
- Pseudo-random stream: [ChaCha20](#) is used to generate a pseudo-random byte stream. For its generation, a fixed 96-bit null-nonce (0x000000000000000000000000) is used, along with a key derived from a shared secret and with a 0x00-byte stream of the desired output size as the message.
- The terms *origin node* and *final node* refer to the initial packet sender and the final packet recipient, respectively.
- The terms *hop* and *node* are sometimes used interchangeably, but a *hop* usually refers to an intermediate node in the route rather than an end node. *origin node* → *hop* → ... → *hop* → *final node*
- The term *processing node* refers to the specific node along the route that is currently processing the forwarded packet.
- The term *peers* refers only to hops that are direct neighbors (in the overlay network): more specifically, *sending peers* forward packets to *receiving peers*.
- Each hop in the route has a variable length *hop_payload*.
 - The variable length *hop_payload* is prefixed with a *bigsize* encoding the length in bytes, excluding the prefix and the trailing HMAC.
- These *hop_payloads* are used to build the combined payload in the packet, also referred to as *hop_payloads* or the mix-header.

Key Generation

A number of encryption and verification keys are derived from the shared secret:

- *rho*: used as key when generating the pseudo-random byte stream that is used to obfuscate the per-hop information

- *mu*: used during the HMAC generation
- *um*: used during error reporting
- *pad*: use to generate random filler bytes for the starting mix-header packet

The key generation function takes a key-type (*rho*=0x72686F, *mu*=0x6d75, *um*=0x756d, or *pad*=0x706164) and a 32-byte secret as inputs and returns a 32-byte key.

Keys are generated by computing an HMAC (with SHA256 as hashing algorithm) using the appropriate key-type (i.e. *rho*, *mu*, *um*, or *pad*) as HMAC-key and the 32-byte shared secret as the message. The resulting HMAC is then returned as the key.

Notice that the key-type does not include a C-style 0x00-termination-byte, e.g. the length of the *rho* key-type is 3 bytes, not 4.

Pseudo Random Byte Stream

The pseudo-random byte stream is used to obfuscate the packet at each hop of the path, so that each hop may only recover the address and HMAC of the next hop. The pseudo-random byte stream is generated by encrypting (using ChaCha20) a 0x00-byte stream, of the required length, which is initialized with a key derived from the shared secret and a 96-bit zero-nonce (0x000000000000000000000000).

The use of a fixed nonce is safe, since the keys are never reused.

Packet Structure

The packet consists of four sections:

- a version byte
- a 33-byte compressed secp256k1 public_key, used during the shared secret generation
- a 1300-byte hop_payloads consisting of multiple, variable length, hop_payload payloads
- a 32-byte hmac, used to verify the packet's integrity

The network format of the packet consists of the individual sections serialized into one contiguous byte-stream and then transferred to the packet recipient. Due to the fixed size of the packet, it need not be prefixed by its length when transferred over a connection.

The overall structure of the packet is as follows:

1. type: onion_packet
2. data:
 - [byte:version]
 - [point:public_key]
 - [1300*byte:hop_payloads]
 - [32*byte:hmac]

For this specification (*version 0*), version has a constant value of 0x00.

The hop_payloads field is a structure that holds obfuscated routing information, and associated HMAC. It is 1300 bytes long and has the following structure:

1. type: hop_payloads

2. data:

- [bigsize:length]
- [length*byte:payload]
- [32*byte:hmac]
- ...
- filler

Where, the length, payload, and hmac are repeated for each hop; and where, filler consists of obfuscated, deterministically-generated padding, as detailed in [Filler Generation](#). Additionally, hop_payloads is incrementally obfuscated at each hop.

Using the payload field, the origin node is able to specify the path and structure of the HTLCs forwarded at each hop. As the payload is protected under the packet-wide HMAC, the information it contains is fully authenticated with each pair-wise relationship between the HTLC sender (origin node) and each hop in the path.

Using this end-to-end authentication, each hop is able to cross-check the HTLC parameters with the payload's specified values and to ensure that the sending peer hasn't forwarded an ill-crafted HTLC.

Since no payload TLV value can ever be shorter than 2 bytes, length values of 0 and 1 are reserved. (0 indicated a legacy format no longer supported, and 1 is reserved for future use).

payload format

This is formatted according to the Type-Length-Value format defined in [BOLT #1](#).

1. tlv_stream: payload

2. types:

1. type: 2 (amt_to_forward)
2. data:
 - [tu64:amt_to_forward]
3. type: 4 (outgoing_cltv_value)
4. data:
 - [tu32:outgoing_cltv_value]
5. type: 6 (short_channel_id)
6. data:
 - [short_channel_id:short_channel_id]
7. type: 8 (payment_data)
8. data:
 - [32*byte:payment_secret]
 - [tu64:total_msat]
9. type: 10 (encrypted_recipient_data)
10. data:
 - [...*byte:encrypted_recipient_data]
11. type: 12 (current_path_key)
12. data:
 - [point:path_key]
13. type: 16 (payment_metadata)

14. data:
 - [...*byte:payment_metadata]
15. type: 18 (total_amount_msat)
16. data:
 - [tu64:total_msat]

short_channel_id is the ID of the outgoing channel used to route the message; the receiving peer should operate the other end of this channel.

amt_to_forward is the amount, in millisatoshis, to forward to the next receiving peer specified within the routing information, or for the final destination.

For non-final nodes, this includes the origin node's computed *fee* for the receiving peer, calculated according to the receiving peer's advertised fee schema (as described in [BOLT #7](#)).

outgoing_cltv_value is the CLTV value that the *outgoing* HTLC carrying the packet should have. Inclusion of this field allows a hop to both authenticate the information specified by the origin node, and the parameters of the HTLC forwarded, and ensure the origin node is using the current cltv_expiry_delta value.

If the values don't correspond, this indicates that either a forwarding node has tampered with the intended HTLC values or that the origin node has an obsolete cltv_expiry_delta value.

The requirements ensure consistency in responding to an unexpected outgoing_cltv_value, whether it is the final node or not, to avoid leaking its position in the route.

Requirements

The creator of encrypted_recipient_data (usually, the recipient of payment):

- MUST create encrypted_data_tlv for each node in the blinded route (including itself).
- MUST include encrypted_data_tlv.payment_relay for each non-final node.
- MUST include exactly one of encrypted_data_tlv.short_channel_id or encrypted_data_tlv.next_node_id for each non-final node.
- MUST set encrypted_data_tlv.payment_constraints for each non-final node and MAY set it for the final node:
 - max_cltv_expiry to the largest block height at which the route is allowed to be used, starting from the final node's chosen max_cltv_expiry height at which the route should expire, adding the final node's min_final_cltv_expiry_delta and then adding encrypted_data_tlv.payment_relay.cltv_expiry_delta at each hop.
 - htlc_minimum_msat to the largest minimum HTLC value the nodes will allow.
- If it sets encrypted_data_tlv.allowed_features:
 - MUST set it to an empty array.
- MUST compute the total fees and CLTV delta of the route as follows and communicate them to the sender:
 - $\text{total_fee_base_msat}(n+1) = (\text{fee_base_msat}(n+1) * 1000000 + \text{total_fee_base_msat}(n) * (1000000 + \text{fee_proportional_millionths}(n+1)) + 1000000 - 1) / 1000000$
 - $\text{total_fee_proportional_millionths}(n+1) = ((\text{total_fee_proportional_millionths}(n) + \text{fee_proportional_millionths}(n+1)) * 1000000 + \text{total_fee_proportional_millionths}(n) * \text{fee_proportional_millionths}(n+1) + 1000000 - 1) / 1000000$

- $\text{total_cltv_delta} = \text{cltv_delta}(0) + \text{cltv_delta}(1) + \dots + \text{cltv_delta}(n) + \text{min_final_cltv_expiry_delta}$
- MUST create the `encrypted_recipient_data` from the `encrypted_data_tlv` as required in [Route Blinding](#).

The writer of the TLV payload:

- For every node inside a blinded route:
 - MUST include the `encrypted_recipient_data` provided by the recipient
 - For the first node in the blinded route:
 - MUST include the `path_key` provided by the recipient in `current_path_key`
 - If it is the final node:
 - MUST include `amt_to_forward`, `outgoing_cltv_value` and `total_amount_msat`.
 - The value set for `outgoing_cltv_value`:
 - MUST use the current block height as a baseline value.
 - if a [random offset](#) was added to improve privacy:
 - SHOULD add the offset to the baseline value.
 - MUST NOT include any other tlv field.
- For every node outside of a blinded route:
 - MUST include `amt_to_forward` and `outgoing_cltv_value`.
 - For every non-final node:
 - MUST include `short_channel_id`
 - MUST NOT include `payment_data`
 - For the final node:
 - MUST NOT include `short_channel_id`
 - if the recipient provided `payment_secret`:
 - MUST include `payment_data`
 - MUST set `payment_secret` to the one provided
 - MUST set `total_msat` to the total amount it will send
 - if the recipient provided `payment_metadata`:
 - MUST include `payment_metadata` with every HTLC
 - MUST not apply any limits to the size of `payment_metadata` except the limits implied by the fixed onion size

The reader:

- If `encrypted_recipient_data` is present:
 - If `path_key` is set in the incoming `update_add_htlc`:
 - MUST return an error if `current_path_key` is present.
 - MUST use that `path_key` as `path_key` for decryption.
 - Otherwise:
 - MUST return an error if `current_path_key` is not present.
 - MUST use that `current_path_key` as the `path_key` for decryption.
 - SHOULD add a random delay before returning errors.
 - MUST return an error if `encrypted_recipient_data` does not decrypt using the `path_key` as described in [Route Blinding](#).

- If `payment_constraints` is present:
 - MUST return an error if:
 - the expiry is greater than `encrypted_recipient_data.payment_constraints.max_cltv_expiry`.
 - the amount is below `encrypted_recipient_data.payment_constraints.htlc_minimum_msat`.
- If `allowed_features` is missing:
 - MUST process the message as if it were present and contained an empty array.
- MUST return an error if:
 - `encrypted_recipient_data.allowed_features.features` contains an unknown feature bit (even if it is odd).
 - `encrypted_recipient_data` contains both `short_channel_id` and `next_node_id`.
 - the payment uses a feature not included in `encrypted_recipient_data.allowed_features.features`.
- If it is not the final node:
 - MUST return an error if the payload contains other tlv fields than `encrypted_recipient_data` and `current_path_key`.
 - MUST return an error if `encrypted_recipient_data` does not contain either `short_channel_id` or `next_node_id`.
 - MUST return an error if `encrypted_recipient_data` does not contain `payment_relay`.
 - MUST use values from `encrypted_recipient_data.payment_relay` to calculate `amt_to_forward` and `outgoing_cltv_value` as follows:
 - $\text{amt_to_forward} = ((\text{amount_msat} - \text{fee_base_msat}) * 1000000 + 1000000 + \text{fee_proportional_millionths} - 1) / (1000000 + \text{fee_proportional_millionths})$
 - $\text{outgoing_cltv_value} = \text{cltv_expiry} - \text{payment_relay.cltv_expiry_delta}$
- If it is the final node:
 - MUST return an error if the payload contains other tlv fields than `encrypted_recipient_data`, `current_path_key`, `amt_to_forward`, `outgoing_cltv_value` and `total_amount_msat`.
 - MUST return an error if `amt_to_forward`, `outgoing_cltv_value` or `total_amount_msat` are not present.
 - MUST return an error if `amt_to_forward` is below what it expects for the payment.
 - MUST return an error if `incoming_cltv_expiry < outgoing_cltv_value`.
 - MUST return an error if `incoming_cltv_expiry < current_block_height + min_final_cltv_expiry_delta`.
- Otherwise (it is not part of a blinded route):
 - MUST return an error if `path_key` is set in the incoming `update_add_htlc` or `current_path_key` is present.
 - MUST return an error if `amt_to_forward` or `outgoing_cltv_value` are not present.
 - if it is not the final node:
 - MUST return an error if:
 - `short_channel_id` is not present,
 - it cannot forward the HTLC to the peer indicated by the channel `short_channel_id`.
 - `incoming amount_msat - fee < amt_to_forward` (where `fee` is the advertised fee as described in [BOLT #7](#))
 - $\text{cltv_expiry} - \text{cltv_expiry_delta} < \text{outgoing_cltv_value}$

- If it is the final node:
 - MUST return an error if:
 - `total_msat` is not present
 - `incoming amount_msat < amt_to_forward`.
 - `incoming cltv_expiry < outgoing_cltv_value`.
 - `incoming cltv_expiry < current_block_height + min_final_cltv_expiry_delta`.

Additional requirements are specified [here](#) for multi-part payments, and [here](#) for blinded payments.

Basic Multi-Part Payments

An HTLC may be part of a larger “multi-part” payment: such “base” atomic multipath payments will use the same `payment_hash` for all paths.

Note that `amt_to_forward` is the amount for this HTLC only: a `total_msat` field containing a greater value is a promise by the ultimate sender that the rest of the payment will follow in succeeding HTLCs; we call these outstanding HTLCs which have the same preimage, an “HTLC set”.

Note that there are two distinct tlv fields that can be used to transmit `total_msat`. The last one, `total_amount_msat`, was introduced with blinded paths for which the `payment_secret` doesn’t make sense.

`payment_metadata` is to be included in every payment part, so that invalid payment details can be detected as early as possible.

Requirements

The writer: - if the invoice offers the `basic_mpp` feature: - MAY send more than one HTLC to pay the invoice. - MUST use the same `payment_hash` on all HTLCs in the set. - SHOULD send all payments at approximately the same time. - SHOULD try to use diverse paths to the recipient for each HTLC. - SHOULD retry and/or re-divide HTLCs which fail. - if the invoice specifies an amount: - MUST set `total_msat` to at least that amount, and less than or equal to twice amount. - otherwise: - MUST set `total_msat` to the amount it wishes to pay. - MUST ensure that the total `amt_to_forward` of the HTLC set which arrives at the payee is equal to or greater than `total_msat`. - MUST NOT send another HTLC if the total `amt_to_forward` of the HTLC set is already greater or equal to `total_msat`. - MUST include `payment_secret`. - otherwise: - MUST set `total_msat` equal to `amt_to_forward`.

The final node: - MUST fail the HTLC if dictated by Requirements under [Failure Messages](#) - Note: “amount paid” specified there is the `total_msat` field. - if it does not support `basic_mpp`: - MUST fail the HTLC if `total_msat` is not exactly equal to `amt_to_forward`. - otherwise, if it supports `basic_mpp`: - MUST add it to the HTLC set corresponding to that `payment_hash`. - SHOULD fail the entire HTLC set if `total_msat` is not the same for all HTLCs in the set. - if the total `amt_to_forward` of this HTLC set is equal to or greater than `total_msat`: - SHOULD fulfill all HTLCs in the HTLC set - otherwise, if the total `amt_to_forward` of this HTLC set is less than `total_msat`: - MUST NOT fulfill any HTLCs in the HTLC set - MUST fail all HTLCs in the HTLC set after some reasonable timeout. - SHOULD wait for at least 60 seconds after the initial HTLC. - SHOULD use `mpp_timeout` for the failure message. - MUST require `payment_secret` for all HTLCs in the set. - if it fulfills any HTLCs in the HTLC set: - MUST fulfill the entire HTLC set.

Rationale

If `basic_mpp` is present it causes a delay to allow other partial payments to combine. The total amount must be sufficient for the desired payment, just as it must be for single payments. But this must be reasonably bounded to avoid a denial-of-service.

Because invoices do not necessarily specify an amount, and because payers can add noise to the final amount, the total amount must be sent explicitly. The requirements allow exceeding this slightly, as it simplifies adding noise to the amount when splitting, as well as scenarios in which the senders are genuinely independent (friends splitting a bill, for example).

Because a node may need to pay more than its desired amount (due to the `htlc_minimum_msat` value of channels in the desired path), nodes are allowed to pay more than the `total_msat` they specified. Otherwise, nodes would be constrained in which paths they can take when retrying payments along specific paths. However, no individual HTLC may be for less than the difference between the total paid and `total_msat`.

The restriction on sending an HTLC once the set is over the agreed total prevents the preimage being released before all the partial payments have arrived: that would allow any intermediate node to immediately claim any outstanding partial payments.

An implementation may choose not to fulfill an HTLC set which otherwise meets the amount criterion (eg. some other failure, or invoice timeout), however if it were to fulfill only some of them, intermediary nodes could simply claim the remaining ones.

Route Blinding

1. subtype: `blinded_path`
2. data:
 - `[scid:dir_or_pubkey:first_node_id]`
 - `[point:first_path_key]`
 - `[byte:num_hops]`
 - `[num_hops*blinded_path_hop:path]`
3. subtype: `blinded_path_hop`
4. data:
 - `[point:blinded_node_id]`
 - `[u16:enclen]`
 - `[enclen*byte:encrypted_recipient_data]`

A blinded path consists of: 1. an initial introduction point (`first_node_id`) 2. an initial key to share a secret with the first node (`first_path_key`) 3. a series of tweaked node ids (`path.blinded_node_id`) 4. a series of binary blobs encrypted to the nodes (`path.encrypted_recipient_data`) to tell them the next hop.

For example, Dave wants Alice to reach him via public node Bob then Carol. He creates a chain of public keys ("path_keys") for Bob, Carol and finally himself, so he can share a secret with each of them. These keys are a simple chain, so each node can derive the next path_key without having to be told explicitly.

From these shared secrets, Dave creates and encrypts three `encrypted_data_tlvs`: 1. `encrypted_data_bob`: For Bob to tell him to forward to Carol 2. `encrypted_data_carol`: For Carol to tell her to forward to him 3. `encrypted_data_dave`: For himself to indicate the path was used, and any metadata he wants.

To mask the node ids, he also derives three blinding factors from the shared secrets, which turn Bob into Bob', Carol into Carol' and Dave into Dave'.

So this is the blinded_path he hands to Alice.

1. first_node_id: Bob
2. first_path_key: the first path key for Bob
3. path: [Bob', encrypted_data_bob], [Carol', encrypted_data_carol], [Dave', encrypted_data_dave]

There are two different ways for Alice to construct an onion which gets to Bob (since he's probably not a direct peer of hers) which are described in the requirements below.

But after Bob the path is always the same: he will send Carol the path_key he derived, along with the onion. She will use the path_key to derive the tweak for the onion (which Alice encrypted for Carol' not Carol) so she can decrypt it, and also to derive the key to decrypt encrypted_data_tlv which will tell her to forward to Dave (and possibly additional restrictions Dave specified).

Requirements

Note that the creator of the blinded path (i.e. the recipient) is creating it for the sender to use to create an onion, and for the intermediate nodes to read the instructions, hence there are two reader sections here.

The writer of a blinded_path:

- MUST create a viable path to itself (N_r) i.e. $N_0 \rightarrow N_1 \rightarrow \dots \rightarrow N_r$.
- MUST set first_node_id to N_0
- MUST create a series of ECDH shared secrets for each node in the route using the following algorithm:
 - $e_0 \leftarrow \{0;1\}^{256}$ (e_0 SHOULD be obtained via CSPRNG)
 - $E_0 = e_0 \cdot G$
 - For every node in the route:
 - let $N_i = k_i \cdot G$ be the node_id (k_i is N_i 's private key)
 - $ss_i = \text{SHA256}(e_i \cdot N_i) = \text{SHA256}(k_i \cdot E_i)$ (ECDH shared secret known only by N_r and N_i)
 - $\rho_i = \text{HMAC256}(\text{"rho"}, ss_i)$ (key used to encrypt encrypted_recipient_data for N_i by N_r)
 - $e_{i+1} = \text{SHA256}(E_i \parallel ss_i) \cdot e_i$ (ephemeral private path key, only known by N_r)
 - $E_{i+1} = \text{SHA256}(E_i \parallel ss_i) \cdot E_i$ (path_key. NB: N_i MUST NOT learn e_i)
- MUST set first_path_key to E_0
- MUST create a series of blinded node IDs B_i for each node using the following algorithm:
 - $B_i = \text{HMAC256}(\text{"blinded_node_id"}, ss_i) \cdot N_i$ (blinded node_id for N_i , private key known only by N_i)
 - MUST set blinded_node_id for each blinded_path_hop in path to B_i
- MAY replace E_{i+1} with a different value, but if it does:
 - MUST set encrypted_data_tlv[i].next_path_key_override to E_{i+1}
- MAY store private data in encrypted_data_tlv[r].path_id to verify that the route is used in the right context and was created by them
- SHOULD add padding data to ensure all encrypted_data_tlv[i] have the same length

- MUST encrypt each `encrypted_data_tlv[i]` with ChaCha20-Poly1305 using the corresponding $\text{\$}\rho_i\text{\$}$ key and an all-zero nonce to produce `encrypted_recipient_data[i]`
- MAY add additional “dummy” hops at the end of the path (which it will ignore on receipt) to obscure the path length.

The reader of the `blinded_path`: - MUST prepend its own onion payloads to reach the `first_node_id` - MUST include the corresponding `encrypted_recipient_data` in each onion payload within path - For the first entry in path: - if it is sending a payment: - SHOULD create an unblinded onion payment to `first_node_id`, and include `first_path_key` as `current_path_key`. - otherwise: - MUST encrypt the first blinded path onion to the first blinded `node_id`. - MUST set `next_path_key_override` in the prior onion payload to `first_path_key`. - For each successive entry in path: - MUST encrypt the onion to the corresponding blinded `node_id`.

The reader of the `encrypted_recipient_data`:

- MUST compute:
 - $\text{\$}\text{ss}_i = \text{SHA256}(k_i * E_i)\text{\$}$ (standard ECDH)
 - $\text{\$}\text{b}_i = \text{HMAC256}(\text{text}\{\text{"blinded_node_id"}\}, \text{ss}_i) * k_i\text{\$}$
 - $\text{\$}\rho_i = \text{HMAC256}(\text{text}\{\text{"rho"}\}, \text{ss}_i)\text{\$}$
- MUST decrypt the `encrypted_recipient_data` field using $\text{\$}\rho_i\text{\$}$ as a key using ChaCha20-Poly1305 and an all-zero nonce key.
- If the `encrypted_recipient_data` field is missing, cannot be decrypted into an `encrypted_data_tlv` or contains unknown even fields:
 - MUST return an error
- If the `encrypted_data_tlv` contains a `next_path_key_override`:
 - MUST use it as the next `path_key`.
- Otherwise:
 - MUST use $\text{\$}E_{i+1} = \text{SHA256}(E_i || \text{ss}_i) * E_i\text{\$}$ as the next `path_key`
- MUST forward the onion and include the next `path_key` in the lightning message for the next node
- If it is the final recipient:
 - MUST ignore the message if the `path_id` does not match the blinded route it created for this purpose

Rationale

Route blinding is a lightweight technique to provide recipient anonymity. It’s more flexible than rendezvous routing because it simply replaces the public keys of the nodes in the route with random public keys while letting senders choose what data they put in the onion for each hop. Blinded routes are also reusable in some cases (e.g. onion messages).

Each node in the blinded route needs to receive $\text{\$}E_i\text{\$}$ to be able to decrypt the onion and the `encrypted_recipient_data` payload.

When concatenating two blinded routes generated by different nodes, the last node of the first route needs to know the first `path_key` of the second route: the `next_path_key_override` field must be used to transmit this information. In theory this method could be used for payments (not just onion messages), but we recommend using an unblinded path to reach the `first_node_id` and using `current_path_key` there: this means that the node can tell it is being used as an introductory point, but also does not require blinded path support on the nodes to reach that point, and gives meaningful errors on the unblinded part of the payment.

The final recipient must verify that the blinded route is used in the right context (e.g. for a specific payment) and was created by them. Otherwise a malicious sender could create different blinded routes to all the nodes that they suspect could be the real recipient and try them until one accepts the message. The recipient can protect against that by storing `$`E_r`$` and the context (e.g. a `payment_hash`), and verifying that they match when receiving the onion. Otherwise, to avoid additional storage cost, it can put some private context information in the `path_id` field (e.g. the `payment_preimage`) and verify that when receiving the onion. Note that it's important to use private information in that case, that senders cannot have access to.

Whenever the introduction point receives a failure from the blinded route, it should add a random delay before forwarding the error. Failures are likely to be probing attempts and message timing may help the attacker infer its distance to the final recipient.

Note that nodes in the blinded route return failures through `update_fail_malformed_htlc` and therefore do not and can not provide timing information via attribution data to the sender.

The padding field can be used to ensure that all `encrypted_recipient_data` have the same length. It's particularly useful when adding dummy hops at the end of a blinded route, to prevent the sender from figuring out which node is the final recipient.

When route blinding is used for payments, the recipient specifies the fees and expiry that blinded nodes should apply to the payment instead of letting the sender configure them. The recipient also adds additional constraints to the payments that can go through that route to protect against probing attacks that would let malicious nodes unblind the identity of the blinded nodes. It should set `payment_constraints.max_cltv_expiry` to restrict the lifetime of a blinded route and reduce the risk that an intermediate node updates its fees and rejects payments (which could be used to unblind nodes inside the route).

Inside `encrypted_recipient_data`: `encrypted_data_tlv`

The `encrypted_recipient_data` is a TLV stream, encrypted for a given blinded node, that may contain the following TLV fields:

1. `tlv_stream`: `encrypted_data_tlv`
2. types:
 1. type: 1 (padding)
 2. data:
 - `[...*byte:padding]`
 3. type: 2 (`short_channel_id`)
 4. data:
 - `[short_channel_id:short_channel_id]`
 5. type: 4 (`next_node_id`)
 6. data:
 - `[point:node_id]`
 7. type: 6 (`path_id`)
 8. data:
 - `[...*byte:data]`
 9. type: 8 (`next_path_key_override`)

- 10. data:
 - [point:path_key]
- 11. type: 10 (payment_relay)
- 12. data:
 - [u16:cltv_expiry_delta]
 - [u32:fee_proportional_millionths]
 - [tu32:fee_base_msat]
- 13. type: 12 (payment_constraints)
- 14. data:
 - [u32:max_cltv_expiry]
 - [tu64:htlc_minimum_msat]
- 15. type: 14 (allowed_features)
- 16. data:
 - [...*byte:features]

Rationale

Encrypted recipient data is created by the final recipient to give to the sender, containing instructions for the node on how to handle the message (it can also be created by the sender themselves: the node forwarding cannot tell). It's used in both payment onions and onion messages onions. See [Route Blinding](#).

Accepting and Forwarding a Payment

Once a node has decoded the payload it either accepts the payment locally, or forwards it to the peer indicated as the next hop in the payload.

Non-strict Forwarding

A node MAY forward an HTLC along an outgoing channel other than the one specified by `short_channel_id`, so long as the receiver has the same node public key intended by `short_channel_id`. Thus, if `short_channel_id` connects nodes A and B, the HTLC can be forwarded across any channel connecting A and B. Failure to adhere will result in the receiver being unable to decrypt the next hop in the onion packet.

Rationale

In the event that two peers have multiple channels, the downstream node will be able to decrypt the next hop payload regardless of which channel the packet is sent across.

Nodes implementing non-strict forwarding are able to make real-time assessments of channel bandwidths with a particular peer, and use the channel that is locally-optimal.

For example, if the channel specified by `short_channel_id` connecting A and B does not have enough bandwidth at forwarding time, then A is able use a different channel that does. This can reduce payment latency by preventing the HTLC from failing due to bandwidth constraints across `short_channel_id`, only to have the sender attempt the same route differing only in the channel between A and B.

Non-strict forwarding allows nodes to make use of private channels connecting them to the receiving node, even if the channel is not known in the public channel graph.

Recommendation

Implementations using non-strict forwarding should consider applying the same fee schedule to all channels with the same peer, as senders are likely to select the channel which results in the lowest overall cost. Having distinct policies may result in the forwarding node accepting fees based on the most optimal fee schedule for the sender, even though they are providing aggregate bandwidth across all channels with the same peer.

Alternatively, implementations may choose to apply non-strict forwarding only to like-policy channels to ensure their expected fee revenue does not deviate by using an alternate channel.

Payload for the Last Node

When building the route, the origin node **MUST** use a payload for the final node with the following values:

- `payment_secret`: set to the payment secret specified by the recipient (e.g. `payment_secret` from a [BOLT #11](#) payment invoice)
- `outgoing_cltv_value`: set to the final expiry specified by the recipient (e.g. `min_final_cltv_expiry_delta` from a [BOLT #11](#) payment invoice)
- `amt_to_forward`: set to the final amount specified by the recipient (e.g. `amount` from a [BOLT #11](#) payment invoice)

This allows the final node to check these values and return errors if needed, but it also eliminates the possibility of probing attacks by the second-to-last node. Such attacks could, otherwise, attempt to discover if the receiving peer is the last one by re-sending HTLCs with different amounts/expiries. The final node will extract its onion payload from the HTLC it has received and compare its values against those of the HTLC. See the [Returning Errors](#) section below for more details.

If not for the above, since it need not forward payments, the final node could simply discard its payload.

Shared Secret

The origin node establishes a shared secret with each hop along the route using Elliptic-curve Diffie-Hellman between the sender's ephemeral key at that hop and the hop's node ID key. The resulting curve point is serialized to the compressed format and hashed using SHA256. The hash output is used as the 32-byte shared secret.

Elliptic-curve Diffie-Hellman (ECDH) is an operation on an EC private key and an EC public key that outputs a curve point. For this protocol, the ECDH variant implemented in `libsecp256k1` is used, which is defined over the `secp256k1` elliptic curve. During packet construction, the sender uses the ephemeral private key and the hop's public key as inputs to ECDH, whereas during packet forwarding, the hop uses the ephemeral public key and its own node ID private key. Because of the properties of ECDH, they will both derive the same value.

Blinding Ephemeral Onion Keys

In order to ensure multiple hops along the route cannot be linked by the ephemeral public keys they see, the key is blinded at each hop. The blinding is done in a deterministic way that allows the sender to compute the corresponding blinded private keys during packet construction.

The blinding of an EC public key is a single scalar multiplication of the EC point representing the public key with a 32-byte blinding factor. Due to the commutative property of scalar multiplication, the blinded private key is the multiplicative product of the input's corresponding private key with the same blinding factor.

The blinding factor itself is computed as a function of the ephemeral public key and the 32-byte shared secret. Concretely, it is the SHA256 hash value of the concatenation of the public key serialized in its compressed format and the shared secret.

Packet Construction

In the following example, it's assumed that a *sending node* (origin node), n_0 , wants to route a packet to a *receiving node* (final node), n_r . First, the sender computes a route $\{n_0, n_1, \dots, n_{r-1}, n_r\}$, where n_0 is the sender itself and n_r is the final recipient. All nodes n_i and n_{i+1} MUST be peers in the overlay network route. The sender then gathers the public keys for n_1 to n_r and generates a random 32-byte sessionkey. Optionally, the sender may pass in *associated data*, i.e. data that the packet commits to but that is not included in the packet itself. Associated data will be included in the HMACs and must match the associated data provided during integrity verification at each hop. The `payment_hash` is commonly used for this purpose.

To construct the onion, the sender initializes the ephemeral private key for the first hop ek_1 to the sessionkey and derives from it the corresponding ephemeral public key epk_1 by multiplying with the secp256k1 base point. For each of the k hops along the route, the sender then iteratively computes the shared secret ss_k and ephemeral key for the next hop ek_{k+1} as follows:

- The sender executes ECDH with the hop's public key and the ephemeral private key to obtain a curve point, which is hashed using SHA256 to produce the shared secret ss_k .
- The blinding factor is the SHA256 hash of the concatenation between the ephemeral public key epk_k and the shared secret ss_k .
- The ephemeral private key for the next hop ek_{k+1} is computed by multiplying the current ephemeral private key ek_k by the blinding factor.
- The ephemeral public key for the next hop epk_{k+1} is derived from the ephemeral private key ek_{k+1} by multiplying with the base point.

Once the sender has all the required information above, it can construct the packet. Constructing a packet routed over r hops requires r 32-byte ephemeral public keys, r 32-byte shared secrets, r 32-byte blinding factors, and r variable length hop_payload payloads. The construction returns a single 1366-byte packet along with the first receiving peer's address.

The packet construction is performed in the reverse order of the route, i.e. the last hop's operations are applied first.

The packet is initialized with 1300 *random* bytes derived from a CSPRNG (ChaCha20). The *pad* key referenced above is used to extract additional random bytes from a ChaCha20 stream, using it as a CSPRNG for this purpose. Once the paddingKey has been obtained, ChaCha20 is used with an all zero nonce, to generate 1300 random bytes. Those random bytes are then used as the starting state of the mix-header to be created.

A filler is generated (see [Filler Generation](#)) using the shared secret.

For each hop in the route, in reverse order, the sender applies the following operations:

- The *rho*-key and *mu*-key are generated using the hop's shared secret.
- *shift_size* is defined as the length of the *hop_payload* plus the bigsize encoding of the length and the length of that HMAC. Thus if the payload length is *l* then the *shift_size* is $1 + l + 32$ for $l < 253$, otherwise $3 + l + 32$ due to the bigsize encoding of *l*.
- The *hop_payloads* field is right-shifted by *shift_size* bytes, discarding the last *shift_size* bytes that exceed its 1300-byte size.
- The bigsize-serialized length, serialized *hop_payload* and *hmac* are copied into the *shift_size* bytes as the beginning of the mix-header.
- The *rho*-key is used to generate a 1300 byte pseudo-random stream which is then applied, with XOR, to the *hop_payloads* field.
- If this is the last hop, i.e. the first iteration, then the tail of the *hop_payloads* field is overwritten with the routing information filler.
- The next HMAC is computed (with the *mu*-key as HMAC-key) over the concatenated *hop_payloads* and associated data.

The resulting final HMAC value is the HMAC that will be used by the first receiving peer in the route.

The packet generation returns a serialized packet that contains the version byte, the ephemeral pubkey for the first hop, the HMAC for the first hop, and the obfuscated *hop_payloads*.

The following Go code is an example implementation of the packet construction:

```
const RoutingInfoSize int32 = 1300

func NewOnionPacket(paymentPath []*btcec.PublicKey, sessionKey *btcec.PrivateKey,
    hopsData []HopData, assocData []byte) (*OnionPacket, error) {

    numHops := len(paymentPath)
    hopSharedSecrets := make([][sha256.Size]byte, numHops)

    // Initialize ephemeral key for the first hop to the session key.
    var ephemeralKey big.Int
    ephemeralKey.Set(sessionKey.D)

    for i := 0; i < numHops; i++ {
        // Perform ECDH and hash the result.
        ecdhResult := scalarMult(paymentPath[i], ephemeralKey)
        hopSharedSecrets[i] = sha256.Sum256(ecdhResult.SerializeCompressed())

        // Derive ephemeral public key from private key.
        ephemeralPrivKey := btcec.PrivKeyFromBytes(btcec.S256(), ephemeralKey.Bytes())
        ephemeralPubKey := ephemeralPrivKey.PubKey()

        // Compute blinding factor.
        sha := sha256.New()
        sha.Write(ephemeralPubKey.SerializeCompressed())
        sha.Write(hopSharedSecrets[i])

        var blindingFactor big.Int
        blindingFactor.SetBytes(sha.Sum(nil))
    }
}
```

```

    // Blind ephemeral key for next hop.
    ephemeralKey.Mul(&ephemeralKey, &blindingFactor)
    ephemeralKey.Mod(&ephemeralKey, btcec.S256().Params().N)
}

// Generate the padding, called "filler strings" in the paper.
filler := generateHeaderPadding(numHops, hopSharedSecrets)

// Allocate and initialize fields to zero-filled slices
var mixHeader [RoutingInfoSize]byte
var nextHmac [32]byte

// Our starting packet needs to be filled out with random bytes, we
// generate some deterministically using the session private key.
paddingKey := generateKey("pad", sessionKey.Serialize())
paddingBytes := generateCipherStream(paddingKey, RoutingInfoSize)
copy(mixHeader[:], paddingBytes)

// Compute the routing information for each hop along with a
// MAC of the routing information using the shared key for that hop.
for i := numHops - 1; i >= 0; i-- {
    rhoKey := generateKey("rho", hopSharedSecrets[i])
    muKey := generateKey("mu", hopSharedSecrets[i])

    hopsData[i].HMAC = nextHmac

    // Shift and obfuscate routing information
    streamBytes := generateCipherStream(rhoKey, RoutingInfoSize)

    rightShift(mixHeader[:], hopDataSize)
    buf := &bytes.Buffer{}
    hopsData[i].Encode(buf)
    copy(mixHeader[:], buf.Bytes())
    xor(mixHeader[:], mixHeader[:], streamBytes[:RoutingInfoSize])

    // These need to be overwritten, so every node generates a correct padding
    if i == numHops-1 {
        copy(mixHeader[len(mixHeader)-len(filler):], filler)
    }

    packet := append(mixHeader[:], assocData...)
    nextHmac = calcMac(muKey, packet)
}

packet := &OnionPacket{
    Version:      0x00,
    EphemeralKey: sessionKey.PubKey(),
    RoutingInfo:  mixHeader,
    HeaderMAC:    nextHmac,
}
return packet, nil
}

```

Onion Decryption

There are two kinds of onion_packet we use:

1. onion_routing_packet in update_add_htlc for payments, which contains a payload TLV (see [Adding an HTLC](#))
2. onion_message_packet in onion_message for messages, which contains an onionmsg_tlv TLV (see [Onion Messages](#))

Those sections specify the associated_data to use, the path_key (if any), the extracted payload format and handling (including how to determine the next peer, if any), and how to handle errors. The processing itself is identical.

Requirements

A reader: - if version is not 0: - MUST abort processing the packet and fail. - if public_key is not a valid pubkey: - MUST abort processing the packet and fail. - if the onion is for a payment: - if hmac has previously been received: - if the preimage is known: - MAY immediately redeem the HTLC using the preimage. - otherwise: - MUST abort processing the packet and fail. - if path_key is specified: - Calculate the blinding_ss as ECDH(path_key, node_privkey). - Either: - Tweak public_key by multiplying by $\$`HMAC256(\text{"blinded\_node\_id"}, \text{blinding\_ss})`$$. - or (equivalently): - Tweak its own node_privkey below by multiplying by $\$`HMAC256(\text{"blinded\_node\_id"}, \text{blinding\_ss})`$$. - Derive the shared secret ss as ECDH(public_key, node_privkey) (see [Shared Secret](#)). - Derive mu as $\$`HMAC256(\text{"mu"}, \text{ss})`$$ (see [Key Generation](#)). - Derive the HMAC as $\$`HMAC256(\text{mu}, \text{hop\_payloads} || \text{associated\_data})`$$. - MUST use a constant time comparison of the computed HMAC and hmac. - If the computed HMAC and hmac differ: - MUST abort processing the packet and fail. - Derive rho as $\$`HMAC256(\text{"rho"}, \text{ss})`$$ (see [Key Generation](#)). - Derive bytestream of twice the length of hop_payloads using rho (see [Pseudo Random Byte Stream](#)). - Set unwrapped_payloads to the XOR of hop_payloads and bytestream. - Remove a bigsize from the front of unwrapped_payloads as payload_length. If that is malformed: - MUST abort processing the packet and fail. - If the payload_length is less than two: - MUST abort processing the packet and fail. - If there are fewer than payload_length bytes remaining in unwrapped_payloads: - MUST abort processing the packet and fail. - Remove payload_length bytes from the front of unwrapped_payloads, as the current payload. - If there are fewer than 32 bytes remaining in unwrapped_payloads: - MUST abort processing the packet and fail. - Remove 32 bytes as next_hmac from the front of unwrapped_payloads. - If unwrapped_payloads is smaller than hop_payloads: - MUST abort processing the packet and fail. - If next_hmac is not all-zero (not the final node): - Derive blinding_tweak as $\$`SHA256(\text{public\_key} || \text{ss})`$$ (see [Blinding Ephemeral Onion Keys](#)). - SHOULD forward an onion to the next peer with: - version set to 0. - public_key set to the incoming public_key multiplied by blinding_tweak. - hop_payloads set to the unwrapped_payloads, truncated to the incoming hop_payloads size. - hmac set to next_hmac. - If it cannot forward: - MUST fail. - Otherwise (all-zero next_hmac): - This is the final destination of the onion.

Rationale

In the case where blinded paths are used, the sender did not actually encrypt this onion for our node_id, but for a tweaked version: we can derive the tweak used from path_key which is given alongside the onion. Then we either tweak our node private key the same way to decrypt the onion, or tweak to the onion ephemeral key which is mathematically equivalent.

Filler Generation

Upon receiving a packet, the processing node extracts the information destined for it from the route information and the per-hop payload. The extraction is done by deobfuscating and left-shifting the field. This would make the field shorter at each hop, allowing an attacker to deduce the route length. For this reason, the field is pre-padded before forwarding. Since the padding is part of the HMAC, the origin node will have to pre-generate an identical padding (to that which each hop will generate) in order to compute the HMACs correctly for each hop. The filler is also used to pad the field-length, in the case that the selected route is shorter than 1300 bytes.

Before deobfuscating the hop_payloads, the processing node pads it with 1300 0x00-bytes, such that the total length is 2*1300. It then generates the pseudo-random byte stream, of matching length, and applies it with XOR to the hop_payloads. This deobfuscates the information destined for it, while simultaneously obfuscating the added 0x00-bytes at the end.

In order to compute the correct HMAC, the origin node has to pre-generate the hop_payloads for each hop, including the incrementally obfuscated padding added by each hop. This incrementally obfuscated padding is referred to as the filler. Keep in mind that while the mix-header is generated in reverse-route order, the filler is generated in route order.

The following example code shows how the filler is generated in Go:

```
const (
    // The mix-header in Lightning is a max of 1300 bytes
    NumMaxHops int = 20
    HopSize    = 65
)

func generateFiller(numHops int, sharedSecrets [][]sharedSecretSize]byte) []byte {
    fillerSize := uint((NumMaxHops + 1) * HopSize)
    filler := make([]byte, fillerSize)

    // The last hop does not obfuscate, it's not forwarding anymore.
    for i := 0; i < numHops-1; i++ {

        // Left-shift the field
        copy(filler[:], filler[HopSize:])

        // Zero-fill the last hop
        copy(filler[len(filler)-HopSize:], bytes.Repeat([]byte{0x00}, HopSize))

        // Generate pseudo-random byte stream
        streamKey := generateKey("rho", sharedSecrets[i])
        streamBytes := generateCipherStream(streamKey, fillerSize)

        // Obfuscate
        xor(filler, filler, streamBytes)
    }

    // Cut filler down to the correct length (numHops+1)*HopSize
    // bytes will be prepended by the packet generation.
    return filler[(NumMaxHops-numHops+2)*HopSize:]
}
```


Note that this example implementation is for demonstration purposes only; the filler can be generated much more efficiently. The last hop need not obfuscate the filler, since it won't forward the packet any further and thus need not extract an HMAC either.

Returning Errors

The onion routing protocol includes a mechanism for returning encrypted error messages to the origin node. The returned error messages may be failures reported by any hop, including the final node. The format of the forward packet is not usable for the return path, since no hop besides the origin has access to the information required for its generation. Note that these error messages are not reliable, as they are not placed on-chain due to the possibility of hop failure.

Return packets are limited to 32768 bytes (32 KiB), leaving room in `update_fail_htlc` for `attribution_data` and future extensions. Earlier versions permitted larger return packets, so intermediate nodes truncate them to their first 32768 bytes instead of rejecting the corresponding message.

Intermediate hops store the shared secret from the forward path and reuse it to authenticate and obfuscate any corresponding return packet during each hop. In addition, each node locally stores data regarding its own sending peer in the route, so it knows where to return-forward any eventual return packets.

Erring node

The node generating the error message builds a *return packet* consisting of the following fields:

1. data:
 - [32*byte:hmac]
 - [u16:failure_len]
 - [failure_len*byte:failuremsg]
 - [u16:pad_len]
 - [pad_len*byte:pad]

Where `hmac` is an HMAC authenticating the remainder of the packet, with a key generated using the above process, with key type `um`, `failuremsg` as defined below, and `pad` as the extra bytes used to conceal length.

The erring node then generates a new key, using the key type `ammag`. This key is then used to generate a pseudo-random stream, which is in turn applied to the packet using XOR.

Error handling for HTLCs with `path_key` is particularly fraught, since differences in implementations (or versions) may be leveraged to de-anonymize elements of the blinded path. Thus the decision turn every error into `invalid_onion_blinding` which will be converted to a normal onion error by the introduction point.

Initialization of `attribution_data`

For the layout of `attribution_data`, see [Removing an HTLC: `update\_fulfill\_htlc`, `update\_fail\_htlc`, and `update\_fail\_malformed\_htlc`](#).

The `htlc_hold_times` field specifies the `htlc` hold time for each hop in units of 100 milliseconds. For example, a value of 3 represents 300 ms. Nodes along the path that lack accurate timing information may report a value of zero.

The erring node puts its hold time at the start of this array and zeroes out the rest. The size of the field is based on the maximum supported number of hops in a route (20).

The field `truncated_hmacs` contains truncated authentication codes series for each hop, with the same um key that is used for hmac in the return packet. Regular 32 byte HMACs are truncated to the first 4 bytes to save space.

In theory this truncation makes it possible for malicious nodes to guess the HMAC. However, game theory is against them because a wrong guess will get them penalized in future pathfinding.

The size of the field is based on the maximum number of hops in a route (20) and the truncated HMAC size (4 bytes). Each hop adds 20 HMACs, one for each possible position that the hop could be at in the path. This is necessary because only the sender knows the position of each hop in the path.

It is not required to store all $20 * 20 = 400$ HMACs. The node in position 0 needs to store an HMAC for every one of the 20 positions. For the node in position 1, the maximum remaining route length is only 19. For that position just 19 HMACs need to be kept, and so on. This makes for a total number of HMACs of $20+19+18+\dots+1 = 210$ HMACs.

The layout of the `hmacs` field is shown below. The actual format is much longer, but for readability the format is described as if the maximum route length would be just three hops.

`hmac_0_2 | hmac_0_1 | hmac_0_0 | hmac_1_1 | hmac_1_0 | hmac_2_0`

`hmac_x_y` is the hmac added by node `x` (counted from the node that is currently handling the failure message) assuming that this node is `y` hops away from the erring node.

Each HMAC is computed from the combination of the following elements, concatenated in the specified order.

- The return packet before applying the pseudo-random byte stream.
- The concatenation of the first `y+1` hold times in `htlc_hold_times`. For example, `hmac_0_2` would cover all three hold times.
- The concatenation of `y` downstream hmacs that correspond to downstream node positions relative to `x`. For example, `hmac_0_2` would cover `hmac_1_1` and `hmac_2_0`.

The erring node stores its 20 HMACs at the start of the array and zeroes out the rest. Strictly speaking the erring node would only need to add the single `hmac_0_0` here, because there is no downstream data to cover. However, for verification efficiency at the origin node, we still require all HMACs to be calculated. The redundant HMACs will cover portions of the zero-initialized data.

Finally a new key is generated, using the key type `ammagext`. This key is then used to generate a pseudo-random stream, which is in turn applied to the `attribution_data` field using XOR.

Requirements

The *erring node*: - MUST construct the return packet such that its total length is no more than 32768 bytes (32 KiB). - MUST set `pad` such that the `failure_len` plus `pad_len` is at least 256. - SHOULD set `pad` such that the `failure_len` plus `pad_len` is equal to 256. Deviating from this may cause older nodes to be unable to parse the

return message. - if option_attribution_data is advertised: - if path_key is not set in the incoming update_add_htlc: - MUST initialize attribution_data and include it in update_fail_htlc

Intermediate nodes

Transformation of the return packet

Generate the node's ammag key, generate the pseudo-random byte streams, and apply the result to obfuscate the return packet. This is then stored as the reason field of the update_htlc_fail message.

This obfuscation step is identical to the obfuscation steps that the erring node carries out.

Transformation of attribution_data

- Shift all existing hold times to the right (4 bytes).
- Shift and prune all existing HMACs.

At each step backwards, one HMAC for every hop can be pruned. When HMACs for all 20 positions are present, and it turns out that there is another hop upstream, each existing HMAC that now corresponds to position 21 due to the preceding hop becomes obsolete.

For the simplified three-hop layout above, the shift/prune operation would apply a transformation that results in:

- | - | - | hmac_0'_1 | hmac_0'_0 | hmac_1'_0

The former hmac_x'_y now becomes hmac_{x+1}_y. The left-most HMAC for each hop is discarded.

Update of attribution_data

- Put the node's hold time at the start of htlc_hold_times. The shift operation above has opened up a slot for that.
- Calculate its own 20 truncated HMACs and put them at the start of hmacs in the newly opened slots.
- Generate the node's ammagext key, generate the pseudo-random byte stream, and apply the result to obfuscate the attribution_data field. This obfuscation step is identical to the obfuscation steps that the erring node carries out.

Requirements

The *intermediate node*: - if the return packet received from downstream is longer than 32768 bytes: - MUST truncate it to its first 32768 bytes before transforming the return packet or updating attribution_data - if option_attribution_data is advertised: - if path_key is not set in the incoming update_add_htlc: - if attribution_data is received from downstream: - MUST transform attribution_data as described above - otherwise: - MUST instantiate an all-zeroes attribution_data block - MUST update attribution_data as described above - MUST transform the return packet as described above. - MUST return-forward the update_htlc_fail message

Origin node

The origin node is able to detect that it's the intended final recipient of the return message, because of course, it was the originator of the corresponding forward packet. When an origin node receives an `update_htlc_fail` message matching a transfer it initiated (i.e. it cannot return-forward the error any further) it generates the `ammag`, `ammagext` and `um` keys for each hop in the route.

It then iteratively decrypts the message, using each hop's `ammag` and `ammagext` keys. At each hop, the following steps are carried out:

For origin nodes supporting `option_attribution_data`:

- Verify the HMAC in `attribution_data` that corresponds to the hop's position in the path using the hop's `um` key. If the HMAC is invalid, processing of the message can stop and the node should penalize this hop for future path selection. This is what makes the failure 'attributable'.

Because HMACs cover all data including HMACs added by downstream nodes, it is not possible for a malicious node to tamper with the message without revealing themselves. Blame is still only assignable to a pair of nodes though, because it is impossible to know whether sender or receiver modified the message. This is true for other failure cases in Lightning too.

When not every path node supports `option_attribution_data`, the origin node will still have attribution data up to the first node downstream without support.

- Record the reported `htlc` hold time for this hop.

The origin node can use this information to score nodes on latency. When a zero hold time is reported, the origin node should distribute any potential latency penalty across multiple nodes. This encourages path nodes to provide timing data to avoid being held responsible for the high latency of other nodes.

For all nodes:

- Compute the HMAC for the return packet, using the hop's `um` key.
- When the computed HMAC matches `hmac` in the return packet, the origin node will know that the current hop is the sender of the failure. They can then parse `failurormsg`.

The association between the forward and return packets is handled outside of this onion routing protocol, e.g. via association with an HTLC in a payment channel.

Requirements

The *origin node*: - once the return message has been decrypted: - SHOULD store a copy of the message. - SHOULD continue decrypting, until the loop has been repeated 27 times (maximum route length of `tlv_payload_type`). - SHOULD use constant `ammag` and `um` keys to obfuscate the route length. - When the failure source cannot be identified from the return packet AND `attribution_data` is present: - SHOULD use `attribution_data` to identify the failure source

Rationale

The requirements for the *origin node* should help hide the payment sender. By continuing decrypting 27 times (dummy decryption cycles after the error is found) the erroring node cannot learn its relative position in the route by performing a timing analysis if the sender were to retry the same route multiple times.

Failure Messages

The failure message encapsulated in `failuremsg` has an identical format as a normal message: a 2-byte type `failure_code` followed by data applicable to that type. The message data is followed by an optional [TLV stream](#).

Below is a list of the currently supported `failure_code` values, followed by their use case requirements.

Notice that the `failure_codes` are not of the same type as other message types, defined in other BOLTs, as they are not sent directly on the transport layer but are instead wrapped inside return packets. The numeric values for the `failure_code` may therefore reuse values, that are also assigned to other message types, without any danger of causing collisions.

The top byte of `failure_code` can be read as a set of flags: * 0x8000 (BADONION): unparsable onion encrypted by sending peer * 0x4000 (PERM): permanent failure (otherwise transient) * 0x2000 (NODE): node failure (otherwise channel) * 0x1000 (UPDATE): channel forwarding parameter was violated

The following `failure_codes` are defined:

1. type: NODE | 2 (temporary_node_failure)

General temporary failure of the processing node.

1. type: PERM | NODE | 2 (permanent_node_failure)

General permanent failure of the processing node.

1. type: PERM | NODE | 3 (required_node_feature_missing)

The processing node has a required feature which was not in this onion.

1. type: BADONION | PERM | 4 (invalid_onion_version)
2. data:
 - [sha256:sha256_of_onion]

The version byte was not understood by the processing node.

1. type: BADONION | PERM | 5 (invalid_onion_hmac)
2. data:
 - [sha256:sha256_of_onion]

The HMAC of the onion was incorrect when it reached the processing node.

1. type: BADONION | PERM | 6 (invalid_onion_key)
2. data:
 - [sha256:sha256_of_onion]

The ephemeral key was unparsable by the processing node.

1. type: UPDATE|7 (temporary_channel_failure)
2. data:
 - [u16:len]
 - [len*byte:channel_update]

The channel from the processing node was unable to handle this HTLC, but may be able to handle it, or others, later.

1. type: PERM|8 (permanent_channel_failure)

The channel from the processing node is unable to handle any HTLCs.

1. type: PERM|9 (required_channel_feature_missing)

The channel from the processing node requires features not present in the onion.

1. type: PERM|10 (unknown_next_peer)

The onion specified a short_channel_id which doesn't match any leading from the processing node.

1. type: UPDATE|11 (amount_below_minimum)
2. data:
 - [u64:htlc_msat]
 - [u16:len]
 - [len*byte:channel_update]

The HTLC amount was below the htlc_minimum_msat of the channel from the processing node.

1. type: UPDATE|12 (fee_insufficient)
2. data:
 - [u64:htlc_msat]
 - [u16:len]
 - [len*byte:channel_update]

The fee amount was below that required by the channel from the processing node.

1. type: UPDATE|13 (incorrect_cltv_expiry)
2. data:
 - [u32:cltv_expiry]
 - [u16:len]
 - [len*byte:channel_update]

The cltv_expiry does not comply with the cltv_expiry_delta required by the channel from the processing node: it does not satisfy the following requirement:

$$\text{cltv_expiry} - \text{cltv_expiry_delta} \geq \text{outgoing_cltv_value}$$

1. type: UPDATE|14 (expiry_too_soon)
2. data:
 - [u16:len]

- [len*byte:channel_update]

The CLTV expiry is too close to the current block height for safe handling by the processing node.

1. type: PERM|15 (incorrect_or_unknown_payment_details)
2. data:
 - [u64:htlc_msat]
 - [u32:height]

The payment_hash is unknown to the final node, the payment_secret doesn't match the payment_hash, the amount for that payment_hash is too low, the CLTV expiry of the htlc is too close to the current block height for safe handling or payment_metadata isn't present while it should be.

The htlc_msat parameter is superfluous, but left in for backwards compatibility. The value of htlc_msat is required to be at least the value specified in the final hop onion payload. It therefore does not have any substantial informative value to the sender (though may indicate the penultimate node took a lower fee than expected). A penultimate hop sending an amount or an expiry that is too low for the htlc is handled through final_incorrect_cltv_expiry and final_incorrect_htlc_amount.

The height parameter is set by the final node to the best known block height at the time of receiving the htlc. This can be used by the sender to distinguish between sending a payment with the wrong final CLTV expiry and an intermediate hop delaying the payment so that the receiver's invoice CLTV delta requirement is no longer met.

Note: Originally PERM|16 (incorrect_payment_amount) and 17 (final_expiry_too_soon) were used to differentiate incorrect htlc parameters from unknown payment hash. Sadly, sending this response allows for probing attacks whereby a node which receives an HTLC for forwarding can check guesses as to its final destination by sending payments with the same hash but much lower values or expiry heights to potential destinations and check the response. Care must be taken by implementations to differentiate the previously non-permanent case for final_expiry_too_soon (17) from the other, permanent failures now represented by incorrect_or_unknown_payment_details (PERM|15).

1. type: 18 (final_incorrect_cltv_expiry)
2. data:
 - [u32:cltv_expiry]

The CLTV expiry in the HTLC is less than the value in the onion.

1. type: 19 (final_incorrect_htlc_amount)
2. data:
 - [u64:incoming_htlc_amt]

The amount in the HTLC is less than the value in the onion.

1. type: UPDATE|20 (channel_disabled)
2. data:
 - [u16:disabled_flags]
 - [u16:len]
 - [len*byte:channel_update]

The channel from the processing node has been disabled. No flags for disabled_flags are currently defined, thus it is currently always two zero bytes.

1. type: 21 (expiry_too_far)

The CLTV expiry in the HTLC is too far in the future.

1. type: PERM|22 (invalid_onion_payload)
2. data:
 - [bigsize:type]
 - [u16:offset]

The decrypted onion per-hop payload was not understood by the processing node or is incomplete. If the failure can be narrowed down to a specific tlv type in the payload, the erring node may include that type and its byte offset in the decrypted byte stream.

1. type: 23 (mpp_timeout)

The complete amount of the multi-part payment was not received within a reasonable time.

1. type: BADONION|PERM|24 (invalid_onion_blinding)
2. data:
 - [sha256:sha256_of_onion]

An error occurred within the blinded path.

Requirements

An *erring node*: - if path_key is set in the incoming update_add_htlc: - MUST return an invalid_onion_blinding error. - if current_path_key is set in the onion payload and it is not the final node: - MUST return an invalid_onion_blinding error. - otherwise: - MUST select one of the above error codes when creating an error message. - MUST include the appropriate data for that particular error type. - if there is more than one error: - SHOULD select the first error it encounters from the list above.

An *erring node* MAY: - if the per-hop payload in the onion is invalid (e.g. it is not a valid tlv stream) or is missing required information (e.g. the amount was not specified): - return an invalid_onion_payload error. - if an otherwise unspecified transient error occurs for the entire node: - return a temporary_node_failure error. - if an otherwise unspecified permanent error occurs for the entire node: - return a permanent_node_failure error. - if a node has requirements advertised in its node_announcement features, which were NOT included in the onion: - return a required_node_feature_missing error.

A *forwarding node* MUST: - if path_key is set in the incoming update_add_htlc: - return an invalid_onion_blinding error. - if current_path_key is set in the onion payload: - return an invalid_onion_blinding error. - otherwise: - select one of the above error codes when creating an error message.

A *forwarding node* MAY, but a *final node* MUST NOT: - if during forwarding to its receiving peer, an otherwise unspecified, transient error occurs in the outgoing channel (e.g. channel capacity reached, too many in-flight HTLCs, etc.): - return a temporary_channel_failure error. - if an otherwise unspecified, permanent error occurs during forwarding to its receiving peer (e.g. channel recently closed): - return a

permanent_channel_failure error. - if the outgoing channel has requirements advertised in its channel_announcement's features, which were NOT included in the onion: - return a required_channel_feature_missing error. - if the receiving peer specified by the onion is NOT known: - return an unknown_next_peer error. - if the HTLC amount is less than the currently specified minimum amount: - report the amount of the outgoing HTLC and the current channel setting for the outgoing channel. - return an amount_below_minimum error. - if the HTLC does NOT pay a sufficient fee: - report the amount of the incoming HTLC and the current channel setting for the outgoing channel. - return a fee_insufficient error. - if the incoming cltv_expiry minus the outgoing_cltv_value is below the cltv_expiry_delta for the outgoing channel: - report the cltv_expiry of the outgoing HTLC and the current channel setting for the outgoing channel. - return an incorrect_cltv_expiry error. - if the cltv_expiry is unreasonably near the present: - report the current channel setting for the outgoing channel. - return an expiry_too_soon error. - if the cltv_expiry is more than max_htlc_cltv in the future: - return an expiry_too_far error. - if the channel is disabled: - report the current channel setting for the outgoing channel. - return a channel_disabled error.

A forwarding node and *a final node* MAY: - if the onion version byte is unknown: - return an invalid_onion_version error. - if the onion HMAC is incorrect: - return an invalid_onion_hmac error. - if the ephemeral key in the onion is unparsable: - return an invalid_onion_key error.

A forwarding node MUST NOT, but *a final node* MUST: - if the payment hash has already been paid: - MAY treat the payment hash as unknown. - MAY succeed in accepting the HTLC. - if the payment_secret doesn't match the expected value for that payment_hash, or the payment_secret is required and is not present: - MUST fail the HTLC. - MUST return an incorrect_or_unknown_payment_details error. - if the amount paid is less than the amount expected: - MUST fail the HTLC. - MUST return an incorrect_or_unknown_payment_details error. - if the payment hash is unknown: - MUST fail the HTLC. - MUST return an incorrect_or_unknown_payment_details error. - if the amount paid is more than twice the amount expected: - SHOULD fail the HTLC. - SHOULD return an incorrect_or_unknown_payment_details error. - Note: this allows the origin node to reduce information leakage by altering the amount while not allowing for accidental gross overpayment. - if the cltv_expiry value is unreasonably near the present: - MUST fail the HTLC. - MUST return an incorrect_or_unknown_payment_details error. - if the cltv_expiry from the final node's HTLC is below outgoing_cltv_value: - MUST return final_incorrect_cltv_expiry error. - if amount_msat from the final node's HTLC is below amt_to_forward: - MUST return a final_incorrect_htlc_amount error. - if it returns a channel_update: - MUST set short_channel_id to the short_channel_id used by the incoming onion.

Rationale

In the case of multiple short_channel_id aliases, the channel_update short_channel_id should refer to the one the original sender is expecting, to both avoid confusion and to avoid leaking information about other aliases (or the real location of the channel UTXO).

The channel_update field used to be mandatory in messages whose failure_code includes the UPDATE flag. However, because nodes applying an update contained in the onion to their gossip data is a massive fingerprinting vulnerability, the channel_update field is no longer mandatory and nodes are expected to transition away from including it. Nodes which do not provide a channel_update are expected to set the channel_update len field to zero.

Some nodes may still use the channel_update for retries of the same payment, however.

Receiving Failure Codes

Requirements

The *origin node*: - MUST ignore any extra bytes in `failurmsg`. - if the *final node* is returning the error: - if the PERM bit is set: - SHOULD fail the payment. - otherwise: - if the error code is understood and valid: - MAY retry the payment. In particular, `final_expiry_too_soon` can occur if the block height has changed since sending, and in this case `temporary_node_failure` could resolve within a few seconds. - otherwise, an *intermediate hop* is returning the error: - if the NODE bit is set: - SHOULD remove all channels connected with the erring node from consideration. - if the PERM bit is NOT set: - SHOULD restore the channels as it receives new `channel_updates` from its peers. - otherwise: - if UPDATE is set, AND the `channel_update` is valid and more recent than the `channel_update` used to send the payment: - MAY consider the `channel_update` when calculating routes to retry the payment which failed - MUST NOT expose the `channel_update` to third-parties in any other context, including applying the `channel_update` to the local network graph, send the `channel_update` to peers as gossip, etc. - SHOULD then retry routing and sending the payment. - MAY use the data specified in the various failure types for debugging purposes.

Successful Payments

The `update_fulfill_htlc` message contains an optional `attribution_data` field, identical to the one used in the failure case described above. With attribution data, the sender can obtain hold times as reported and committed to by each hop along the payment path.

In addition, the final node can return data to the origin node in an optional `fulfillment_payload` field. It is the success counterpart of the failure message that an erring node returns in the reason field of `update_fail_htlc`. The final node builds the plaintext as a `fulfillment_payload_tlv`s [TLV stream](#):

1. `tlv_stream`: `fulfillment_payload_tlv`s
2. types:
 1. type: 1 (padding)
 2. data:
 - [...*byte:padding]

padding conceals the length of the conveyed data and is ignored by the origin. Apart from padding, no record types are defined yet.

The final node generates a key using key type `fulfillment` and its shared secret with the origin node, then encrypts this plaintext with ChaCha20-Poly1305 using that key and an all-zero nonce, as for `encrypted_recipient_data` in route blinding. The ciphertext and its 16-byte tag are the `fulfillment_payload`. Each intermediate node then obfuscates the `fulfillment_payload` with its `ammag` key, exactly as it obfuscates a failure return packet; the final node does not, as the ciphertext is already confidential and forms the innermost layer. The `attribution_data` HMACs added by intermediate nodes, which already cover the hold times, now additionally cover the `fulfillment_payload`. For each intermediate HMAC, the covered fields are, in order: the `fulfillment_payload` as received from the downstream node, before the current node applies its own `ammag` obfuscation; the hold times; and the downstream HMACs. The final node's `attribution_data` HMACs do not cover the `fulfillment_payload`. With no payload present, there is nothing extra for intermediate nodes to cover. The final node still initializes and updates `attribution_data`, and the origin still verifies the final node's `attribution_data` HMACs for its reported hold time.

A successful payment is relayed back with `update_fulfill_htlc`, so a `fulfillment_payload` propagates back through every hop, even a blinded path: a blinded final node can originate it, and blinded hops obfuscate and relay it. The origin can still recover it, because it shares a secret with every hop, blinded ones included, computed from each hop's public key (the blinded public key for a blinded hop) when it built the onion. Returning a `fulfillment_payload` is therefore not gated by the `path_key` condition that applies to `attribution_data`.

The origin knows the number of hops in the path, so the final node is its last hop: the origin peels each intermediate hop's `ammag` obfuscation, including obfuscation for blinded hops whose blinded public keys it used when constructing the route, and decrypts the `fulfillment_payload` with the last hop's `fulfillment` key, one operation per hop. The origin verifies the final node's `attribution_data` HMACs without a `fulfillment_payload` input. For each intermediate hop, the origin removes that hop's `ammag` obfuscation from the `fulfillment_payload` and verifies the hop's `attribution_data` HMACs over the resulting `fulfillment_payload` bytes. A recipient that conceals its position with dummy hops in a blinded path must therefore originate the `fulfillment_payload` as though it were the last hop, applying the obfuscation for the concealed hops itself, so that this decoding succeeds without revealing its position.

The two integrity checks cover different extents. The Poly1305 tag is verified end-to-end and detects tampering by any hop, blinded or not. The `attribution_data` HMACs are contributed only by hops with `path_key` not set, so they provide per-hop attribution where available, as in the failure case. This split is intentional: `attribution_data` is per-hop data for scoring path nodes, which the origin cannot do for blinded hops it cannot identify and whose hold times would aid de-anonymization, whereas a `fulfillment_payload` is end-to-end data that reveals nothing about the intermediate hops.

Like return packets, `fulfillment_payload` is limited to 32768 bytes (32 KiB) to leave room for `attribution_data` and future extensions. Because `fulfillment_payload` was introduced with this limit, receiving a larger payload is a protocol violation as specified in [BOLT #2](#).

Requirements

`attribution_data` is initialized, transformed, updated and verified exactly as in the failure case.

The *final node*: - if `option_attribution_data` is advertised: - MAY include a `fulfillment_payload`, even when reached through a blinded path - if it includes a `fulfillment_payload`: - MUST ensure the `fulfillment_payload` is no more than 32768 bytes (32 KiB) - MUST set `fulfillment_payload_tlvs` to a serialized TLV stream - MUST pad with padding such that the serialized `fulfillment_payload_tlvs` stream is at least 256 bytes and a multiple of 256 bytes, excluding the 16-byte Poly1305 tag. The size calculation includes the type and length bytes of the padding record; this can require padding to the next multiple of 256 bytes - SHOULD pad such that the serialized `fulfillment_payload_tlvs` stream is exactly 256 bytes if possible - MUST encrypt the plaintext with ChaCha20-Poly1305 as above - MUST NOT include the `fulfillment_payload` in its `attribution_data` HMACs - MUST still initialize and update `attribution_data` for its reported hold time

An *intermediate node*: - if it receives a `fulfillment_payload` from the downstream node: - MUST obfuscate it with its `ammag` key and include it in the outgoing `update_fulfill_htlc`, as it transforms a return packet - MUST include the `fulfillment_payload` as received from the downstream node in its `attribution_data` HMACs

The *origin node*: - MUST recover a fulfillment_payload by peeling each intermediate hop's ammag obfuscation and decrypting it with the last hop's fulfillment key - MUST verify the final node's attribution_data HMACs as in the failure case, except that the HMAC input excludes the fulfillment_payload, and process its reported hold time - MUST remove an intermediate hop's ammag obfuscation from the fulfillment_payload before verifying that hop's attribution_data HMACs over the resulting fulfillment_payload bytes - MUST ignore a fulfillment_payload whose Poly1305 tag is invalid, or whose decrypted fulfillment_payload_tlv stream has malformed lengths, duplicate or unordered types, or unknown even types - MUST ignore padding records in a valid fulfillment_payload_tlv stream

Onion Messages

Onion messages allow peers to use existing connections to query for invoices (see [BOLT 12](#)). Like gossip messages, they are not associated with a particular local channel. Like HTLCs, they use [onion messages](#) protocol for end-to-end encryption.

Onion messages use the same form as HTLC onion_packet, with a slightly more flexible format: instead of 1300 byte payloads, the payload length is implied by the total length (minus 66 bytes for the header and trailing bytes). The onionmsg_payloads themselves are the same as the hop_payloads format, except there is no "legacy" length: a 0 length would mean an empty onionmsg_payload.

Onion messages are unreliable: in particular, they are designed to be cheap to process and require no storage to forward. As a result, there is no error returned from intermediary nodes.

For consistency, all onion messages use [Route Blinding](#).

The onion_message Message

1. type: 513 (onion_message) (option_onion_messages)
2. data:
 - [point:path_key]
 - [u16:len]
 - [len*byte:onion_message_packet]
3. type: onion_message_packet
4. data:
 - [byte:version]
 - [point:public_key]
 - [...*byte:onionmsg_payloads]
 - [32*byte:hmac]
5. type: onionmsg_payloads
6. data:
 - [bigsize:length]
 - [length*u8:onionmsg_tlv]
 - [32*byte:hmac]
 - ...
 - filler

The onionmsg_tlv itself is a TLV: an intermediate node expects an encrypted_recipient_data which it can decrypt into an encrypted_data_tlv using the path_key which it is handed along with the onion message.

Field numbers 64 and above are reserved for payloads for the final hop.

1. tlv_stream: onionmsg_tlv
2. types:
 1. type: 2 (reply_path)
 2. data:
 - [blinded_path:path]
 3. type: 4 (encrypted_recipient_data)
 4. data:
 - [...*byte:encrypted_recipient_data]
 5. type: 64 (invoice_request)
 6. data:
 - [tlv_invoice_request:invreq]
 7. type: 66 (invoice)
 8. data:
 - [tlv_invoice:inv]
 9. type: 68 (invoice_error)
 10. data:
 - [tlv_invoice_error:inverr]

Requirements

The creator of encrypted_recipient_data (usually, the recipient of the onion):

- MUST create the encrypted_recipient_data from the encrypted_data_tlv as required in [Route Blinding](#).
- MUST NOT include payment_relay or payment_constraints in any encrypted_data_tlv
- MUST include either next_node_id or short_channel_id in the encrypted_data_tlv for each non-final node.

The writer:

- MUST set the onion_message_packet version to 0.
- MUST construct the onion_message_packet onionmsg_payloads as detailed above using Sphinx.
- MUST NOT use any associated_data in the Sphinx construction.
- SHOULD set onion_message_packet len to 1366 or 32834.
- SHOULD retry via a different path if it expects a response and doesn't receive one after a reasonable period.
- For the non-final nodes' onionmsg_tlv:
 - MUST NOT set fields other than encrypted_recipient_data.
- For the final node's onionmsg_tlv:
 - if the final node is permitted to reply:
 - MUST set reply_path path_key to the initial path key for the first_node_id
 - MUST set reply_path first_node_id to the unblinded node id of the first node in the reply path.
 - For every reply_path path:
 - MUST set blinded_node_id to the blinded node id to encrypt the onion hop for.

- MUST set `encrypted_recipient_data` to a valid encrypted `encrypted_data_tlv` stream which meets the requirements of the `onionmsg_tlv` when used by the recipient.
- MAY use `path_id` to contain a secret so it can recognize use of this `reply_path`.
- otherwise:
 - MUST NOT set `reply_path`.

The reader:

- if it advertises `option_onion_messages_only_channels`:
 - MUST NOT accept onion messages from peers without an established channel.
- otherwise:
 - SHOULD accept onion messages from peers without an established channel.
- MAY rate-limit messages by dropping them.
- MUST decrypt `onion_message_packet` using an empty `associated_data`, and `path_key`, as described in [Onion Decryption](#) to extract an `onionmsg_tlv`.
- If decryption fails, the result is not a valid `onionmsg_tlv`, or it contains unknown even types:
 - MUST ignore the message.
- if `encrypted_data_tlv` contains `allowed_features`:
 - MUST ignore the message if:
 - `encrypted_data_tlv.allowed_features.features` contains an unknown feature bit (even if it is odd).
 - the message uses a feature not included in `encrypted_data_tlv.allowed_features.features`.
- if it is not the final node according to the onion encryption:
 - if the `onionmsg_tlv` contains other tlv fields than `encrypted_recipient_data`:
 - MUST ignore the message.
 - if the `encrypted_data_tlv` contains `path_id`:
 - MUST ignore the message.
 - otherwise:
 - if `next_node_id` is present:
 - the *next peer* is the peer with that node id.
 - otherwise, if `short_channel_id` is present and corresponds to an announced `short_channel_id` or a local alias for a channel:
 - the *next peer* is the peer at the other end of that channel.
 - otherwise:
 - MUST ignore the message.
 - SHOULD forward the message using `onion_message` to the *next peer*.
 - if it forwards the message:
 - MUST set `path_key` in the forwarded `onion_message` to the next `path_key` as calculated in [Route Blinding](#).
- otherwise (it is the final node):
 - if `path_id` is set and corresponds to a path the reader has previously published in a `reply_path`:
 - if the onion message is not a reply to that previous onion:
 - MUST ignore the onion message

- otherwise (unknown or unset `path_id`):
 - if the onion message is a reply to an onion message which contained a `path_id`:
 - MUST respond (or not respond) exactly as if it did not send the initial onion message.
- if the `onionmsg_tlv` contains more than one payload field:
 - MUST ignore the message.
- if it wants to send a reply:
 - MUST create an onion message using `reply_path`.
 - MUST send the reply via `onion_message` to the node indicated by the `first_node_id`, using `reply_path` `path_key` to send along `reply_path` path.

Rationale

Accepting onion messages from channelless peers is the baseline: a node that advertises `option_onion_messages` is expected to accept onion messages from any connected peer. This is a SHOULD rather than a MUST because acceptance is always best effort: a node may still drop messages due to rate-limiting.

A node that wants to restrict acceptance to channel peers (e.g. to limit the DoS surface of free bandwidth/CPU for any node that can connect) advertises `option_onion_messages_only_channels` as an explicit opt-out. A restrictive node should not silently drop these messages while advertising only the baseline: it has to advertise the bit, which makes the bit a reliable signal. A sender can therefore detect it and simply not attempt a relay it knows will fail, rather than sending a message that is silently dropped.

Care must be taken that replies are only accepted using the exact `reply_path` given, otherwise probing is possible. That means checking both ways: non-replies don't use the reply path, and replies always use the reply path.

The requirement to discard messages with `onionmsg_tlv` fields which are not strictly required ensures consistency between current and future implementations. Even odd fields can be a problem since they are parsed (and thus may be rejected!) by nodes which understand them, and ignored by those which don't.

All onion messages are blinded, even though this overhead is not always necessary (33 bytes here, the 16-byte MAC for each `encrypted_data_tlv` in the onion). This blinding allows nodes to use a path provided by others without knowing its contents. Using it universally simplifies implementations a little, and makes it more difficult to distinguish onion messages.

len allows larger messages to be sent than the standard 1300 bytes allowed for an HTLC onion, but this should be used sparingly as it reduces the anonymity set, hence the recommendation that it either looks like an HTLC onion, or if larger, be a fixed size.

Onion messages don't explicitly require a channel, but for spam-reduction a node may choose to ratelimit such peers, especially messages it is asked to forward.

`max_htlc_cltv` Selection

This `max_htlc_cltv` value is defined as 2016 blocks, based on historical value deployed by Lightning implementations.

Test Vector

Returning Errors

The test vectors use the following parameters:

```
pubkey[0] = 0x02eec7245d6b7d2ccb30380bfbe2a3648cd7a942653f5aa340edcea1f283686619
htlc hold time[0] = 1
```

```
pubkey[1] = 0x0324653eac434488002cc06bbfb7f10fe18991e35f9fe4302dbea6d2353dc0ab1c
htlc hold time[1] = 2
```

```
pubkey[2] = 0x027f31ebc5462c1fdce1b737ecff52d37d75dea43ce11c74d25aa297165faa2007
htlc hold time[2] = 3
```

```
pubkey[3] = 0x032c0b7cf95324a07d05398b240174dc0c2be444d96b159aa6c7f7b1e668680991
htlc hold time[3] = 4
```

```
pubkey[4] = 0x02edabbd16b41c8371b92ef2f04c1185b4f03b6dcd52ba9b78d9d7c89c8f221145
htlc hold time[4] = 5
```

```
nhops = 5  
sessionkey = 0x4141414141414141414141414141414141414141414141414141414141414141
```

```
failure_source = node 4
failure_message = `incorrect_or_unknown_payment_details`
  htlc_msat = 100
  height = 800000
  tlvs {
    type = 34001
    value = [128, 128, ..., 128] (300 bytes)
  }
}
```

The following is an in-depth trace of an example of error message creation:

[illegible]

ammag_key = 2f36bb8822e1f0d04c27b7d8bb7d7dd586e032a3218b8d414afbba6f169a4d68

stream =

e9c975b07c9a374ba64fd9be3aae955e917d34d1fa33f2e90f53bbf4394713c6a8c9b16ab5f12fd45edd73c1b0c8b330
02df376801ff58aaa94000bf8a86f92620f343baef38a580102395ae3abf9128d1047a0736ff9b83d456740ebbb4aeb3
aa9737f18fb4afb4aa074fb26c4d702f4296888550a3bde8c05247e045b866baef0499f079fdaeeff6538f31d44deaf
ffdf3afa2fb4ca9082b8f1c465371a9894dd8c243fb4847e004f5256b3e90e2edde4c9fb3082ddfe4d1e734cacd96ef
0706bf63c9984e22cd98851bcccd1c3494351feb458c9c6af41c0044bea3c47552b1d992ae542b17a2d0bba1a096c78d
169034ecb55b6e3a7263c26017f033031228833c1daefc0dedb8cf7c3e37c9c37ebfe42f3225c326e8bcfd338804c145
b16e34e4f5984bc119af09d471a61f39e9e389c4120cadabc5d9b7b1355a8cce050ca8ad72f642fc26919927b347808
bade4b1c321b08bc363f20745ba2f97f0ced2996a232f55ba28fe7dfa70a9ab0433a085388f25cce8d53de6a2fbd7546
377d6ede9027ad173ba1f95767461a3689ef405ab608a21086165c64b02c1782b04a6dba2361a7784603069124e12f2f
6dcb1ec7612a4fbf94c0e14631a2bef6190c3d5f35e0c4b32aa85201f449d830fd8f782ec758b0910428e3ec3ca1dba3
b6c7d89f69e1ee1b9df3dfbbf6d361e1463886b38d52e8f43b73a3bd48c6f36f5897f514b93364a31d49d1d506340b13
15883d425cb36f4ea553430d538fd6f3596d4afc518db2f317dd051abc0d4bfb0a7870c3db70f19fe78d6604bbf088fc
b4613f54e67b038277fedcd9680eb97bdfc3be1ab2cbcbaf6d25b8a7ac34d8c190f98d3064ecd3b95b8895157c6a37f
31ef4de094b2cb9dbf8ff1f419ba0ecac1bb13df0253b826bec2ccca1e745dd3b3e7cc6277ce284d649e7b828572773
5ff4ef6cca6c18e2714f4e2a1ac67b25213d3bb49763b3b94e7ebf72507b71fb2fe0329666477ee7cb7ebd6b88ad5add
8b217188b1ca0fa13de1ec09cc674346875105be6e0e0d6c8928eb0df23c39a639e04e4aedf535c4e093f08b2c905a14
f25c0c0fe47a5a1535ab9eae0d9d67bdd79de13a08d59ee05385c7ea4af1ad3248e61dd22f8990e9e99897d653dd7b1b
1433a6d464ea9f74e377f2d8ce99ba7dbc753297644234d25ecb5bd528e2e2082824681299ac30c05354baaa9c3967d8
6d7c07736f87fc0f63e5036d47235d7ae12178ced3ae36ee5919c093a02579e4fc9edad2c446c656c790704bfc8e2c49
1a42500aa1d75c8d4921ce29b753f883e17c79b09ea324f1f32ddf1f3284cd70e847b09d90f6718c42e5c94484cc9cbb
0df659d255630a3f5a27e7d5dd14fa6b974d1719aa98f01a20fb4b7b1c77b42d57fab3c724339d459ee4a1c6b5d3bd4e
08624c786a257872acc9ad3ff62222f2265a658d9f2a007229a5293b67ec91c84c4b4407c228434bad8a815ca9b256c7
76bd2c9f

error packet for node 4:

146e94a9086dbbed6a0ab6932d00c118a7195dbf69b7d7a12b0e6956fc54b5e0a989f165b5f12fd45edd73a5b0c48630
ff5be69500d3d82a29c0803f0a0679a6a073c33a6fb8250090a3152eba3f11a85184fa87b67f1b0354d6f48e3b342e33
2a17b7710f342f342a87cf32eccdf0afc2160808d58abb5e5840d2c760c538e63a6f841970f97d2e6fe5b8739dc45e2f
7f5f532f227bcc2988ab0f9cc6d3f12909cd5842c37bc8c7608475a5ebbe10626d5ecc1f3388ad5f645167b44a4d166f
87863fe34918cea25c18059b4c4d9cb414b59f6bc50c1cea749c80c43e2344f5d23159122ed4ab9722503b212016470d
9610b46c35dbeebaf2e342e09770b38392a803bc9d2e7c8d6d384ffcbef74943fe3f64afb2a543a6683c7db3088441c5
31eeb4647518cb41992f8954f1269fb969630944928c2d2b45593731b5da0c4e70d04a0a57afe4af42e99912fbb4f888
3a5ecb9cb29b883cb6bfa0f4db2279ff8c6d2b56a232f55ba28fe7dfa70a9ab0433a085388f25cce8d53de6a2fbd7546
377d6ede9027ad173ba1f95767461a3689ef405ab608a21086165c64b02c1782b04a6dba2361a7784603069124e12f2f
6dcb1ec7612a4fbf94c0e14631a2bef6190c3d5f35e0c4b32aa85201f449d830fd8f782ec758b0910428e3ec3ca1dba3
b6c7d89f69e1ee1b9df3dfbbf6d361e1463886b38d52e8f43b73a3bd48c6f36f5897f514b93364a31d49d1d506340b13
15883d425cb36f4ea553430d538fd6f3596d4afc518db2f317dd051abc0d4bfb0a7870c3db70f19fe78d6604bbf088fc
b4613f54e67b038277fedcd9680eb97bdfc3be1ab2cbcbaf6d25b8a7ac34d8c190f98d3064ecd3b95b8895157c6a37f
31ef4de094b2cb9dbf8ff1f419ba0ecac1bb13df0253b826bec2ccca1e745dd3b3e7cc6277ce284d649e7b828572773
5ff4ef6cca6c18e2714f4e2a1ac67b25213d3bb49763b3b94e7ebf72507b71fb2fe0329666477ee7cb7ebd6b88ad5add
8b217188b1ca0fa13de1ec09cc674346875105be6e0e0d6c8928eb0df23c39a639e04e4aedf535c4e093f08b2c905a14
f25c0c0fe47a5a1535ab9eae0d9d67bdd79de13a08d59ee05385c7ea4af1ad3248e61dd22f8990e9e99897d653dd7b1b
1433a6d464ea9f74e377f2d8ce99ba7dbc753297644234d25ecb5bd528e2e2082824681299ac30c05354baaa9c3967d8
6d7c07736f87fc0f63e5036d47235d7ae12178ced3ae36ee5919c093a02579e4fc9edad2c446c656c790704bfc8e2c49
1a42500aa1d75c8d4921ce29b753f883e17c79b09ea324f1f32ddf1f3284cd70e847b09d90f6718c42e5c94484cc9cbb
0df659d255630a3f5a27e7d5dd14fa6b974d1719aa98f01a20fb4b7b1c77b42d57fab3c724339d459ee4a1c6b5d3bd4e
08624c786a257872acc9ad3ff62222f2265a658d9f2a007229a5293b67ec91c84c4b4407c228434bad8a815ca9b256c7
76bd2c9f

attribution data for node 4:

d77d0711b5f71d1d1be56bd88b3bb7ebc1792bb739ea7ebc1bc3b031b8bc2df3a50e25aeb99f47d7f7ab39e24187d3f4
df9c4333463b053832ee9ac07274a5261b8b2a01fc09ce9ea7cd04d7b585dfb8cf5958e3f3f2a4365d1ec0dfd1d83c6a6
221b5b7d1ff30156a2289a1d3ee559e7c7256bda444bb8e046f860e00b3a59a85e1e1a43de215fd5e6bf646a5deab97b
1912c934e31b1cfd344764d6ca7e14ea7b3f2a951aba907c964c0f5d19a44e6d1d7279637321fa598adde927b3087d23
8f8b426ecde500d318617cdb7a56e6ce3520fc95be41a549973764e4dc483853ecc313947709f1b5199cb077d46e701f

a633e11d3e13b03e9212c115ca6fa004b2f3dd912814693b705a561a06da54cdf603677a3abecdc22c7358c2de3cef77
1b366a568150aecc86ad1990bb0f4e2865933b03ea0df87901bff467908273dc6cea31cbab0e2b8d398d10b001058c2
59ed221b7b55762f4c7e49c8c11a45a107b7a2c605c26dc5b0b10d719b1c844670102b2b6a36c43fe4753a78a483fc39
166ae28420f112d50c10ee64ca69569a2f690712905236b7c2cb7ac8954f02922d2d918c56d42649261593c47b14b324
a65038c3c5be8d3c403ce0c8f19299b1664bf077d7cf1636c4fb9685a8e58b7029fd0939fa07925a60bed339b23f9732
93598f595e75c8f9d455d7cebe4b5e23357c8bd47d66d6628b39427e37e0aecbabf46c11be6771f7136e108a143ae9ba
fba0fc47a51b6c7deef4cba54bae906398ee3162a41f2191ca386b628bde7e1dd63d1611aa01a95c456df337c763cb8c
3a81a6013aa633739d8cd554c688102211725e6adad165adc1bcd429d020c51b4b25d2117e8bb27eb0cc7020f9070d4a
d19ac31a76ebdf5f9246646aeadbfb9a3f1d75bd8237961e786302516a1a781780e8b73f58dc06f307e58bd0eb1d8f5c
9111f01312974c1dc777a6a2d3834d8a2a40014e9818d0685cb3919f6b3b788ddc640b0ff9b1854d7098c7dd6f35196e
902b26709640bc87935a3914869a807e8339281e9cedaaca99474c3e7bdd35050bb998ab4546f9900904e0e39135e861
ff7862049269701081ebce32e4cca992c6967ff0fd239e38233eaf614af31e186635e9439ec5884d798f9174da6ff569
d68ed5c092b78bd3f880f5e88a7a8ab36789e1b57b035fb6c32a6358f51f83e4e5f46220bcad072943df8bd9541a61b7
dae8f30fa3dd5fb39b1fd9a0b8e802552b78d4ec306ecee15bfe6da14b29ba6d19ce5be4dd478bca74a52429cd5309d4
04655c3dec85c252

forwarding error packet

shared_secret = 21e13c2d7cfe7e18836df50872466117a295783ab8aab0e7ecc8c725503ad02d

ammag_key = cd9ac0e09064f039fa43a31dea05f5fe5f6443d40a98be4071af4a9d704be5ad

stream =

617ca1e4624bc3f04fece3aa5a2b615110f421ec62408d16c48ea6c1b7c33fe7084a2bd9d4652fc5068e5052bf6d0aca
e2176018a3d8c75f37842712913900263cff92f39f3c18aa1f4b20a93e70fc429af7b2b1967ca81a761d40582daf0eb4
9cef66e3d6fbca0218d3022d32e994b41c884a27c28685ef1eb14603ea80a204b2f2f474b6ad5e71c6389843e3611ebe
afc62390b717ca53b3670a33c517ef28a659c251d648bf4c966a4ef187113ec9848bf110816061ca4f2f68e76ceb88bd
6208376460b916fb2ddeb77a65e8f88b2e71a2cbf4ea4958041d71c17d05680c051c3676fb0dc8108e5d78fb1e2c44d7
9a202e9d14071d536371ad47c39a05159e8d6c41d17a1e858faaaf572623aa23a38ffc73a4114cb1ab1cd7f906c6bd4e
21b29694f9830d12e8ad4b1ac320b3d5bfb4e534f02cefe9a983d66939581412acb1927eb93e8ed73145cddf24266bdc
c95923ecb38c8c9c5f4465335b0f18bf9f2d03fa02d57f258db27983d94378bc796cbe7737180dd7e39a36e461ebcb7e
c82b6dcd9d3f209381f7b3a23e798c4f92e13b4bb972ee977e24f4b83bb59b577c210e1a612c2b035c8271d9bc1fb91
5776ac6560315b124465866830473aa238c35089cf2adb9c6e9f05ab113c1d0a4a18ba0cb9951b928c0358186532c36d
4c3daa65657be141cc22e326f88e445e898893fd5f0a7dd231ee5bc972077b1e12a8e382b75d4b557e895a2adc757f2e
451e33e0ae3fb54566ee09155da6ada818aa4a4a2546832a8ba22f0ef9ec6a1c78e03a7c29cb126bcfa81aea61cd8b07
ab9f4e5e0ad0d9a3a0c66d2d0a00cc05884d183a68e816e76b75842d55895f5b91c5c1b9f7052763aae8a647aa079921
4275b6e781f0816fc9fffb802a0101eb5a2de6b3375d3e3478f892b2de7f1900d8ca9bf188fcba89fc49d03c38fa2587a
8ae119abfb295b15fa11cb188796bedc4fdcfceef296e44fbfa7e84569cc6346389a421782e40a298e1e2b6f9cae3103c
3f39d24541e4ab7b61daf1a5f2fe936a59d87cccdaf7c226acc451ceec3e81bc4828b4925feae3526d5e2bf93bd5f4
fdc0e069010aea1ae7e0d480d438918598896b776bf08fea124f91b3a13414b56934857707902612fc97b0e5d02cbe6e
5901ad304c7e8656390efccf3a1b22e18a2935181b78d5d2c89540ede8b0e6194d0d3945780bf577f622cb12deedbf82
10eb1450c298b9ee19f3c7082aabc2bdbd3384f3539dc3766978567135549df0d48287735854a6098fa40a9e48eaa27e
0d159beb65dd871e4c0b3ffa65f0375f0d3253f582193135ece60d5b9d8ba6739d87964e992cbec674b728d9eaaed59
5462c41d15fb497d4baa062368005d13fc99e1402563117a6c140c10363b05196a4cbb6b84ae807d62c748485c15e331
6841e4a98c3aac81e3bc996b4baca77eac8cdeb99cc7c5eb85c907cef4abe15cbe87fdc5dc2326019196235ac2059
34fcf8e3

error packet for node 3:

7512354d6a26781d25e65539772ba049b7ed7c530bf75ab7ef80cf974b978a07a1c3dabc61940011585323f70fa98cfa
1d4c868da30b1f751e44a72d9b3f79809c8c51c9f0843daa8fe83587844fedeachb7348362003b31922cbb4d6169b2087
b6f8d192d9cfe5363254cd1fde24641bde9e422f170c3eb146f194c48a459ae2889d706dc654235fa9dd20307ea54091
d09970bf956c067a3bcc05af03c41e01af949a131533778bf6ee3b546caf2eabe9d53d0fb2e8cc952b7e0f5326a69ed2
e58e088729a1d85971c6b2e129a5643f3ac43da031e655b27081f10543262cf9d72d6f64d5d96387ac0d43da3e3a03da
0c309af121dcf3e99192efa754eab6960c256fffd4c546208e292e0ab9894e3605db098dc16b40f17c320aa4a0e42fc8b
105c22f08c9bc6537182c24e32062c6cd6d7ec7062a0c2c2ecdac1588c82185cdc61d874ee916a7873ac54cddf929354
f307e870011704a0e9fbc5c7802d6140134028aca0e78a7e2f3d9e5c7e49e20c3a56b624bfea51196ec9e88e4e56be38
ff56031369f45f1e03be826d44a182f270c153ee0d9f8cf9f1f4132f33974e37c7887d5b857365c873cb218cbf20d4be
3abdb2a2011b14add0a5672e01e5845421cf6dd6faca1f2f443757aae575c53ab797c2227ecdab03882bbb4599318ce
fafa72fa0c9a0f5a51d13c9d0e5d25bfcfb0154ed25895260a9df8743ac188714a3f16960e6e2ff663c08bffd41743d

50960ea2f28cda0bc3bd4a180e297b5b41c700b674cb31d99c7f2a1445e121e772984abff2bbe3f42d757ceeda3d03fb
1ffe710aecabda21d738b1f4620e757e57b123dbc3c4aa5d9617dfa72f4a12d788ca596af14bea583f502f16fdc13a5e
739afb0715424af2767049f6b9aa107f69c5da0e85f6d8c5e46507e14616d5d0b797c3dea8b74a1b12d4e47ba7f57f09
d515f6c7314543f78b5e85329d50c5f96ee2f55bbe0df742b4003b24ccb4598a64413ee4807dc7f2a9c0b92424e4ae1
b418a3cdf02ea4da5c3b12139348aa7022cc8272a3a1714ee3e4ae111cfff1bdf62c503c80bdf27b2feaea0d5ab8fe0
0f9cec66e570b00fd24b4a2ed9a5f6384f148a4d6325110a41ca5659ebc5b98721d298a52819b6fb150f273383f1c575
4d320be428941922da790e17f482989c365c078f7f3ae100965e1b38c052041165295157e1a7c5b7a57671b842d4d85a
7d971323ad1f45e17a16c4656d889fc75c12fc3d8033f598306196e29571e414281c5da19c12605f48347ad5b4648e37
1757cbe1c40adb93052af1d6110cfbf611af5c8fc682b7e2ade3bfcab85c7717d19fc9f97964ba6025aebbc91a6671e2
59949dcf40984342118de1f6b514a7786bd4f6598ffbe1604cef476b2a4cb1343db608aca09d1d38fc23e98ee9c65e7f
6023a8d1e61fd4f34f753454bd8e858c8ad6be6403edc599c220e03ca917db765980ac781e758179cd93983e9c1e769e
4241d47c

attribution data for node 3:

1571e10db7f8aa9f8e7e99caaf9c892e106c817df1d8e3b7b0e39d1c48f631e473e17e205489dd7b3c634cac3be0825c
bf01418cd46e83c24b8d9c207742db9a0f0e5bcd888086498159f08080ba7bf36dee297079eb841391ccd3096da76461
e314863b6412efe0ffe228d51c6097db10d3edb2e50ea679820613bfe9db11ba02920ab4c1f2a79890d997f1fc022f3a
b78f0029cc6de0c90be74d55f4a99bf77a50e20f8d076fe61776190a61d2f41c408871c0279309cba3b60fcdc7efc4a0
e90b47cb4a418fc78f362ecc7f15ebbc9f854c09c7be300ebc1a40a69d4c7cb7a19779b6905e82bec221a709c1dab8c
bdcde7b527aca3f54bde651aa9f3f2178829cee3f1c0b9292758a40cc63bd998fcd0d3ed4bdcaf1023267b8f8e44130a
63ad15f76145936552381eabb6d684c0a3af6ba8efcf207cebaea5b7acdbb63f8e7221102409d10c23f0514dc9f4d0ef
b2264161a193a999a23e992632710580a0d320f676d367b9190721194514457761af05207cdab2b6328b1b3767each36
a7ef4f7bd2e16762d13df188e0898b7410f62459458712a44bf594ae662fd89eb300abb6952ff8ad40164f2bcd7f86db
5c7650b654b79046de55d51aa8061ce35f867a3e8f5bf98ad920be827101c64fb871d86e53a4b3c0455bfac578416821
8aa72cbee86d9c750a9fa63c363a8b43d7b7f4b2762516706a306f0aa3be1ec788b5e13f8b24837e53ac414f211e11c7a
093cd9653dfa5fba4e377c79adfa5e841e2ddb6afcf054fc715c05ddc6c8fc3e1ee3406e1ffcecb2df77dc2f02652614d1
bfcfaddebaa53ba919c7051034e2c7b7cfaabdf89f26e7f8e3f956d205dfab747ad0cb505b85b54a68439621b25832cb
c2898919d0cd7c0a64cfd235388982dd4dd68240cb668f57e1d2619a656ed326f8c92357ee0d9acead3c20008bc5f04c
a8059b55d77861c6d04dfc57cfba57315075acbe1451c96cf28e1e328e142890248d18f53b5d3513ce574dea7156cf59
6fdb3d909095ec287651f9cf1bcdc791c5938a5dd9b47e84c004d24ab3ae74492c7e8dcc1da15f65324be2672947ec82
074cac8ce2b925bc555facbbf1b55d63ea6fba6a785c97d4caf2e1dad9551b7f66c31caae5ebc7c0047e892f201308f
cf452c588be0e63d89152113d87bf0dbd01603b4cdc7f0b724b0714a9851887a01f709408882e18230fe810b9fafa58a
666654576d8eba3005f07221f55a6193815a672e5db56204053bc4286fa3db38250396309fd28011b5708a26a2d76c4a
333b69b6bfd272fb

forwarding error packet

shared_secret = 3a6b412548762f0dbccce5c7ae7bb8147d1caf9b5471c34120b30bc9c04891cc

ammag_key = 1bf08df8628d452141d56adfd1b25c1530d7921c23cecfc749ac03a9b694b0d3

stream =

6149f48b5a7e8f3d6f5d870b7a698e204cf64452aab4484ff1dee671fe63fd4b5f1b78ee2047dfa61e3d576b149bedaf
83058f85f06a3172a3223ad6c4732d96b32955da7d2feb4140e58d86fc0f2eb5d9d1878e6f8a7f65ab9212030e8e9155
73ebbd7f35e1a430890be7e67c3fb4bbf2def662fa625421e7b411c29ebe81ec67b77355596b05cc155755664e59c16e
21410aabe53e80404a615f44ebb31b365ca77a6e91241667b26c6cad24fb2324cf64e8b9dd6e2ce65f1f098cfd1ef41b
a2d4c7def0ff165a0e7c84e7597c40e3dffe97d417c144545a0e38ee33ebaae12cc0c14650e453d46bfc48c0514f3547
73435ee89b7b2810606eb73262c77a1d67f3633705178d79a1078c3a01b5fad9651feb63603d19dec3a00c1f69af2d
ab2595931ca50d8280758b1cc91ba2dc43dbbc3d91bf25c08b46c2ecef7a32cec64d4b61ee3a629ef563afe058b71e71
bcb69033948bc8728c5ebe65ec596e4f305b9fc159d53f723dfc95b57f3d51717f1c89af97a6d587e89e62efcc92198a
1b2bd66e2d875505ea4046c04389f8cb0ee98f0af03af2652e2f3d9a9c48430f2891a4d9b16e7d18099e4a3dd334c24a
ba1e2450792c2f22092c170da549d43a440021e699bd6c20d8bbf1961100a01ebccce06a4609f5ad93066287ac68294c
fa9ea7cea03a508983b134a9f0118b16409a61c06aaa95897d2067cb7cd59123f3e2ccf0e16091571d616c44818f118b
b7835a679f5c0eea8cf1bd5479882b2c2a341ec26dbe5da87b3d37d66b1fbd176f71ab203a3b6eaf7f214d579e7d0e4a
3e59089ebd26ba04a62403ae7a793516ec16d971d51c5c0107a917d1a70221e6de16edca7cb057c7d06902b5191f298a
a4d478a0c3a6260c257eae504ebbf2b591688e6f3f77af770b6f566ae9868d2f26c12574d3bf9323af59f0fe0072ff94
ae597c2aa6fbcfb0831989e02f9d3d1b9fd6dd97f509185d9ecbf272e38bd621ee94b97af8e1cd43853a8f6aa6e83725
85c71bf88246d064ade524e1e0bd8496b620c4c2d3ae06b6b064c97536aaf8d515046229f72bee8aa398cd0cc21afd54
49595016bef4c77cb1e2e9d31fe1ca3ffde06515e6a4331ccc84edf702e5777b10fc844faf17601a4be3235931f6fec

4582a8d247c1d6e4773f8fb6de320cf02bbb1767192782dc550d8e266e727a2aa2a414b816d1826ea46af7170153719
3c22bbcc0123d7ff5a23b0aa8d7967f36fef27b14fe1866ff3ab215eb29e07af49e19174887d71da7e7fe1b7aa1b3c80
5c063e0fafedf125fa6c57e38cce33a3f7bb35fd8a9f0950de3c22e49743c05f40bc55f960b8a8b5e2fde4bb229f1255
38438de418cb318d13968532499118cb7dcaaf8b6d635ac4001273bdafd12c8ea0702fb2f0dac81dbaaf68c1c3226638
2b293fa3951cb952ed5c1bdc41750cdcbc0bd62c51bb685616874e251f031a929c06faef5bfbcb0857f815ae20620b823f
0abecfb5

error packet for node 2:

145bc1c63058f7204abbd2320d422e69fb1b3801a14312f81e5e29e6b5f4774cfed8a25241d3dfb7466e749c1b326155
9e49090853612e07bd669dfb5f4c54162fa504138dabd6ebcf0db8017840c35f12a2cfb84f89cc7c8959a6d51815b1d2
c5136cedec2e4106bb5f2af9a21bd0a02c40b44ded6e6a90a145850614fb1b0eef2a03389f3f2693bc8a755630fc81ff
f1d87a147052863a71ad5aeb8770537f333e07d841761ec448257f948540d8f26b1d5b66f86e073746106dfdbb86ac9
475acf59d95ece037fba360670d924dce53aaa74262711e62a8fc9eb70cd8618fbedae22853d3053c7f10b1a6f75369d
7f73c419baa7dbf9f1fc5895362dcc8b6bd60cca4943ef7143956c91992119bccbe1666a20b7de8a2ff30a46112b53a6
bb79b763903ecbd1f1f74952fb1d8eb0950c504df31fe702679c23b463f82a921a2c931500ab08e686cfff2d87258d25
4fb17843959cccd265a57ba26c740f0f231bb76df932b50c12c10be90174b37d454a3f8b284c849e86578a6182c4a7b2
e47dd57d44730a1be9fec4ad07287a397e28dce4fda57e9cdfdb2eb5afdf0d38ef19d982341d18d07a556bb16c1416f4
80a396f278373b8fd9897023a4ac506e65cf4c306377730f9c8ca63cf47565240b59c4861e52f1dab84d938e96fb3182
0064d534aca05fd3d2600834fe4caea98f2a748eb8f200af77bd9fbf46141952b9ddda66ef0ebea17ea1e7bb5bce65b6
e71554c56dd0d4e14f4cf74c77a150776bf31e7419756c71e7421dc22efe9cf01de9e19fc8808d5b525431b944400db1
21a77994518d6025711cb25a18774068bba7faaa16d8f65c91bec8768848333156dcb4a08dfbbd9fef392da3e4de13d4
d74e83a7d6e46cfe530ee7a6f711e2caf8ad5461ba8177b2ef0a518baf9058ff9156e6aa7b08d938bd8d1485a787809d
7b4c8aed97be880708470cd2b2cdf8e2f13428cc4b04ef1f2acbc9562f3693b948d0aa94b0e6113cfa684f8e4a67dc4
31dfb835726874bef1de36f273f52ee694ec46b0700f77f8538067642a552968e866a72a3f2031ad116663ac17b172b4
46c5bc705b84777363a9a3fdc6443c07b2f4ef58858122168d4ebbaee920cefc312e1cea870ed6e15eec046ab2073bbf
08b0a3366f55cfc6ad4681a12ab0946534e7b6f90ea8992d530ec3daa6b523b3cf03101c60cadd914f30dec932c1ef43
41b5a8efac3c921e203574cfe0f1f83433fddb8ccfd273f7c3cab7bc27efe3bb61fdccd5146f1185364b9b621e7fb2b7
4b51f5ee6be72ab6ff46a6359dc2c855e61469724c1dbeb273df9d2e1c1fb74891239c0019dc12d5c7535f7238f963b7
61d7102b585372cf021b64c4fc85bfb3161e59d2e298bba44cfcd34d6859d9dba9dc6271e5047d525468c814f2ae43847
4b0a977273036da1a2292f88fcfb89574a6bdca1185b40f8aa54026d5926725f99ef028da1be892e3586361efe15f4a1
48ff1bc9

attribution data for node 2:

34e34397b8621ec2f2b54dbe6c14073e267324cd60b152bce76aec8729a6ddfeb61bc263be4b57bd592aae604a32bea6
9afe6ef4a6b573c26b17d69381ec1fc9b5aa769d148f2f1f8b5377a73840bb6dc641f68e356323d766ffff0aaca5039fe
7fc27038195844951a97d5a5b26698a4ca1e9cd4bca1fccaa0aac5fee91b18977d2ad0e399ba159733fc98f6e96898ebc
39bf0028c9c81619233bab6fad0328aa183a635fac20437fa6e00e899b2527c3697a8ab7342e42d55a679b176ab76671
fcd480a9894cb897fa6af0a45b917a162bed6c491972403185df7235502f7ada65769d1bfb12d29f10e25b0d3cc08bbf
6de8481ac5c04df32b4533b4f764c2aefb7333202645a629fb16e4a208e9045dc36830759c852b31dd613d8b2b10bbea
d1ed4eb60c85e8a4517deba5ab53e39867c83c26802beee2ee545bdd713208751added5fc0eb2bc89a5aa2dec18ee37
dac39f22a33b60cc1a369d24de9f3d2d8b63c039e248806de4e36a47c7a0aed30edd30c3d62debd1ad82bf7aedd7ede
c413850d91c261e12beec7ad1586a9ad25b2db62c58ca17119d61dcc4f3e5c4520c42a8e384a45d8659b338b3a08f9e1
23a1d3781f5fc97564ccff2c1d97f06fa0150cfa1e20eacabefb0c339ec109336d207cc63d9170752fc58314c43e6d4a
528fd0975afa85f3aa186ff1b6b8cb12c97ed4ace295b0ef5f075f0217665b8bb180246b87982d10f43c9866b2287810
6f5214e99188781180478b07764a5e12876ddcb709e0a0a8dd42cf004c695c6fc1669a6fd0e4a1ca54b024d0d80eac49
2a9e5036501f36fb25b72a054189294955830e43c18e55668337c8c6733abb09fc2d4ade18d5a853a2b82f7b4d77151a
64985004f1d9218f2945b63c56fdebd1e96a2a7e49fa70acb4c39873947b83c191c10e9a8f40f60f3ad5a2be47145c22
ea59ed3f5f4e61cb069e875fb67142d281d784bf925cc286eacc2c43e94d08da4924b83e58dbf2e43fa625bdd620eba6
d9ce960fff17d14ed1f2dbee7d08eceb540fdc75ff06dabc767267658fad8ce99e2a3236e46d2deedcb51c3c6f8158935
7edebac9772a70b3d910d83cd1b9ce6534a011e9fa557b891a23b5d88afcc0d9856c6dabeab25eea55e9a248182229e4
927f268fe5431672fccce52f434ca3d27d1a2136bae5770bb36920df12fbc01d0e8165610efa04794f414c1417f1d4059
435c5385bfe2de83ce0e238d6fd2dbd3c0487c69843298577bfa480fe2a16ab2a0e4bc712cd8b5a14871cda61c993b68
35303d9043d7689a

forwarding error packet

shared_secret = a6519e98832a0b179f62123b3567c106db99ee37bef036e783263602f3488fae

ammag_key = 59ee5867c5c151daa31e36ee42530f429c433836286e63744f2020b980302564

stream =
0f10c86f05968dd91188b998ee45dcddfbbf89fe9a99aa6375c42ed5520a257e048456fe417c15219ce39d921555956ae
2ff795177c63c819233f3bcb9b8b28e5ac6e33a3f9b87ca62dff43f4cc4a2755830a3b7e98c326b278e2bd31f4a9973e
e99121c62873f5bfb2d159d3d48c5851e3b341f9f6634f51939188c3b9fff45feeb11160bb39ce3332168b8e744a92107
db575ace7866e4b8f390f1edc4acd726ed106555900a0832575c3a7ad11bb1fe388ff32b99bcf2a0d0767a83cf293a22
0a983ad014d404bfa20022d8b369fe06f7ecc9c74751dcda0ff39d8bca74bf9956745ba4e5d299e0da8f68a9f660040b
eac03e795a046640cf8271307a8b64780b0588422f5a60ed7e36d60417562938b400802dac5f87f267204b6d5bcfd8a0
5b221ec294d883271b06ca709042ed5dbb64a7328d2796195eab4512994fb88919c73b3e5dd7bf68b2136d34cff39b3b
e266b71e004509bf975a240800bb8ae5eed248423a991ae80ef751b2d03b67fb93ffdd7969d5b500fe446a4fffb4cd04d
0767a5d367ebd3f8f260f38ae1e9d9f9a7bd1a99ca1e10ee36bd241f06fc2b481c9b7450d9c9704204666807783264a0
e93468e22db4dc4a7a4db2963dddf4366d08e225cf94848aac794bcecb7e850113e38cc3647a03a5dffa3442b1bb58b1d
e7fa7f436feb4d7c23cbd2de6d55d4025fcd383cc9d49c0b130e2fd5a9097c216683c842f898a8a2159761cca9aa1c81
8194e3b7bea6da6652d5189f3b6b0ca1d5398b6d14e311d9c7f00399c29e94deb98496f4cd97c5d7d6a65cab3791f60
d728d6422a422c0cff5f7dfd4ce2d7e8d38dd71ae18763acc832c57275497f61b2620cca13cc64c0c48353f3817016f9
1448d6fc1cc451ee1f4a429e43292bbcd54fcd807e2c47675bac1781d9d81e9e6dc69028d428f5ee261750f626bcaf41
6a0e7badadf73fe1922207ae6c5209d16849e4a108f4a6f38694075f55177105ac4c2b97f6a474b94c03257d8d12b019
6e2905d914b8c2213a1b9dc9608e1a2a1e03fe0820a813275de83be5e9734875787a9e006eb8574c23ddd49e2347d1ec
fcedf3caa0a5dd45666368525b48ac14225d6422f82dbf59860ee4dc78e845d3c57668ce9b9e7a8d012491cef242078b
458a956ad67c360fb6d8b86ab201d6217e49b55fa02a1dea2dbe88d0b08d30670d1b93c35cc5e41e088fccb267e41d61
51cf8560496e1beee680744d9dabb383a957466b4dc3e2bce7b135211da483d998a22fa687cc609641126c5dee3ed
87291067916b5b065f40582163291d48e81ecd975d0d6fd52a31754f8ef15e43a560bd30ea5bf21915bd2e7007e607ab
bc6261edc8430cc7f789675b1fe83e807c5c475bd5178eba2fc40674706b0a68c6a428e5dec36e413e653c6db1178923
ff87e2389a78bf9e93b713de4f4753f9f9d6a361369b609e1970c91ff9bd191c472e0bf2e8681412260ad0ef5855dc39
f2084d45

error packet for node 1:

1b4b09a935ce7af95b336baae307f2b400e3a7e808d9b4cf421cc4b3955620acb69dcd656128dae8857adbd4e6b37fb
b1be9c1f2f02e61e9e59a630c4c77cf383cb37b07413aa4de2f2fbf5b40ae40a91a8f4c6d74aeacef1bb1be4ecbc26ec
2c824d2bc45db4b9098e732a769788f1cff3f5b41b0d25c132d40dc5ad045ef0043b15332ca3c5a09de2cdb17455a0f8
2a8f20da08346282823dab062cddb2111e238528141d69de13de6d83994fbc711e3e269df63a12d3a4177c5c149150eb
4dc2f589cd8acabcbdbba14dec3b0dada12d663b36176cd3c257c5460bab93981ad99f58660efa9b31d7e63b399153296
95b3fa60e0a3bdb93e7e29a54ca6a8f360d3848866198f9c3da3ba958e7730847fe1e6478ce8597848d3412b4ae48b06
e05ba9a104e648f6eaf183226b5f63ed2e68f77f7e38711b393766a6fab7921b03eba82b5d7cb78e34dc961948d6161e
add7cf5d95d9c56df2ff5faa6ccf85eacdc9ff2fc3abafe41c365a5bd14fd486d6b5e2f24199319e7813e02e798877ff
e31a70ae2398d9e31b9e3727e6c1a3c0d995c67d37bb6e72e9660aaaa9232670f382add2edd468927e3303b614267254
6997fe105583e7c5a3c4c2b599731308b5416e6c9a3f3ba55b181ad0439d3535356108b059f2cb8742eed7a58d4eba9f
e79eaa77c34b12aff1abdae93197aab0e74cb271269ca464b3b06aef1d6573df5e1224179616036b368677f2647937
6681b772d3760e871d99efd34cca5cd6beca95190d967da820b21e5bec60082ea46d776b0517488c84f26d12873912d1
f68fafd67bcf4c298e43cfa754959780682a2db0f75f95f0598c0d04fd014c50e4beb86a9e37d95f2bba7e5065ae052d
c306555bca203d104c44a538b438c9762de299e1c4ad30d5b4a6460a76484661fc907682af202cd69b9a4473813b2fdc
1142f1403a49b7e69a650b7cde9ff133997dcc6d43f049ecac5fce097a21e2bce49c810346426585e3a5a18569b4cddd
5ff6bdec66d0b69fcbc5ab3b137b34cc8aefb8b850a764df0e685c81c326611d901c392a519866e132bbb73234f6a358
ba284fbabf21aa3605cacbaf9d0c901390a98b7a7dac9d4f0b405f7291c88b2ff45874241c90ac6c5fc895a440453c34
4d3a365cb929f9c91b9e39cb98b142444aae03a6ae8284c77eb04b0a163813d4c21883df3c0f398f47bf127b5525f222
107a2d8fe55289f0cfd3f4bbad6c5387b0594ef8a966afc9e804ccaf75fe39f35c6446f7ee076d433f2f8a44dba1515a
cc78e589fa8c71b0a006fe14feebd51d0e0aa4e51110d16759eee86192eee90b34432130f387e0ccd2ee71023f1f641c
ddb571c690107e08f592039fe36d81336a421e89378f351e633932a2f5f697d25b620ffb8e84bb6478e9bd229bf3b164
b48d754ae97bd23f19e3c56b3bcdaaeb3bd7fc02ec02066b324cb72a09b6b43dec1097f49d69d3c138ce6f1a6402898
baf7568c

attribution data for node 1:

74a4ea61339463642a2182758871b2ea724f31f531aa98d80f1c3043febca41d5ee52e8b1e127e61719a0d078db89097
48d57839e58424b91f063c4fbc8a221bef261140e66a9b596ca6d420a973ad54fef30646ae53ccf0855b61f291a81e0e
c6dc0f6bf69f0ca0e5889b7e23f577ba67d2a7d6a2aa91264ab9b20630ed52f8ed56cc10a869807cd1a4c2cd802d8433
fee5685d6a04edb0bff248a480b93b01904bed3bb31705d1ecb7332004290cc0cd9cc2f7907cf9db28eec02985301668
f53fbc28c3e095c8f3a6cd8cab28e5e442fd9ba608b8b12e098731bbfda755393bd403c62289093b40390b2bae337fc8
7d2606ca028311d73a9ffbdffef56020c735ada30f54e577c6a9ec515ae2739290609503404b118d7494499ecf0457d7

5015bb60a16288a4959d74cf5ac5d8d6c113de39f748a418d2a7083b90c9c0a09a49149fd1f2d2cde4412e5aa2421eca
6fd4f6fe6b2c362ff37d1a0608c931c7ca3b8fefcfcd4c44ef9c38357a0767b14f83cb49bd1989fb3f8e2ab202ac98bd8
439790764a40bf309ea2205c1632610956495720030a25dc7118e0c868fdfa78c3e9ecce58215579a0581b3bafdb7dbb
e53be9e904567fdc0ce1236aab5d22f1ebc18997e3ea83d362d891e04c5785fd5238326f767bce499209f8db211a50e1
402160486e98e7235cf397dbb9ae19fd9b79ef589c821c6f99f28be33452405a003b33f4540fe0a41dfcc286f4d7cc10
b70552ba7850869abadcd4bb7f256823face853633d6e2a999ac9fcd259c71d08e266db5d744e1909a62c0db673745ad
9585949d108ab96640d2bc27fb4acac7fa8b170a30055a5ede90e004df9a44bdc29aeb4a6bec1e85dde1de6aaf01c6a5
d12405d0bec22f49026cb23264f8c04b8401d3c2ab6f2e109948b6193b3bec27adfe19fb8afb8a92364d6fc5b219e873
7d583e7ff3a4bcb75d53edda3bf3f52896ac36d8a877ad9f296ea6c045603fc62ac4ae41272bde85ef7c3b3fd3538aac
fd5b025fefbe277c2906821ecb20e6f75ea479fa3280f9100fb0089203455c56b6bc775e5c2f0f58c63edd63fa3eec0b
40da4b276d0d41da2ec0ead865a98d12bc694e23d8eaadd2b4d0ee88e9570c88fb878930f492e036d27998d593e47763
927ff7eb80b188864a3846dd2238f7f95f4090ed399ae95deaeb37abca1cf37c397cc12189affb42dca46b4ff6988eb8
c060691d155302d448f50ff70a794d97c0408f8cee9385d6a71fa412e36edcb22dbf433db9db4779f27b682ee17fc05e
70c8e794b9f7f6d1

forwarding error packet

shared_secret = 53eb63ea8a3fec3b3cd433b85cd62a4b145e1dda09391b348c4e1cd36a03ea66

ammag_key = 3761ba4d3e726d8abb16cba5950ee976b84937b61b7ad09e741724d7dee12eb5

stream =

3699fd352a948a05f604763c0bca2968d5eaca2b0118602e52e59121f050936c8dd90c24df7dc8cf8f1665e39a6c75e9
e2c0900ea245c9ed3b0008148e0ae18bbfaea0c711d67eade980c6f5452e91a06b070bbde68b5494a92575c114660fb5
3cf04bf686e67ffa4a0f5ae41a59a39a8515cb686db553d25e71e7a97cc2febca55df2711b6209c502b2f8827b13d3a
d2f491c45a0cafe7b4d8d8810e805dee25d676ce92e0619b9c206f922132d806138713a8f69589c18c3fdc5acee41c12
34b17ecab96b8c56a46787bba2c062468a13919afc18513835b472a79b2c35f9a91f38eb3b9e998b1000cc4a0dbd62ac
1a5cc8102e373526d7e8f3c3a1b4bf2f8a3947fe350cb89f73aa1bb054edfa9895c0fc971c2b5056dc8665902b51fce
d6dff80c4d247db977c15a710ce280fbd0ae3ca2a245b1c967aeb5a1a4a441c50bc9ceb33ca64b5ca93bd8b50060520f
35a54a148a4112e8762f9d0b0f78a7f46a5f06c7a4b0845d020eb505c9e527aabab71009289a6919520d32af1f9f51ce
4b3655c6f1aae1e26a16dc9aae55e9d4a6f91d4ba76e96fcb851161da3fc39d0d97ce30a5855c75ac2f613ff36a24801
bcbd33f0ce4a3572b9a2fca21efb3b07897fc07ee71e8b1c0c6f8dbb7d2c4ed13f11249414fc94047d1a4a0be94d45db
56af4c1a3bf39c9c5aa18209eaebb9e025f670d4c8cc1ee598c912db154eaa3d0c93cb3957e126c50486bf98c852ad32
6b5f80a19df6b2791f3d65b8586474f4c5dcb2aca0911d2257d1bb6a1e9fc1435be879e75d23290f9feb93ed40baeaca
1c399fc91fb1da3e5f0f5d63e543a8d12fe6f7e654026d3a118ab58cb14bef9328d4af254215eb1f639828cc6405a3ab
02d90bb70a798787a52c29b3a28fc67b0908563a65f08112abd4e9115cb01db09460c602aba3ddc375569dc3abe42c61
c5ea7feb39ad8b05d8e2718e68806c0e1c34b0bc85492f985f8b3e76197a50d63982b780187078f5c59ebd814faeffc
7b2c6ee39d4f9c8c45fb5f685756c563f4b9d028fe7981b70752f5a31e44ba051ab40f3604c8596f1e95dc9b0911e7ed
e63d69b5eeecd245fbecbfc233cf6eba842c0fec795a5adeab2100b1a1bc62c15046d48ec5709da4af64f59a2e552ddb
dcda1f543bb4b687e79f2253ff0cd9ba4e6bfae8e510e5147273d288fd4336dbd0b6617bf0ef71c0b4f1f9c1dc999c17
ad32fe196b1e2b27baf4d59bba8e5193a9595bd786be00c32bae89c5dbed1e994fddffbec49d0e2d270bcc1068850e5d
7e7652e274909b3cf5e3bc6bf64def0bbeac974a76d835e9a10bdd7896f27833232d907b7405260e3c986569bb8fdd65
a55b020b91149f27bda9e63b4c2cc5370bcc81ef044a68c40c1b178e4265440334cc40f59ab5f82a022532805bfa6592
57c8d8ab9b4aef6abbd05de284c2eb165ef35737e3d387988c566f7b1ca0b1fc3e7b4ed991b77f23775e1c36a09a9913
84a33b78

error packet for node 0:

2dd2f49c1f5af0fcad371d96e8cddbdc5096dc309c1d4e110f955926506b3c03b44c192896f45610741c85ed4074212
537e0c118d472ff3a559ae244acd9d783c65977765c5d4e00b723d00f12475aafafff7b31c1be5a589e6e25f8da2959
107206dd42bbcb43438129ce6cce2b6b4ae63edc76b876136ca5ea6cd1c6a04ca86eca143d15e53ccdc9e23953e49dc2
f87bb11e5238cd6536e57387225b8fff3bf5f3e686fd08458ffe0211b87d64770db9353500af9b122828a006da754cf9
79738b4374e146ea79dd93656170b89c98c5f2299d6e9c0410c826c721950c780486cd6d5b7130380d7eaff994a8503a
8fef3270ce94889fe996da66ed121741987010f785494415ca991b2e8b39ef2df6bde98efd2aec7d251b2772485194c8
368451ad49c2354f9d30d95367bde316fec6cbdddc7dc0d25e99d3075e13d3de0822669861dafcd29de74eac48b64411
987285491f98d78584d0c2a163b7221ea796f9e8671b2bb91e38ef5e18aaf32c6c02f2fb690358872a1ed28166172631
a82c2568d23238017188ebbd48944a147f6cdb3690d5f88e51371cb70adf1fa02afe4ed8b581afc8bcc5104922843a55
d52acde09bc9d2b71a663e178788280f3c3eae127d21b0b95777976b3eb17be40a702c244d0e5f833ff49dae6403ff44
b131e66df8b88e33ab0a58e379f2c34bf5113c66b9ea8241fc7aa2b1fa53cf4ed3cdd91d407730c66fb039ef3a36d405
0dde37d34e80bcfe02a48a6b14ae28227b1627b5ad07608a7763a531f2ffc96dff850e8c583461831b19feffc783bc1b

eab6301f647e9617d14c92c4b1d63f5147ccda56a35df8ca4806b8884c4aa3c3cc6a174fdc2232404822569c01aba686c1df5eccc059ba97e9688c8b16b70f0d24eacfdba15db1c71f72af1b2af85bd168f0b0800483f115eccc9b02adf03bd4a88eab03e43ce342877af2b61f9d3d85497cd1c6b96674f3d4f07f635bb26add1e36835e321d70263b1c04234e222124dad30fffb9f2a138e3ef453442df1af7e566890aedee568093aa922dd62db188aa8361c55503f8e2c2e6ba93de744b55c15260f15ec8e69bb01048ca1fa7bbbd26975bde80930a5b95054688a0ea73af0353cc84b997626a987cc06a517e18f91e02908829d4f4efc011b9867bd9bfe04c5f94e4b9261d30cc39982eb7b250f12aee2a4cce0484ff34eebba89bc6e35bd48d3968e4ca2d77527212017e202141900152f2fd8af0ac3aa456aae13276a13b9b9492a9a636e18244654b3245f07b20eb76b8e1cea8c55e5427f08a63a16b0a633af67c8e48ef8e53519041c9138176eb14b8782c6c2ee76146b8490b97978ee73cd0104e12f483be5a4af414404618e9f6633c55dda6f22252cb793d3d16fae4f0e1431434e7acc8fa2c009d4f6e345ade172313d558a4e61b4377e31b8ed4e28f7cd13a7fe3f72a409bc3bdabfe0ba47a6d861e21f64d2fac706dab18b3e546df4

attribution data for node 0:

84986c936d26bfd3bb2d34d3ec62cfd63e0032fdb3d9d75f3e5d456f73dffa7e35aab1db4f1bd3b98ff585caf004f656c51037a3f4e810d275f3f6aea0c8e3a125ebee5f374b6440bcb9bb2955ebf706f42be9999a62ed49c7a81fc73c0b4a16419fd6d334532f40bf179dd19afec21bd8519d5e6ebc3802501ef373bc378eee1f14a6fc5fab5b697c91ce31d5922199d1b0ad5ee12176aacafc7c81d54bc5b8fb7e63f3bfd40a3b6e21f985340cbd1c124c7f85f0369d1aa86ebc66def417107a7861131c8bcd73e8946f4fb54bfac87a2dc15bd7af642f32ae583646141e8875ef81ec9083d7e32d5f135131eab7a43803360434100ff67087762bbe3d6afe2034f5746b8c50e0c3c20dd62a4c174c38b1df7365dccebc7f24f19406649fbf48981448abe5c858bbd4bef6eb983ae7a23e9309fb33b5e7c0522554e88ca04b1d65fc190947dead8c0ccd32932976537d869b5ca53ed4945bccafab2a014ea4cbdc6b0250b25be66ba0afff2ff19c0058c68344fd1b9c472567147525b13b1bc27563e61310110935cf89fda0e34d0575e2389d57bdf2869398ca2965f64a6f04e1d1c2edf2082b97054264a47824dd1a9691c27902b39d57ae4a94dd6481954a9bd1b5cff4ab29ca221fa2bf9b28a362c9661206f896fc7cec563fb80aa5eacbc26c09fa4ef7a981e63028a9c4dac12f82ccb5bea090d56bbb1a4c431e315d9a169299224a8dbd099fb67ea61dfc604edf8a18ee742550b636836bb552dabb28820221bf8546331f32b0c143c1c89310c4fa2e1e0e895ce1a1eb0f43278fdb528131a3e32bfffef0c6de9006418f5309cba773ca38b6ad8507cc59445ccc0257506ebc16a4c01d4cd97e03fcf7a2049fea0db28447858f73b8e9f9e98b391b136c9dc510288630a1f0af93b26a8891b857bfe4b818af99a1e011e6dbaa53982d29cf74ae7dffe45545279f19931708ed3eede5e82280eab908e8eb80abff3f1f023ab66869297b40da8496861dc455ac3abe1efa8a6f9e2c4eda48025d43a486a3f26f269743eaa30d6f0e1f48db6287751358a41f5b07aee0f098862e3493731fe2697acce734f004907c6f11eef189424fee52cd30ad708707eaf2e441f52bcf3d0c5440c1742458653c0c8a27b5ade784d9e09c8b47f1671901a29360e7e5e94946b9c75752a1a8d599d2a3e14ac81b84d42115cd688c8383a64fc6e7e1dc5568bb4837358ebe63207a4067af66b2027ad2ce8fb7ae3a452d40723a51fdf9f9c9913e8029a222cf81d12ad41e58860d75deb6de30ad

Returning success

A successful payment without a fulfillment_payload using the parameters above would result in the following attribution data values:

attribution data for node 4:

d77d0711b5f71d1d1be56bd88b3bb7ebc1792bb739ea7ebc1bc3b031b8bc2df3a50e25aeb99f47d7f7ab39e24187d3f4df9c4333463b05

attribution data for node 3:

1571e10db7f8aa9f8e7e99caaf9c892e106c817df1d8e3b7b0e39d1c48f631e473e17e205489dd7b3c634cac3be0825cbf01418cd46e83c2

attribution data for node 2:

34e34397b8621ec2f2b54dbe6c14073e267324cd60b152bce76aec8729a6ddefb61bc263be4b57bd592aae604a32bea69afe6ef4a6b573c

attribution data for node 1:

74a4ea61339463642a2182758871b2ea724f31f531aa98d80f1c3043febca41d5ee52e8b1e127e61719a0d078db8909748d57839e58424b

attribution data for node 0:

84986c936d26bfd3bb2d34d3ec62cfd63e0032fdb3d9d75f3e5d456f73dffa7e35aab1db4f1bd3b98ff585caf004f656c51037a3f4e810d

A successful payment with a fulfillment_payload using the same route parameters contains the following fulfillment_payload_tlvs records:

```
type  = 1 (`padding`)  
value = 245 zero bytes
```

```
type  = 65537  
value = 070809
```

These two records serialize to 256 bytes: 247 bytes for the padding record and 9 bytes for the type 65537 record. With the 16-byte Poly1305 tag, this results in a 272-byte fulfillment_payload.

The following are the fulfillment_payload and attribution data values at each hop:

fulfillment payload for node 4:

b5b1fed711469adc0da510494818ce54624b537e7a8aba4b190403d90439798a694bdca0d0d53c18487d9940664a8416eb855290621e1

attribution data for node 4:

d77d0711b5f71d1d1be56bd88b3bb7ebc1792bb739ea7ebc1bc3b031b8bc2df3a50e25aeb99f47d7f7ab39e24187d3f4df9c4333463b05

fulfillment payload for node 3:

d4cd5f33730d592c4249f3e31233af0572bf729218ca375ddd8aa518b3fa466d6101f77904b013dd4ef3c912d9278edc09923288c1c6df3

attribution data for node 3:

1571e10db7f8aa9f8e7e99caaf9c892e106c817df1d8e3b7b0e39d1c48f631e473e17e205489dd7b3c634cac3be0825cbf01418cd46e83c2

fulfillment payload for node 2:

b584abb82973d6112d1474e8685a21253e4936c0b27e7f122c5443694d99bb263e1a8f9724f7cc7b50ce9e79cdbc63738a97bd0d31acee4

attribution data for node 2:

34e34397b8621ec2f2b54dbe6c14073e267324cd60b152bce76aec8729a6ddefb61bc263be4b57bd592aae604a32bea69afe6ef4a6b573c

fulfillment payload for node 1:

ba9463d72ce55bc83c9ccd70861ffdf8c5b1a9291be4d9257016ae3c6d3becc6765fe07333369e629ef7475898e535dda560281a4dcf2654

attribution data for node 1:

74a4ea61339463642a2182758871b2ea724f31f531aa98d80f1c3043febca41d5ee52e8b1e127e61719a0d078db8909748d57839e58424b

fulfillment payload for node 0:

8c0d9ee20671d1cdca98bb4c8dd5d490105b63021afcb90b22f33f1d9d6b7faafb86ec57ec4b56ad11e122bb0289403447a0b814ef8aefb

attribution data for node 0:

84986c936d26bfd3bb2d34d3ec62cfdb63e0032fdb3d9d75f3e5d456f73dffa7e35aab1db4f1bd3b98ff585caf004f656c51037a3f4e810d

References

Authors

[FIXME:]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

BOLT #7: P2P Node and Channel Discovery

This specification describes simple node discovery, channel discovery, and channel update mechanisms that do not rely on a third-party to disseminate the information.

Node and channel discovery serve two different purposes:

- Node discovery allows nodes to broadcast their ID, host, and port, so that other nodes can open connections and establish payment channels with them.
- Channel discovery allows the creation and maintenance of a local view of the network's topology, so that a node can discover routes to desired destinations.

To support channel and node discovery, three *gossip messages* are supported:

- For node discovery, peers exchange `node_announcement` messages, which supply additional information about the nodes. There may be multiple `node_announcement` messages, in order to update the node information.
- For channel discovery, peers in the network exchange `channel_announcement` messages containing information regarding new channels between the two nodes. They can also exchange `channel_update` messages, which update information about a channel. There can only be one valid `channel_announcement` for any channel, but at least two `channel_update` messages are expected.

Table of Contents

- [Definition of short_channel_id](#)
- [The announcement_signatures Message](#)
- [The channel_announcement Message](#)
- [The node_announcement Message](#)
- [The channel_update Message](#)
- [Query Messages](#)
- [Rebroadcasting](#)
- [HTLC Fees](#)
- [Pruning the Network View](#)
- [Recommendations for Routing](#)
- [References](#)

Definition of short_channel_id

The `short_channel_id` is the unique description of the funding transaction. It is constructed as follows: 1. the most significant 3 bytes: indicating the block height 2. the next 3 bytes: indicating the transaction index within the block 3. the least significant 2 bytes: indicating the output index that pays to the channel.

The standard human readable format for `short_channel_id` is created by printing the above components, in the order: block height, transaction index, and output index. Each component is printed as a decimal number, and separated from each other by the small letter x. For example, a `short_channel_id` might be written as 539268x845x1, indicating a channel on the output 1 of the transaction at index 845 of the block at height 539268.

Rationale

The `short_channel_id` human readable format is designed so that double-clicking or double-tapping it will select the entire ID on most systems. Humans prefer decimal when reading numbers, so the ID components are written in decimal. The small letter `x` is used since on most fonts, the `x` is visibly smaller than decimal digits, making it easy to visibly group each component of the ID.

The `announcement_signatures` Message

This is a direct message between the two endpoints of a channel and serves as an opt-in mechanism to allow the announcement of the channel to the rest of the network. It contains the necessary signatures, by the sender, to construct the `channel_announcement` message.

1. type: 259 (`announcement_signatures`)
2. data:
 - [`channel_id:channel_id`]
 - [`short_channel_id:short_channel_id`]
 - [`signature:node_signature`]
 - [`signature:bitcoin_signature`]

The willingness of the initiating node to announce the channel is signaled during channel opening by setting the `announce_channel` bit in `channel_flags` (see [BOLT #2](#)).

Requirements

The `announcement_signatures` message is created by constructing a `channel_announcement` message, corresponding to the newly confirmed channel funding transaction, and signing it with the secrets matching an endpoint's `node_id` and `bitcoin_key`.

A node: - If the `open_channel` message has the `announce_channel` bit set AND a shutdown message has not been sent: - After `channel_ready` has been sent and received AND the funding transaction has enough confirmations to ensure that it won't be reorganized: - MUST send `announcement_signatures` for the funding transaction. - After `splice_locked` has been sent and received AND the splice transaction has enough confirmations to ensure that it won't be reorganized: - MUST send `announcement_signatures` for the matching splice transaction. - Otherwise: - MUST NOT send the `announcement_signatures` message. - Upon reconnection (once the above timing requirements have been met): - If it has NOT previously received `announcement_signatures` for the funding transaction: - MUST send its own `announcement_signatures` message. - If it receives `announcement_signatures` for the funding transaction: - MUST respond with its own `announcement_signatures` message. - If it has NOT previously received `announcement_signatures` for a splice transaction: - MUST SET the `announcement_signatures` bit in the `retransmit_flags` of `my_current_funding_locked`. - If the `announcement_signatures` bit is set in the *remote* `retransmit_flags`: - MUST retransmit its `announcement_signatures` message.

A recipient node: - If the `short_channel_id` doesn't match one of its funding transactions: - SHOULD send a warning. - If the `node_signature` OR the `bitcoin_signature` is NOT correct: - MAY send a warning and close the connection, or send an error and fail the channel. - If it has sent AND received a valid `announcement_signatures` message: - If the funding transaction has at least 6 confirmations: - SHOULD queue the `channel_announcement` message for its peers. - If it has not sent `channel_ready`: - SHOULD defer handling

the announcement_signatures until after it has sent channel_ready. - If it has not sent splice_locked for the transaction matching this short_channel_id: - SHOULD defer handling the announcement_signatures until after it has sent splice_locked.

Rationale

Channels must not be announced before the funding transaction has enough confirmations, because a blockchain reorganization would otherwise invalidate the short_channel_id.

When splicing is used, a channel_announcement is generated for every splice transaction once both sides have sent splice_locked. This lets the network know that the transaction spending a currently active channel is a splice and not a closing transaction, and this channel can still be used with its updated short_channel_id.

The channel_announcement Message

This gossip message contains ownership information regarding a channel. It ties each on-chain Bitcoin key to the associated Lightning node key, and vice-versa. The channel is not practically usable until at least one side has announced its fee levels and expiry, using channel_update.

Proving the existence of a channel between node_1 and node_2 requires:

1. proving that the funding transaction pays to bitcoin_key_1 and bitcoin_key_2
2. proving that node_1 owns bitcoin_key_1
3. proving that node_2 owns bitcoin_key_2

Assuming that all nodes know the unspent transaction outputs, the first proof is accomplished by a node finding the output given by the short_channel_id and verifying that it is indeed a P2WSH funding transaction output for those keys specified in [BOLT #3](#).

The last two proofs are accomplished through explicit signatures: bitcoin_signature_1 and bitcoin_signature_2 are generated for each bitcoin_key and each of the corresponding node_ids are signed.

It's also necessary to prove that node_1 and node_2 both agree on the announcement message: this is accomplished by having a signature from each node_id (node_signature_1 and node_signature_2) signing the message.

1. type: 256 (channel_announcement)
2. data:
 - [signature:node_signature_1]
 - [signature:node_signature_2]
 - [signature:bitcoin_signature_1]
 - [signature:bitcoin_signature_2]
 - [u16:len]
 - [len*byte:features]
 - [chain_hash:chain_hash]
 - [short_channel_id:short_channel_id]
 - [point:node_id_1]
 - [point:node_id_2]
 - [point:bitcoin_key_1]

◦ [point:bitcoin_key_2]

Requirements

The origin node: - MUST set chain_hash to the 32-byte hash that uniquely identifies the chain that the channel was opened within: - for the *Bitcoin blockchain*: - MUST set chain_hash value (encoded in hex) equal to 6fe28c0ab6f1b372c1a6a246ae63f74f931e8365e15a089c68d6190000000000. - When announcing a channel creation: - MUST set short_channel_id to refer to the confirmed funding transaction, as specified in [BOLT #2](#). - When announcing a splice transaction: - MUST set short_channel_id to refer to the confirmed splice transaction for which splice_locked has been sent and received, as specified in [BOLT #2](#). - SHOULD keep relaying payments that use the short_channel_ids of its previous channel_announcements. - SHOULD send a new channel_update using the short_channel_id that matches the latest channel_announcement. - Note: the corresponding output MUST be a P2WSH, as described in [BOLT #3](#). - MUST set node_id_1 and node_id_2 to the public keys of the two nodes operating the channel, such that node_id_1 is the lexicographically-lesser of the two compressed keys sorted in ascending lexicographic order. - MUST set bitcoin_key_1 and bitcoin_key_2 to node_id_1 and node_id_2's respective funding_pubkeys. - MUST compute the double-SHA256 hash h of the message, beginning at offset 256, up to the end of the message. - Note: the hash skips the 4 signatures but hashes the rest of the message, including any future fields appended to the end. - MUST set node_signature_1 and node_signature_2 to valid signatures of the hash h (using node_id_1 and node_id_2's respective secrets). - MUST set bitcoin_signature_1 and bitcoin_signature_2 to valid signatures of the hash h (using bitcoin_key_1 and bitcoin_key_2's respective secrets). - MUST set features based on what features were negotiated for this channel, according to [BOLT #9](#). - MUST set len to the minimum length required to hold the features bits it sets. - If the funding transaction has less than 6 confirmations: - MUST NOT send channel_announcement.

The receiving node: - MUST verify the integrity AND authenticity of the message by verifying the signatures. - if node_id_1 is not lexicographically less than node_id_2: - SHOULD send a warning. - MAY close the connection. - MUST ignore the message. - if there is an unknown even bit in the features field: - MUST NOT attempt to route messages through the channel. - if the short_channel_id's output does NOT correspond to a P2WSH (using bitcoin_key_1 and bitcoin_key_2, as specified in [BOLT #3](#)) OR the output is spent: - MUST ignore the message. - if the specified chain_hash is unknown to the receiver: - MUST ignore the message. - if the short_channel_id's output does NOT have at least 6 confirmations: - MAY accept the message if the output is close to 6 confirmations, in case the receiving node hasn't received the latest block(s) yet. - otherwise: - SHOULD ignore the message. - otherwise: - if bitcoin_signature_1, bitcoin_signature_2, node_signature_1 OR node_signature_2 are invalid OR NOT correct: - SHOULD send a warning. - MAY close the connection. - MUST ignore the message. - otherwise: - if node_id_1 OR node_id_2 are blacklisted: - SHOULD ignore the message. - otherwise: - if the transaction referred to was NOT previously announced as a channel: - SHOULD queue the message for rebroadcasting. - MAY choose NOT to for messages longer than the minimum expected length. - if it has previously received a valid channel_announcement, for the same transaction, in the same block, but for a different node_id_1 or node_id_2: - SHOULD blacklist the previous message's node_id_1 and node_id_2, as well as this node_id_1 and node_id_2 AND forget any channels connected to them. - otherwise: - SHOULD store this channel_announcement. - once its funding output has been spent OR reorganized out: - SHOULD forget a channel after a 72-block delay. - SHOULD NOT rebroadcast this channel_announcement to its peers.

Rationale

Both nodes are required to sign to indicate they are willing to route other payments via this channel (i.e. be part of the public network); requiring their Bitcoin signatures proves that they control the channel.

The blacklisting of conflicting nodes disallows multiple different announcements. Such conflicting announcements should never be broadcast by any node, as this implies that keys have leaked.

While channels should not be advertised before they are sufficiently deep, the requirement against rebroadcasting only applies if the transaction has not moved to a different block.

In order to avoid storing excessively large messages, yet still allow for reasonable future expansion, nodes are permitted to restrict rebroadcasting (perhaps statistically).

New channel features are possible in the future: backwards compatible (or optional) features will have *odd* feature bits, while incompatible features will have *even* feature bits ([“It’s OK to be odd!”](#)).

A delay of 72 blocks is used when forgetting a channel after detecting that it has been spent: this can allow a new `channel_announcement` to propagate to indicate that this channel was spliced and not closed. Thanks to this delay, payments can still be relayed on the channel while the splice transaction is waiting for enough confirmations.

The node_announcement Message

This gossip message allows a node to indicate extra data associated with it, in addition to its public key. To avoid trivial denial of service attacks, nodes not associated with an already known channel are ignored.

1. type: 257 (node_announcement)
2. data:
 - [signature:signature]
 - [u16:flen]
 - [flen*byte:features]
 - [u32:timestamp]
 - [point:node_id]
 - [3*byte:rgb_color]
 - [32*byte:alias]
 - [u16:addrlen]
 - [addrlen*byte:addresses]

`timestamp` allows for the ordering of messages, in the case of multiple announcements. `rgb_color` and `alias` allow intelligence services to assign nodes colors like black and cool monikers like ‘IRATEMONK’ and ‘WISTFULTOLL’.

`addresses` allows a node to announce its willingness to accept incoming network connections: it contains a series of address descriptors for connecting to the node. The first byte describes the address type and is followed by the appropriate number of bytes for that type.

The following address descriptor types are defined:

- 1: ipv4; data = [4:ipv4_addr] [2:port] (length 6)

- 2: ipv6; data = [16:ipv6_addr] [2:port] (length 18)
- 3: Deprecated (length 12). Used to contain Tor v2 onion services.
- 4: Tor v3 onion service; data = [35:onion_addr] [2:port] (length 37)
 - version 3 ([prop224](#)) onion service addresses; Encodes: [32:32_byte_ed25519_pubkey] || [2:checksum] || [1:version], where checksum = sha3(".onion checksum" || pubkey || version)[:2].
- 5: DNS hostname; data = [1:hostname_len] [hostname_len:hostname] [2:port] (length up to 258)
 - hostname bytes MUST be ASCII characters.
 - Non-ASCII characters MUST be encoded using Punycode: <https://en.wikipedia.org/wiki/Punycode>

Requirements

The origin node: - MUST set timestamp to be greater than that of any previous node_announcement it has previously created. - MAY base it on a UNIX timestamp. - MUST set signature to the signature of the double-SHA256 of the entire remaining packet after signature (using the key given by node_id). - MAY set alias AND rgb_color to customize its appearance in maps and graphs. - Note: the first byte of rgb_color is the red value, the second byte is the green value, and the last byte is the blue value. - MUST set alias to a valid UTF-8 string, with any alias trailing-bytes equal to 0. - SHOULD fill addresses with an address descriptor for each public network address that expects incoming connections. - MUST set addr_len to the number of bytes in addresses. - MUST place address descriptors in ascending order. - SHOULD NOT place any zero-typed address descriptors anywhere. - SHOULD use placement only for aligning fields that follow addresses. - MUST NOT create an address descriptor with port equal to 0. - SHOULD ensure ipv4_addr AND ipv6_addr are routable addresses. - MUST set features according to [BOLT #9](#) - SHOULD set f_len to the minimum length required to hold the features bits it sets. - SHOULD not announce a Tor v2 onion service. - MUST NOT announce more than one type 5 DNS hostname.

The receiving node: - if node_id is NOT a valid compressed public key: - SHOULD send a warning. - MAY close the connection. - MUST NOT process the message further. - if signature is NOT a valid signature (using node_id of the double-SHA256 of the entire message following the signature field, including any future fields appended to the end): - SHOULD send a warning. - MAY close the connection. - MUST NOT process the message further. - if features field contains *unknown even bits*: - SHOULD NOT connect to the node. - Unless paying a [BOLT #11](#) invoice which does not have the same bit(s) set, MUST NOT attempt to send payments to the node. - MUST NOT route a payment *through* the node. - SHOULD ignore the first address descriptor that does NOT match the types defined above. - if addr_len is insufficient to hold the address descriptors of the known types: - SHOULD send a warning. - MAY close the connection. - if port is equal to 0: - SHOULD ignore that address descriptor. - if node_id is NOT previously known from a channel_announcement message, OR if timestamp is NOT greater than the last-received node_announcement from this node_id: - SHOULD ignore the message. - otherwise: - if timestamp is greater than the last-received node_announcement from this node_id: - SHOULD queue the message for rebroadcasting. - MAY choose NOT to queue messages longer than the minimum expected length. - MAY use rgb_color AND alias to reference nodes in interfaces. - SHOULD insinuate their self-signed origins. - SHOULD ignore Tor v2 onion services. - if more than one type 5 address is announced: - SHOULD ignore the additional data. - MUST not forward the node_announcement.

Rationale

New node features are possible in the future: backwards compatible (or optional) ones will have *odd* feature *bits*, incompatible ones will have *even* feature *bits*. These will be propagated normally; incompatible feature bits here refer to the nodes, not the node_announcement message itself.

New address types may be added in the future; as address descriptors have to be ordered in ascending order, unknown ones can be safely ignored. Additional fields beyond addresses may also be added in the future—with optional padding within addresses, if they require certain alignment.

Security Considerations for Node Aliases

Node aliases are user-defined and provide a potential avenue for injection attacks, both during the process of rendering and during persistence.

Node aliases should always be sanitized before being displayed in HTML/Javascript contexts or any other dynamically interpreted rendering frameworks. Similarly, consider using prepared statements, input validation, and escaping to protect against injection vulnerabilities and persistence engines that support SQL or other dynamically interpreted querying languages.

- [Stored and Reflected XSS Prevention](#)
- [DOM-based XSS Prevention](#)
- [SQL Injection Prevention](#)

Don't be like the school of [Little Bobby Tables](#).

The channel_update Message

After a channel has been initially announced, each side independently announces the fees and minimum expiry delta it requires to relay HTLCs through this channel. Each uses the 8-byte channel shortid that matches the channel_announcement and the 1-bit channel_flags field to indicate which end of the channel it's on (origin or final). A node can do this multiple times, in order to change fees.

Note that the channel_update gossip message is only useful in the context of *relaying* payments, not *sending* payments. When making a payment A -> B -> C -> D, only the channel_updates related to channels B -> C (announced by B) and C -> D (announced by C) will come into play. When building the route, amounts and expiries for HTLCs need to be calculated backward from the destination to the source. The exact initial value for amount_msat and the minimal value for cltv_expiry, to be used for the last HTLC in the route, are provided in the payment request (see [BOLT #11](#)).

1. type: 258 (channel_update)
2. data:
 - [signature:signature]
 - [chain_hash:chain_hash]
 - [short_channel_id:short_channel_id]
 - [u32:timestamp]
 - [byte:message_flags]
 - [byte:channel_flags]
 - [u16:cltv_expiry_delta]

- [u64:htlc_minimum_msat]
- [u32:fee_base_msat]
- [u32:fee_proportional_millionths]
- [u64:htlc_maximum_msat]

The `channel_flags` bitfield is used to indicate the direction of the channel: it identifies the node that this update originated from and signals various options concerning the channel. The following table specifies the meaning of its individual bits:

Bit Position	Name	Meaning
0	direction	Direction this update refers to.
1	disable	Disable the channel.

The `message_flags` bitfield is used to provide additional details about the message:

Bit Position	Name
0	must_be_one
1	dont_forward

The `node_id` for the signature verification is taken from the corresponding `channel_announcement`: `node_id_1` if the least-significant bit of flags is 0 or `node_id_2` otherwise.

Requirements

The origin node: - MUST NOT send a created `channel_update` before `channel_ready` has been received. - MAY create a `channel_update` to communicate the channel parameters to the channel peer, even though the channel has not yet been announced (i.e. the `announce_channel` bit was not set or the `channel_update` is sent before the peers exchanged [announcement signatures](#)). - MUST set the `short_channel_id` to either an alias it has received from the peer, or the real channel `short_channel_id`. - MUST set `dont_forward` to 1 in `message_flags` - MUST NOT forward such a `channel_update` to other peers, for privacy reasons. - Note: such a `channel_update`, one not preceded by a `channel_announcement`, is invalid to any other peer and would be discarded. - MUST set signature to the signature of the double-SHA256 of the entire remaining packet after signature, using its own `node_id`. - MUST set `chain_hash` AND `short_channel_id` to match the 32-byte hash AND 8-byte channel ID that uniquely identifies the channel specified in the `channel_announcement` message. - if the origin node is `node_id_1` in the message: - MUST set the direction bit of `channel_flags` to 0. - otherwise: - MUST set the direction bit of `channel_flags` to 1. - MUST set `htlc_maximum_msat` to the maximum value it will send through this channel for a single HTLC. - MUST set this to less than or equal to the channel capacity. - MUST set this to less than or equal to `max_htlc_value_in_flight_msat` it received from the peer. - MUST set this to greater than or equal to `htlc_minimum_msat`. - MUST set `must_be_one` in `message_flags` to 1. - MUST set bits in `channel_flags` and `message_flags` that are not assigned a meaning to 0. - MAY create and send a `channel_update` with the `disable` bit set to 1, to signal a channel's temporary unavailability (e.g. due to a loss of connectivity) OR permanent unavailability (e.g. prior to an on-chain settlement). - MAY send a subsequent `channel_update` with the `disable` bit set to 0 to re-enable the channel. - MUST set `timestamp` to greater than 0, AND to greater than any previously-sent `channel_update` for this `short_channel_id`. - SHOULD base `timestamp` on a UNIX timestamp. - MUST set `cltv_expiry_delta` to the number of blocks it will subtract from an incoming HTLC's `cltv_expiry`. - MUST set `htlc_minimum_msat` to the minimum HTLC value (in

millisatoshi) that the channel peer will accept. - MUST set `htlc_minimum_msat` to less than or equal to `htlc_maximum_msat`. - MUST set `fee_base_msat` to the base fee (in millisatoshi) it will charge for any HTLC. - MUST set `fee_proportional_millionths` to the amount (in millionths of a satoshi) it will charge per transferred satoshi. - SHOULD NOT create redundant `channel_updates` - If it creates a new `channel_update` with updated channel parameters: - SHOULD keep accepting the previous channel parameters for 10 minutes

The receiving node: - if the `short_channel_id` does NOT match a previous `channel_announcement`: - MUST ignore `channel_updates` that do NOT correspond to one of its own channels. - if the channel output has been spent: - MUST ignore `channel_updates`, unless they have the `disable` bit set to 1. - SHOULD NOT rebroadcast `channel_updates` to its peers, unless they have the `disable` bit set to 1. - SHOULD accept `channel_updates` for its own channels (even if non-public), in order to learn the associated origin nodes' forwarding parameters. - if signature is not a valid signature, using `node_id` of the double-SHA256 of the entire message following the signature field (including unknown fields following `fee_proportional_millionths`): - SHOULD send a warning and close the connection. - MUST NOT process the message further. - if the specified `chain_hash` value is unknown (meaning it isn't active on the specified chain): - MUST ignore the channel update. - if the timestamp is equal to the last-received `channel_update` for this `short_channel_id` AND `node_id`: - if the fields below timestamp differ: - MAY blacklist this `node_id`. - MAY forget all channels associated with it. - if the fields below timestamp are equal: - SHOULD ignore this message - if timestamp is lower than that of the last-received `channel_update` for this `short_channel_id` AND for `node_id`: - SHOULD ignore the message. - otherwise: - if the timestamp is unreasonably far in the future: - MAY discard the `channel_update`. - otherwise: - SHOULD queue the message for rebroadcasting. - MAY choose NOT to for messages longer than the minimum expected length. - if `htlc_maximum_msat` < `htlc_minimum_msat`: - SHOULD ignore this channel during route considerations. - if `htlc_maximum_msat` is greater than channel capacity: - MAY blacklist this `node_id` - SHOULD ignore this channel during route considerations. - otherwise: - SHOULD consider the `htlc_maximum_msat` when routing.

Rationale

The timestamp field is used by nodes for pruning `channel_updates` that are either too far in the future or have not been updated in two weeks; so it makes sense to have it be a UNIX timestamp (i.e. seconds since UTC 1970-01-01). This cannot be a hard requirement, however, given the possible case of two `channel_updates` within a single second.

It is assumed that more than one `channel_update` message changing the channel parameters in the same second may be a DoS attempt, and therefore, the node responsible for signing such messages may be blacklisted. However, a node may send a same `channel_update` message with a different signature (changing the nonce in signature signing), and hence fields apart from signature are checked to see if the channel parameters have changed for the same timestamp. It is also important to note that ECDSA signatures are malleable. So, an intermediate node who received the `channel_update` message can rebroadcast it just by changing the `s` component of signature with `-s`. This should however not result in the blacklist of the `node_id` from where the message originated.

The recommendation against redundant `channel_updates` minimizes spamming the network, however it is sometimes inevitable. For example, a channel with a peer which is unreachable will eventually cause a `channel_update` to indicate that the channel is disabled, with another update re-enabling the channel when the peer reestablishes contact. Because gossip messages are batched and replace previous ones, the result may be a single seemingly-redundant update.

When a node creates a new `channel_update` to change its channel parameters, it will take some time to propagate through the network and payers may use older parameters. It is recommended to keep accepting older parameters for at least 10 minutes to improve payment latency and reliability.

The `must_be_one` field in `message_flags` was previously used to indicate the presence of the `htlc_maximum_msat` field. This field must now always be present, so `must_be_one` is a constant value, and ignored by receivers.

Query Messages

Understanding of messages used to be indicated with the `gossip_queries` feature bit; now these messages are universally supported, that feature has now been slightly repurposed. Not offering this feature means a node is not worth querying for gossip: either they do not store the entire gossip map, or they are only connected to a single peer (this one).

There are several messages which contain a long array of `short_channel_ids` (called `encoded_short_ids`) so we include an encoding byte which allows for different encoding schemes to be defined in the future, if they provide benefit.

Encoding types: * 0: uncompressed array of `short_channel_id` types, in ascending order. * 1: Previously used for zlib compression, this encoding MUST NOT be used.

This encoding is also used for arrays of other types (timestamps, flags, ...), and specified with an `encoded_` prefix. For example, `encoded_timestamps` is an array of timestamps with a 0 prefix.

Query messages can be extended with optional fields that can help reduce the number of messages needed to synchronize routing tables by enabling:

- timestamp-based filtering of `channel_update` messages: only ask for `channel_update` messages that are newer than the ones you already have.
- checksum-based filtering of `channel_update` messages: only ask for `channel_update` messages that carry different information from the ones you already have.

Nodes can signal that they support extended gossip queries with the `gossip_queries_ex` feature bit.

The `query_short_channel_ids/reply_short_channel_ids_end` Messages

1. type: 261 (`query_short_channel_ids`)
2. data:
 - [chain_hash:chain_hash]
 - [u16:len]
 - [len*byte:encoded_short_ids]
 - [query_short_channel_ids_tlvs:tlvs]
3. tlv_stream: `query_short_channel_ids_tlvs`
4. types:
 1. type: 1 (`query_flags`)
 2. data:
 - [byte:encoding_type]
 - [...*byte:encoded_query_flags]

encoded_query_flags is an array of bitfields, one bigsize per bitfield, one bitfield for each short_channel_id. Bits have the following meaning:

Bit Position	Meaning
0	Sender wants channel_announcement
1	Sender wants channel_update for node 1
2	Sender wants channel_update for node 2
3	Sender wants node_announcement for node 1
4	Sender wants node_announcement for node 2

Query flags must be minimally encoded, which means that one flag will be encoded with a single byte.

1. type: 262 (reply_short_channel_ids_end)
2. data:
 - [chain_hash:chain_hash]
 - [byte:full_information]

This is a general mechanism which lets a node query for the channel_announcement and channel_update messages for specific channels (identified via short_channel_ids). This is usually used either because a node sees a channel_update for which it has no channel_announcement or because it has obtained previously unknown short_channel_ids from reply_channel_range.

Requirements

The sender: - SHOULD NOT send this to a peer which does not offer gossip_queries. - MUST NOT send query_short_channel_ids if it has sent a previous query_short_channel_ids to this peer and not received reply_short_channel_ids_end. - MUST set chain_hash to the 32-byte hash that uniquely identifies the chain that the short_channel_ids refer to. - MUST set the first byte of encoded_short_ids to the encoding type. - MUST encode a whole number of short_channel_ids to encoded_short_ids - MAY send this if it receives a channel_update for a short_channel_id for which it has no channel_announcement. - SHOULD NOT send this if the channel referred to is not an unspent output. - MAY include an optional query_flags. If so: - MUST set encoding_type, as for encoded_short_ids. - Each query flag is a minimally-encoded bigsize. - MUST encode one query flag per short_channel_id.

The receiver: - if the first byte of encoded_short_ids is not a known encoding type: - MAY send a warning. - MAY close the connection. - if encoded_short_ids does not decode into a whole number of short_channel_id: - MAY send a warning. - MAY close the connection. - if it has not sent reply_short_channel_ids_end to a previously received query_short_channel_ids from this sender: - MAY send a warning. - MAY close the connection. - if the incoming message includes query_short_channel_ids_tlvs: - if encoding_type is not a known encoding type: - MAY send a warning. - MAY close the connection. - if encoded_query_flags does not decode to exactly one flag per short_channel_id: - MAY send a warning. - MAY close the connection. - MUST respond to each known short_channel_id: - if the incoming message does not include encoded_query_flags: - with a channel_announcement and the latest channel_update for each end - MUST follow with any node_announcements for each channel_announcement - otherwise: - We define query_flag for the Nth short_channel_id in encoded_short_ids to be the Nth bigsize of the decoded encoded_query_flags. - if bit 0 of query_flag is set: - MUST reply with a channel_announcement - if bit 1 of query_flag is set and it has

received a `channel_update` from `node_id_1`: - MUST reply with the latest `channel_update` for `node_id_1` - if bit 2 of `query_flag` is set and it has received a `channel_update` from `node_id_2`: - MUST reply with the latest `channel_update` for `node_id_2` - if bit 3 of `query_flag` is set and it has received a `node_announcement` from `node_id_1`: - MUST reply with the latest `node_announcement` for `node_id_1` - if bit 4 of `query_flag` is set and it has received a `node_announcement` from `node_id_2`: - MUST reply with the latest `node_announcement` for `node_id_2` - SHOULD NOT wait for the next outgoing gossip flush to send these. - SHOULD avoid sending duplicate `node_announcements` in response to a single `query_short_channel_ids`. - MUST follow these responses with `reply_short_channel_ids_end`. - if does not maintain up-to-date channel information for `chain_hash`: - MUST set `full_information` to 0. - otherwise: - SHOULD set `full_information` to 1.

Rationale

Future nodes may not have complete information; they certainly won't have complete information on unknown `chain_hash` chains. While this `full_information` field (previously and confusingly called `complete`) cannot be trusted, a 0 does indicate that the sender should search elsewhere for additional data.

The explicit `reply_short_channel_ids_end` message means that the receiver can indicate it doesn't know anything, and the sender doesn't need to rely on timeouts. It also causes a natural ratelimiting of queries.

The `query_channel_range` and `reply_channel_range` Messages

1. type: 263 (`query_channel_range`)
2. data:
 - [`chain_hash:chain_hash`]
 - [`u32:first_blocknum`]
 - [`u32:number_of_blocks`]
 - [`query_channel_range_tlvs:tlvs`]
3. `tlv_stream`: `query_channel_range_tlvs`
4. types:
 1. type: 1 (`query_option`)
 2. data:
 - [`bigsize:query_option_flags`]

`query_option_flags` is a bitfield represented as a minimally-encoded `bigsize`. Bits have the following meaning:

Bit Position	Meaning
0	Sender wants timestamps
1	Sender wants checksums

Though it is possible, it would not be very useful to ask for checksums without asking for timestamps too: the receiving node may have an older `channel_update` with a different checksum, asking for it would be useless. And if a `channel_update` checksum is actually 0 (which is quite unlikely) it will not be queried.

1. type: 264 (`reply_channel_range`)
2. data:
 - [`chain_hash:chain_hash`]
 - [`u32:first_blocknum`]
 - [`u32:number_of_blocks`]

- [byte:sync_complete]
 - [u16:len]
 - [len*byte:encoded_short_ids]
 - [reply_channel_range_tlvs:tlvs]
3. tlv_stream: reply_channel_range_tlvs
4. types:
1. type: 1 (timestamps_tlv)
 2. data:
 - [byte:encoding_type]
 - [...*byte:encoded_timestamps]
 3. type: 3 (checksums_tlv)
 4. data:
 - [...*channel_update_checksums:checksums]

For a single channel_update, timestamps are encoded as:

1. subtype: channel_update_timestamps
2. data:
 - [u32:timestamp_node_id_1]
 - [u32:timestamp_node_id_2]

Where: * timestamp_node_id_1 is the timestamp of the channel_update for node_id_1, or 0 if there was no channel_update from that node. * timestamp_node_id_2 is the timestamp of the channel_update for node_id_2, or 0 if there was no channel_update from that node.

For a single channel_update, checksums are encoded as:

1. subtype: channel_update_checksums
2. data:
 - [u32:checksum_node_id_1]
 - [u32:checksum_node_id_2]

Where: * checksum_node_id_1 is the checksum of the channel_update for node_id_1, or 0 if there was no channel_update from that node. * checksum_node_id_2 is the checksum of the channel_update for node_id_2, or 0 if there was no channel_update from that node.

The checksum of a channel_update is the CRC32C checksum as specified in [RFC3720](#) of this channel_update without its signature and timestamp fields.

This allows querying for channels within specific blocks.

Requirements

The sender of query_channel_range: - SHOULD NOT send this to a peer which does not offer gossip_queries. - MUST NOT send this if it has sent a previous query_channel_range to this peer and not received all reply_channel_range replies. - MUST set chain_hash to the 32-byte hash that uniquely identifies the chain that it wants the reply_channel_range to refer to - MUST set first_blocknum to the first block it wants to know channels for - MUST set number_of_blocks to 1 or greater. - MAY append an additional query_channel_range_tlv, which specifies the type of extended information it would like to receive.

The receiver of query_channel_range: - if it has not sent all reply_channel_range to a previously received query_channel_range from this sender: - MAY send a warning. - MAY close the connection. - MUST respond with one or more reply_channel_range: - MUST set chain_hash equal to that of query_channel_range, - MUST limit number_of_blocks to the maximum number of blocks whose results could fit in encoded_short_ids - MAY split block contents across multiple reply_channel_range. - the first reply_channel_range message: - MUST set first_blocknum less than or equal to the first_blocknum in query_channel_range - MUST set first_blocknum plus number_of_blocks greater than first_blocknum in query_channel_range. - successive reply_channel_range message: - MUST have first_blocknum equal or greater than the previous first_blocknum. - MUST set sync_complete to false if this is not the final reply_channel_range. - the final reply_channel_range message: - MUST have first_blocknum plus number_of_blocks equal or greater than the query_channel_range first_blocknum plus number_of_blocks. - MUST set sync_complete to true.

If the incoming message includes query_option, the receiver MAY append additional information to its reply: - if bit 0 in query_option_flags is set, the receiver MAY append a timestamps_tlv that contains channel_update timestamps for all short_channel_ids in encoded_short_ids - if bit 1 in query_option_flags is set, the receiver MAY append a checksums_tlv that contains channel_update checksums for all short_channel_ids in encoded_short_ids

Rationale

A single response might be too large for a single packet, so multiple replies may be required. We want to allow a peer to store canned results for (say) 1000-block ranges, so replies can exceed the requested range. However, we require that each reply be relevant (overlapping the requested range).

By insisting that replies be in increasing order, the receiver can easily determine if replies are done: simply check if first_blocknum plus number_of_blocks equals or exceeds the first_blocknum plus number_of_blocks it asked for.

The addition of timestamp and checksum fields allow a peer to omit querying for redundant updates.

The gossip_timestamp_filter Message

1. type: 265 (gossip_timestamp_filter)
2. data:
 - [chain_hash:chain_hash]
 - [u32:first_timestamp]
 - [u32:timestamp_range]

This message allows a node to constrain future gossip messages to a specific range. A node which wants any gossip messages has to send this, otherwise no gossip messages would be received.

Note that this filter replaces any previous one, so it can be used multiple times to change the gossip from a peer.

Requirements

The sender: - MUST set `chain_hash` to the 32-byte hash that uniquely identifies the chain that it wants the gossip to refer to. - If the receiver does not offer `gossip_queries`: - SHOULD set `first_timestamp` to `0xFFFFFFFF` and `timestamp_range` to 0.

The receiver: - SHOULD send all gossip messages whose timestamp is greater or equal to `first_timestamp`, and less than `first_timestamp` plus `timestamp_range`. - MAY wait for the next outgoing gossip flush to send these. - SHOULD send gossip messages as it generates them regardless of timestamp. - Otherwise (relayed gossip): - SHOULD restrict future gossip messages to those whose timestamp is greater or equal to `first_timestamp`, and less than `first_timestamp` plus `timestamp_range`. - If a `channel_announcement` has no corresponding `channel_updates`: - MUST NOT send the `channel_announcement`. - If the funding output of the `channel_announcement` has been spent: - SHOULD NOT send the `channel_announcement`. - Otherwise: - MUST consider the timestamp of the `channel_announcement` to be the timestamp of a corresponding `channel_update`. - MUST consider whether to send the `channel_announcement` after receiving the first corresponding `channel_update`. - If a `channel_announcement` is sent: - MUST send the `channel_announcement` prior to any corresponding `channel_updates` and `node_announcements`.

Rationale

Since `channel_announcement` doesn't have a timestamp, we generate a likely one. If there's no `channel_update` then it is not sent at all, which is most likely in the case of pruned channels.

Otherwise the `channel_announcement` is usually followed immediately by a `channel_update`. Ideally we would specify that the first (oldest) `channel_update`'s timestamp is to be used as the time of the `channel_announcement`, but new nodes on the network will not have this, and further would require the first `channel_update` timestamp to be stored. Instead, we allow any update to be used, which is simple to implement.

In the case where the `channel_announcement` is nonetheless missed, `query_short_channel_ids` can be used to retrieve it.

Nodes can use `timestamp_filter` to reduce their gossip load when they have many peers (eg. setting `first_timestamp` to `0xFFFFFFFF` after the first few peers, in the assumption that propagation is adequate). This assumption of adequate propagation does not apply for gossip messages generated directly by the node itself, so they should ignore filters.

Requirements

A node: - MUST NOT relay any gossip messages it did not generate itself, unless explicitly requested.

Rebroadcasting

Requirements

A receiving node: - upon receiving a new `channel_announcement` or a `channel_update` or `node_announcement` with an updated timestamp: - SHOULD update its local view of the network's topology accordingly. - after applying the changes from the announcement: - if there are no channels associated with the corresponding origin node: - MAY purge the origin node from its set of known nodes. - otherwise: - SHOULD update the

appropriate metadata AND store the signature associated with the announcement. - Note: this will later allow the node to rebuild the announcement for its peers.

A node: - MUST not send gossip it did not generate itself, until it receives `gossip_timestamp_filter`. - SHOULD flush outgoing gossip messages once every 60 seconds, independently of the arrival times of the messages. - Note: this results in staggered announcements that are unique (not duplicated). - SHOULD NOT forward gossip messages to peers who sent networks in `init` and did not specify the `chain_hash` of this gossip message. - MAY re-announce its channels regularly. - Note: this is discouraged, in order to keep the resource requirements low.

Rationale

Once the gossip message has been processed, it's added to a list of outgoing messages, destined for the processing node's peers, replacing any older updates from the origin node. This list of gossip messages will be flushed at regular intervals; such a store-and-delayed-forward broadcast is called a *staggered broadcast*. Also, such batching forms a natural rate limit with low overhead.

HTLC Fees

Requirements

The origin node: - SHOULD accept HTLCs that pay a fee equal to or greater than: $\text{fee_base_msat} + (\text{amount_to_forward} * \text{fee_proportional_millionths} / 1000000)$ - SHOULD accept HTLCs that pay an older fee, for some reasonable time after sending `channel_update`. - Note: this allows for any propagation delay.

Pruning the Network View

Requirements

A node: - SHOULD monitor the funding transactions in the blockchain, to identify channels that are being closed. - if the funding output of a channel is spent and received 72 block confirmations: - SHOULD be removed from the local network view AND be considered closed. - if the announced node no longer has any associated open channels: - MAY prune nodes added through `node_announcement` messages from their local view. - Note: this is a direct result of the dependency of a `node_announcement` being preceded by a `channel_announcement`.

Recommendation on Pruning Stale Entries

Requirements

A node: - if the timestamp of the latest `channel_update` in either direction is older than two weeks (1209600 seconds): - MAY prune the channel. - MAY ignore the channel. - Note: this is an individual node policy and MUST NOT be enforced by forwarding peers, e.g. by closing channels when receiving outdated gossip messages.

Rationale

Several scenarios may result in channels becoming unusable and its endpoints becoming unable to send updates for these channels. For example, this occurs if both endpoints lose access to their private keys and can neither sign `channel_updates` nor close the channel on-chain. In this case, the channels are unlikely to be part of a computed route, since they would be partitioned off from the rest of the network; however, they would remain in the local network view and would be forwarded to other peers indefinitely.

The oldest `channel_update` is used to prune the channel since both sides need to be active in order for the channel to be usable. Doing so prunes channels even if one side continues to send fresh `channel_updates` but the other node has disappeared.

Recommendations for Routing

When calculating a route for an HTLC, both the `cltv_expiry_delta` and the fee need to be considered: the `cltv_expiry_delta` contributes to the time that funds will be unavailable in the event of a worst-case failure. The relationship between these two attributes is unclear, as it depends on the reliability of the nodes involved.


If a route is computed by simply routing to the intended recipient and summing the `cltv_expiry_deltas`, then it's possible for intermediate nodes to guess their position in the route. Knowing the CLTV of the HTLC, the surrounding network topology, and the `cltv_expiry_deltas` gives an attacker a way to guess the intended recipient. Therefore, it's highly desirable to add a random offset to the CLTV that the intended recipient will receive, which bumps all CLTVs along the route.

In order to create a plausible offset, the origin node MAY start a limited random walk on the graph, starting from the intended recipient and summing the `cltv_expiry_deltas`, and use the resulting sum as the offset. This effectively creates a *shadow route extension* to the actual route and provides better protection against this attack vector than simply picking a random offset would.

Other more advanced considerations involve diversification of route selection, to avoid single points of failure and detection, and balancing of local channels.

Routing Example

Consider four nodes:



Each advertises the following `cltv_expiry_delta` on its end of every channel:

1. A: 10 blocks
2. B: 20 blocks
3. C: 30 blocks
4. D: 40 blocks

C also uses a `min_final_cltv_expiry_delta` of 18 (the default) when requesting payments.

Also, each node has a set fee scheme that it uses for each of its channels:

1. A: 100 base + 1000 millionths
2. B: 200 base + 2000 millionths
3. C: 300 base + 3000 millionths
4. D: 400 base + 4000 millionths

The network will see eight `channel_update` messages:

1. A->B: `cltv_expiry_delta` = 10, `fee_base_msat` = 100, `fee_proportional_millionths` = 1000
2. A->D: `cltv_expiry_delta` = 10, `fee_base_msat` = 100, `fee_proportional_millionths` = 1000
3. B->A: `cltv_expiry_delta` = 20, `fee_base_msat` = 200, `fee_proportional_millionths` = 2000
4. D->A: `cltv_expiry_delta` = 40, `fee_base_msat` = 400, `fee_proportional_millionths` = 4000
5. B->C: `cltv_expiry_delta` = 20, `fee_base_msat` = 200, `fee_proportional_millionths` = 2000
6. D->C: `cltv_expiry_delta` = 40, `fee_base_msat` = 400, `fee_proportional_millionths` = 4000
7. C->B: `cltv_expiry_delta` = 30, `fee_base_msat` = 300, `fee_proportional_millionths` = 3000
8. C->D: `cltv_expiry_delta` = 30, `fee_base_msat` = 300, `fee_proportional_millionths` = 3000

B->C. If B were to send 4,999,999 millisatoshi directly to C, it would neither charge itself a fee nor add its own `cltv_expiry_delta`, so it would use C's requested `min_final_cltv_expiry_delta` of 18. Presumably it would also add a *shadow route* to give an extra CLTV of 42. Additionally, it could add extra CLTV deltas at other hops, as these values represent a minimum, but chooses not to do so here, for the sake of simplicity:

- `amount_msat`: 4999999
- `cltv_expiry`: `current-block-height` + 18 + 42
- `onion_routing_packet`:
 - `amt_to_forward` = 4999999
 - `outgoing_cltv_value` = `current-block-height` + 18 + 42

A->B->C. If A were to send 4,999,999 millisatoshi to C via B, it needs to pay B the fee it specified in the B->C `channel_update`, calculated as per [HTLC Fees](#):

$$\text{fee_base_msat} + (\text{amount_to_forward} * \text{fee_proportional_millionths} / 1000000)$$

$$200 + (4999999 * 2000 / 1000000) = 10199$$

Similarly, it would need to add B->C's `channel_update` `cltv_expiry_delta` (20), C's requested `min_final_cltv_expiry_delta` (18), and the cost for the *shadow route* (42). Thus, A->B's `update_add_htlc` message would be:

- `amount_msat`: 5010198
- `cltv_expiry`: `current-block-height` + 20 + 18 + 42
- `onion_routing_packet`:
 - `amt_to_forward` = 4999999
 - `outgoing_cltv_value` = `current-block-height` + 18 + 42

B->C's `update_add_htlc` would be the same as B->C's direct payment above.

A->D->C. Finally, if for some reason A chose the more expensive route via D, A->D's `update_add_htlc` message would be:

- `amount_msat`: 5020398
- `cltv_expiry`: `current-block-height + 40 + 18 + 42`
- `onion_routing_packet`:
 - `amt_to_forward` = 4999999
 - `outgoing_cltv_value` = `current-block-height + 18 + 42`

And D->C's `update_add_htlc` would again be the same as B->C's direct payment above.



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Paying People

Lightning payments are still HTLCs under the hood, but humans don't paste onion packets into a wallet. They scan invoices, or they fetch offers.

BOLT 11 is the classic invoice. It's a bech32 string (`lnbc...` on mainnet) that encodes amount, payment hash, expiry, a description, a fallback on-chain address, and routing hints for unannounced channels. Almost every Lightning wallet you've used is speaking BOLT 11.

BOLT 12 is offers. Instead of a one-shot invoice, a receiver publishes a reusable offer. The payer requests an invoice over the Lightning network itself (onion messages), gets a fresh invoice back, and pays it. Offers can be static, they can request payer information, and they can be used for refunds. It's the more recent, more flexible cousin of BOLT 11.

If you're building a merchant flow, start with BOLT 11 — it's everywhere — then read BOLT 12 to see where the protocol is going.

In this chapter

- BOLT 11 — Invoice protocol
- BOLT 12 — Offers

BOLT # 11: Invoice Protocol for Lightning Payments

A simple, extendable, QR-code-ready protocol for requesting payments over Lightning.

Table of Contents

- [Encoding Overview](#)
- [Human-Readable Part](#)
- [Data Part](#)
 - [Tagged Fields](#)
 - [Feature Bits](#)

- [Payer / Payee Interactions](#)
 - [Payer / Payee Requirements](#)
- [Implementation](#)
- [Examples](#)
- [Authors](#)

Encoding Overview

The format for a Lightning invoice uses [bech32 encoding](#), which is already used for Bitcoin Segregated Witness. It can be simply reused for Lightning invoices even though its 6-character checksum is optimized for manual entry, which is unlikely to happen often given the length of Lightning invoices.

If a URI scheme is desired, the current recommendation is to either use 'lightning:' as a prefix before the BOLT-11 encoding (note: not 'lightning:/''), or for fallback to Bitcoin payments, to use 'bitcoin:', as per BIP-21, with the key 'lightning' and the value equal to the BOLT-11 encoding.

Requirements

A writer: - MUST encode the payment request in Bech32 (see BIP-0173) - SHOULD use upper case for QR codes (see BIP-0173) - MAY exceed the 90-character limit specified in BIP-0173.

A reader: - MUST parse the address as Bech32, as specified in BIP-0173 (also without the character limit). - Note: this includes handling uppercase as specified by BIP-0173 - if the checksum is incorrect: - MUST fail the payment.

Human-Readable Part

The human-readable part of a Lightning invoice consists of two sections: 1. prefix: ln + BIP-0173 currency prefix (e.g. lnbc for Bitcoin mainnet, lntb for Bitcoin testnet, lntbs for Bitcoin signet, and lnbcrt for Bitcoin regtest) 1. amount: optional number in that currency, followed by an optional multiplier letter. The unit encoded here is the 'social' convention of a payment unit – in the case of Bitcoin the unit is 'bitcoin' NOT satoshis.

The following multiplier letters are defined:

- m (milli): multiply by 0.001
- u (micro): multiply by 0.000001
- n (nano): multiply by 0.000000001
- p (pico): multiply by 0.000000000001

Requirements

A writer: - MUST encode prefix using the currency required for successful payment. - if a specific minimum amount is required for successful payment: - MUST include that amount. - MUST encode amount as a positive decimal integer with no leading 0s. - If the p multiplier is used the last decimal of amount MUST be 0. - SHOULD use the shortest representation possible, by using the largest multiplier or omitting the multiplier.

A reader: - if it does NOT understand the prefix: - MUST fail the payment. - if the amount is empty: - SHOULD indicate to the payer that amount is unspecified. - otherwise: - if amount contains a non-digit OR is followed by

anything except a multiplier (see table above): - MUST fail the payment. - if the multiplier is present: - MUST multiply amount by the multiplier value to derive the amount required for payment. - if multiplier is p and the last decimal of amount is not 0: - MUST fail the payment.

Rationale

The amount is encoded into the human readable part, as it's fairly readable and a useful indicator of how much is being requested.

Donation addresses often don't have an associated amount, so amount is optional in that case. Usually a minimum payment is required for whatever is being offered in return.

The p multiplier would allow to specify sub-millisatoshi amounts, which cannot be transferred on the network, since HTLCs are denominated in millisatoshis. Requiring a trailing 0 decimal ensures that the amount represents an integer number of millisatoshis.

Note that non-largest multipliers have been encountered in the wild, and as such invoice parsers should handle them.

Data Part

The data part of a Lightning invoice consists of multiple sections:

1. timestamp: seconds-since-1970 (35 bits, big-endian)
2. zero or more tagged parts
3. signature: compact ECDSA/secp256k1 signature of the above (520 bits: 64-byte R || S + 1-byte recovery id)

Requirements

A writer: - MUST set timestamp to the number of seconds since Midnight 1 January 1970, UTC in big-endian. - MUST set signature to a valid compact ECDSA signature over secp256k1 of the SHA-256 hash of: the human-readable part (as UTF-8 bytes) concatenated with the data part (excluding the signature), with 0 bits appended to pad to a byte boundary. The signature is encoded as 64 bytes (R || S), followed by a trailing 1-byte recovery id in {0,1,2,3}.

A reader: - MUST check that the signature is valid (see the n tagged field specified below).

Rationale

signature covers an exact number of bytes even though the SHA2 standard actually supports hashing in bit boundaries, because it's not widely implemented. The recovery ID allows public-key recovery, so the identity of the payee node can be implied.

Tagged Fields

Each Tagged Field is of the form:

1. type (5 bits)
2. data_length (10 bits, big-endian)

3. data (data_length x 5 bits)

Note that the maximum length of a Tagged Field's data is constricted by the maximum value of data_length. This is 1023 x 5 bits, or 639 bytes.

Currently defined tagged fields are:

- p (1): data_length 52. 256-bit SHA256 payment_hash. Preimage of this provides proof of payment.
- s (16): data_length 52. This 256-bit secret prevents forwarding nodes from probing the payment recipient.
- d (13): data_length variable. Short description of purpose of payment (UTF-8), e.g. '1 cup of coffee' or 'ナンセンス 1杯'
- m (27): data_length variable. Additional metadata to attach to the payment. Note that the size of this field is limited by the maximum hop payload size. Long metadata fields reduce the maximum route length.
- n (19): data_length 53. 33-byte public key of the payee node
- h (23): data_length 52. 256-bit description of purpose of payment (SHA256). This is used to commit to an associated description that is over 639 bytes, but the transport mechanism for the description in that case is transport specific and not defined here.
- x (6): data_length variable. expiry time in seconds (big-endian). Default is 3600 (1 hour) if not specified.
- c (24): data_length variable. min_final_cltv_expiry_delta to use for the last HTLC in the route. Default is 18 if not specified.
- f (9): data_length variable, depending on version. Fallback on-chain address: for Bitcoin, this starts with a 5-bit version and contains a witness program or P2PKH or P2SH address.
- r (3): data_length variable. One or more entries containing extra routing information for a private route; there may be more than one r field
 - pubkey (264 bits)
 - short_channel_id (64 bits)
 - fee_base_msat (32 bits, big-endian)
 - fee_proportional_millionths (32 bits, big-endian)
 - cltv_expiry_delta (16 bits, big-endian)
- 9 (5): data_length variable. One or more 5-bit values containing features supported or required for receiving this payment. See [Feature Bits](#).

Requirements

A writer: - MUST include exactly one p field. - MUST include exactly one s field. - MUST set payment_hash to the SHA2 256-bit hash of the payment_preimage that will be given in return for payment. - MUST include either exactly one d or exactly one h field. - if d is included: - MUST set d to a valid UTF-8 string. - SHOULD use a complete description of the purpose of the payment. - if h is included: - MUST make the preimage of the hashed description in h available through some unspecified means. - SHOULD use a complete description of the purpose of the payment. - MAY include one x field. - if x is included: - MUST use the minimum data_length possible, i.e. no leading 0 field-elements. - SHOULD include one c field (min_final_cltv_expiry_delta). - MUST set c to the minimum cltv_expiry it will accept for the last HTLC in the route. - MUST use the minimum data_length possible, i.e. no leading 0 field-elements. - MAY include one n field. (Otherwise performing public-key recovery is required) - MUST set n to the public key used to create the signature. - MAY include one or more f fields. - for Bitcoin payments: - MUST set an f field to a valid witness version and program, OR to 17 followed by a public key hash, OR to 18 followed by a script hash. - if there is

NOT a public channel associated with its public key: - MUST include at least one `r` field. - `r` field MUST contain one or more ordered entries, indicating the forward route from a public node to the final destination. - Note: for each entry, the `pubkey` is the node ID of the start of the channel; `short_channel_id` is the short channel ID field to identify the channel; and `fee_base_msat`, `fee_proportional_millionths`, and `cltv_expiry_delta` are as specified in [BOLT #7](#). - MAY include more than one `r` field to provide multiple routing options. - if `9` contains non-zero bits: - MUST use the minimum `data_length` possible to encode the non-zero bits with no 0 field-elements at the start. - otherwise: - MUST omit the `9` field altogether. - MUST pad field data to a multiple of 5 bits, using 0s. - if a writer offers more than one of any field type, it: - MUST specify the most-preferred field first, followed by less-preferred fields, in order.

A reader: - MUST skip over `f` fields that use an unknown version. - MUST fail the payment if any field with fixed `data_length` (`p`, `h`, `s`, `n`) does not have the correct length (52, 52, 52, 53). - MUST fail the payment if neither a `d` field nor a `h` field is present, or if both are present. - if the `9` field contains unknown *odd* bits that are non-zero: - MUST ignore the bit. - if the `9` field contains unknown *even* bits that are non-zero: - MUST fail the payment. - SHOULD indicate the unknown bit to the user. - MUST check that the SHA2 256-bit hash in the `h` field exactly matches the hashed description. - if a valid `n` field is provided: - MUST use the `n` field to validate the signature instead of performing public-key recovery. - If the signature is not compliant with the low-S standard rule [low-S](#): - MUST fail the payment - otherwise: - MUST perform ECDSA public-key recovery and accept both high-S and low-S signatures. - if a valid `s` field is not provided: - MUST fail the payment. - otherwise: - MUST use the `s` field as [payment_secret](#) - if the `c` field (`min_final_cltv_expiry_delta`) is not provided: - MUST use an expiry delta of at least 18 when making the payment - if an `m` field is provided: - MUST use that as [payment_metadata](#) - if a `c`, `x`, or `9` field is provided which has a non-minimal `data_length` (i.e. begins with 0 field elements): - SHOULD treat the invoice as invalid.

Rationale

The type-and-length format allows future extensions to be backward compatible. `data_length` is always a multiple of 5 bits, for easy encoding and decoding.

The `d` field allows inline descriptions, but may be insufficient for complex orders. Thus, the `h` field allows a summary: though the method by which the description is served is as-yet unspecified and will probably be transport dependent.

The `m` field allows metadata to be attached to the payment. This supports applications where the recipient doesn't keep any context for the payment.

The `n` field can be used to explicitly specify the destination node ID, instead of requiring public-key recovery.

The `x` field gives warning as to when a payment will be refused: mainly to avoid confusion. The default was chosen to be reasonable for most payments and to allow sufficient time for on-chain payment, if necessary.

The `c` field allows a way for the destination node to require a specific minimum CLTV expiry for its incoming HTLC. Destination nodes may use this to require a higher, more conservative value than the default one (depending on their fee estimation policy and their sensitivity to time locks). Note that remote nodes in the route specify their required `cltv_expiry_delta` in the `channel_update` message, which they can update at all times.

The `f` field allows on-chain fallback; however, this may not make sense for tiny or time-sensitive payments. It's possible that new address forms will appear; thus, multiple `f` fields (in an implied preferred order) help with transition, and `f` fields with versions 19-31 will be ignored by readers.

The `r` field allows limited routing assistance: as specified, it only allows minimum information to use private channels, however, it could also assist in future partial-knowledge routing.

Security Considerations for Payment Descriptions

Payment descriptions are user-defined and provide a potential avenue for injection attacks: during the processes of both rendering and persistence.

Payment descriptions should always be sanitized before being displayed in HTML/JavaScript contexts (or any other dynamically interpreted rendering frameworks). Implementers should be extra perceptive to the possibility of reflected XSS attacks, when decoding and displaying payment descriptions. Avoid optimistically rendering the contents of the payment request until all validation, verification, and sanitization processes have been successfully completed.

Furthermore, consider using prepared statements, input validation, and/or escaping, to protect against injection vulnerabilities in persistence engines that support SQL or other dynamically interpreted querying languages.

- [Stored and Reflected XSS Prevention](#)
- [DOM-based XSS Prevention](#)
- [SQL Injection Prevention](#)

Don't be like the school of [Little Bobby Tables](#).

Feature Bits

Feature bits allow forward and backward compatibility, and follow the *it's ok to be odd* rule. Features appropriate for use in the `9` field are marked in [BOLT 9](#).

The field is big-endian. The least-significant bit is numbered 0, which is *even*, and the next most significant bit is numbered 1, which is *odd*.

Requirements

A writer: - MUST set the `9` field to a feature vector compliant with the [BOLT 9 origin node requirements](#).

A reader: - if the feature vector does not set all known, transitive feature dependencies: - MUST NOT attempt the payment. - if the `basic_mpp` feature is offered in the invoice: - MAY pay using [Basic multi-part payments](#). - otherwise: - MUST NOT use [Basic multi-part payments](#).

Payer / Payee Interactions

These are generally defined by the rest of the Lightning BOLT series, but it's worth noting that [BOLT #4](#) specifies that the payee SHOULD accept up to twice the expected amount, so the payer can make payments harder to track by adding small variations.

- [illegible]

- [illegible]

Please send \$3 for a cup of coffee to the same peer, within one minute

Inbc2500u1pvjluezsp5zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zygspp5qqqsyqcyq5rqwzqfqqqsyqcyq

Breakdown:

- lnbtc: prefix, Lightning on Bitcoin mainnet
- 2500u: amount (2500 micro-bitcoin)
- 1: Bech32 separator
- pvj1uez: timestamp (1496314658)
- s: payment secret...
- p: payment hash...
- d: short description
 - q5: data_length ($q = 0, 5 = 20; 0 * 32 + 20 == 20$)
 - xysxxatsyp3k7enxv4js: '1 cup coffee'
- x: expiry time
 - qz: data_length ($q = 0, z = 2; 0 * 32 + 2 == 2$)
 - pu: 60 seconds ($p = 1, u = 28; 1 * 32 + 28 == 60$)
- 9: features
 - qr: data_length ($q = 0, r = 3; 0 * 32 + 3 == 3$)
 - sgq: b100000100000000
- uk0rl77nj30yxdy8j9vdx85fkpmdla2087ne0xh8nhedh8w27kyke0lp53ut353s06fv3qfegext0eh0ymjpf39tuven09sam30g4v
signature
- fna3rh: Bech32 checksum
- Signature breakdown:
 - e59e3ffbd3945e4334879158d31e89b076dff54f3fa7979ae79df2db9dcaf5896cbfe1a478b8d2307e92c88139464cb7
hex of signature data (32-byte r, 32-byte s)
 - 1 (int) recovery flag contained in signature

- [illegible]

Please send 0.0025 BTC for a cup of nonsense (ナンセンス 1杯) to the same peer, within one minute

Inbc2500u1pvjluezsp5zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zygspp5qqqsyqcyq5rqwzqfqqqsyqcyq

Breakdown:

- [illegible]

Now send \$24 for an entire list of things (hashed)

[illegible]

Breakdown:

- lnbc: prefix, Lightning on Bitcoin mainnet
- 20m: amount (20 milli-bitcoin)
- 1: Bech32 separator

- [illegible]

The same, on testnet, with a fallback address `mk2QpYatsKicvFVuTAQLBryyccRXMUaGHP`

[illegible]

Breakdown:

- lntb: prefix, Lightning on Bitcoin testnet
- 20m: amount (20 milli-bitcoin)
- 1: Bech32 separator
- pvj luez: timestamp (1496314658)
- h: tagged field: hash of description...
- s: payment secret...
- p: payment hash...
- f: tagged field: fallback address
 - pp: data_length ($p = 1; 1 * 32 + 1 == 33$)
 - 3 = 17, so P2PKH address
 - x9et2e20v6pu37c5d9vax37wxq72un98: 160-bit P2PKH address
- 9: features...
- dj545axuxtnfemtpwkc45hx9d2ft7x04mt8q7y6t0k2dge9e7h8kpy9p34ytysl j3yu569aalz2xdk8xkd7ltxql d94u8h2esmsmac
signature
- ghe9k8: Bech32 checksum

- lnb: prefix, Lightning on Bitcoin mainnet
- 20m: amount (20 milli-bitcoin)
- 1: Bech32 separator
- pvjluez: timestamp (1496314658)
- s: payment secret...
- p: payment hash...
- h: tagged field: hash of description...
- f: tagged field: fallback address
 - pp: data_length ($p = 1; 1 * 32 + 1 == 33$)
 - 3 = 17, so P2PKH address
 - qjmp7lwpagxun9pygexvgpjdc4j85f: 160-bit P2PKH address
- r: tagged field: route information
 - 9y: data_length ($9 = 5, y = 4; 5 * 32 + 4 == 164$)
 - q20q82gphp2nflc7jtzrcazrra7wwgzxqc8u7754cdlpfrmcgae92qgzqvzq2ps8pqqqqqqppqqqqq9qqqvpeuqafqx
 - pubkey:
 - 029e03a901b85534ff1e92c43c74431f7ce72046060fcf7a95c37e148f78c77255
 - short_channel_id: 66051x263430x1800
 - fee_base_msat: 1 millisatoshi
 - fee_proportional_millionths: 20
 - cltv_expiry_delta: 3
 - pubkey:
 - 039e03a901b85534ff1e92c43c74431f7ce72046060fcf7a95c37e148f78c77255
 - short_channel_id: 197637x395016x2314
 - fee_base_msat: 2 millisatoshi
 - fee_proportional_millionths: 30
 - cltv_expiry_delta: 4
- 9: features...
- dfjcdk6w3ak5pca9hwfwh63zrrz06wwfya0ydlzpgzxkn5xagsqz7x9j4jwe7yj7vaf2k9lqsd45kts2fd0fkr28am0u4w95tt2n

signature

- `lnbc`: prefix, Lightning on Bitcoin mainnet
- `20m`: amount (20 milli-bitcoin)
- `1`: Bech32 separator
- `pvjluez`: timestamp (1496314658)

[illegible]

- lnb: prefix, Lightning on bitcoin mainnet
- 9678785340p: amount (9678785340 pico-bitcoin = 967878534 milli satoshi)
- 1: Bech32 separator
- pwmna7l: timestamp (1572468703)
- s: payment secret...
- p: payment hash
 - p5: data_length ($p = 1, 5 = 20; 1 * 32 + 20 == 52$)
 - gc3xfm08u9qy06djf8dfflhugl6p7lgza6dsjxq454gxhj9t7a0s: payment hash
462264ede7e14047e9b249da94fefc47f41f7d02ee9b091815a5506bc8abf75f

- d: short description
 - 8d: data_length ($8 = 7, d = 13; 7 * 32 + 13 == 237$)
 - gfkx7cmtwd68yetpd5s9xar0wfjn5gpc8qhrsdfq24f5ggrxdaezqsnvda3kkum5wfjkzmfqf3jkgem9wgsyuctwdus9xgrc
'Blockstream Store: 88.85 USD for Blockstream Ledger Nano S x 1, "Back In My Day" Sticker x 2, "I Got Lightning Working" Sticker x 2 and 1 more items'
- x: expiry time
 - qy: data_length ($q = 0, y = 2; 0 * 32 + 4 == 4$)
 - jw5q: 604800 seconds ($j = 18, w = 14, 5 = 20, q = 0; 18 * 32^3 + 14 * 32^2 + 20 * 32 + 0 == 604800$)
- c: min_final_cltv_expiry_delta
 - qp: data_length ($q = 0, p = 1; 0 * 32 + 1 == 1$)
 - 2: min_final_cltv_expiry_delta = 10
- r: tagged field: route information
 - zj: data_length ($z = 2, j = 18; 2 * 32 + 18 == 82$)
 - q0gxwkzc8w6323m55m4jyxcjwmy7stt9hwkwe2qxm8zpsgg7jcuwz87fcqqeuqqyqqqqqlgqqqqn3qq9q:
 - pubkey: 03d06758583bb5154774a6eb221b1276c9e82d65bbaceca806d90e20c108f4b1c7
 - short_channel_id: 589390x3312x1
 - fee_base_msat = 1000
 - fee_proportional_millionths = 2500
 - cltv_expiry_delta = 40
- 9: features...
- rvgkpnmps664wgkp43l22qsgdw4ve24aca4nymnxdldnp8vh9v2sdxlu5ywdxefs fvm0fq3sesf08uf6q9a2ke0hc9j6z6wlgx5z5k
signature
- u2v9wz: Bech32 checksum

- 2a25dxl5hrntdtn6zvydt7d66hyzsyhqs4wdynavys42xgl6sgx9c4g7me86a27t07mdtfry458rtjr0v92cnmswpsjscgt2vcse3s
signature
- z3uapa: Bech32 checksum

- [illegible]

Please send 0.01 BTC with payment metadata 0x01fafaf0

lnbc10m1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfgyppqdp9wpslujmt9de6zqmt9w3skgct5v

Breakdown:

- lnb: prefix, Lightning on Bitcoin mainnet
- 10m: amount (10 milli-bitcoin)
- 1: Bech32 separator
- pvj luez: timestamp (1496314658)
- p: payment hash
 - p5: data_length ($p = 1, 5 = 20; 1 * 32 + 20 == 52$)
 - qqqsycyq5rqwzfqqqsyqcyq5rqwzfqqqsyqcyq5rqwzfypq: payment hash
0001020304050607080900010203040506070809000102030405060708090102
- d: short description
 - p9: data_length ($p = 1, 9 = 5; 1 * 32 + 5 == 37$)
 - wpshjmt9de6zqmt9w3skgct5vysxjmnnd9jx2: ‘payment metadata inside’
- m: metadata
 - q8: data_length ($q = 0, 8 = 7; 0 * 32 + 7 == 7$)
 - q8a04uq: 0x01fafaf0
- s: payment secret
 - p5: data_length ($p = 1, 5 = 20; 1 * 32 + 20 == 52$)
 - zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zygs:
0x11
- 9: features
 - q2: data_length ($q = 0, 2 = 10; 0 * 32 + 10 == 10$)
 - gqqqqqqsgq: [b01000] = 8 + 14 + 48
- 7hf8he7ecf7n4ffphs6awl9t6676rrclv9ckg3d3ncn7fct63p6s365duk5wrk202cfy3aj5xnnp5gs3vrdrvruverwwq7yzhkf5a3x signature
- d05wj c: Bech32 checksum

Public-key recovery with high-S signature

lnbc1pvjluezsp5zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zyg3zygspp5qqqsyqcyq5rqlwzqfqqqsyqcyq5rqlwz

Breakdown:

- [illegible]

Examples of Invalid Invoices

Same, but adding invalid unknown feature 100

lnbc25m1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfppdq5vdlhkven9v5sxyetpdeessp5zyg?

Breakdown:

- `1nbc`: prefix, Lightning on Bitcoin mainnet
- `25m`: amount (25 milli-bitcoin)
- `1`: Bech32 separator
- `pvj1uezt`: timestamp (1496314658)

- [illegible]

Bech32 checksum is invalid.

Inbc2500u1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqypqdpquwpc4curk03c9wlrswe78q4e

Malformed bech32 string (no 1)

pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqpqdpquwpc4curk03c9wlrswe78q4eyqc7d8d0xc

Malformed bech32 string (mixed case)

LNBC2500u1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfgyppgdpquwpc4curk03c9wlrswe78q4

Signature is not recoverable.

lnbc2500u1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqypqdq5xysxxatsyp3k7enxv4jsqxzpus

String is too short.

lnbc1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfypqdpl2pkx2ctnv5sxxmmwwd5kgetjypeh2

Invalid multiplier

lnbc2500x1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfypqdq5xysxxatsyp3k7enxv4jsxqzpus

Invalid sub-millisatoshi precision.

lnbc25000000001p1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfypqdq5xysxxatsyp3k7enxv4jsx

Missing required s field.

lnbc20m1pvjluezpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqypqhpf58yjmdan79s6qqdhdzgynm4zw

Non canonical signature (high-S) with 'n' field defined

lnbc25m1p70xwfzpp5qqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfqqqsyqcyq5rqwzqfypqdpl2pkx2ctnv5sxxmmwwd5kgetj.

Authors

[FIXME:]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

BOLT #12: Negotiation Protocol for Lightning Payments

Table of Contents

- [Limitations of BOLT 11](#)
- [Payment Flow Scenarios](#)
- [Encoding](#)
- [Signature calculation](#)
- [Offers](#)
- [Invoice Requests](#)
- [Invoices](#)
- [Invoice Errors](#)
- [Payer Proofs](#)

Limitations of BOLT 11

The BOLT 11 invoice format has proven popular but has several limitations:

1. The entangling of bech32 encoding makes it awkward to send in other forms (e.g. inside the lightning network itself).
2. The signature applying to the entire invoice makes it impossible to prove an invoice without revealing its entirety.
3. Fields cannot generally be extracted for external use: the h field was a boutique extraction of the d field only.
4. The lack of the 'it's OK to be odd' rule makes backward compatibility harder.
5. The 'human-readable' idea of separating amounts proved fraught: p was often mishandled, and amounts in pico-bitcoin are harder than the modern satoshi-based counting.
6. Developers found the bech32 encoding to have an issue with extensions, which means we want to replace or discard it anyway.
7. The payment_secret designed to prevent probing by other nodes in the path was only useful if the invoice remained private between the payer and payee.
8. Invoices must be given per user and are actively dangerous if two payment attempts are made for the same user.

Payment Flow Scenarios

Here we use “user” as shorthand for the individual user’s lightning node and “merchant” as the shorthand for the node of someone who is selling or has sold something.

There are two basic payment flows supported by BOLT 12:

The general user-pays-merchant flow is: 1. A merchant publishes an *offer*, such as on a web page or a QR code. 2. Every user requests a unique *invoice* over the lightning network using an *invoice_request* message, which contains the offer fields. 3. The merchant replies with the *invoice*. 4. The user makes a payment to the merchant as indicated by the invoice.

The merchant-pays-user flow (e.g. ATM or refund): 1. The merchant publishes an *invoice_request* which includes an amount it wishes to send to the user. 2. The user sends an *invoice* over the lightning network for the amount in the *invoice_request*, using a (possibly temporary) *invoice_node_id*. 3. The merchant confirms the *invoice_node_id* to ensure it’s about to pay the correct person, and makes a payment to the invoice.

Payment Proofs and Payer Proofs

Note that the normal lightning “proof of payment” can only demonstrate that an invoice was paid (by showing the preimage of the *payment_hash*), not who paid it. The merchant can claim an invoice was paid, and once revealed, anyone can claim they paid the invoice, too.[1]

Providing a key in *invoice_request* allows the payer to prove that they were the one to request the invoice. In addition, the Merkle construction of the BOLT 12 invoice signature allows the user to reveal invoice fields in case of a dispute selectively.

Encoding

Each of the forms documented here are in [TLV](#) format.

The supported ASCII encoding is the human-readable prefix, followed by a 1, followed by a bech32-style data string of the TLVs in order, optionally interspersed with + (for indicating additional data is to come). There is no checksum, unlike bech32m.

Requirements

Writers of a bolt12 string: - MUST either use all lowercase or all UPPERCASE. - SHOULD use uppercase for QR codes. - SHOULD use lower case otherwise. - MAY use +, optionally followed by whitespace, to separate large bolt12 strings.

Readers of a bolt12 string: - MUST handle strings which are all lowercase, or all uppercase. - if it encounters a + followed by zero or more whitespace characters between two bech32 characters: - MUST remove the + and whitespace.

Rationale

The use of bech32 is arbitrary but already exists in the bitcoin world. We currently omit the six-character trailing checksum: QR codes have their own checksums anyway, and errors don't result in loss of funds, simply an invalid offer (or inability to parse).

The use of + (which is ignored) allows use over limited text fields like Twitter:

ln01xxxxxxxxx+

yyyyyyyyyyyyy+

zzzzz

See [format-string-test.json](#).

Signature Calculation

All signatures are created as per [BIP-340](#) and tagged as recommended there. Thus we define $H(\text{tag}, \text{msg})$ as $\text{SHA256}(\text{SHA256}(\text{tag}) \parallel \text{SHA256}(\text{tag}) \parallel \text{msg})$, and $\text{SIG}(\text{tag}, \text{msg}, \text{key})$ as the signature of $H(\text{tag}, \text{msg})$ using key.

Each form is signed using one or more *signature TLV elements*: TLV types 240 through 1000 (inclusive). For these, the tag is "lightning" || messagename || fieldname, and msg is the Merkle-root; "lightning" is the literal 9-byte ASCII string, messagename is the name of the TLV stream being signed (i.e. "invoice_request", "invoice" or "payer_proof") and the fieldname is the TLV field containing the signature (e.g. "signature").

The formulation of the Merkle tree is similar to that proposed in [BIP-341](#), with each TLV leaf paired with a nonce leaf to avoid revealing adjacent nodes in proofs.

The Merkle tree's leaves are, in TLV-ascending order for each tlv: 1. The $H(\text{"LnLeaf"}, \text{tlv})$. 2. The $H(\text{"LnNonce"} \parallel \text{first-tlv}, \text{tlv-type})$ where first-tlv is the numerically-first TLV entry in the stream, and tlv-type is the "type" field (1-9 bytes) of the current tlv.

The Merkle tree inner nodes are $H(\text{"LnBranch"}, \text{lesser-SHA256} \parallel \text{greater-SHA256})$; this ordering means proofs are more compact since left/right is inherently determined.

If there is not exactly a power of 2 leaves, then the tree depth will be uneven, with the deepest tree on the lowest-order leaves.

e.g. consider the encoding of an invoice signature with TLVs TLV0, TLV1, and TLV2 (of types 0, 1, and 2 respectively):

```
L1=H("LnLeaf", TLV0)
L1nonce=H("LnNonce" || TLV0, 0)
L2=H("LnLeaf", TLV1)
L2nonce=H("LnNonce" || TLV0, 1)
L3=H("LnLeaf", TLV2)
L3nonce=H("LnNonce" || TLV0, 2)
```

Assume $L1 < L1nonce$, $L2 > L2nonce$ and $L3 > L3nonce$.

L1 L1nonce

L2 L2nonce

L3 L3nonce

$$\begin{array}{ccc}
 \begin{array}{c} \backslash \quad / \\ v \quad v \end{array} & \begin{array}{c} \backslash \quad / \\ v \quad v \end{array} & \begin{array}{c} \backslash \quad / \\ v \quad v \end{array} \\
 L1A=H("LnBranch", L1 || L1nonce) & L2A=H("LnBranch", L2nonce || L2) & L3A=H("LnBranch", L3nonce || L3)
 \end{array}$$

Assume $L1A < L2A$:

$$\begin{array}{ccc}
 \begin{array}{c} L1A \quad L2A \\ \backslash \quad / \\ v \quad v \end{array} & & \begin{array}{c} L3A=H("LnBranch", L3nonce || L3) \\ | \\ v \\ L3A=H("LnBranch", L3nonce || L3) \end{array} \\
 L1A2A=H("LnBranch", L1A || L2A) & &
 \end{array}$$

Assume $L1A2A > L3A$:

$$\begin{array}{ccc}
 L1A2A=H("LnBranch", L1A || L2A) & & L3A \\
 \backslash & & / \\
 v & & v \\
 \text{Root}=H("LnBranch", L3A || L1A2A)
 \end{array}$$

Signature = SIG("lightninginvoicesignature", Root, nodekey)

Offers

Offers are a precursor to an invoice_request: readers will request an invoice (or multiple) based on the offer. An offer can be much longer-lived than a particular invoice, so it has some different characteristics; in particular the amount can be in a non-lightning currency. It's also designed for compactness to fit inside a QR code easily.

Note that the non-signature TLV elements get mirrored into invoice_request and invoice messages, so they each have specific and distinct TLV ranges.

The human-readable prefix for offers is `lno`.

TLV Fields for Offers

1. tlv_stream: offer
2. types:
 1. type: 2 (offer_chains)
 2. data:
 - [...*chain_hash:chains]
 3. type: 4 (offer_metadata)
 4. data:
 - [...*byte:data]
 5. type: 6 (offer_currency)
 6. data:
 - [...*utf8:iso4217]
 7. type: 8 (offer_amount)
 8. data:
 - [tu64:amount]
 9. type: 10 (offer_description)
 10. data:
 - [...*utf8:description]
 11. type: 12 (offer_features)

- 12. data:
 - [...*byte:features]
- 13. type: 14 (offer_absolute_expiry)
- 14. data:
 - [tu64:seconds_from_epoch]
- 15. type: 16 (offer_paths)
- 16. data:
 - [...*blinded_path:paths]
- 17. type: 18 (offer_issuer)
- 18. data:
 - [...*utf8:issuer]
- 19. type: 20 (offer_quantity_max)
- 20. data:
 - [tu64:max]
- 21. type: 22 (offer_issuer_id)
- 22. data:
 - [point:id]

Requirements For Offers

A writer of an offer: - MUST NOT set any TLV fields outside the inclusive ranges: 1 to 79 and 1000000000 to 1999999999. - if the chain for the invoice is not solely bitcoin: - MUST specify offer_chains the offer is valid for. - otherwise: - SHOULD omit offer_chains, implying that bitcoin is only chain. - if a specific minimum offer_amount is required for successful payment: - MUST set offer_amount to the amount expected (per item). - MUST set offer_amount greater than zero. - if the currency for offer_amount is that of all entries in chains: - MUST specify offer_amount in multiples of the minimum lightning-payable unit (e.g. milli-satoshis for bitcoin). - otherwise: - MUST specify offer_currency iso4217 as an ISO 4217 three-letter code. - MUST specify offer_amount in the currency unit adjusted by the ISO 4217 exponent (e.g. USD cents). - MUST set offer_description to a complete description of the purpose of the payment. - otherwise: - MUST NOT set offer_amount - MUST NOT set offer_currency - MAY set offer_description - MAY set offer_metadata for its own use. - if it supports bolt12 offer features: - MUST set offer_features.features to the bitmap of bolt12 features. - if the offer expires: - MUST set offer_absolute_expiry seconds_from_epoch to the number of seconds after midnight 1 January 1970, UTC that invoice_request should not be attempted. - if it is connected only by private channels: - MUST include offer_paths containing one or more paths to the node from publicly reachable nodes. - otherwise: - MAY include offer_paths. - if it includes offer_paths: - MAY set offer_issuer_id. - otherwise: - MUST set offer_issuer_id to the node's public key to request the invoice from. - if it sets offer_issuer: - SHOULD set it to identify the issuer of the invoice clearly. - if it includes a domain name: - SHOULD begin it with either user@domain or domain - MAY follow with a space and more text - if it can supply more than one item for a single invoice: - if the maximum quantity is known: - MUST set that maximum in offer_quantity_max. - MUST NOT set offer_quantity_max to 0. - otherwise: - MUST set offer_quantity_max to 0. - otherwise: - MUST NOT set offer_quantity_max.

A reader of an offer: - if the offer contains any TLV fields outside the inclusive ranges: 1 to 79 and 1000000000 to 1999999999: - MUST NOT respond to the offer. - if offer_features contains unknown *odd* bits that are non-zero: - MUST ignore the bit. - if offer_features contains unknown *even* bits that are non-zero: - MUST NOT respond to the offer. - SHOULD indicate the unknown bit to the user. - if offer_chains is not set: - if the node does not accept bitcoin invoices: - MUST NOT respond to the offer - otherwise: (offer_chains is set): - if the

node does not accept invoices for at least one of the chains: - MUST NOT respond to the offer - if offer_amount is set and offer_description is not set: - MUST NOT respond to the offer. - if offer_amount is set and is not greater than zero: - MUST NOT respond to the offer. - if offer_currency is set and offer_amount is not set: - MUST NOT respond to the offer. - if neither offer_issuer_id nor offer_paths are set: - MUST NOT respond to the offer. - if num_hops is 0 in any blinded_path in offer_paths: - MUST NOT respond to the offer. - if it uses offer_amount to provide the user with a cost estimate: - MUST take into account the currency units for offer_amount: - offer_currency field if set - otherwise, the minimum lightning-payable unit (e.g. milli-satoshis for bitcoin). - MUST warn the user if the received invoice_amount differs significantly from that estimate. - if the current time is after offer_absolute_expiry: - MUST NOT respond to the offer. - if it chooses to send an invoice request, it sends an onion message: - if offer_paths is set: - MUST send the onion message via any path in offer_paths to the final onion_msg_hop.blinded_node_id in that path - otherwise: - MUST send the onion message to offer_issuer_id - MAY send more than one invoice request onion message at once.

Rationale

The entire offer is reflected in the invoice_request, both for completeness (so all information will be returned in the invoice), and so that the offer node can be stateless. This makes offer_metadata particularly useful, since it can contain an authentication cookie to validate the other fields.

Because offer fields are copied into the invoice request (and then the invoice), they require distinct ranges. The range 1-79 is the normal range, with another billion for self-assigning experimental ranges.

A signature is unnecessary, and makes for a longer string (potentially limiting QR code use on low-end cameras); if the offer has an error, no invoice will be given since the request includes all the non-signature fields.

The offer_issuer_id can be omitted for brevity, if offer_paths is set, as each of the final blinded_node_id in the paths can serve as a valid public key for the destination.

Because offer_amount can be in a different currency (using the offer_currency field) it is merely a guide: the issuer will convert it into a number of millisatoshis for invoice_amount at the time they generate an invoice, or the invoice request can specify the exact amount in invreq_amount, but the issuer may then reject it if it disagrees.

offer_quantity_max is allowed to be 1, which seems useless, but useful in a system which bases it on available stock. It would be painful to have to special-case the “only one left” offer generation.

Offers can be used to simply send money without expecting anything in return (tips, kudos, donations, etc), which means the description field is optional (the offer_issuer field is very useful for this case!); if you are charging for something specific, the description is vital for the user to know what it was they paid for.

An empty offer_chains (present but with zero entries) is explicitly invalid because it would make invoice requests impossible. The payer cannot set invreq_chain to “one of offer_chains” when there are no chains listed. Rejecting such offers early provides clear feedback rather than leaving implementations to fail at the invoice request stage.

Invoice Requests

Invoice Requests are a request for an invoice; the human-readable prefix for invoice requests is lnr.

There are two similar-looking uses for invoice requests, which are almost identical from a workflow perspective, but are quite different from a user's point of view.

One is a response to an offer; this contains the `offer_issuer_id` or `offer_paths` and all other offer details, and is generally received over an onion message: if it's valid and refers to a known offer, the response is generally to reply with an invoice using the `reply_path` field of the onion message.

The second case is publishing an invoice request without an offer, such as via QR code. It contains neither `offer_issuer_id` nor `offer_paths`, setting the `invreq_payer_id` (and possibly `invreq_paths`) instead, as it is the one paying: the other offer fields are filled by the creator of the `invoice_request`, forming a kind of offer-to-send-money.

Note: the `invreq_metadata` is numbered 0 (not in the 80-159 range for other `invreq` fields) making it the "numerically-first TLV entry" for [Signature Calculation](#). This ensures that merkle leaves are unguessable, allowing a future compact representation to hide fields while still allowing signature validation.

TLV Fields for `invoice_request`

1. `tlv_stream`: `invoice_request`
2. types:
 1. type: 0 (`invreq_metadata`)
 2. data:
 - [...*byte:blob]
 3. type: 2 (`offer_chains`)
 4. data:
 - [...*chain_hash:chains]
 5. type: 4 (`offer_metadata`)
 6. data:
 - [...*byte:data]
 7. type: 6 (`offer_currency`)
 8. data:
 - [...*utf8:iso4217]
 9. type: 8 (`offer_amount`)
 10. data:
 - [tu64:amount]
 11. type: 10 (`offer_description`)
 12. data:
 - [...*utf8:description]
 13. type: 12 (`offer_features`)
 14. data:
 - [...*byte:features]
 15. type: 14 (`offer_absolute_expiry`)
 16. data:
 - [tu64:seconds_from_epoch]
 17. type: 16 (`offer_paths`)
 18. data:
 - [...*blinded_path:paths]
 19. type: 18 (`offer_issuer`)

- 20. data:
 - [...*utf8:issuer]
- 21. type: 20 (offer_quantity_max)
- 22. data:
 - [tu64:max]
- 23. type: 22 (offer_issuer_id)
- 24. data:
 - [point:id]
- 25. type: 80 (invreq_chain)
- 26. data:
 - [chain_hash:chain]
- 27. type: 82 (invreq_amount)
- 28. data:
 - [tu64:msat]
- 29. type: 84 (invreq_features)
- 30. data:
 - [...*byte:features]
- 31. type: 86 (invreq_quantity)
- 32. data:
 - [tu64:quantity]
- 33. type: 88 (invreq_payer_id)
- 34. data:
 - [point:key]
- 35. type: 89 (invreq_payer_note)
- 36. data:
 - [...*utf8:note]
- 37. type: 90 (invreq_paths)
- 38. data:
 - [...*blinded_path:paths]
- 39. type: 91 (invreq_bip_353_name)
- 40. data:
 - [u8:name_len]
 - [name_len*byte:name]
 - [u8:domain_len]
 - [domain_len*byte:domain]
- 41. type: 240 (signature)
- 42. data:
 - [bip340sig:sig]

Requirements for Invoice Requests

The writer: - if it is responding to an offer: - MUST copy all fields from the offer (including unknown fields). - if offer_chains is set: - MUST set invreq_chain to one of offer_chains unless that chain is bitcoin, in which case it SHOULD omit invreq_chain. - otherwise: - if it sets invreq_chain it MUST set it to bitcoin. - MUST set signature.sig as detailed in [Signature Calculation](#) using the invreq_payer_id. - if offer_amount is not present: - MUST specify invreq_amount. - otherwise: - MAY omit invreq_amount. - if it sets invreq_amount: - MUST specify invreq_amount.msat as greater or equal to amount expected by offer_amount (and, if present,

offer_currency and invreq_quantity). - MUST set invreq_payer_id to a transient public key. - MUST remember the secret key corresponding to invreq_payer_id. - if offer_quantity_max is present: - MUST set invreq_quantity to greater than zero. - if offer_quantity_max is non-zero: - MUST set invreq_quantity less than or equal to offer_quantity_max. - otherwise: - MUST NOT set invreq_quantity - otherwise (not responding to an offer): - MUST set offer_description to a complete description of the purpose of the payment. - MUST set (or not set) offer_absolute_expiry and offer_issuer as it would for an offer. - MUST set invreq_payer_id (as it would set offer_issuer_id for an offer). - MUST set invreq_paths as it would set (or not set) offer_paths for an offer. - MUST NOT include signature, offer_metadata, offer_chains, offer_amount, offer_currency, offer_features, offer_quantity_max, offer_paths or offer_issuer_id - if the chain for the invoice is not solely bitcoin: - MUST specify invreq_chain the offer is valid for. - MUST set invreq_amount. - MUST NOT set any non-signature TLV fields outside the inclusive ranges: 0 to 159 and 1000000000 to 2999999999 - MUST set invreq_metadata to an unpredictable series of bytes. - if it sets invreq_amount: - MUST set msat in multiples of the minimum lightning-payable unit (e.g. milli-satoshis for bitcoin) for invreq_chain (or for bitcoin, if there is no invreq_chain). - if it supports bolt12 invoice request features: - MUST set invreq_features.features to the bitmap of features. - if it received the offer from which it constructed this invoice_request using BIP 353 resolution: - MUST include invreq_bip_353_name with, - name set to the post-~~£~~, pre-@ part of the BIP 353 HRN, - domain set to the post-@ part of the BIP 353 HRN.

The reader: - MUST reject the invoice request if invreq_payer_id or invreq_metadata are not present. - MUST reject the invoice request if any non-signature TLV fields are outside the inclusive ranges: 0 to 159 and 1000000000 to 2999999999 - if invreq_features contains unknown *odd* bits that are non-zero: - MUST ignore the bit. - if invreq_features contains unknown *even* bits that are non-zero: - MUST reject the invoice request. - MUST reject the invoice request if signature is not correct as detailed in [Signature Calculation](#) using the invreq_payer_id. - if num_hops is 0 in any blinded_path in invreq_paths: - MUST reject the invoice request. - if offer_issuer_id is present, and invreq_metadata is identical to a previous invoice_request: - MAY simply reply with the previous invoice. - otherwise: - MUST NOT reply with a previous invoice. - if offer_issuer_id or offer_paths are present (response to an offer): - MUST reject the invoice request if the offer fields do not exactly match a valid, unexpired offer. - if offer_paths is present: - MUST ignore the invoice_request if it did not arrive via one of those paths. - otherwise: - MUST ignore any invoice_request if it arrived via a blinded path. - if offer_quantity_max is present: - MUST reject the invoice request if there is no invreq_quantity field. - if offer_quantity_max is non-zero: - MUST reject the invoice request if invreq_quantity is zero, OR greater than offer_quantity_max. - otherwise: - MUST reject the invoice request if there is an invreq_quantity field. - if offer_amount is present: - MUST calculate the *expected amount* using the offer_amount: - if offer_currency is not the invreq_chain currency, convert to the invreq_chain currency. - if invreq_quantity is present, multiply by invreq_quantity.quantity. - if invreq_amount is present: - MUST reject the invoice request if invreq_amount.msat is less than the *expected amount*. - MAY reject the invoice request if invreq_amount.msat greatly exceeds the *expected amount*. - otherwise (no offer_amount): - MUST reject the invoice request if it does not contain invreq_amount. - SHOULD send an invoice in response using the onionmsg_tlv reply_path. - otherwise (no offer_issuer_id or offer_paths, not a response to our offer): - MUST reject the invoice request if any of the following are present: - offer_chains, offer_features or offer_quantity_max. - MUST reject the invoice request if invreq_amount is not present. - MAY use offer_amount (or offer_currency) for informational display to user. - if it sends an invoice in response: - MUST use invreq_paths if present, otherwise MUST use invreq_payer_id as the node id to send to. - if invreq_chain is not present: - MUST reject the invoice request if bitcoin is not a supported chain. - otherwise: - MUST reject the invoice request if invreq_chain.chain is not a supported chain. - if invreq_bip_353_name is present: - MUST reject the invoice request if name or domain contain any bytes which are not 0-9, a-z, A-Z, -, _ or ..

Rationale

`invreq_metadata` might typically contain information about the derivation of the `invreq_payer_id`. This should not leak any information (such as using a simple BIP-32 derivation path); a valid system might be for a node to maintain a base payer key and encode a 128-bit tweak here. The `payer_id` would be derived by tweaking the base key with `SHA256(payer_base_pubkey || tweak)`. It's also the first entry (if present), ensuring an unpredictable nonce for hashing.

`invreq_payer_note` allows you to compliment, taunt, or otherwise engrave graffiti into the invoice for all to see.

Users can give a tip (or obscure the amount sent) by specifying an `invreq_amount` in their invoice request, even though the offer specifies an `offer_amount`. The recipient will only accept this if the invoice request amount exceeds the amount it's expecting (i.e. its `offer_amount` after any currency conversion, multiplied by `invreq_quantity`, if any).

Non-offer-response invoice requests are currently required to explicitly state the `invreq_amount` in the chain currency, so `offer_amount` and `offer_currency` are redundant (but may be informative for the payer to know how the sender claims `invreq_amount` was derived).

The requirement to use `offer_paths` if present, ensures a node does not reveal it is the source of an offer if it is asked directly. Similarly, the requirement that the correct path is used for the offer ensures that cannot be made to reveal that it is the same node that created some other offer.

Invoices

Invoices are a payment request, and when the payment is made, the payment preimage can be combined with the invoice to form a cryptographic receipt.

The recipient sends an invoice in response to an `invoice_request` using the `onion_message` invoice field.

1. `tlv_stream`: invoice
2. `types`:
 1. `type`: 0 (`invreq_metadata`)
 2. `data`:
 - `[...*byte:blob]`
 3. `type`: 2 (`offer_chains`)
 4. `data`:
 - `[...*chain_hash:chains]`
 5. `type`: 4 (`offer_metadata`)
 6. `data`:
 - `[...*byte:data]`
 7. `type`: 6 (`offer_currency`)
 8. `data`:
 - `[...*utf8:iso4217]`
 9. `type`: 8 (`offer_amount`)
 10. `data`:
 - `[tu64:amount]`
 11. `type`: 10 (`offer_description`)

12. data:
 ▪ [...*utf8:description]

13. type: 12 (offer_features)

14. data:
 ▪ [...*byte:features]

15. type: 14 (offer_absolute_expiry)

16. data:
 ▪ [tu64:seconds_from_epoch]

17. type: 16 (offer_paths)

18. data:
 ▪ [...*blinded_path:paths]

19. type: 18 (offer_issuer)

20. data:
 ▪ [...*utf8:issuer]

21. type: 20 (offer_quantity_max)

22. data:
 ▪ [tu64:max]

23. type: 22 (offer_issuer_id)

24. data:
 ▪ [point:id]

25. type: 80 (invreq_chain)

26. data:
 ▪ [chain_hash:chain]

27. type: 82 (invreq_amount)

28. data:
 ▪ [tu64:msat]

29. type: 84 (invreq_features)

30. data:
 ▪ [...*byte:features]

31. type: 86 (invreq_quantity)

32. data:
 ▪ [tu64:quantity]

33. type: 88 (invreq_payer_id)

34. data:
 ▪ [point:key]

35. type: 89 (invreq_payer_note)

36. data:
 ▪ [...*utf8:note]

37. type: 90 (invreq_paths)

38. data:
 ▪ [...*blinded_path:paths]

39. type: 91 (invreq_bip_353_name)

40. data:
 ▪ [u8:name_len]
 ▪ [name_len*byte:name]
 ▪ [u8:domain_len]
 ▪ [domain_len*byte:domain]

- 41. type: 160 (invoice_paths)
- 42. data:
 - [...*blinded_path:paths]
- 43. type: 162 (invoice_blindedpay)
- 44. data:
 - [...*blinded_payinfo:payinfo]
- 45. type: 164 (invoice_created_at)
- 46. data:
 - [tu64:timestamp]
- 47. type: 166 (invoice_relative_expiry)
- 48. data:
 - [tu32:seconds_from_creation]
- 49. type: 168 (invoice_payment_hash)
- 50. data:
 - [sha256:payment_hash]
- 51. type: 170 (invoice_amount)
- 52. data:
 - [tu64:msat]
- 53. type: 172 (invoice_fallbacks)
- 54. data:
 - [...*fallback_address:fallbacks]
- 55. type: 174 (invoice_features)
- 56. data:
 - [...*byte:features]
- 57. type: 176 (invoice_node_id)
- 58. data:
 - [point:node_id]
- 59. type: 240 (signature)
- 60. data:
 - [bip340sig:sig]
- 3. subtype: blinded_payinfo
- 4. data:
 - [u32:fee_base_msat]
 - [u32:fee_proportional_millionths]
 - [u16:cltv_expiry_delta]
 - [u64:htlc_minimum_msat]
 - [u64:htlc_maximum_msat]
 - [u16:flen]
 - [flen*byte:features]
- 5. subtype: fallback_address
- 6. data:
 - [byte:version]
 - [u16:len]
 - [len*byte:address]

Invoice Features

Bits	Description	Name
16	Multi-part-payment support	MPP/compulsory
17	Multi-part-payment support	MPP/optional

The 'MPP support' invoice feature indicates that the payer **MUST** (16) or **MAY** (17) use multiple part payments to pay the invoice.

Some implementations may not support MPP (e.g. for small payments), or may (due to capacity limits on a single channel) require it.

Requirements

A writer of an invoice: - **MUST** set `invoice_created_at` to the number of seconds since Midnight 1 January 1970, UTC when the invoice was created. - **MUST** set `invoice_amount` to the minimum amount it will accept, in units of the minimal lightning-payable unit (e.g. milli-satoshis for bitcoin) for `invreq_chain`. - if the invoice is in response to an `invoice_request`: - **MUST** copy all non-signature fields from the invoice request (including unknown fields). - if `invreq_amount` is present: - **MUST** set `invoice_amount` to `invreq_amount` - otherwise: - **MUST** set `invoice_amount` to the *expected amount*. - **MUST** set `invoice_payment_hash` to the SHA256 hash of the `payment_preimage` that will be given in return for payment. - if `offer_issuer_id` is present: - **MUST** set `invoice_node_id` to the `offer_issuer_id` - otherwise, if `offer_paths` is present: - **MUST** set `invoice_node_id` to the final `blinded_node_id` on the path it received the invoice request - **MUST** specify exactly one signature TLV element: `signature`. - **MUST** set `sig` to the signature using `invoice_node_id` as described in [Signature Calculation](#). - if it requires multiple parts to pay the invoice: - **MUST** set `invoice_features.features` bit `MPP/compulsory` - or if it allows multiple parts to pay the invoice: - **MUST** set `invoice_features.features` bit `MPP/optional` - if the expiry for accepting payment is not 7200 seconds after `invoice_created_at`: - **MUST** set `invoice_relative_expiry.seconds_from_creation` to the number of seconds after `invoice_created_at` that payment of this invoice should not be attempted. - if it accepts onchain payments: - **MAY** specify `invoice_fallbacks` - **SHOULD** specify `invoice_fallbacks` in order of most-preferred to least-preferred if it has a preference. - for the bitcoin chain, it **MUST** set each `fallback_address` with `version` as a valid witness version and `address` as a valid witness program - **MUST** include `invoice_paths` containing one or more paths to the node. - **MUST** specify `invoice_paths` in order of most-preferred to least-preferred if it has a preference. - **MUST** include `invoice_blindedpay` with exactly one `blinded_payinfo` for each `blinded_path` in `paths`, in order. - **MUST** set `features` in each `blinded_payinfo` to match `encrypted_data_tlv.allowed_features` (or empty, if no `allowed_features`). - **SHOULD** ignore any payment which does not use one of the paths. - **MUST NOT** set any non-signature TLV fields outside the inclusive ranges: 0 to 239 and 1000000000 to 3999999999.

A reader of an invoice: - **MUST** reject the invoice if `invoice_amount` is not present. - **MUST** reject the invoice if `invoice_created_at` is not present. - **MUST** reject the invoice if `invoice_payment_hash` is not present. - **MUST** reject the invoice if `invoice_node_id` is not present. - if `invreq_chain` is not present: - **MUST** reject the invoice if bitcoin is not a supported chain. - otherwise: - **MUST** reject the invoice if `invreq_chain.chain` is not a supported chain. - if `invoice_features` contains unknown *odd* bits that are non-zero: - **MUST** ignore the bit. - if `invoice_features` contains unknown *even* bits that are non-zero: - **MUST** reject the invoice. - if `invoice_relative_expiry` is present: - **MUST** reject the invoice if the current time since 1970-01-01 UTC is greater than `invoice_created_at` plus `seconds_from_creation`. - otherwise: - **MUST** reject the invoice if the current time since 1970-01-01 UTC is greater than `invoice_created_at` plus 7200. - **MUST** reject the invoice if

invoice_paths is not present or is empty. - MUST reject the invoice if num_hops is 0 in any blinded_path in invoice_paths. - MUST reject the invoice if invoice_blindedpay is not present. - MUST reject the invoice if invoice_blindedpay does not contain exactly one blinded_payinfo per invoice_paths.blinded_path. - For each invoice_blindedpay.payinfo: - MUST NOT use the corresponding invoice_paths.path if payinfo.features has any unknown even bits set. - MUST reject the invoice if this leaves no usable paths. - if the invoice is a response to an invoice_request: - MUST reject the invoice if all fields in ranges 0 to 159 and 1000000000 to 2999999999 (inclusive) do not exactly match the invoice request. - if offer_issuer_id is present (invoice_request for an offer): - MUST reject the invoice if invoice_node_id is not equal to offer_issuer_id - otherwise, if offer_paths is present (invoice_request for an offer without id): - MUST reject the invoice if invoice_node_id is not equal to the final blinded_node_id it sent the invoice request to. - otherwise (invoice_request without an offer): - MAY reject the invoice if it cannot confirm that invoice_node_id is correct, out-of-band. - MUST reject the invoice if signature is not a valid signature using invoice_node_id as described in [Signature Calculation](#). - SHOULD prefer to use earlier invoice_paths over later ones if it has no other reason for preference. - if invoice_features contains the MPP/compulsory bit: - MUST pay the invoice via multiple separate blinded paths. - otherwise, if invoice_features contains the MPP/optional bit: - MAY pay the invoice via multiple separate payments. - otherwise: - MUST NOT use multiple parts to pay the invoice. - if invreq_amount is present: - MUST reject the invoice if invoice_amount is not equal to invreq_amount - otherwise: - SHOULD confirm authorization if invoice_amount.msat is not within the amount range authorized. - for the bitcoin chain, if the invoice specifies invoice_fallbacks: - MUST ignore any fallback_address for which version is greater than 16. - MUST ignore any fallback_address for which address is less than 2 or greater than 40 bytes. - MUST ignore any fallback_address for which address does not meet known requirements for the given version - if invreq_paths is present: - MUST reject the invoice if it did not arrive via one of those paths. - otherwise, neither offer_issuer_id nor offer_paths are present (not derived from an offer): - MUST reject the invoice if it arrived via a blinded path. - otherwise (derived from an offer): - MUST reject the invoice if it did not arrive via invoice_request onionmsg_tlv reply_path.

Rationale

Because the messaging layer is unreliable, it's quite possible to receive multiple requests for the same offer. As it's the caller's responsibility to make invreq_metadata both unpredictable and unique, the writer doesn't have to check all the fields are duplicates before simply returning a previous invoice. Note that such caching is optional, and should be carefully limited when e.g. currency conversion is involved, or if the invoice has expired.

The invoice duplicates fields rather than committing to the previous invreq. This flattened format simplifies storage at some space cost, as the payer need only remember the invoice for any refunds or proof.

The reader of the invoice cannot trust the invoice correctly reflects the invreq fields, hence the requirements to check that they are correct, although allowance is made for simply sending an unrequested invoice directly.

Note that the recipient of the invoice can determine the expected amount from either the offer it received, or the invreq it sent, so often already has authorization for the expected amount.

The default invoice_relative_expiry of 7200 seconds, which is generally a sufficient time for payment, even if new channels need to be opened.

Blinded paths provide an equivalent to payment_secret and payment_metadata used in BOLT 11. Even if invoice_node_id or invreq_payer_id is public, we force the use of blinding paths to keep these features. If the

recipient does not care about the added privacy offered by blinded paths, they can create a path of length 1 with only themselves.

Rather than provide detailed per-hop-payinfo for each hop in a blinded path, we aggregate the fees and CLTV deltas. This avoids trivially revealing any distinguishing non-uniformity which may distinguish the path.

In the case of an invoice where there was no offer (just an invoice request), the payer needs to ensure that the invoice is from the intended payment recipient. This is the basis for the suggestion to confirm the `invoice_node_id` for this case.

Raw invoices (not based on an `invoice_request`) are generally not supported, though an implementation is allowed to support them, and we may define the behavior in the future. The redundant requirement to check `invreq_chain` explicitly is a nod to this: if the invoice is a response to an invoice request, that field must have existed due to the invoice request requirements, and we also require it to be mirrored here.

Invoice Errors

Informative errors can be returned in an onion message `invoice_error` field (via the onion `reply_path`) for either `invoice_request` or `invoice`.

TLV Fields for `invoice_error`

1. `tlv_stream`: `invoice_error`
2. `types`:
 1. `type`: 1 (`erroneous_field`)
 2. `data`:
 - `[tu64:tlv_fieldnum]`
 3. `type`: 3 (`suggested_value`)
 4. `data`:
 - `[...*byte:value]`
 5. `type`: 5 (`error`)
 6. `data`:
 - `[...*utf8:msg]`

Requirements

A writer of an `invoice_error`: - MUST set `error` to an explanatory string. - MAY set `erroneous_field` to a specific field number in the `invoice` or `invoice_request` which had a problem. - if it sets `erroneous_field`: - MAY set `suggested_value`. - if it sets `suggested_value`: - MUST set `suggested_value` to a valid field for that `tlv_fieldnum`. - otherwise: - MUST NOT set `suggested_value`.

A reader of an `invoice_error`: FIXME!

Rationale

Usually an error message is sufficient for diagnostics, however future enhancements may make automated handling useful.

In particular, we could allow non-offer-response `invoice_requests` to omit `invreq_amount` in the future and use offer fields to indicate alternate currencies. (“I will send you 10c!”). Then the sender of the invoice would have to guess how many msat that was, and could use the `invoice_error` to indicate if the recipient disagreed with the conversion so the sender can send a new invoice.

Payer Proofs

Payer proofs are proofs of invoice payment; the human-readable prefix for payer proofs is `lnp`.

The payer proof is based on the invoice TLV stream, with the exception that `invreq_metadata` cannot be included. Various other invoice fields may be omitted for privacy: numbers corresponding to (but not identical to) their position in the invoice are included, as well as the minimal hashes for missing merkle branches, to allow verification of the invoicing node’s signature.

To prove that this `payer_proof` was created by someone who has the secret key used to request the invoice in the first place, it includes a signature using the `invreq_payer_id`: this signs the proof fields and invoice fields.

TLV Fields for `payer_proof`

1. `tlv_stream`: `payer_proof`
2. `types`:
 1. `type`: 2 (`offer_chains`)
 2. `data`:
 - `[...*chain_hash:chains]`
 3. `type`: 4 (`offer_metadata`)
 4. `data`:
 - `[...*byte:data]`
 5. `type`: 6 (`offer_currency`)
 6. `data`:
 - `[...*utf8:iso4217]`
 7. `type`: 8 (`offer_amount`)
 8. `data`:
 - `[tu64:amount]`
 9. `type`: 10 (`offer_description`)
 10. `data`:
 - `[...*utf8:description]`
 11. `type`: 12 (`offer_features`)
 12. `data`:
 - `[...*byte:features]`
 13. `type`: 14 (`offer_absolute_expiry`)
 14. `data`:
 - `[tu64:seconds_from_epoch]`
 15. `type`: 16 (`offer_paths`)
 16. `data`:
 - `[...*blinded_path:paths]`
 17. `type`: 18 (`offer_issuer`)
 18. `data`:
 - `[...*utf8:issuer]`

19. type: 20 (offer_quantity_max)
20. data:
 ▪ [tu64:max]
21. type: 22 (offer_issuer_id)
22. data:
 ▪ [point:id]
23. type: 80 (invreq_chain)
24. data:
 ▪ [chain_hash:chain]
25. type: 82 (invreq_amount)
26. data:
 ▪ [tu64:msat]
27. type: 84 (invreq_features)
28. data:
 ▪ [...*byte:features]
29. type: 86 (invreq_quantity)
30. data:
 ▪ [tu64:quantity]
31. type: 88 (invreq_payer_id)
32. data:
 ▪ [point:key]
33. type: 89 (invreq_payer_note)
34. data:
 ▪ [...*utf8:note]
35. type: 90 (invreq_paths)
36. data:
 ▪ [...*blinded_path:paths]
37. type: 91 (invreq_bip_353_name)
38. data:
 ▪ [u8:name_len]
 ▪ [name_len*byte:name]
 ▪ [u8:domain_len]
 ▪ [domain_len*byte:domain]
39. type: 160 (invoice_paths)
40. data:
 ▪ [...*blinded_path:paths]
41. type: 162 (invoice_blindedpay)
42. data:
 ▪ [...*blinded_payinfo:payinfo]
43. type: 164 (invoice_created_at)
44. data:
 ▪ [tu64:timestamp]
45. type: 166 (invoice_relative_expiry)
46. data:
 ▪ [tu32:seconds_from_creation]
47. type: 168 (invoice_payment_hash)

48. data:
 ▪ [sha256:payment_hash]
 49. type: 170 (invoice_amount)
 50. data:
 ▪ [tu64:msat]
 51. type: 172 (invoice_fallbacks)
 52. data:
 ▪ [...*fallback_address:fallbacks]
 53. type: 174 (invoice_features)
 54. data:
 ▪ [...*byte:features]
 55. type: 176 (invoice_node_id)
 56. data:
 ▪ [point:node_id]
 57. type: 240 (signature)
 58. data:
 ▪ [bip340sig:sig]
 59. type: 241 (proof_signature)
 60. data:
 ▪ [bip340sig:sig]
 61. type: 1001 (proof_preimage)
 62. data:
 ▪ [32*byte:preimage]
 63. type: 1002 (proof_omitted_tlvs)
 64. data:
 ▪ [...*bigsize:missing]
 65. type: 1003 (proof_missing_hashes)
 66. data:
 ▪ [...*sha256:hashes]
 67. type: 1004 (proof_leaf_hashes)
 68. data:
 ▪ [...*sha256:hashes]
 69. type: 1005 (proof_note)
 70. data:
 ▪ [...*utf8:note]

Requirements

A writer of a payer_proof: - MUST NOT include invreq_metadata. - MUST include invreq_payer_id, invoice_payment_hash, invoice_node_id, signature and (if present) invoice_features from the invoice. - MUST include proof_preimage containing the payment_preimage returned from successful payment of this invoice. - For each non-signature TLV in the invoice in ascending-type order: - If the field is to be included in the payer_proof: - MUST copy it into the payer_proof. - MUST append the nonce (H("LnNonce" || TLV0,type)) to proof_leaf_hashes. - otherwise, if the TLV type is not zero: - MUST append a *marker number* to proof_omitted_tlvs - If the previous TLV type was included: - The *marker number* is that previous tlv type, plus one. - Otherwise, if proof_omitted_tlvs is empty: - The *marker number* is 1. - Otherwise, if the last proof_omitted_tlvs entry is 239: - The *marker number* is 1000000000. - Otherwise: - The *marker number* is one

greater than the last `proof_omitted_tlvs` entry. - If `proof_omitted_tlvs` is empty: - MAY omit `proof_omitted_tlvs` from the `payer_proof`. - MAY include an annotation on the proof in `proof_note`. - MUST populate `proof_missing_hashes` with the merkle hash of the omitted branch of each internal node that has exactly one branch entirely omitted, in post-order depth-first smallest-to-largest TLV order. - MUST copy signature into the `payer_proof`. - MUST set `proof_signature` as detailed in [Signature Calculation](#) using the `invreq_payer_id`.

A reader of a `payer_proof`: - MUST reject the `payer_proof` if: - `invreq_payer_id`, `invoice_payment_hash`, `invoice_node_id`, `signature`, `proof_preimage`, `proof_missing_hashes`, `proof_leaf_hashes` or `proof_signature` are missing. - `SHA256(proof_preimage)` does not equal `invoice_payment_hash`. - `proof_omitted_tlvs` are not in strict ascending order (no duplicates). - `proof_omitted_tlvs` contains 0. - `proof_omitted_tlvs` contains number outside both ranges 1 to 239 and 1000000000 to 3999999999. - `proof_omitted_tlvs` contains the number of an included TLV field. - `proof_omitted_tlvs` contains a number which is not: - one greater than an included TLV number, - one greater than the previous `proof_omitted_tlvs` entry, or 0 if it is the first number, or - 1000000000, if the previous `proof_omitted_tlvs` entry is 239. - `proof_leaf_hashes` does not contain exactly one hash for each field in ranges 1 to 239 and 1000000000 to 3999999999. - There are not exactly enough `proof_missing_hashes` to reconstruct the merkle tree root using the `proof_omitted_tlvs` values (with 0 implied as the first omitted TLV). - `signature` is not a valid signature using `invoice_node_id` as described in [Signature Calculation](#) (with messagename "invoice") of the reconstructed merkle-root of the invoice (i.e. without fields 1001 through 999999999 inclusive). - `proof_signature` is not a valid signature using `invreq_payer_id` as described in [Signature Calculation](#).

Rationale

Using the invoice as a base enshrines information about the payment including important offer and `invoice_request` fields. However, many fields are not useful (such as payment paths), or may compromise privacy (such as `invreq_payer_note` containing delivery address information), so being able to elide them while still allowing signature validation is vital.

We disallow including `invreq_metadata`: that is the hashing nonce, thus allowing brute-force of omitted fields.

`invreq_payer_id` is the key whose signature we have to attach to the proof, and `invoice_node_id` and `signature` are needed to validate the original invoice. `invoice_features` may indicate additional details in the future which would require additional fields to be in the proof. Note that `invoice_amount` is not compulsory, though it would probably be very useful in most cases.

The requirement to include minimal hashes (rather than one for every unknown leaf) minimizes the size, especially when many consecutive fields are omitted. As the exact TLV types of omitted TLVs are unimportant (as long as ordering is maintained), we renumber them to be minimal, as further obfuscation of values.

The proof fields are outside the established offer, invoice request and invoice TLV ranges, and above the signature range (240-1000), so they are committed to by the `proof_signature`.


The optional `proof_note` field allows a challenge-response system to be implemented: someone requiring proof can ask for a signature with a particular note. It can also be missing.

Example for Payer Proofs

Consider a trivial TLV construct (not a valid invoice) which we are trying to prove, with the following fields:

0 - Omitted 10 - Omitted 20 - Omitted 30 - Omitted 40 - Included 50 - Omitted 60 - Omitted 240 - Omitted (signature field)

Here is the full signature Merkle tree, with omitted nodes marked with (o):



Note that the signature TLV 240 is not included in the merkle tree.

proof_leaf_hashes contains the nonce hashes for the present non-signature TLVs:

1. $H(\text{"LnNonce"} \parallel \text{TLV0}, 40)$

Since four adjacent nodes (0, 10 20 and 30) are omitted, we can (and must) simply provide the hash of the node above them, marked with an asterisk.

Thus, proof_missing_hashes contains the following hashes in order:

1. Merkle of $H(\text{"LnLeaf"}, \text{TLV50})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 50)$
2. Merkle of $H(\text{"LnLeaf"}, \text{TLV60})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 60)$
3. Merkle of (Merkle of (Merkle of $H(\text{"LnLeaf"}, \text{TLV0})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 0)$) (Merkle of $H(\text{"LnLeaf"}, \text{TLV10})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 10)$)) and (Merkle of (Merkle of $H(\text{"LnLeaf"}, \text{TLV20})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 20)$) (Merkle of $H(\text{"LnLeaf"}, \text{TLV30})$ and $H(\text{"LnNonce"} \parallel \text{TLV0}, 30)$))

In this example the correct proof_missing_hashes order is not ascending TLV order: the omitted subtree containing TLVs 0, 10, 20, and 30 is emitted after the omitted TLVs 50 and 60, because it is the missing sibling of their parent's parent.

The proof_omitted_tlvs array is based on the omitted tlvs: [0, 10, 20, 30, 50, 60]. It uses the minimal values which hide the real field numbers without changing their order, 0 is implied (as it's always omitted), giving an array of [1, 2, 3, 41, 42].

The algorithm for creating proof_missing_hashes is most easily implemented in a recursive fashion, traversing smallest-to-largest TLV (left-to-right in the above representation). When you need to combine two hashes where one side is entirely omitted and the other is not, append that hash to proof_missing_hashes.

Reconstruction is the exact opposite: when you need to combine a hash where one side is entirely omitted and the other is not, pull a hash from proof_missing_hashes. If there are insufficient proof_missing_hashes, or it isn't empty when you have completed the merkle tree, the number of proof_missing_hashes was incorrect.

See the [Payer Proof Test Vectors](#) for more examples.

FIXME: Possible future extensions:

1. The offer can require delivery info in the `invoice_request`.
2. An offer can be updated: the response to an `invoice_request` is another offer, perhaps with a signature from the original `offer_issuer_id`
3. Any empty TLV fields can mean the value is supposed to be known by other means (i.e. transport-specific), but is still hashed for sig.
4. We could upgrade to allow multiple offers in one `invreq` and `invoice`, to make a shopping list.
5. All-zero `offer_id` == gratuitous payment.
6. Streaming invoices?
7. Re-add recurrence.
8. Re-add `invoice_replace` for requesting replacement of a (stuck-payment) invoice with a new one.
9. Allow non-offer `invoice_request` with alternate currencies?
10. Add `offer_quantity_unit` to indicate stepping for quantity (e.g. 100 grams).

[1] https://www.youtube.com/watch?v=4SYc_fIMnMQ

When Channels Close

Channels are great until they aren't. A peer disappears, a revocation secret gets broadcast, an HTLC is about to expire. Then you're on-chain, and the chain doesn't care about your encrypted transport.

BOLT 05 is the playbook for that. It tells an implementation how to watch for commitment transactions, when to broadcast a penalty, how to resolve HTLCs on-chain, and how to sweep outputs so you actually get your coins back. It's less about messages and more about timeouts, scripts, and not racing yourself.

Read this after you understand commitment transactions in BOLT 03. The on-chain rules only make sense once you know what those transactions look like.

Most of the time a cooperative close from BOLT 02 is enough. BOLT 05 is for the rest of the time — which is exactly when you need a spec.

In this chapter

- BOLT 05 — On-chain closing and recovery

BOLT #5: Recommendations for On-chain Transaction Handling

Abstract

Lightning allows for two parties (a local node and a remote node) to conduct transactions off-chain by giving each of the parties a *cross-signed commitment transaction*, which describes the current state of the channel (basically, the current balance). This *commitment transaction* is updated every time a new payment is made and is spendable at all times.

There are three ways a channel can end:

1. The good way (*mutual close*): at some point the local and remote nodes agree to close the channel. They generate a *closing transaction* (which is similar to a commitment transaction, but without any pending payments) and publish it on the blockchain (see [BOLT #2: Channel Close](#)).
2. The bad way (*unilateral close*): something goes wrong, possibly without evil intent on either side. Perhaps one party crashed, for instance. One side publishes its *latest commitment transaction*.
3. The ugly way (*revoked transaction close*): one of the parties deliberately tries to cheat, by publishing an *outdated commitment transaction* (presumably, a prior version, which is more in its favor).

Because Lightning is designed to be trustless, there is no risk of loss of funds in any of these three cases; provided that the situation is properly handled. The goal of this document is to explain exactly how a node should react when it encounters any of the above situations, on-chain.

Table of Contents

- [General Nomenclature](#)
- [Commitment Transaction](#)
- [Failing a Channel](#)
- [Mutual Close Handling](#)
- [Unilateral Close Handling: Local Commitment Transaction](#)
 - [HTLC Output Handling: Local Commitment, Local Offers](#)
 - [HTLC Output Handling: Local Commitment, Remote Offers](#)
- [Unilateral Close Handling: Remote Commitment Transaction](#)
 - [HTLC Output Handling: Remote Commitment, Local Offers](#)
 - [HTLC Output Handling: Remote Commitment, Remote Offers](#)
- [Revoked Transaction Close Handling](#)
 - [Penalty Transactions Weight Calculation](#)
- [Generation of HTLC Transactions](#)
- [General Requirements](#)
 - [Limitations of v3 transactions](#)
- [Appendix A: Expected Weights](#)
 - [Expected Weight of the `to\_local` Penalty Transaction Witness](#)
 - [Expected Weight of the `offered\_htlc` Penalty Transaction Witness](#)
 - [Expected Weight of the `accepted\_htlc` Penalty Transaction Witness](#)
- [Authors](#)

General Nomenclature

Any unspent output is considered to be *unresolved* and can be *resolved* as detailed in this document. Usually this is accomplished by spending it with another *resolving* transaction. Although, sometimes simply noting the output for later wallet spending is sufficient, in which case the transaction containing the output is considered to be its own *resolving* transaction.

Outputs that are *resolved* are considered *irrevocably resolved* once the remote's *resolving* transaction is included in a block at least 100 deep, on the most-work blockchain. 100 blocks is far greater than the longest known Bitcoin fork and is the same wait time used for confirmations of miners' rewards (see [Reference Implementation](#)).

Requirements

A node: - once it has broadcast a funding transaction OR sent a commitment signature for a commitment transaction that contains an HTLC output: - until all outputs are *irrevocably resolved*: - MUST monitor the blockchain for transactions that spend any output that is NOT *irrevocably resolved*. - MUST *resolve* all outputs, as specified below. - MUST be prepared to resolve outputs multiple times, in case of blockchain reorganizations. - upon the funding transaction being spent, if the channel is NOT already closed: - MAY send a descriptive error. - SHOULD fail the channel. - SHOULD ignore invalid transactions.

Rationale

Once a local node has some funds at stake, monitoring the blockchain is required to ensure the remote node does not close unilaterally.

Invalid transactions (e.g. bad signatures) can be generated by anyone, (and will be ignored by the blockchain anyway), so they should not trigger any action.

Commitment Transaction

The local and remote nodes each hold a *commitment transaction*. Each of these commitment transactions has up to seven types of outputs:

1. *local node's main output*: Zero or one output, to pay to the *local node's* `delayed_pubkey`.
2. *remote node's main output*: Zero or one output, to pay to the *remote node's* `delayed_pubkey`.
3. *local node's anchor output*: one output paying to the *local node's* `funding_pubkey`.
4. *remote node's anchor output*: one output paying to the *remote node's* `funding_pubkey`.
5. *shared anchor*: one output spendable by anyone (pay-to-anchor).
6. *local node's offered HTLCs*: Zero or more pending payments (*HTLCs*), to pay the *remote node* in return for a payment preimage.
7. *remote node's offered HTLCs*: Zero or more pending payments (*HTLCs*), to pay the *local node* in return for a payment preimage.

To incentivize the local and remote nodes to cooperate, an `OP_CHECKSEQUENCEVERIFY` relative timeout encumbers the *local node's outputs* (in the *local node's commitment transaction*) and the *remote node's outputs* (in the *remote node's commitment transaction*). So for example, if the local node publishes its commitment transaction, it will have to wait to claim its own funds, whereas the remote node will have immediate access to its own funds. As a consequence, the two commitment transactions are not identical, but they are (usually) symmetrical.

See [BOLT #3: Commitment Transaction](#) for more details.

Failing a Channel

Although closing a channel can be accomplished in several ways, the most efficient is preferred.

Various error cases involve closing a channel. The requirements for sending error messages to peers are specified in [BOLT #1: The error Message](#).

Requirements

A node: - if a *local commitment transaction* has NOT ever contained a `to_local` or HTLC output: - MAY simply forget the channel. - otherwise: - if the *current commitment transaction* does NOT contain `to_local` or other HTLC outputs: - MAY simply wait for the remote node to close the channel. - until the remote node closes: - MUST NOT forget the channel. - otherwise: - if it has received a valid `closing_signed` message that includes a sufficient fee: - SHOULD use this fee to perform a *mutual close*. - otherwise: - if the node knows or assumes its channel state is outdated: - MUST NOT broadcast its *last commitment transaction*. - otherwise: - MUST broadcast the *last commitment transaction*, for which it has a signature, to perform a *unilateral close*. - MUST spend any `to_local_anchor` or `shared_anchor` output, providing sufficient fees as incentive to include the commitment transaction in a block. For `to_local_anchor`, special care must be taken when spending to a third-party, because this re-introduces the pinning vulnerability that was addressed by adding the CSV delay to the non-anchor outputs. - SHOULD use [replace-by-fee](#) or other mechanism on the spending transaction if it proves insufficient for timely inclusion in a block.

Rationale

Since `dust_limit_satoshis` is supposed to prevent creation of uneconomic outputs (which would otherwise remain forever, unspent on the blockchain), all commitment transaction outputs MUST be spent.

In the early stages of a channel, it's common for one side to have little or no funds in the channel; in this case, having nothing at stake, a node need not consume resources monitoring the channel state.

There exists a bias towards preferring mutual closes over unilateral closes, because outputs of the former are unencumbered by a delay and are directly spendable by wallets. In addition, mutual close fees tend to be less exaggerated than those of commitment transactions (or in the case of commitments using anchor outputs, the commitment transaction may require a child transaction to cause it to be mined). So, the only reason not to use the mutual close transaction would be if the fee offered was too small for it to be processed.

Mutual Close Handling

A closing transaction *resolves* the funding transaction output.

In the case of a mutual close, a node need not do anything else, as it has already agreed to the output, which is sent to its specified scriptpubkey (see [BOLT #2: Closing initiation: shutdown](#)).

Unilateral Close Handling: Local Commitment Transaction

This is the first of two cases involving unilateral closes. In this case, a node discovers its *local commitment transaction*, which *resolves* the funding transaction output.

However, a node cannot claim funds from the outputs of a unilateral close that it initiated, until the `OP_CHECKSEQUENCEVERIFY` delay has passed (as specified by the remote node's `to_self_delay` field). Where relevant, this situation is noted below.

Requirements

A node: - upon discovering its *local commitment transaction*: - SHOULD spend the `to_local` output to a convenient address. - MUST wait until the `OP_CHECKSEQUENCEVERIFY` delay has passed (as specified by the

remote node's `to_self_delay` field) before spending the output. - Note: if the output is spent (as recommended), the output is *resolved* by the spending transaction, otherwise it is considered *resolved* by the commitment transaction itself. - MAY ignore the `to_remote` output. - Note: No action is required by the local node, as `to_remote` is considered *resolved* by the commitment transaction itself. - MUST handle HTLCs offered by itself as specified in [HTLC Output Handling: Local Commitment, Local Offers](#). - MUST handle HTLCs offered by the remote node as specified in [HTLC Output Handling: Local Commitment, Remote Offers](#).

Rationale

Spending the `to_local` output avoids having to remember the complicated witness script, associated with that particular channel, for later spending.

The `to_remote` output is entirely the business of the remote node, and can be ignored.

HTLC Output Handling: Local Commitment, Local Offers

Each HTLC output can only be spent by either the *local offerer*, by using the HTLC-timeout transaction after it's timed out, or the *remote recipient*, if it has the payment preimage.

There can be HTLCs which are not represented by any outputs: either because they were trimmed as dust, or because the transaction has only been partially committed.

The HTLC output has *timed out* once the height of the latest block is equal to or greater than the HTLC `cltv_expiry`.

Requirements

A node: - if the commitment transaction HTLC output is spent using the payment preimage, the output is considered *irrevocably resolved*: - MUST extract the payment preimage from the transaction input witness. - if the commitment transaction HTLC output has *timed out* and hasn't been *resolved*: - MUST *resolve* the output by spending it using the HTLC-timeout transaction. - once the resolving transaction has reached reasonable depth: - MUST fail the corresponding incoming HTLC (if any). - MUST resolve the output of that HTLC-timeout transaction. - SHOULD resolve the HTLC-timeout transaction by spending it to a convenient address. - Note: if the output is spent (as recommended), the output is *resolved* by the spending transaction, otherwise it is considered *resolved* by the HTLC-timeout transaction itself. - MUST wait until the `OP_CHECKSEQUENCEVERIFY` delay has passed (as specified by the remote node's `open_channel to_self_delay` field) before spending that HTLC-timeout output. - for any committed HTLC that does NOT have an output in this commitment transaction: - if the payment preimage is known: - MUST fulfill the corresponding incoming HTLC (if any). - otherwise: - once the commitment transaction has reached reasonable depth: - MUST fail the corresponding incoming HTLC (if any). - if no *valid* commitment transaction contains an output corresponding to the HTLC: - MAY fail the corresponding incoming HTLC sooner.

Rationale

The payment preimage either serves to prove payment (when the offering node originated the payment) or to redeem the corresponding incoming HTLC from another peer (when the offering node is forwarding the payment). Once a node has extracted the payment, it no longer cares about the fate of the HTLC-spending transaction itself.

In cases where both resolutions are possible (e.g. when a node receives payment success after timeout), either interpretation is acceptable; it is the responsibility of the recipient to spend it before this occurs.

The local HTLC-timeout transaction needs to be used to time out the HTLC (to prevent the remote node fulfilling it and claiming the funds) before the local node can back-fail any corresponding incoming HTLC, using `update_fail_htlc` (presumably with reason `permanent_channel_failure`), as detailed in [BOLT #2](#). If the incoming HTLC is also on-chain, a node must simply wait for it to timeout: there is no way to signal early failure.

There are several reasons a committed HTLC may not have an output in the confirmed commitment transaction: the HTLC may be smaller than `dust_limit_satoshis`, the HTLC may not have been added to the commitment transaction yet, or the HTLC may have already been failed or fulfilled. In any case, if the payment preimage is known for the HTLC, the upstream HTLC needs to be fulfilled to avoid loss of funds.

If the payment preimage is not known for the missing HTLC, the correct action depends on the possibility of a blockchain reorganization that swaps out the confirmed commitment transaction for one with the HTLC present. If the HTLC is too small to appear in *any commitment transaction*, such a reorganization is not possible, and the HTLC can be safely failed immediately. Otherwise, a reorganization delay is required before failing the incoming HTLC. The requirement that the incoming HTLC be failed before its own timeout still applies as an upper bound.

HTLC Output Handling: Local Commitment, Remote Offers

Each HTLC output can only be spent by the recipient, using the HTLC-success transaction, which it can only populate if it has the payment preimage. If it doesn't have the preimage (and doesn't discover it), it's the offerer's responsibility to spend the HTLC output once it's timed out.

There are several possible cases for an offered HTLC:

1. The offerer is NOT irrevocably committed to it. The recipient will usually not know the preimage, since it will not forward HTLCs until they're fully committed. So using the preimage would reveal that this recipient is the final hop; thus, in this case, it's best to allow the HTLC to time out.
2. The offerer is irrevocably committed to the offered HTLC, but the recipient has not yet committed to an outgoing HTLC. In this case, the recipient can either forward or timeout the offered HTLC.
3. The recipient has committed to an outgoing HTLC, in exchange for the offered HTLC. In this case, the recipient must use the preimage, once it receives it from the outgoing HTLC; otherwise, it will lose funds by sending an outgoing payment without redeeming the incoming payment.

Requirements

A local node: - if it receives (or already possesses) a payment preimage for an unresolved HTLC output that it has been offered AND for which it has committed to an outgoing HTLC: - MUST *resolve* the output by spending it, using the HTLC-success transaction. - MUST NOT reveal its own preimage when it's not the final recipient. [Preimage-Extraction](#) - MUST resolve the output of that HTLC-success transaction. - otherwise: - if the *remote node* is NOT irrevocably committed to the HTLC: - MUST NOT *resolve* the output by spending it. - SHOULD resolve that HTLC-success transaction output by spending it to a convenient address. - MUST wait until the `OP_CHECKSEQUENCEVERIFY` delay has passed (as specified by the *remote node's* `open_channel's` `to_self_delay` field), before spending that HTLC-success transaction output.

If the output is spent (as is recommended), the output is *resolved* by the spending transaction, otherwise it's considered *resolved* by the HTLC-success transaction itself.

If it's NOT otherwise resolved, once the HTLC output has expired, it is considered *irrevocably resolved*.

Unilateral Close Handling: Remote Commitment Transaction

The *remote node's* commitment transaction *resolves* the funding transaction output.

There are no delays constraining node behavior in this case, so it's simpler for a node to handle than the case in which it discovers its local commitment transaction (see [Unilateral Close Handling: Local Commitment Transaction](#)).

Requirements

A local node: - upon discovering a *valid* commitment transaction broadcast by a *remote node*: - if possible: - MUST handle each output as specified below. - MAY take no action in regard to the associated `to_remote`, which is simply a P2WPKH output to the *local node*. - Note: `to_remote` is considered *resolved* by the commitment transaction itself. - MAY take no action in regard to the associated `to_local`, which is a payment output to the *remote node*. - Note: `to_local` is considered *resolved* by the commitment transaction itself. - MUST handle HTLCs offered by itself as specified in [HTLC Output Handling: Remote Commitment, Local Offers](#) - MUST handle HTLCs offered by the remote node as specified in [HTLC Output Handling: Remote Commitment, Remote Offers](#) - otherwise (it is NOT able to handle the broadcast for some reason): - MUST inform the user of potentially lost funds.

Rationale

There may be more than one valid, *unrevoked* commitment transaction after a signature has been received via `commitment_signed` and before the corresponding `revoke_and_ack`. As such, either commitment may serve as the *remote node's* commitment transaction; hence, the local node is required to handle both.

In the case of data loss, a local node may reach a state where it doesn't recognize all of the *remote node's* commitment transaction HTLC outputs. It can detect the data loss state, because it has signed the transaction, and the commitment number is greater than expected. It can derive its own `remotepubkey` for the transaction, in order to salvage its own funds. Note: in this scenario, the node will be unable to salvage the HTLCs.

HTLC Output Handling: Remote Commitment, Local Offers

Each HTLC output can only be spent by either the *local offerer*, after it's timed out, or by the *remote recipient*, by using the HTLC-success transaction if it has the payment preimage.

There can be HTLCs which are not represented by any outputs: either because the outputs were trimmed as dust, or because the remote node has two *valid* commitment transactions with differing HTLCs.

The HTLC output has *timed out* once the depth of the latest block is equal to or greater than the HTLC `cltv_expiry`.

Requirements

A local node: - if the commitment transaction HTLC output is spent using the payment preimage: - MUST extract the payment preimage from the HTLC-success transaction input witness. - Note: the output is considered *irrevocably resolved*. - if the commitment transaction HTLC output has *timed out* AND NOT been *resolved*: - MUST *resolve* the output, by spending it to a convenient address. - for any committed HTLC that does NOT have an output in this commitment transaction: - if the payment preimage is known: - MUST fulfill the corresponding incoming HTLC (if any). - otherwise: - once the commitment transaction has reached reasonable depth: - MUST fail the corresponding incoming HTLC (if any). - if no *valid* commitment transaction contains an output corresponding to the HTLC: - MAY fail the corresponding incoming HTLC sooner.

Rationale

If the commitment transaction belongs to the *remote* node, the only way for it to spend the HTLC output (using a payment preimage) is for it to use the HTLC-success transaction.

The payment preimage either serves to prove payment (when the offering node is the originator of the payment) or to redeem the corresponding incoming HTLC from another peer (when the offering node is forwarding the payment). After a node has extracted the payment, it no longer need be concerned with the fate of the HTLC-spending transaction itself.

In cases where both resolutions are possible (e.g. when a node receives payment success after timeout), either interpretation is acceptable: it's the responsibility of the recipient to spend it before this occurs.

Once it has timed out, the local node needs to spend the HTLC output (to prevent the remote node from using the HTLC-success transaction) before it can back-fail any corresponding incoming HTLC, using `update_fail_htlc` (presumably with reason `permanent_channel_failure`), as detailed in [BOLT #2](#). If the incoming HTLC is also on-chain, a node simply waits for it to timeout, as there's no way to signal early failure.

There are several reasons a committed HTLC may not have an output in the confirmed commitment transaction: the HTLC may be smaller than `dust_limit_satoshis`, the HTLC may not have been added to the commitment transaction yet, or the HTLC may have already been failed or fulfilled. In any case, if the payment preimage is known for the HTLC, the upstream HTLC needs to be fulfilled to avoid loss of funds.

If the payment preimage is not known for the missing HTLC, the correct action depends on the possibility of a blockchain reorganization that swaps out the confirmed commitment transaction for one with the HTLC present. If the HTLC is too small to appear in *any commitment transaction*, such a reorganization is not possible, and the HTLC can be safely failed immediately. Otherwise, a reorganization delay is required before failing the incoming HTLC. The requirement that the incoming HTLC be failed before its own timeout still applies as an upper bound.

HTLC Output Handling: Remote Commitment, Remote Offers

The remote HTLC outputs can only be spent by the local node if it has the payment preimage. If the local node does not have the preimage (and doesn't discover it), it's the remote node's responsibility to spend the HTLC output once it's timed out.

There are actually several possible cases for an offered HTLC:

1. The offerer is not irrevocably committed to it. In this case, the recipient usually won't know the preimage, since it won't forward HTLCs until they're fully committed. As using the preimage would reveal that this recipient is the final hop, it's best to allow the HTLC to time out.
2. The offerer is irrevocably committed to the offered HTLC, but the recipient hasn't yet committed to an outgoing HTLC. In this case, the recipient can either forward it or wait for it to timeout.
3. The recipient has committed to an outgoing HTLC in exchange for an offered HTLC. In this case, the recipient must use the preimage, if it receives it from the outgoing HTLC; otherwise, it will lose funds by sending an outgoing payment without redeeming the incoming one.

Requirements

A local node: - if it receives (or already possesses) a payment preimage for an unresolved HTLC output that it was offered AND for which it has committed to an outgoing HTLC: - MUST *resolve* the output by spending it to a convenient address. - otherwise: - if the remote node is NOT irrevocably committed to the HTLC: - MUST NOT *resolve* the output by spending it.

If not otherwise resolved, once the HTLC output has expired, it is considered *irrevocably resolved*.

Revoked Transaction Close Handling

If any node tries to cheat by broadcasting an outdated commitment transaction (any previous commitment transaction besides the most current one), the other node in the channel can use its revocation private key to claim all the funds from the channel's original funding transaction.

Requirements

Once a node discovers a commitment transaction for which *it* has a revocation private key, the funding transaction output is *resolved*.

A local node: - MUST NOT broadcast a commitment transaction for which *it* has exposed the `per_commitment_secret`. - MAY take no action regarding the *local node's main output*, as this is a simple P2WPKH output to itself. - Note: this output is considered *resolved* by the commitment transaction itself. - MUST *resolve* the *remote node's main output* by spending it using the revocation private key. - MUST *resolve* the *remote node's offered HTLCs* in one of three ways: * spend the *commitment tx* using the payment revocation private key. * spend the *commitment tx* using the payment preimage (if known). * spend the *HTLC-timeout tx*, if the remote node has published it. - MUST *resolve* the *local node's offered HTLCs* in one of three ways: * spend the *commitment tx* using the payment revocation private key. * spend the *commitment tx* once the HTLC timeout has passed. * spend the *HTLC-success tx*, if the remote node has published it. - MUST *resolve* the *remote node's HTLC-timeout transaction* by spending it using the revocation private key. - MUST *resolve* the *remote node's HTLC-success transaction* by spending it using the revocation private key. - SHOULD extract the payment preimage from the transaction input witness, if it's not already known. - if `SIGHASH_SINGLE|SIGHASH_ANYONECANPAY` is used for HTLC transactions (i.e. `option_anchors` or `zero_fee_commitments` applies): - MAY use a single transaction to *resolve* all the outputs. - If `zero_fee_commitments` applies: - MUST NOT exceed the 10kB limit of v3 transactions, otherwise the transaction may not reach miners (as explained in the section about [v3 limitations](#) below). - if confirmation doesn't happen before reaching `security_delay` blocks from expiry: - SHOULD *resolve* revoked outputs in their own, separate penalty transactions. A previous penalty transaction

claiming multiple revoked outputs at once may be blocked from confirming because of a transaction pinning attack. - otherwise: - MAY use a single transaction to *resolve* all the outputs. - MUST handle its transactions being invalidated by HTLC transactions.

Rationale

A single transaction that resolves all the outputs will be under the standard size limit because of the 483 HTLC-per-party limit (see [BOLT #2](#)).

Note: when `SIGHASH_SINGLE|SIGHASH_ANYONECANPAY` is used for HTLC transactions (i.e. `option_anchors` or `zero_fee_commitments`), the cheating node can pin spends of its HTLC-timeout/HTLC-success outputs. Using a single penalty transaction for all revoked outputs is thus unsafe as it could be blocked to propagate long enough for the *local node's to_local output's* relative locktime to expire and the cheating party escaping the penalty on this output. Though this situation doesn't prevent faithful punishment of the second-level revoked output if the pinning transaction confirms.

The `security_delay` is a fixed-point relative to the absolute expiration of the revoked output at which the punishing node must broadcast a single-spend transaction for the revoked output and actively fee-bump it until its confirmation. The exact value of `security_delay` is left as a matter of node policy, though we recommend 18 blocks (similar to incoming HTLC deadline).

Penalty Transactions Weight Calculation

There are three different scripts for penalty transactions, with the following witness weights (details of weight computation are in [Appendix A](#)):

```
to_local_penalty_witness: 160 bytes
offered_htlc_penalty_witness: 243 bytes
accepted_htlc_penalty_witness: 249 bytes
```

The penalty *txinput* itself takes up 41 bytes and has a weight of 164 bytes, which results in the following weights for each input:

```
to_local_penalty_input_weight: 324 bytes
offered_htlc_penalty_input_weight: 407 bytes
accepted_htlc_penalty_input_weight: 413 bytes
```

The rest of the penalty transaction takes up $4+1+1+8+1+34+4=53$ bytes of non-witness data: assuming it has a pay-to-witness-script-hash (the largest standard output script), in addition to a 2-byte witness header.

In addition to spending these outputs, a penalty transaction may optionally spend the commitment transaction's `to_remote` output (e.g. to reduce the total amount paid in fees). Doing so requires the inclusion of a P2WPKH witness and an additional *txinput*, resulting in an additional $108 + 164 = 272$ bytes.

In the worst case scenario, the node holds only incoming HTLCs, and the HTLC-timeout transactions are not published, which forces the node to spend from the commitment transaction.

With a maximum standard weight of 400000 bytes, the maximum number of HTLCs that can be swept in a single transaction is as follows:

```
max_num_htlcs = (400000 - 324 - 272 - (4 * 53) - 2) / 413 = 966
```

Thus, 483 bidirectional HTLCs (containing both `to_local` and `to_remote` outputs) can be resolved in a single penalty transaction. Note: even if the `to_remote` output is not swept, the resulting `max_num_htlcs` is 967; which yields the same unidirectional limit of 483 HTLCs.

Generation of HTLC Transactions

If `option_anchors` and `zero_fee_commitments` do not apply, then HTLC-timeout and HTLC-success transactions are signed with `SIGHASH_ALL` and are complete transactions with (hopefully!) reasonable fees and must be used directly.

Otherwise, `SIGHASH_SINGLE|SIGHASH_ANYONECANPAY` MUST be used on the HTLC signatures received from the peer, as this allows HTLC transactions to be combined with other transactions. The local signature MUST use `SIGHASH_ALL`, otherwise anyone can attach additional inputs and outputs to the tx.

If `SIGHASH_SINGLE|SIGHASH_ANYONECANPAY` is used, then the HTLC-timeout and HTLC-success transactions are signed with the input and output having the same value. This means they have a zero fee and MUST be combined with other inputs to arrive at a reasonable fee.

Requirements

A node which broadcasts an HTLC-success or HTLC-timeout transaction for a commitment transaction: - if `SIGHASH_SINGLE|SIGHASH_ANYONECANPAY` is used: - MUST combine it with inputs contributing sufficient fee to ensure timely inclusion in a block. - MAY combine it with other transactions.

General Requirements

A node: - upon discovering a transaction that spends a funding transaction output which does not fall into one of the above categories (mutual close, unilateral close, or revoked transaction close): - MUST warn the user of potentially lost funds. - Note: the existence of such a rogue transaction implies that its private key has leaked and that its funds may be lost as a result. - MAY simply monitor the contents of the most-work chain for transactions. - Note: on-chain HTLCs should be sufficiently rare that speed need not be considered critical. - MAY monitor (valid) broadcast transactions (a.k.a the mempool). - Note: watching for mempool transactions should result in lower latency HTLC redemptions.

Limitations of v3 transactions

When `zero_fee_commitments` applies, we use v3 transactions, also known as TRUC transactions, which have different policy rules than v2 transactions.

Implementers should read [BIP-431](#) to make sure they don't create transactions that aren't relayed by bitcoin nodes.

The main limitation compared to v2 transactions (but not the only one) is that transactions are restricted to 10kB (instead of 400kB): when batching transactions (e.g. HTLC transactions), nodes MUST NOT exceed that limit, otherwise the transaction will not reach miners and funds may be at risk.

Appendix A: Expected Weights

Expected Weight of the to_local Penalty Transaction Witness

As described in [BOLT #3](#), the witness for this transaction is:

```
<sig> 1 { OP_IF <revocationpubkey> OP_ELSE to_self_delay OP_CSV OP_DROP <local_delayedpubkey>
OP_ENDIF OP_CHECKSIG }
```

The *expected weight* of the to_local penalty transaction witness is calculated as follows:

```
to_local_script: 83 bytes
- OP_IF: 1 byte
  - OP_DATA: 1 byte (revocationpubkey length)
  - revocationpubkey: 33 bytes
- OP_ELSE: 1 byte
  - OP_DATA: 1 byte (delay length)
  - delay: 8 bytes
  - OP_CHECKSEQUENCEVERIFY: 1 byte
  - OP_DROP: 1 byte
  - OP_DATA: 1 byte (local_delayedpubkey length)
  - local_delayedpubkey: 33 bytes
- OP_ENDIF: 1 byte
- OP_CHECKSIG: 1 byte
```

```
to_local_penalty_witness: 160 bytes
- number_of_witness_elements: 1 byte
- revocation_sig_length: 1 byte
- revocation_sig: 73 bytes
- one_length: 1 byte
- witness_script_length: 1 byte
- witness_script (to_local_script)
```

Expected Weight of the offered_htlc Penalty Transaction Witness

The *expected weight* of the offered_htlc penalty transaction witness is calculated as follows (some calculations have already been made in [BOLT #3](#)):

```
offered_htlc_script: 133 bytes
```

```
offered_htlc_penalty_witness: 243 bytes
- number_of_witness_elements: 1 byte
- revocation_sig_length: 1 byte
- revocation_sig: 73 bytes
- revocation_key_length: 1 byte
- revocation_key: 33 bytes
- witness_script_length: 1 byte
- witness_script (offered_htlc_script)
```

Expected Weight of the accepted_htlc Penalty Transaction Witness

The *expected weight* of the accepted_htlc penalty transaction witness is calculated as follows (some calculations have already been made in [BOLT #3](#)):

accepted_htlc_script: 139 bytes

accepted_htlc_penalty_witness: 249 bytes

- number_of_witness_elements: 1 byte
- revocation_sig_length: 1 byte
- revocation_sig: 73 bytes
- revocationpubkey_length: 1 byte
- revocationpubkey: 33 bytes
- witness_script_length: 1 byte
- witness_script (accepted_htlc_script)

Authors

[FIXME:]



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Finding the Network

Gossip tells you about nodes you’ve already heard of. It doesn’t tell you how to find the first one.

BOLT 10 is DNS bootstrap. Lightning nodes can publish SRV and node records under a seed domain so a new peer can resolve a handful of entry points, connect, and start receiving gossip. It’s a small document, and it’s optional in spirit — you can always paste a peer address by hand — but it’s how a node joins the network without a phone book.

If you’re running a seed, this BOLT is the contract you have with the rest of the network. If you’re writing a client, it’s the one-time lookup before BOLT 07 takes over.

In this chapter

- BOLT 10 — DNS bootstrap

BOLT #10: DNS Bootstrap and Assisted Node Location

Overview

This specification describes a node discovery mechanism based on the Domain Name System (DNS). Its purpose is twofold:

- Bootstrap: providing the initial node discovery for nodes that have no known contacts in the network
- Assisted Node Location: supporting nodes in discovery of the current network address of previously known peers

A domain name server implementing this specification is referred to as a *DNS Seed* and answers incoming DNS queries of type A, AAAA, or SRV, as specified in RFCs 1035¹, 3596², and 2782³, respectively. The DNS server is authoritative for a subdomain, referred to as a *seed root domain*, and clients may query it for subdomains.

The subdomains consist of a number of dot-separated *conditions* that further narrow the desired results.

Table of Contents

- [DNS Seed Queries](#)
 - [Query Semantics](#)
- [Reply Construction](#)
- [Policies](#)
- [Examples](#)
- [References](#)
- [Authors](#)

DNS Seed Queries

A client MAY issue queries using the A, AAAA, or SRV query types, specifying conditions for the desired results the seed should return.

Queries distinguish between *wildcard* queries and *node* queries, depending on whether the l-key is set or not.

Query Semantics

The conditions are key-value pairs: the key is a single-letter, while the remainder of the key-value pair is the value. The following key-value pairs MUST be supported by a DNS seed:

- r: realm byte
 - used to specify what realm the returned nodes must support
 - default value: 0 (Bitcoin)
- a: address types
 - a bitfield that uses the types from [BOLT #7](#) as bit index
 - used to specify what address types should be returned for SRV queries
 - MAY only be used for SRV queries
 - default value: 6 (i.e. 2 || 4, since bit 1 and bit 2 are set for IPv4 and IPv6, respectively)
- l: node_id
 - a bech32-encoded node_id of a specific node
 - used to ask for a single node instead of a random selection
 - default value: null
- n: number of desired reply records
 - default value: 25

Conditions are passed in the DNS seed query as individual, dot-separated subdomain components.

For example, a query for `r0.a2.n10.lseed.bitcoinstats.com` would imply: return 10 (n10) IPv4 (a2) records for nodes supporting Bitcoin (r0).

Requirements

The DNS seed: - MUST evaluate the conditions from the *seed root domain* by 'going up-the-tree', i.e. evaluating right-to-left in a fully qualified domain name. - E.g. to evaluate the above case: first evaluate n10, then a2, and finally r0. - if a condition (key) is specified more than once: - MUST discard any earlier value for that condition

AND use the new value instead. - E.g. for `n5.r0.a2.n10.lseed.bitcoinstats.com`, the result is: ~~n10~~, a2, r0, n5. - SHOULD return results that match all conditions. - if it does NOT implement filtering by a given condition: - MAY ignore the condition altogether (i.e. the seed filtering is best effort only). - for A and AAAA queries: - MUST return only nodes listening on the default port 9735, as defined in [BOLT #1](#). - for SRV queries: - MAY return nodes that are listening on non-default ports, since SRV records return a *(hostname,port)*-tuple. - upon receiving a *wildcard* query: - MUST select a random subset of up to n IPv4 or IPv6 addresses of nodes that are listening for incoming connections. - upon receiving a *node* query: - MUST select the record matching the `node_id`, if any, AND return all addresses associated with that node.

Querying clients: - MUST NOT rely on any given condition being met by the results.

Reply Construction

The results are serialized in a reply with a query type matching the client's query type. For example, A, AAAA, and SRV queries respectively result in A, AAAA, and SRV replies. Additionally, replies may be augmented with additional records (e.g. to add A or AAAA records matching the returned SRV records).

For A and AAAA queries, the reply contains the domain name and the IP address of the results.

The domain name MUST match the domain in the query, in order not to be filtered by intermediate resolvers.

For SRV queries, the reply consists of *(virtual hostnames, port)*-tuples. A virtual hostname is a subdomain of the seed root domain that uniquely identifies a node in the network. It is constructed by prepending the `node_id` condition to the seed root domain.

The DNS seed: - MAY additionally return the corresponding A and AAAA records that indicate the IP address for the SRV entries in the additional section of the reply. - MAY omit these additional records upon detecting a repeated query. - Reason: due to the large size of the resulting reply, the reply may be dropped by intermediate resolvers. - if no entries match all the conditions: - MUST return an empty reply.

Policies

The DNS seed: - MUST NOT return replies with a TTL less than 60 seconds. - MAY filter nodes from its local views for various reasons, including faulty nodes, flaky nodes, or spam prevention. - MUST reply to random queries (i.e. queries to the seed root domain and to the `_nodes._tcp.` alias for SRV queries) with *random and unbiased* samples from the set of all known good nodes, in accordance with the Bitcoin DNS Seed policy^{[4](#)}.

Examples

Querying for AAAA records:

```
$ dig lseed.bitcoinstats.com AAAA
lseed.bitcoinstats.com. 60      IN      AAAA    2a02:aa16:1105:4a80:1234:1234:37c1:9c9
```

Querying for SRV records:

```
$ dig lseed.bitcoinstats.com SRV
lseed.bitcoinstats.com. 59      IN      SRV      10 10 6331
ln1qwktp6jxltmppyhl578eax6fcjc2m807qalr76a5gfm7k9qqfjwy4mctz.lseed.bitcoinstats.com.
lseed.bitcoinstats.com. 59      IN      SRV      10 10 9735
ln1qv2w3tledmzczw227nnkqrlltvmtydl8gu4w4d70g9td7avke6nmz2tdefqp.lseed.bitcoinstats.com.
```



```
lseed.bitcoinstats.com. 59 IN SRV 10 10 9735
ln1qtnyymv99pqf0r9cuexvvqtxrlgejuecf8myfsa96vcpflgll5cqmr2xsu.lseed.bitcoinstats.com.
lseed.bitcoinstats.com. 59 IN SRV 10 10 4280
ln1qdfvlysfpyh96apy3w3qdwl8jjkdhnu6a689ka540tnde6gnx86cf7ga2d.lseed.bitcoinstats.com.
lseed.bitcoinstats.com. 59 IN SRV 10 10 4281
ln1qwf789tclpe4n34649xrqlxt97whsvfk5pm07ggqms3vrjwdj3cu6332zs.lseed.bitcoinstats.com.
```

Querying for the A for the first virtual hostname from the previous example:

```
$ dig ln1qwktp6jxltmpphyl578eax6fcjc2m807qalr76a5gfm7k9qqfjwy4mctz.lseed.bitcoinstats.com A
ln1qwktp6jxltmpphyl578eax6fcjc2m807qalr76a5gfm7k9qqfjwy4mctz.lseed.bitcoinstats.com. 60 IN A
139.59.143.87
```

Querying for only IPv4 nodes (a2) via seed filtering:

```
$dig a2.lseed.bitcoinstats.com SRV
a2.lseed.bitcoinstats.com. 59 IN SRV 10 10 9735
ln1q2jy22cg2nckgxttj8txmamwe9rtw325v4m04ug2dm9sxlrh9cagrrpy86.lseed.bitcoinstats.com.
a2.lseed.bitcoinstats.com. 59 IN SRV 10 10 9735
ln1qfrkq32xayuq63anmc2zp5vtd2jxafhdzzudmuws0hvxshsgd2zd7jsqv7f.lseed.bitcoinstats.com.
```

Querying for only IPv6 nodes (a4) supporting Bitcoin (r0) via seed filtering:

```
$dig r0.a4.lseed.bitcoinstats.com SRV
r0.a4.lseed.bitcoinstats.com. 59 IN SRV 10 10 9735
ln1qwx3prnvmxuwnaqhzwrrpwy4pjf5m8fv4m8kcjkdvyryzmlcmj5dakwrx.lseed.bitcoinstats.com.
r0.a4.lseed.bitcoinstats.com. 59 IN SRV 10 10 9735
ln1qwr7x7q2gvj7kwzrz7urqq9x7mq0lf9xn6svs8dn7q8gu5q4e852znqj3j7.lseed.bitcoinstats.com.
```

References

- [RFC 1035 - Domain Names](#)
- [RFC 3596 - DNS Extensions to Support IP Version 6](#)
- [RFC 2782 - A DNS RR for specifying the location of services \(DNS SRV\)](#)
- [Expectations for DNS Seed operators](#)

Authors

[FIXME: Insert Author List]



This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

Newer Work

Lightning's original channels use a 2-of-2 OP_CHECKMULTISIG funding output. It works. It also looks like a Lightning channel on-chain, and it predates Taproot.

Simple Taproot Channels move funding into a Taproot output. Cooperative closes look like a single-key spend. The commitment format changes, musig2 shows up, and a few BOLT 02 / BOLT 03 messages grow extra fields. The rest of the protocol — invoices, onions, gossip — stays the same.

This document is newer than the numbered BOLTs, and implementations are still catching up. Read it after you're comfortable with BOLT 02 and BOLT 03, because it's a variation on those two, not a replacement for the rest of the book.

Treat it as a look at where channel construction is heading, not as required reading for a first implementation.

In this chapter

- Simple Taproot Channels

Extension BOLT XX: Simple Taproot Channels

Authors: * Olaoluwa Osuntokun roasbeef@lightning.engineering * Eugene Siegel eugene@lightning.engineering

Created: 2022-04-20

Table of Contents

- [Introduction](#)
 - [Abstract](#)
 - [Motivation](#)
 - [Preliminaries](#)
 - [Pay-To-Taproot-Outputs](#)
 - [Tapscript Tree Semantics](#)
 - [BIP 86](#)
 - [Taproot Key Path Spends](#)
 - [Taproot Script Path Spends](#)
 - [MuSig2 & Associated Changes](#)
 - [Key Aggregation](#)
 - [Nonce Generation](#)
 - [Nonce Handling](#)
 - [Signing](#)
 - [Nothing Up My Sleeves Points](#)
 - [Design Overview](#)
 - [Specification](#)
 - [Feature Bits](#)
 - [New TLV Types](#)
 - [Channel Funding](#)
 - [open_channel Extensions](#)
 - [accept_channel Extensions](#)
 - [funding_created Extensions](#)
 - [funding_signed Extensions](#)
 - [channel_ready Extensions](#)
 - [RBF Cooperative Close](#)
 - [shutdown Extensions](#)
 - [JIT \(Just-In-Time\) Nonce Pattern](#)
 - [closing_complete Extensions](#)

- [closing_sig Extensions](#)
- [RBF Nonce Rotation Protocol](#)
- [Channel Operation](#)
 - [commitment_signed Extensions](#)
 - [revoke_and_ack Extensions](#)
- [Message Retransmission](#)
 - [channel_reestablish Extensions](#)
- [Funding Transactions](#)
- [Commitment Transactions](#)
 - [To Local Outputs](#)
 - [To Remote Outputs](#)
 - [Anchor Outputs](#)
- [HTLC Scripts & Transactions](#)
 - [Offered HTLCs](#)
 - [Accepted HTLCs](#)
 - [HTLC Second Level Transactions](#)
 - [HTLC-Success Transactions](#)
 - [HTLC-Timeout Transactions](#)
- [Appendix](#)
- [Footnotes](#)
- [Test Vectors](#)
- [Acknowledgements](#)

Introduction

Abstract

The activation of the Taproot soft-fork suite enables a number of updates to the Lightning Network, allowing developers to improve the privacy, security, and flexibility of the system. This document specifies extensions to BOLTs 2, 3, and 5 which describe new, Taproot based channels to update to the Lightning Network to take advantage of *some* of these new capabilities. Namely, we mechanically translate the current funding and commitment design to utilize musig2 and the new tapscript tree capabilities.

Motivation

The activation of Taproot grants protocol developers with a number of new tools including: schnorr, musig2, scriptless scripts, PTLs, merkalized script trees and more. While it's technically possible to craft a *single* update to the Lightning Network to take advantage of *all* these new capabilities, we instead propose a step-wise update process, with each step layering on top of the prior with new capabilities. While the ultimate realization of a more advanced protocol may be delayed as a result of this approach, packing up smaller incremental updates will be easier to implement and review, and may also hasten the timeline to deploy initial taproot based channels.

In this document, we start by revising the most fundamental component of a Lightning Channel: the funding output. By first porting the funding output to Segwit V1 (P2TR) musig2 based output, we're able to re-anchor all channels in the network with the help of dynamic commitments. Once those channels are re-anchored,

dynamic commitments allows developers to ship incremental changes to the commitment transaction, HTLC structure, and overall channel structure – all without additional on-chain transactions.

By restricting the surface area of the *initial* update, we aim to expedite the initial roll-out while also providing a solid base for future updates without blocking off any interesting design paths. Implementing and shipping a constrained update also gives implementers an opportunity to get more familiar with the new tooling and capabilities, before tackling more advanced protocol designs.

Preliminaries

In this section we lay out some preliminary concepts, protocol flows, and notation that'll be used later in the core channel specification.

Pay-To-Taproot-Outputs

The Taproot soft fork package (BIPs [340](#), [341](#), and [342](#)) introduced a new [Segwit](#) witness version: v1. This new witness version utilizes a new standardized public key script template:

```
OP_1 <taproot_output_key>
```

The `taproot_output_key` is a 32-byte x-only secp256k1 public key, with all digital signatures based off of the schnorr signature scheme described in [BIP 340](#).

A `taproot_output_key` commits to an internal key, and optional script root via the following mapping:

```
taproot_output_key = taproot_internal_key + tagged_hash("TapTweak", taproot_internal_key || script_root)*G
```

It's important to note that while `taproot_output_key` is serialized as a 32-byte public key, in order to properly spend the script path, the parity (even or odd) of the output public key must be remembered.

The `taproot_internal_key` is also a BIP 340 public key, ideally derived anew for each output. The `tagged_hash` scheme is described in BIP 340, we briefly define the function as:

```
tagged_hash(tag, msg) = SHA256(SHA256(tag) || SHA256(tag) || msg)
```

Tapscript Tree Semantics

The `script_root` is the root of a Tapscript tree as defined in BIP 341. A tree is composed of `tap_leaves` and `tap_branches`.

A `tap_leaf` is a two tuple of (`leaf_version`, `leaf_script`). A `tap_leaf` is serialized as:

```
leaf_version || compact_size_of(leaf_script) || leaf_script
```

where `compact_size_of` is the variable length integer used in the Bitcoin P2P protocol.

The digest of a `tap_leaf` is computed as:

```
tagged_hash("TapLeaf", leaf_encoding)
```

where `leaf_encoding` is serialized as:

```
leaf_version || compact_size_of(leaf_script) || leaf_script
```

A tap_branch can commit to either a tap_leaf or tap_branch. Before hashing, a tap_branch sorts the two tap_node arguments based on lexicographical ordering:

```
tagged_hash("TapBranch", sort(node_1, node_2))
```

A tapscript tree is constructed by hashing each pair of leaves into a tap_branch, then combining these branches with each other until only a single hash remains. There're no strict requirements at the consensus level for tree construction, allowing systems to "weight" the tree by placing items that occur more often at the top of the tree. A simple algorithm that attempts to [produce a balanced script tree can be found here](#).

BIP 86

A taproot_output_key doesn't necessarily *need* to commit to a script_root, unless the script_root is being revealed, a normal key (skipping the tweaking operation) can be used. In many cases it's also nice to prove to a 3rd party that the output can only be spent with a top-level signature and not also a script path.

[BIP 86](#) defines a taproot output key derivation scheme, wherein only the internal key is committed to without a script root:

```
taproot_output_key = taproot_internal_key + tagged_hash("TapTweak", taproot_internal_key)*G
```

We use BIP 86 whenever we want to ensure that a script path spend isn't possible, and an output can only be spent via a key path.

We recommend that in any instances where a normal unencumbered output is used (cooperative closures, etc) BIP 86 is used as well.

Taproot Key Path Spends

A taproot key path spend a transaction to spend a taproot output without revealing the script root, which may or may not exist.

The witness of a keypath spend is simply a 64 or 65 byte schnorr signature. A signature is 64 bytes if the new SIGHASH_DEFAULT alias for SIGHASH_ALL is being used. Otherwise a single byte for the sighash type is used as normal:

```
<key_path_sig>
```

If an output commits to a script root (or uses BIP 86), a private key will need to be tweaked in order to generate a valid signature, and negated if the resulting output key has an odd y coordinate (see BIP 341 for details).

For any revocation path based on a key spend (the HTLC revocation paths), the tap tweak value **MUST** be known in order to spend the out unilaterally. As a result, it's recommend that for a given revoked state, the tap tweak value is also stored.

Taproot Script Path Spends

A script path spend is executed when a tapscript leaf is revealed. A script path spends has three components: a control block, the script leaf being revealed, and the valid witness.

A control block has the following serialization:

```
(output_key_y_parity | leaf_version) || internal_key || inclusion_proof
```

The first segment uses a spare bit of the `leaf_version` to encode the parity of the output key, which is checked during control block verification. The `internal_key` is just that. The `inclusion_proof` is the series of sibling hashes in the path from the revealed leaf all the way up to the root of the tree. In practice, if one hashes the revealed leaf, and each 32-byte hash of the inclusion proof together, they'll reach the script root if the proof was valid.

If only a single leaf is committed to in a tree, then the `inclusion_proof` will be absent.

The final witness stack to spend a script path output is:

```
<witness1> ... <witnessN> <leaf_script> <control_block>
```

Note that whenever a script path spend is required, in order to ensure an output is unilaterally spendable by a party, the finalized control block **MUST** either be stored as is, or the components needed to reconstruct the control block are persisted.

SIGHASH_ALL vs SIGHASH_DEFAULT

Where ever applicable, signers **SHOULD** use the `SIGHASH_DEFAULT` sighash over `SIGHASH_ALL`. When using taproot semantics, both algorithm refer to the very same digest, but `SIGHASH_DEFAULT` allows us to save a byte for any signature, as omitting the explicit sighash byte denotes that the default signing algorithm was used.

MuSig2 & Associated Changes

Musig2 is a multi-signature scheme based on schnorr signatures. It allows N parties to construct an aggregated public key, and then generate a single multi-signature that can only be generated with all the parties of the aggregated key. In practice, we use this to allow the funding output of channel to look just like any other P2TR output.

A musig2 signing flow has 4 stages:

1. First, each party generates a public nonce, and exchanges it with every other party.
2. Next, public keys are exchanged by each party, and combined into a single aggregated public key.
3. Next, each party generates a *partial* signature and exchanges it with every other party.
4. Finally, once all partial signatures are known, they can be combined into a valid schnorr Signature, which can be verified using BIP 340.

Steps 1&2 may be performed out of order, or concurrently. In our case only two parties exist, so as soon as one party knows both partial signatures, they can be combined into a final signature.

Thought this document musig2 refers to the finalized (v1.0.0) [BIP-0327](<https://github.com/bitcoin/bips/blob/master/bip-0327.mediawiki>) specification.

Key Aggregation

We refer to the algorithm of KeyAgg algorithm [bip-musig2 v1.0.0](#) for key aggregation. In order to avoid sending extra ordering information, we always assume that both keys are sorted first with the KeySort algorithm before aggregating ($\text{KeyAgg}(\text{KeySort}(p_1, p_2))$).

Nonce Generation

A musig2 secret nonce is the concatenation of two random, 32-byte integers.

A musig2 public nonce is technically the concatenation of two public keys, each representing the EC-point corresponding to its secret integer, thus resulting in a 66-byte value:

`point_1 || point_2`

For nonce generation, we refer to the NonceGen algorithm defined in the musig2 BIP. We only specify the set of required arguments, though a signer can also opt to specify the additional argument, in order to strengthen their nonce.

Counter Based Nonce Generation & JIT Nonce Strengthening

Each new commitment state requires fresh nonces from both parties, as each commitment state is signed using a fresh musig2 session. Nonces sent by either party in the `commitment_signed` message can be generated Just-In-Time (JIT), thus requiring no additional state. These nonces can also be further strengthened by mixing in the commitment sighash information into the nonce generation algorithm.

Upon channel creation, and subsequent re-establishment, a fresh set of verification nonces are exchanged by each party. This is the nonce the party will use to *verify* a new incoming commitment state sent by the remote party. In order to force close a channel, the holder of a verification nonce must use that same nonce to counter sign the commitment transaction with the other half of the musig2 partial signature. Rather than force an implementation to retain additional signing state (the verification nonce) to avoid holding a “hot” (fully broadcast able) commitment on disk, we instead recommend a counter based approach.

In this scenario are nonce generated via a counter is deemed to be safe as:

1. Each commitment state only has a single signature exchanged to counter sign the commitment transaction.
2. Each commitment state has a unique number which is currently encoded as a 48-bit integer across the sequence and lock time of the commitment transaction.
3. The shachain scheme is used today to generate a fresh nonce-like value for the revocation scheme of today’s penalty based channels.
4. A verification nonce is only used to *sign* a local party’s commitment transaction when they go to broadcast a force close transaction.
5. The remote party never sees the unaggregated partial signature the local party will use to broadcast, only the aggregated signature is seen on chain.

6. As a result of all the above, a verification nonce will never be used to sign the exact same commitment transaction, thereby avoiding nonce reuse.

As a result, we recommend the following counter based scheme to generate verification nonce. With this scheme, signing can remain stateless, as given the commitment height/ number of the state to sign, and the original shachain root, the verification nonce sent by that state can be deterministically reproduced:

1. Given the shachain root used to generate revocation pre-images, `shachain_root`, derive the sha256 hash of this value as: `shachain_root_hash = sha256(shachain_root)`.
2. Derive a *new* shachain root to be used to generate musig2 secret nonces via a HMAC invocation as: `musig2_shachain_root = hmac(msg, shachain_root_hash)`, where `msg` is any string that can serve to uniquely bind the produced secret to this dedicated context. A recommended value is the ASCII string `taproot-rev-root` concatenated with the `funding_txid` (which allows deriving distinct deterministic shachains when splicing).
3. Given a commitment height/ number (`N`), the verification nonce to send to the remote party can be derived by obtaining the `N`th shachain leaf preimage `k_i`. The verification nonce to be derived by calling the `musig2.NonceGen` algorithm with the required values, and the `rand'` value set to `k_i`.

Nonce Handling

For commitment transactions, Due to the asymmetric state of the current revocation channels, two sets of public nonces need to be exchanged by each party: one for the local commitment, and one for the remote commitment.

The nonce for the local commitment **MUST** be sent a priori, thus enabling the counterparty to create a valid partial signature.

The nonce for the remote commitment *could* be sent a priori, too, but for the purposes of simplicity and disambiguation, **MUST** be sent alongside the partial signature. To that end, we will create a new data type that will be reused in multiple TLV extensions described further down in this document.

The exception to this is the co-op close flow: as there's only a single message to sign, we only require a single pair of nonces. These nonces can also be sent right as the co-op flow begins.

Note that if nonces for commitment updates are not generated using the counter based scheme described above, then the nonces **MUST** be persisted to disk. As a result, in order to minimize additional signing state, it is recommended that the shachain counter nonce scheme is used.

Nonce Terminology

At all times, there exist *two* nonces associated with a given commitment state: the signing nonce, and the verification nonce.

A peer's signing nonce is used to generate new commitments for the remote party. Each time a party signs a commitment, they send another signing nonce to the remote party.

A peer's verification nonce is used to verify incoming commitments sent by the remote party. Once a peer verifies and revoke a new commitment, another verification nonce is sent to the remote party.

Both sides will maintain a `musig2` sessions for their local commitment and the commitment of the remote party. A peer's local commitment uses their verification nonce, and the signing nonce of the remote party. The remote commitment of the channel peer uses the peer's signing nonce, and their verification nonce.

In order to sign and broadcast a new commitment to force close a channel, a peer uses their *verification* nonce to generate their final signature.

Signing

Once public nonces have been exchanged, the `NonceAgg` algorithm is used to combine the public nonces into the aggregate nonce which will be a part of the final signature.

The `Sign` method is used to generate a valid partial signature, with the `PartialSigVerify` algorithm used to verify each signature.

Once all partial signatures are known, the `PartialSigAgg` algorithm is used to aggregate the partial signature into a final, valid [BIP 340](#) Schnorr signature.

Partial Signature Encoding

A standard `musig2` partial signature is simply the encoding of the `s` value, which is a 32-byte-element modulo the order of the group.

Throughout this specification, we will be using a custom type, `PartialSignatureWithNonce`, comprised of the aforementioned `s` value, along with the 66-byte-encoding of the public nonce used to sign remote commitments:

```
s || point_1 || point_2
```

Nothing Up My Sleeves Points

Whenever we want to ensure that a given P2TR can *only* be spent via a script path, we utilize a Nothing Up My Sleeve (NUMS) point. A NUMS point is an EC point that no one knows the private key to. If no one knows the private key, then it can't be used for key path signing, forcing the script path to always be taken.

We refer to the `simple_taproot_nums` as the following value:

```
02dca094751109d0bd055d03565874e8276dd53e926b44e3bd1bb6bf4bc130a279
```

The value was [generated using this tool](#), with the seed phrase "Lightning Simple Taproot".

Design Overview

With the preliminaries out of the way, we provide a brief overview of the Simple Taproot Channel design.

The multisig output becomes a single P2TR key, with `musig2` key aggregation used to combine two keys into one.

For commitment transactions, we inherit the anchor output semantics, meaning there are two outputs used for CFPF, with all other scripts inheriting a 1 CSV restriction.

The local output of the commitment transaction uses a script-path based revocation scheme in order to ensure that the information needed by 3rd parties to sweep the anchor outputs is always revealed on chain.

The remote output of the commitment transaction uses the `combined_funding_key` as the top-level internal key, and then commits to a normal remote script.

Anchor outputs use the `local_delayedpubkey` and the `remotepubkey` of both parties as the top-level internal key committing to the script `16 CSV`. Unless active HTLCs exist on the commitment transaction, if a local or remote output is missing from the commitment, its corresponding anchor must not be present, either.

All HTLC scripts use the revocation key as the top-level internal key. This allows the revocation to be executed without additional on-chain bytes, and also reduces the amount of state that nodes need to keep in order to properly handle channel breaches.

Specification

Feature Bits

Inheriting the structure put forth in BOLT 9, we define a new feature bit to be placed in the `init` message, and the `node_announcement` message.

Bits	Name	Description	Context	Dependencies	Link
80/81	<code>option_simple_taproot</code>	Node supports simple taproot channels	IN	<code>option_channel_type</code> , <code>option_simple_close</code>	TODO(roasbeef): link
180/181	<code>option_simple_taproot_staging</code>	Node supports simple taproot channels	IN	<code>option_channel_type</code>	TODO(roasbeef): link

Note that we allocate *two* pairs of feature bits: one the final version of this protocol proposal, and the higher bits (+100) for preliminary experimental deployments before this protocol extension has been finalized. This +100 feature bit is thus refereed to as the staging feature bit.

The Context column decodes as follows: * I: presented in the `init` message. * N: presented in the `node_announcement` messages * C: presented in the `channel_announcement` message. * C-: presented in the `channel_announcement` message, but always odd (optional). * C+: presented in the `channel_announcement` message, but always even (required). * 9: presented in [BOLT 11](#) invoices.

The `option_simple_taproot` feature bit also becomes a defined channel type feature bit for explicit channel negotiation.

It's important to note that given the early version of this channel type *cannot* be announced on the public network (gossip protocol changes are required), the taproot channels type cannot be used as an interchangeable default channel type. As a result, this channel type SHOULD only be used with *explicit* channel

negotiation. Otherwise, two parties that advertise the feature bit would not be able to open publicly advertised channels.

Throughout this document, we assume that `option_simple_taproot` was negotiated, and also the `option_simple_taproot` channel type is used.

New TLV Types

Note that these TLV types exist across different messages, but their type IDs are always the same.

`partial_signature_with_nonce`

- type: 2
- data:
 - [32*byte: `partial_signature`]
 - [66*byte: `public_nonce`]

`next_local_nonce`

- type: 4
- data:
 - [66*byte: `public_nonce`]

`partial_signature`

- type: 6
- data:
 - [32*byte: `partial_signature`]

`shutdown_nonce`

- type: 8
- data:
 - [66*byte: `public_nonce`]

`next_local_nonces`


- type: 22
- data:
 - [...*`nonce_entry`: `entries`]

where `nonce_entry` is: * [32*byte: `funding_txid`] * [66*byte: `public_nonce`]

Channel Funding

`n_a_L`: Alice's local secret nonce

`n_B_R`: Bob's remote public nonce



The `open_channel` and `accept_channel` messages are extended to specify a new TLV type that houses the musig2 public nonces.

We add `option_simple_taproot` to the set of defined channel types.

open_channel Extensions

1. `tlv_stream`: `open_channel_tlvs`
2. `types`:
 1. `type`: 4 (`next_local_nonce`)
 2. `data`:
 - `[66*byte:public_nonce]`

Requirements

The sending node:

- MUST specify the `next_local_nonce` field.
- MUST use the NonceGen algorithm defined in `bip-musig2` to generate `next_local_nonce` to ensure it generates nonces in a safe manner.
- MUST not set the `announce_channel` bit.

The receiving node MUST fail the channel if:

- the message doesn't include a `next_local_nonce` value.
- the specified public nonce cannot be parsed as two compressed `secp256k1` points

Rationale

For the initial Taproot channel type (`option_simple_taproot`), musig2 nonces must be exchanged in the first round trip (`open->`, `<-accept`). This is done to ensure that the initiator is able to generate a valid musig2 partial signature at the next step once the transaction to be signed is fully specified.

The `next_local_nonce` field will be used to verify incoming commitment signatures. Each incoming signature will carry the newly generated nonce used to sign the commitment transaction. At force close broadcast time, the verification nonce is then used to sign the local party's commitment transaction.

accept_channel Extensions

1. `tlv_stream`: `accept_channel_tlvs`
2. `types`:
 1. `type`: 4 (`next_local_nonce`)
 2. `data`:
 - `[66*byte:public_nonce]`

Requirements

The sender:

- MUST set `next_local_nonce` to the `musig2` public nonce used to sign local commitments as specified by the `NonceGen` algorithm of `bip-musig2`.

The recipient:

- MUST reject the channel if `next_local_nonce` is absent, or cannot be parsed as two compressed `secp256k1` points

funding_created Extensions

1. `tlv_stream`: `funding_created_tlvs`
2. `types`:
 1. `type`: 2 (`partial_signature_with_nonce`)
 2. `data`:
 - `[98*byte: partial_signature || public_nonce]`

Requirements

The sender:

- MUST set the original, non-TLV signature field to a 0-byte-array of length 64.
- MUST sort the exchanged `funding_pubkeys` using the `KeySort` algorithm from `bip-musig2`.
- MUST compute the aggregated `musig2` public key from the sorted `funding_pubkeys` using the `KeyAgg` algorithm from `bip-musig2`.
- MUST generate a unique nonce to combine with the `next_local_nonce` previously received in `accept_channel` using `NonceAgg` from `bip-musig2`.
- MUST use the generated secret nonce and the calculated aggregate nonce to construct a `musig2` partial signature for the sender's remote commitment using the `Sign` algorithm from `bip-musig2`.
- MUST include the partial signature and the public counterpart of the generated nonce in the `partial_signature_with_nonce` field.

The recipient:

- MUST fail the channel if `partial_signature_with_nonce` is absent.
- MUST fail the channel if `next_local_nonce` is absent.
- MUST compute the aggregate nonce from:
 - the `next_local_nonce` field the recipient previously sent in `accept_channel`

- the `public_nonce` included as part of the `partial_signature_with_nonce` field
- MUST verify the `partial_signature_with_nonce` field using the `PartialSigVerifyInternal` algorithm of `bip-musig2`:
 - if the partial signature is invalid, MUST fail the channel

funding_signed Extensions

1. `tlv_stream`: `funding_signed_tlvs`
2. `types`:
 1. `type`: 2 (`partial_signature_with_nonce`)
 2. `data`:
 - `[98*byte: partial_signature || public_nonce]`

Requirements

The sender:

- MUST set the original, non-TLV signature field to a 0-byte-array of length 64.
- MUST sort the exchanged `funding_pubkeys` using the `KeySort` algorithm from `bip-musig2`.
- MUST compute the aggregated `musig2` public key from the sorted `funding_pubkeys` using the `KeyAgg` algorithm from `bip-musig2`.
- MUST generate a unique nonce to combine with the `next_local_nonce` previously received in `open_channel` using `NonceAgg` from `bip-musig2`.
- MUST use the generated secret nonce and the calculated aggregate nonce to construct a `musig2` partial signature for the sender's remote commitment using the `Sign` algorithm from `bip-musig2`.
- MUST include the partial signature and the public counterpart of the generated nonce in the `partial_signature_with_nonce` field.

The recipient:

- MUST fail the channel if `partial_signature_with_nonce` is absent.
- MUST compute the aggregate nonce from:
 - the `next_local_nonce` field the recipient previously sent in `open_channel`
 - the `public_nonce` included as part of the `partial_signature_with_nonce` field
- MUST verify the `partial_signature_with_nonce` field using the `PartialSigVerifyInternal` algorithm of `bip-musig2`:
 - if the partial signature is invalid, MUST fail the channel

channel_ready Extensions

We add a new TLV field to the channel_ready message:

1. tlv_stream: channel_ready_tlvs
2. types:
 1. type: 4 (next_local_nonce)
 2. data:
 - [66*byte: public_nonce]

Similar to the next_per_commitment_point, by sending the next_local_nonce value in this message, we ensure that the remote party has our public nonce which is required to generate a new commitment signature.

Requirements

The sender:

- MUST set next_local_nonce to a fresh, unique musig2 nonce as specified by bip-musig2

The recipient:

- MUST fail the channel if next_local_nonce is absent.

RBF Cooperative Close

For simple taproot channels, the cooperative close protocol uses closing_complete and closing_sig messages (as specified in BOLT #2) with extensions for MuSig2 signature generation and RBF (Replace-By-Fee) iterations.

shutdown Extensions

For taproot channels, the shutdown message includes a single nonce:

1. type: 38 (shutdown)
2. data:
 - ...
 - [shutdown_tlvs:tlvs]
3. tlv_stream: shutdown_tlvs
4. types:
 1. type: 8 (shutdown_nonce)
 2. data:
 - [66*byte:public_nonce]

The shutdown_nonce represents the sender's "close nonce" - the nonce they will use when sending closing_sig messages.

Additional nonces are provided just-in-time (JIT) with signatures:

- closing_complete uses PartialSigWithNonce which includes the sender's closer nonce

- `closing_sig` uses just `PartialSig` (no nonce) since the receiver already knows the nonce

Requirements

For taproot channels:

A sending node: - MUST set the `shutdown_nonce` to a valid `musig2` public nonce. - MUST use this nonce when sending `closing_sig` messages.

A receiving node: - MUST verify that the `shutdown_nonce` value is a valid `musig2` public nonce. - MUST store this nonce for use when verifying the peer's `closing_sig` messages.

JIT (Just-In-Time) Nonce Pattern

Final nonces are delivered just-in-time with signatures using an asymmetric pattern:

- `closing_complete` uses `PartialSigWithNonce` (98 bytes) to bundle the signature with the next closer nonce
- `closing_sig` uses `PartialSig` (32 bytes) with a separate `NextCloseeNonce` field

With this pattern, we exchange the set of `closee_nonces` up front in the shutdown message. When either party wants to initiate a new update, they send `closing_complete` which includes their JIT `closer_nonce`. Sending this signatures serves to consume the `closee_nonce` of the remote party, so then they respond with `closing_sig`, they include a new `next_closee_nonce`.

`closing_complete` Extensions

For taproot channels, the `closing_complete` message uses `PartialSigWithNonce` to support `MuSig2` signatures with JIT nonce delivery:

1. `tlv_stream`: `closing_tlvs`
2. `types`:
 1. type: 5 (`closer_no_closee`)
 2. data:
 - `[98*byte:partial_sig_with_nonce]`
 3. type: 6 (`no_closer_closee`)
 4. data:
 - `[98*byte:partial_sig_with_nonce]`
 5. type: 7 (`closer_and_closee`)
 6. data:
 - `[98*byte:partial_sig_with_nonce]`

Note: The TLV type numbers (5, 6, 7) are different from the non-taproot types (1, 2, 3) to distinguish taproot signatures. Each `partial_sig_with_nonce` contains:

- 32 bytes: `MuSig2` partial signature
- 66 bytes: The sender's closer nonce used to produce the partial signature

Requirements

The sender of `closing_complete` (aka. “the closer”):

- For taproot channels (when `option_simple_taproot` was negotiated):
 - MUST provide `PartialSigWithNonce` (98 bytes) in the appropriate TLV field based on output presence
 - MUST compute the aggregated musig2 public key from the sorted `funding_pubkeys` using the `KeyAgg` algorithm from `bip-musig2`
 - For the first closing transaction:
 - MUST use the `shutdown_nonce` received from the peer as their closee nonce
 - MUST use a locally generated nonce as the closer nonce
 - For RBF iterations:
 - MUST use the nonce extracted from the peer’s previous `closing_sig` message as their closee nonce
 - MUST use a fresh locally generated nonce as the closer nonce
 - MUST generate the partial signature using the `Sign` algorithm from `bip-musig2`

The receiver of `closing_complete` (aka. “the closee”):

- For taproot channels:
 - MUST verify that signature fields contain `PartialSigWithNonce` (98 bytes)
 - MUST extract the partial signature (first 32 bytes) and the sender’s closer nonce (remaining 66 bytes)
 - MUST compute the aggregate nonce by combining the closer nonce received with:
 - For the first closing transaction:
 - the `shutdown_nonce` the receiver previously sent (their closee nonce)
 - For RBF iterations:
 - the `next_closee_nonce` the receiver sent in their previous `closing_sig` message
 - MUST verify the partial signature using the `PartialSigVerifyInternal` algorithm of `bip-musig2`
 - MUST store (in memory) the extracted nonce for potential future RBF iterations

closing_sig Extensions

For taproot channels, the `closing_sig` message uses `PartialSig` with a separate next nonce field:

1. `tlv_stream`: `closing_tlvs`
2. types:
 1. type: 5 (`closer_no_closee`)
 2. data:
 - `[32*byte:partial_sig]`
 3. type: 6 (`no_closer_closee`)
 4. data:
 - `[32*byte:partial_sig]`
 5. type: 7 (`closer_and_closee`)
 6. data:
 - `[32*byte:partial_sig]`
 7. type: 22 (`next_closee_nonce`)

8. data:

- [66*byte:public_nonce]

Note: Unlike `closing_complete`, `closing_sig` separates the signature from the next nonce (which will be used for the *next* RBF attempt signature):

- The partial signature (32 bytes) is sent in the appropriate TLV field based on output presence
- The `next_closee_nonce` (66 bytes) is sent separately in TLV type 22
- The receiver (the closer) already knows the current closee signing nonce from either the shutdown message or the previous `next_closee_nonce`

Requirements

The sender of `closing_sig`:

- For taproot channels:
 - MUST provide a 32-byte partial signature in the appropriate TLV field based on output presence
 - For the first closing transaction:
 - MUST use the `shutdown_nonce` sent in their own shutdown message as their closee nonce
 - For RBF iterations:
 - MUST use the `next_closee_nonce` sent in their own previous `closing_sig` message as their closee nonce
 - MUST use the nonce provided in the `PartialSigWithNonce` message of the received `closing_complete` message as the closer nonce.
 - MUST generate the partial signature using the `Sign` algorithm from `bip-musig2`
 - MUST include `next_closee_nonce` for potential future RBF iterations
 - MUST generate `next_closee_nonce` using the `NonceGen` algorithm from `bip-musig2` if included

The receiver of `closing_sig`:

- For taproot channels:
 - MUST verify that signature fields contain 32-byte partial signatures
 - MUST use the appropriate nonces for verification:
 - The sender's closee nonce (from their shutdown message or previous `closing_sig`)
 - The receiver's own closer nonce included in `closing_complete`
 - MUST verify the partial signature using `PartialSigVerifyInternal` from `bip-musig2`
 - MUST combine both partial signatures using `PartialSigAgg` to create the final schnorr signature
 - MUST broadcast the closing transaction with the aggregated signature
 - SHOULD store `next_closee_nonce` for potential future RBF iterations

RBF Nonce Rotation Protocol

For taproot channels supporting RBF cooperative close, nonces are delivered just-in-time (JIT) with signatures using an asymmetric pattern:

1. **Initial Exchange:** Nonces are exchanged in shutdown messages
 - Each party sends their "closee nonce" - the nonce they'll use when signing as the closee
 - This nonce is used by the remote party when they act as closer

2. RBF Iterations: Subsequent nonces are delivered with signatures

- `closing_complete` (from closer): Uses `PartialSigWithNonce` (98 bytes)
 - Bundles the partial signature with the sender's closer nonce
 - The bundling is necessary because the closee doesn't know this nonce yet
- `closing_sig` (from closee): Uses `PartialSig` (32 bytes) + `NextCloseeNonce` field
 - Separates the signature from the next closee nonce because this nonce is for the *next* RBF attempt, not the one currently being signed

3. Asymmetric Roles:

- **Closer:** The party sending `closing_complete`, proposing a new fee
- **Closee:** The party sending `closing_sig`, accepting the fee proposal
- Each party alternates between these roles during RBF iterations

Nonce State Management

With this JIT nonce approach for coop close, an implementation only needs to store (in memory) the current closee nonce for the remote and local party. When either side is ready to sign, it'll generate a new closer nonce, and send that along with the sig.


Security Considerations

- Nonces MUST be generated using cryptographically secure randomness
- Previously used nonces MUST NOT be reused for different closing transactions

Channel Operation

`n_a_L`: Alice's local secret nonce

`n_B_R`: Bob's remote public nonce



commitment_signed Extensions

A new TLV stream is added to the commitment_signed message:

1. tlv_stream: commitment_signed_tlvs
2. types:
 1. type: 2 (partial_signature_with_nonce)
 2. data:
 - [98*byte: partial_signature || public_nonce]

Requirements

The sender:

- MUST set the original, non-TLV signature field to a 0-byte-array of length 64.
- MUST retrieve the musig2 public key previously aggregated from the sorted funding_pubkeys using the KeyAgg algorithm from bip-musig2.
- MUST generate a unique nonce to combine with the next_local_nonce previously received in the latest of channel_ready, channel_reestablish, or revoke_and_ack using NonceAgg from bip-musig2.
- MUST use the generated secret nonce and the calculated aggregate nonce to construct a musig2 partial signature for the sender's remote commitment using the Sign algorithm from bip-musig2.
- MUST include the partial signature and the public counterpart of the generated nonce in the partial_signature_with_nonce field.
- MUST compute each HTLC signature according to [BIP 342](#), as a BIP 340 Schnorr signature (non-partial).
 - Each HTLC signature must be packed as a 64-byte value within the existing htlc_signature field.

The recipient:

- MUST fail the channel if signature is non-zero.
- MUST fail the channel if partial_signature_with_nonce is absent.
- MUST compute the aggregate nonce from:
 - the next_local_nonce field the recipient previously sent in the latest of channel_ready, channel_reestablish, or revoke_and_ack
 - the public_nonce included as part of the partial_signature_with_nonce field
- MUST verify the partial_signature_with_nonce field using the PartialSigVerifyInternal algorithm of bip-musig2:
 - if the partial signature is invalid, MUST fail the channel

- MUST fail the channel if *any* of the received HTLC signatures is invalid according to BIP 342.

revoke_and_ack Extensions

A new TLV stream is added to the revoke_and_ack message:

1. tlv_stream: revoke_and_ack_tlvs
2. types:
 1. type: 22 (next_local_nonces)
 2. data:
 - [...*nonce_entry: nonces]

Similar to sending the next_per_commitment_point, we also send the *next* musig2 nonces, after we revoke a state. Sending these nonces allows the remote party to propose another state transition as soon as the message is received.

Requirements

The sender:

- MUST use the musig2.NonceGen algorithm to generate unique nonces to send in the next_local_nonces field.
- MUST include one entry for each active commitment transaction, indexed by the commitment transaction's funding_txid.

The recipient:

- MUST fail the channel if next_local_nonces is absent.
- MUST fail the channel if an active commitment's funding_txid is missing in next_local_nonces.
- If the local nonce generation is non-deterministic and the recipient co-signs commitments only upon pending broadcast:
 - MUST **securely** store the local nonce.

Message Retransmission

channel_reestablish Extensions

We add a new TLV field to the channel_reestablish message:

1. tlv_stream: channel_reestablish_tlvs
2. types:
 1. type: 22 (next_local_nonces)
 2. data:
 - [...*nonce_entry: nonces]

Similar to the next_per_commitment_point, by sending the next_local_nonces value in this message, we ensure that the remote party has our public nonces, which are required to generate new commitment signatures.

Requirements

The sender:

- MUST set `next_local_nonces` to fresh, unique musig2 nonces as specified by bip-musig2.
- MUST include one entry for each active commitment transaction, indexed by the commitment transaction's `funding_txid`.

The recipient:

- MUST fail the channel if `next_local_nonces` is absent, or cannot be parsed.
- MUST fail the channel if an active commitment's `funding_txid` is missing in `next_local_nonces`.

A node:

- If ALL of the following conditions apply:
 - It has previously sent a `commitment_signed` message
 - It never processed the corresponding `revoke_and_ack` message
 - It decides to retransmit the exact `update_` messages from the last sent `commitment_signed`
- THEN it must regenerate the partial signature using the newly received `next_local_nonces`

Splice Coordination

Splicing allows parties to modify the funding output of an existing channel without closing it. During splice operations, multiple commitment transactions may exist concurrently, each requiring its own MuSig2 nonce coordination. The `next_local_nonces` field enables this coordination by mapping transaction IDs to their respective nonces.

Splice Nonce Management

During splice negotiation:

- Each splice transaction MUST have a unique TXID as the key in the `next_local_nonces` map
- Parties MUST include nonces for all active commitment transactions in their `next_local_nonces` map, indexed by their `funding_txid`

Requirements for Splice Coordination

When a splice is initiated:

- The initiating party MUST generate a fresh nonce for the splice transaction
- Both parties MUST add the splice TXID and corresponding nonce to their `next_local_nonces` map
- The nonce MUST be communicated in the next `revoke_and_ack` or `channel_reestablish` message

When multiple splices are pending:

- Each splice MUST have a distinct TXID and nonce pair
- Nonces MUST NOT be reused across different splice transactions

- The `next_local_nonces` map MUST contain multiple entries during concurrent splice operations

Backward Compatibility

Nodes that do not support splicing will simply use a nonce map with a single entry indexed by the commitment transaction's `funding_txid`.

Funding Transactions

For our Simple Taproot Channels, `musig2` is used to generate a single aggregated public key.

The resulting funding output script takes the form: `OP_1 funding_key`

where:

- `funding_key = combined_funding_key + tagged_hash("TapTweak", combined_funding_key)*G`
- `combined_funding_key = musig2.KeyAgg(musig2.KeySort(pubkey1, pubkey2))`

The funding key is derived via the recommendation in BIP 341 that states: "If the spending conditions do not require a script path, the output key should commit to an unspendable script path instead of having no script path. This can be achieved by computing the output key point as $Q = P + \text{int}(\text{hashTapTweak}(\text{bytes}(P)))G$ ". We refer to BIP 86 for the specifics of this derivation.

By using the `musig2` sorting routine, we inherit the lexicographical key ordering from the base segwit channels. Throughout the lifetime of a channel any keys aggregated via `musig2` also are assumed to use the `KeySort` routine

Commitment Transactions

To Local Outputs

For the simple taproot commitment type, we use the same flow of revocation or delay, while re-using the NUMS point to ensure that the internal key is always revealed (script path always taken). This ensures that the anchor outputs can *always* be spent given chain revealed information. For the remote party, the `remotepubkey` will be revealed once the remote party sweeps. For the local party, we ensure that the `local_delayedpubkey` is revealed with an extra noop data push.

The new output has the following form:

- `OP_1 to_local_output_key`
- where:
 - `to_local_output_key = taproot_nums_point + tagged_hash("TapTweak", taproot_nums_point || to_delay_script_root)*G`
 - `to_delay_script_root = tapscript_root([to_delay_script, revoke_script])`
 - `to_delay_script` is the delay script: ``` OP_CHECKSIGVERIFY OP_CHECKSEQUENCEVERIFY`
 - `revoke_script` is the revoke script: `<local_delayedpubkey> OP_DROP <revocation_pubkey> OP_CHECKSIG`

The parity (even or odd) of the y-coordinate of the derived `to_local_output_key` MUST be known in order to derive a valid control block.

The `tapscript_root` routine constructs a valid taproot commitment according to BIP 341+342. Namely, a leaf version of `0xc0` MUST be used.

In the case of a commitment breach, the `to_delay_script_root` can be used along side `<revocationpubkey>` to derive the private key needed to sweep the top-level key spend path. The control block can be crafted as such:

```
revoke_control_block = (output_key_y_parity | 0xc0) || taproot_nums_point || revoke_script
```

A valid witness is then:

```
<revoke_sig> <revoke_script> <revoke_control_block>
```

In the case of a routine force close, the script path must be revealed so the broadcaster can sweep their funds after a delay. The control block to spend is only 33 bytes, as it just includes the internal key (along with the y-parity bit and leaf version):

```
delay_control_block = (output_key_y_parity | 0xc0) || taproot_nums_point || to_delay_script
```

A valid witness is then:

```
<local_delayed_sig> <to_delay_script> <delay_control_block>
```

As with base channels, the `nSequence` field must be set to `to_self_delay`.

To Remote Outputs

As we inherit the anchor output, semantics we want to ensure that the remote party can unilaterally sweep their funds after the 1 block CSV delay. In order to achieve this property, we'll utilize a NUMS point (`simple_taproot_nums`) By using this point as the internal key, we ensure that the remote party isn't able to by pass the CSV delay.

Using a NUMS key has a key benefit: the static internal key allows the remote party to scan for their output on chain, which is useful for various recovery scenarios.

The `to_remote` output has the following form:

- `OP_1 to_remote_output_key`
- where:
 - `taproot_nums_point = 0245b18183a06ee58228f07d9716f0f121cd194e4d924b037522503a7160432f15`
 - `to_remote_output_key = taproot_nums_point + tagged_hash("TapTweak", taproot_nums_point || to_remote_script_root)*G`
 - `to_remote_script_root = tapscript_root([to_remote_script])`
 - `to_remote_script` is the remote script: `<remotepubkey> OP_CHECKSIGVERIFY 1 OP_CHECKSEQUENCEVERIFY`

This output can be swept by the remote party with the following witness:

```
<remote_sig> <to_remote_script> <to_remote_control_block>
```

where `to_remote_control_block` is:

```
(output_key_y_parity | 0xc0) || combined_funding_key
```

The sequence field of the input MUST also be set to 1.

Anchor Outputs

For simple taproot channels (option_simple_taproot) we'll maintain the same form as base segwit channels, but instead will utilize local_delayedpubkey and the remotepubkey rather than the multi-sig keys as that's no longer revealed due to musig2.

An anchor output has the following form:

- OP_1 anchor_output_key
- where:
 - anchor_internal_key = remotepubkey/local_delayedpubkey
 - anchor_output_key = anchor_internal_key + tagged_hash("TapTweak", anchor_internal_key || anchor_script_root)
 - anchor_script_root = tapscript_root([anchor_script])
 - anchor_script: OP_16 OP_CHECKSEQUENCEVERIFY

This output can be swept by anyone after 16 blocks with the following witness:

<anchor_script> <anchor_control_block>

where anchor_control_block is:

(output_key_y_parity | 0xc0) || anchor_internal_key

nSequence needs to be set to 16.

HTLC Scripts & Transactions

Offered HTLCs

We retain the same second-level HTLC mappings as base channels. The main modifications are using the revocation public key as the internal key, and splitting the timeout and success paths into individual script leaves.

An offered HTLC has the following form:

- OP_1 offered_htlc_key
- where:
 - offered_htlc_key = revocation_pubkey + tagged_hash("TapTweak", revocation_pubkey || htlc_script_root)
 - htlc_script_root = tapscript_root([htlc_timeout, htlc_success])
 - htlc_timeout: <local_htlcpubkey> OP_CHECKSIGVERIFY <remote_htlcpubkey> OP_CHECKSIG
 - htlc_success: OP_SIZE 32 OP_EQUALVERIFY OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY <remote_htlcpubkey> OP_CHECKSIGVERIFY 1 OP_CHECKSEQUENCEVERIFY

In order to spend a offered HTLC, via either script path, an inclusion_proof must be specified along with the control block. This inclusion_proof is simply the tap_leaf hash of the path *not* taken.

Accepted HTLCs

Accepted HTLCs inherit a similar format:

- OP_1 accepted_htlc_key
- where:
 - accepted_htlc_key = revocation_pubkey + tagged_hash("TapTweak", revocation_pubkey || htlc_script_root)
 - htlc_script_root = tapscript_root([htlc_timeout, htlc_success])
 - htlc_timeout: <remote_htlcpubkey> OP_CHECKSIGVERIFY 1 OP_CHECKSEQUENCEVERIFY OP_VERIFY <cltv_expiry> OP_CHECKLOCKTIMEVERIFY
 - htlc_success: OP_SIZE 32 OP_EQUALVERIFY OP_HASH160 <RIPEMD160(payment_hash)> OP_EQUALVERIFY <local_htlcpubkey> OP_CHECKSIGVERIFY <remote_htlcpubkey> OP_CHECKSIG

Note that there is no OP_CHECKSEQUENCEVERIFY in the offered HTLC's timeout path nor in the accepted HTLC's success path. This is not needed here as these paths require a remote signature that commits to the sequence, which needs to be set at 1.

In order to spend an accepted HTLC, via either script path, an inclusion_proof must be specified along with the control block. This inclusion_proof is simply the tap_leaf hash of the path *not* taken.

HTLC Second Level Transactions

For the second level transactions, we retain the existing structure: a 2-of-2 multi-spend going to a CLTV delayed output. Like the normal HTLC transactions, we also place the revocation key as the internal key, allowing easy breach sweeps with each party only retaining the specific script root.

An HTLC-Success or HTLC-Timeout is a one-in-one-out transaction signed using the SIGHASH_SINGLE | SIGHASH_ANYONECANPAY flag. These transactions always have *zero* fees attached, forcing them to be aggregated with each other and a change input.

HTLC-Success Transactions

A HTLC-Success transaction has the following structure:

- version: 2
- locktime: 0
- input:
 - txid: commitment_tx
 - vout: htlc_index
 - sequence: 1
 - witness: <remotehtlcsig> <localhtlcsig> <preimage> <htlc_success_script> <control_block>
- output:
 - value: htlc_value
 - script:
 - OP_1 htlc_success_key

- where:
 - `htlc_success_key = revocation_pubkey + tagged_hash("TapTweak", revocation_pubkey || htlc_script_root)`
 - `htlc_script_root = tapscript_root([htlc_success])`
 - `htlc_success:`

```
<local_delayedpubkey> OP_CHECKSIGVERIFY
<to_self_delay> OP_CHECKSEQUENCEVERIFY
```

HTLC-Timeout Transactions

A HTLC-Timeout transaction has the following structure:

- version: 2
- locktime: `cltv_expiry`
- input:
 - txid: `commitment_tx`
 - vout: `htlc_index`
 - sequence: 1
 - witness: `<remotehtlcsig> <localhtlcsig> <htlc_timeout_script> <control_block>`
- output:
 - value: `htlc_value`
 - script:
 - `OP_1 htlc_timeout_key`
 - where:
 - `htlc_timeout_key = revocation_pubkey + tagged_hash("TapTweak", revocation_pubkey || htlc_script_root)`
 - `htlc_script_root = tapscript_root([htlc_timeout])`
 - `htlc_timeout:`

```
<local_delayedpubkey> OP_CHECKSIGVERIFY
<to_self_delay> OP_CHECKSEQUENCEVERIFY
```

Appendix

Footnotes

Test Vectors

All test vectors are derived deterministically from a single 32-byte seed:

`000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f`. Private keys are produced via `SHA256(seed || label)` where `label` is a fixed ASCII string for each key (e.g. "local-funding", "remote-funding", etc.). The per-commitment point is derived by computing `SHA256(seed || "local-per-commit-secret")` and using the result as the per-commitment secret scalar, then multiplying by the generator to obtain the per-commitment point.

The vectors cover two areas:

1. **Script vectors:** the decomposed tapscript trees for every output type (funding, to_local, to_remote, anchors, offered/accepted HTLCs on both local and remote commits, and second-level HTLC transactions). Each entry includes the raw leaf scripts, leaf hashes, tapscript root, internal key, output key, and the resulting pkScript.
2. **Transaction vectors:** full serialized commitment transactions and their HTLC resolution transactions for three scenarios — a simple commitment with no HTLCs, a commitment with five untrimmed HTLCs, and a commitment with the same HTLCs at a higher fee rate causing some to be trimmed.

Channel parameters used: funding_amount = 10,000,000 sat, dust_limit = 354 sat (taproot dust), csv_delay = 144 blocks, commit_height = 42, fee_per_kw as specified per test case.

Key Derivation Labels

The following labels are used with `SHA256(seed || label)` to derive each private key:

Key	Label
local_funding_privkey	"local-funding"
remote_funding_privkey	"remote-funding"
local_payment_basepoint_secret	"local-payment-basepoint"
remote_payment_basepoint_secret	"remote-payment-basepoint"
local_delayed_payment_basepoint_secret	"local-delayed-payment-basepoint"
remote_revocation_basepoint_secret	"remote-revocation-basepoint"
local_htlc_basepoint_secret	"local-htlc-basepoint"
remote_htlc_basepoint_secret	"remote-htlc-basepoint"
local_per_commit_secret	"local-per-commit-secret"

MuSig2 Nonce Generation

MuSig2 partial signatures require deterministic nonces so that test vectors are reproducible across implementations. The signing randomness for each party's JIT nonces is derived as `SHA256(seed || label)` using:

Party	Label
Local signer	"local-signing-rand"
Remote signer	"remote-signing-rand"

The resulting 32-byte hash is used as the nonce randomness input to the MuSig2 signing session.

Because different MuSig2 implementations (e.g., libsecp256k1-zkp vs btcd) may produce different nonces from the same randomness input, the test vectors include both the raw secret nonces and the derived public nonces so that implementations can inject them directly rather than re-derive them.

Each transaction test case includes the following nonce fields:

Field	Size	Description
local_nonce	66 bytes	Local party's public nonce for this commitment tx
remote_nonce	66 bytes	Remote party's public nonce for this commitment tx
local_sec_nonce	97 bytes	Local party's secret nonce (two 32-byte scalars + 33-byte pubkey)
remote_sec_nonce	97 bytes	Remote party's secret nonce (two 32-byte scalars + 33-byte pubkey)

All nonces in a given test case correspond to the **same commitment transaction** (the local party's commitment, as produced by ForceClose). Specifically:

- `local_nonce` / `local_sec_nonce`: the verification nonce that the local party contributes to the MuSig2 session for their own commitment. This nonce is generated during the initial nonce exchange (or nonce rotation after a prior commitment update) and sent to the remote party.
- `remote_nonce` / `remote_sec_nonce`: the JIT signing nonce that the remote party generates when producing their partial signature on the local party's commitment transaction.

To replay the MuSig2 signing from the test vectors:

1. Aggregate the two public nonces using `NonceAgg` from BIP-327.
2. Use each party's secret nonce with `Sign` to reproduce their partial signature.
3. Combine both partial signatures to obtain the final 64-byte Schnorr signature, which should match the commitment transaction witness.

Commitment Transaction Signatures

The `remote_partial_sig` field in each transaction test case is a 32-byte MuSig2 partial signature scalar (hex-encoded). This is the remote party's contribution to the combined Schnorr signature on the commitment transaction. The commitment transaction witness contains the final aggregated 64-byte Schnorr signature produced by combining both parties' partial signatures.

HTLC Resolution Transactions

Each `htlc_descs` entry describes a second-level HTLC transaction that spends from one of the commitment transaction's HTLC outputs. The entries are ordered by their corresponding commitment transaction output index (BIP 69 ordering), which matches the order of `htlc_signature` messages exchanged during the commitment signing protocol.

The `remote_partial_sig_hex` for HTLC transactions is a 64-byte Schnorr signature (not a MuSig2 partial sig), since HTLC second-level transactions use a regular tapscript spend path rather than a keyspend MuSig2 path. These Schnorr signatures use BIP-340 deterministic nonce derivation with zero auxrand (32 zero bytes). Implementations must use the same nonce derivation to reproduce matching signatures — passing no auxiliary randomness (or explicitly zero bytes) to the BIP-340 `Sign` algorithm.

The `resolution_tx_hex` contains the fully serialized second-level transaction with a complete witness stack. For HTLC-success transactions, the witness layout is:

```
<remote_htlcsig> <local_htlcsig> <payment_preimage> <htlc_success_script> <control_block>
```

For HTLC-timeout transactions, the witness layout is:

```
<remote_htlcsig> <local_htlcsig> <htlc_timeout_script> <control_block>
```

HTLC Trimming

Taproot channels use zero-fee HTLC transactions, meaning the HTLC output value on the commitment transaction equals the HTLC amount directly (no fee is deducted from the output). As a result, HTLC trimming depends solely on whether the HTLC amount falls below the `dust_limit_satoshis`, not on the fee rate. The “commitment tx with some HTLCs trimmed” test case uses `dust_limit = 2500 sat` to trim the three smallest test HTLCs (1000, 2000, and 2000 sats), leaving only the 3000 and 4000 sat HTLCs on the commitment transaction.

Implementations SHOULD regenerate these vectors from the seed using the derivation method described above and verify that all intermediate values (keys, scripts, leaf hashes, transactions) match.

```
{
  "params": {
    "seed": "000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f",
    "funding_amount_satoshis": 10000000,
    "dust_limit_satoshis": 354,
    "csv_delay": 144,
    "commit_height": 42,
    "nums_point": "02dca094751109d0bd055d03565874e8276dd53e926b44e3bd1bb6bf4bc130a279",
    "keys": {
      "local_funding_privkey":
        "20ae2d254ab29afd3dcbf8744a5b88d06070f55a4bd5532483a093ac4db91277",
      "local_funding_pubkey":
        "03b7203dec7c13896b66ff1f58b24f84458c441720a12b5a57426397e22f0a8c78b",
      "remote_funding_privkey":
        "f0c5500a9dbd7cdcd46ced7bdeb937d4dcfb90f9b9357626e7ee54ab024c3df0",
      "remote_funding_pubkey":
        "02956e6845a6f346f97c5e028c0f8ab38a76b0124fd7184deab60f682b3e657fdb",
      "local_payment_basepoint_secret":
        "277975b5b081a9cbc4834e066d7bb494e4fde4f7637257dd3d312a0ae7cb7754",
      "local_payment_basepoint":
        "03955b6085296cbd2447a1dde0f7e273e19b83e83de1814993b1517aaf193b7f33",
      "remote_payment_basepoint_secret":
        "f1cd3a5ca44b52baf4eacb849fbf06e75aace97477b8bfe31d2b814dbbb562b1",
      "remote_payment_basepoint":
        "03595f2ef2a51d2250a21077dbea4a7fc3ce550f10676996bf63719e2a71d1f4c9",
      "local_delayed_payment_basepoint_secret":
        "83ccf0b638c514db5ebefdc6cbf901505e2bb20edb2bb7248ce1a51523325f9b",
      "local_delayed_payment_basepoint":
        "02ae68d8ff4c59864c03a42bbff6c07f9ae18047e0daa9bc40d07c410f9a0f7899",
      "remote_revocation_basepoint_secret":
        "36c4175b91cff9731a63d1472b5b1c4cf3e7b688e87d5fb806b2e8350484e68d",
      "remote_revocation_basepoint":
        "02c354121ef71922b5cb32fa685c08ac0014b558f96e28f383c45eb28b7da264c3",
      "local_htlc_basepoint_secret":
        "786eb5024e4851bea3ddc6e40036c81b1efcf50eed440eedefe5245bde6fc14",
```


[illegible]

```
,
"local_anchor": {
  "scripts": {
    "sweep": "60b2"
  },
  "leaf_hashes": {
    "sweep": "2b88a8f3f52386d61d5b3f2d822df659c35214d7360ed05352ad7ddc1ab03912"
  },
  "tapscript_root": "2b88a8f3f52386d61d5b3f2d822df659c35214d7360ed05352ad7ddc1ab03912",
  "internal_key": "0315ec0138eb42f1ab4603042123988d53c854e89d1d87aa4dbb97a57482029c05",
  "output_key": "02f67ab012701705f3203d132f909a6810ef18c5da4c11d986cb50818803b8344e",
  "pkscript": "5120f67ab012701705f3203d132f909a6810ef18c5da4c11d986cb50818803b8344e"
},
"remote_anchor": {
  "scripts": {
    "sweep": "60b2"
  },
  "leaf_hashes": {
    "sweep": "2b88a8f3f52386d61d5b3f2d822df659c35214d7360ed05352ad7ddc1ab03912"
  },
  "tapscript_root": "2b88a8f3f52386d61d5b3f2d822df659c35214d7360ed05352ad7ddc1ab03912",
  "internal_key": "03595f2ef2a51d2250a21077dbea4a7fc3ce550f10676996bf63719e2a71d1f4c9",
  "output_key": "021249c50576fdf914caa14f9221370b986df520bdbc73f57d5056a86ee03e5ac4",
  "pkscript": "51201249c50576fdf914caa14f9221370b986df520bdbc73f57d5056a86ee03e5ac4"
},
"offered_htlc_local_commit": {
  "scripts": {
    "success":
      "82012088a914b8bcb07f6344b42ab04250c86a6e8b75d3fdbbc688202deba21cf03c42362c9f912094f62ba045a040a2060",

    "timeout":
      "2071e82ef65d5c667159036bfcf662cac2f6c41e38323d148bbbd00fdcd923739ead202deba21cf03c42362c9f912094f62"
  },
  "leaf_hashes": {
    "success": "cd4b7ba74d132998f2bcea85f76082f5018e614c86f27f2631b6569c4914320f",
    "timeout": "dd0bd08b3df902c399f5493a682f6c50c476c89e233ba454e89a234d2d16ffe3"
  },
  "tapscript_root": "f36c8bd45002c5264cfce9944211e7bc6ea974a6b90cf99a87812d18acf28a2a",
  "internal_key": "03d4c77088d346bce67c13bbbf82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
  "output_key": "033e5c3be9f4ce7ae07c28ad5e0eb0ab617c06eeb82b8d6ef10a5bf561848df5f0",
  "pkscript": "51203e5c3be9f4ce7ae07c28ad5e0eb0ab617c06eeb82b8d6ef10a5bf561848df5f0"
},
"offered_htlc_remote_commit": {
  "scripts": {
    "success":
      "82012088a914b8bcb07f6344b42ab04250c86a6e8b75d3fdbbc688202deba21cf03c42362c9f912094f62ba045a040a2060",

    "timeout":
      "2071e82ef65d5c667159036bfcf662cac2f6c41e38323d148bbbd00fdcd923739ead202deba21cf03c42362c9f912094f62"
  },
  "leaf_hashes": {
    "success": "cd4b7ba74d132998f2bcea85f76082f5018e614c86f27f2631b6569c4914320f",
    "timeout": "dd0bd08b3df902c399f5493a682f6c50c476c89e233ba454e89a234d2d16ffe3"
  },
}
```

```
"tapscript_root": "f36c8bd45002c5264cfce9944211e7bc6ea974a6b90cf99a87812d18acf28a2a",
"internal_key": "03d4c77088d346bce67c13bbbf82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
"output_key": "033e5c3be9f4ce7ae07c28ad5e0eb0ab617c06eeb82b8d6ef10a5bf561848df5f0",
"pkscript": "51203e5c3be9f4ce7ae07c28ad5e0eb0ab617c06eeb82b8d6ef10a5bf561848df5f0"
},
"accepted_htlc_local_commit": {
  "scripts": {
    "success":
      "82012088a914b8bcb07f6344b42ab04250c86a6e8b75d3fdbbc688202deba21cf03c42362c9f912094f62ba045a040a2060",

    "timeout":
      "2071e82ef65d5c667159036bfcf662cac2f6c41e38323d148bbbd00fdcd923739ead51b26902f401b1"
  },
  "leaf_hashes": {
    "success": "69192ca730d4480044ade8741b8bd0845a32880aebaf58bc6f9186f8d2be8cbf",
    "timeout": "4da43c795365bf757ed1e9656d12ea744b4cf52b01719a3ea94e6569115623f0"
  },
  "tapscript_root": "1a990caa4bb0ed41ceb19e7466fcea5d9b31e3da968f348f6223201c5831d0a3",
  "internal_key": "03d4c77088d346bce67c13bbbf82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
  "output_key": "029aadbdd9aff986e5ea086cf53ae062972d33d0a5c7f5fb986dafec7fa6d7e6ea",
  "pkscript": "51209aadbdd9aff986e5ea086cf53ae062972d33d0a5c7f5fb986dafec7fa6d7e6ea"
},
"accepted_htlc_remote_commit": {
  "scripts": {
    "success":
      "82012088a914b8bcb07f6344b42ab04250c86a6e8b75d3fdbbc688202deba21cf03c42362c9f912094f62ba045a040a2060",

    "timeout":
      "2071e82ef65d5c667159036bfcf662cac2f6c41e38323d148bbbd00fdcd923739ead51b26902f401b1"
  },
  "leaf_hashes": {
    "success": "69192ca730d4480044ade8741b8bd0845a32880aebaf58bc6f9186f8d2be8cbf",
    "timeout": "4da43c795365bf757ed1e9656d12ea744b4cf52b01719a3ea94e6569115623f0"
  },
  "tapscript_root": "1a990caa4bb0ed41ceb19e7466fcea5d9b31e3da968f348f6223201c5831d0a3",
  "internal_key": "03d4c77088d346bce67c13bbbf82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
  "output_key": "029aadbdd9aff986e5ea086cf53ae062972d33d0a5c7f5fb986dafec7fa6d7e6ea",
  "pkscript": "51209aadbdd9aff986e5ea086cf53ae062972d33d0a5c7f5fb986dafec7fa6d7e6ea"
},
"second_level_htlc_success": {
  "scripts": {
    "success":
      "2015ec0138eb42f1ab4603042123988d53c854e89d1d87aa4dbb97a57482029c05ad029000b2"
  },
  "leaf_hashes": {
    "success": "dbf0400e9c7c57f30b6ad0b0677e396b5a002cbf050d873c8925b966048e6a62"
  },
  "tapscript_root": "dbf0400e9c7c57f30b6ad0b0677e396b5a002cbf050d873c8925b966048e6a62",
  "internal_key": "03d4c77088d346bce67c13bbbf82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
  "output_key": "02df20bcec43daa75161f7d013254e401812e0fee8bc3369220b6a33672fc18ba0",
  "pkscript": "5120df20bcec43daa75161f7d013254e401812e0fee8bc3369220b6a33672fc18ba0"
},
"second_level_htlc_timeout": {
  "scripts": {
    "success":
      "2015ec0138eb42f1ab4603042123988d53c854e89d1d87aa4dbb97a57482029c05ad029000b2"
```

```

    },
    "leaf_hashes": {
      "success": "dbf0400e9c7c57f30b6ad0b0677e396b5a002cbf050d873c8925b966048e6a62"
    },
    "tapscript_root": "dbf0400e9c7c57f30b6ad0b0677e396b5a002cbf050d873c8925b966048e6a62",
    "internal_key": "03d4c77088d346bce67c13bbb82ca112588f4b1c9595a1f8af3be9b2f95a109a0",
    "output_key": "02df20bcec43daa75161f7d013254e401812e0fee8bc3369220b6a33672fc18ba0",
    "pkscript": "5120df20bcec43daa75161f7d013254e401812e0fee8bc3369220b6a33672fc18ba0"
  },
  },
  "transactions": [
    {
      "name": "simple commitment tx with no HTLCs",
      "local_balance_msat": 7000000000,
      "remote_balance_msat": 3000000000,
      "fee_per_kw": 15000,
      "htlcs": null,
      "local_sec_nonce":
        "22a453171ba4a634da1addcf660d63d8e23fb63169a2a7206f4e23290e4cc59bb3c3011fe3c31cb4f1192a2df56c2e52350",
      "remote_sec_nonce":
        "ccdaea6955c7bd9fcb082dc67e32809a5bdd3a3edf770cbb990116c45f4b006e4b56aa81f8e555d6bd1906b159af16d0ac3",
      "local_nonce":
        "025f2272ea289c5fe9d52d411f5a50a6d4882341bf0ecb201d5675850a2ba0b09d025bc489cf67752134ba81f8d7f1146d7",
      "remote_nonce":
        "02d324627074522af8cf4287caf1e073a3493550b99aed2697e58f476ec402e272039c25fc616207e15917b7145cefc4c9",
      "remote_partial_sig": "3fa93659d4c2d590eadbd422595a37597ed58607e026a91f4e6e19329134a931",
      "expected_commitment_tx_hex":
        "020000000001015474cba49124ab0c4327c244bb2907059585c4af3fa5f3469701534120fec0170000000000c5fe178004",
      "htlc_descs": null
    },
    {
      "name": "commitment tx with five HTLCs untrimmed",
      "local_balance_msat": 6988000000,
      "remote_balance_msat": 3000000000,
      "fee_per_kw": 644,
      "htlcs": [
        {
          "incoming": true,
          "amount_msat": 1000000,
          "expiry": 500,
          "preimage": "0000000000000000000000000000000000000000000000000000000000000000"
        },
        {
          "incoming": true,
          "amount_msat": 2000000,
          "expiry": 501,
          "preimage": "0101010101010101010101010101010101010101010101010101010101010101"
        },
        {
          "incoming": false,
          "amount_msat": 2000000,

```

```

    "expiry": 502,
    "preimage": "0202020202020202020202020202020202020202020202020202020202020202"
  },
  {
    "incoming": false,
    "amount_msat": 3000000,
    "expiry": 503,
    "preimage": "0303030303030303030303030303030303030303030303030303030303030303"
  },
  {
    "incoming": true,
    "amount_msat": 4000000,
    "expiry": 504,
    "preimage": "0404040404040404040404040404040404040404040404040404040404040404"
  }
],
"local_sec_nonce":
  "22a453171ba4a634da1addcf660d63d8e23fb63169a2a7206f4e23290e4cc59bb3c3011fe3c31cb4f1192a2df56c2e52350",
"remote_sec_nonce":
  "8c40a8ab26f1bd7d58010e991c72b0f102df6d0f050df28f56cefdafb9d06479958ea0a72180184e338b1b77fdf20690d5c",
"local_nonce":
  "025f2272ea289c5fe9d52d411f5a50a6d4882341bf0ecb201d5675850a2ba0b09d025bc489cf67752134ba81f8d7f1146d7",
"remote_nonce":
  "038a4018a074b5ddfc1424551e871bd739259c4e209f47177eb67c70cfbdb1e57c03c75684ad3a42d8c86a3eddb83b8160c",
"remote_partial_sig": "46efde50f08c128aa6472bbd50ea156fb9bde7b013f5e042f700729c94053613",
"expected_commitment_tx_hex":
  "0200000000001015474cba9124ab0c4327c244bb2907059585c4af3fa5f3469701534120fec0170000000000c5fe1780094",
"htlc_descs": [
  {
    "remote_partial_sig_hex":
      "fc97f7c9fb97e1e48792b0ed174704cd98c886d368ede5adeb6288f0b350c8f88d076488973d0656da72d072e03eb9eb8b31",
    "resolution_tx_hex":
      "020000000000101ec4c0a34c981864f9badcb8383bbe42ec6b32e68c2aa1a7c7c2e8422adde673f02000000000100000001"
  },
  {
    "remote_partial_sig_hex":
      "c99d8d1ca721d1d4b9796cde49698fd46bb3fadd2bc7215de291ae9bd85cd9b2ab9bf40e36c752b698098b0986abe92b3c",
    "resolution_tx_hex":
      "020000000000101ec4c0a34c981864f9badcb8383bbe42ec6b32e68c2aa1a7c7c2e8422adde673f03000000000100000001"
  },
  {
    "remote_partial_sig_hex":
      "b361ed8b70a09f4128fe610db649abcaafbc9b8600136e6f1b9735b4781cd65681faa32c29fef2a1485cc569c655b31b8eb",
    "resolution_tx_hex":
      "020000000000101ec4c0a34c981864f9badcb8383bbe42ec6b32e68c2aa1a7c7c2e8422adde673f04000000000100000001"
  }
],

```

[illegible]

```

"local_sec_nonce":
  "22a453171ba4a634da1addcf660d63d8e23fb63169a2a7206f4e23290e4cc59bb3c3011fe3c31cb4f1192a2df56c2e52350

"remote_sec_nonce":
  "8c5bd39820b3e20d65bf72e530078bf5e8ba057c1654d0a15778a0d3225e0233eb8d7b1d3574e40e57a2cd12dc3b33193c0

"local_nonce":
  "025f2272ea289c5fe9d52d411f5a50a6d4882341bf0ecb201d5675850a2ba0b09d025bc489cf67752134ba81f8d7f1146d7

"remote_nonce":
  "03fd9fa808377737b105f7df362ed513e3946f2bb49dfbca5c2ce2be138ff0607502e4c73701eae82afa7a01993f6232164

"remote_partial_sig": "3e454598e0661188da0e4cf1b806b13c627adb7ab38bab27418cc976769771c3",
"expected_commitment_tx_hex":
  "020000000001015474cba49124ab0c4327c244bb2907059585c4af3fa5f3469701534120fec0170000000000c5fe1780064

"htlc_descs": [
  {
    "remote_partial_sig_hex":
      "7058caa9a075344e095d95736bb6a00c5b7259f439e60e7e1c6cd641a20d6a25c13d2930ebe51a69ffea78daa12b0723717

    "resolution_tx_hex":
      "020000000001018c47e10e0d210da9e50d73b1adda7a598f79ced25806dd0fa6807bb78780dbbb02000000000100000001b

  },
  {
    "remote_partial_sig_hex":
      "818132325cc01e441876615f30c3df5c27df1722db49005dff3856226a41f34c776f98a35dd74fe2f863ea23ac611e06d0a

    "resolution_tx_hex":
      "020000000001018c47e10e0d210da9e50d73b1adda7a598f79ced25806dd0fa6807bb78780dbbb03000000000100000001a

  }
]
}

```

Acknowledgements

The commitment and HTLC scripts are heavily based off of t-bast’s [Lightning Is Getting Taprooty Scriptless-Scripty](#) document.

AJ Towns proposed a more ambitious [“one shot” upgrade which includes PTLCS](#) and a modified revocation scheme which this proposal takes inspiration from.

Arik and Wilmer Paulino for suggesting the “JIT nonce” approach used in this specification when sending musig2 partial signatures.

Conclusion

Thank you for reading the Lightning Book of BOLTs. The specs will keep moving — new feature bits, new scripts, new ways to pay — and the canonical text lives in [lightning/bolts](#), not in this compilation.

I didn't write the protocol. The credit for every message type, every script, and every hard-won MUST belongs to the original BOLT authors and the people who still review pull requests there. This book is a reading order and a set of signposts.

If you implement something, or you find a contradiction, take it upstream. That's how Lightning stays one network instead of a pile of dialects.

See you on the next commit.